Architecting Emergent Realism: A Framework for Advanced Text-to-Image Synthesis in Dynamic Anime Contexts

Abstract

This report outlines a comprehensive framework for advancing text-to-image (T2I) synthesis, specifically targeting the complex, dynamic, and physics-defying visual language of franchises like *Dragon Ball Z* (DBZ). Traditional T2I models, reliant on correlational mimicry, fail to capture the spatio-temporal coherence of character transformations and the volumetric nature of energy constructs. We propose a paradigm shift from static prompting to the design of dynamic, agentic architectures. This framework is built on five strategic pillars: (1) **Causal Control**, leveraging Causal AI principles and Iterative Semantic Deepening (ISD) to enforce logical consistency in generative sequences; (2) **Perceptual Fidelity Anchoring**, integrating physics-based rendering concepts (BRDFs, subsurface scattering) and volumetric primitives (NeRF/3DGS) to simulate, rather than imitate, phenomena like ki auras; (3) **Reflexive Meta-Learning**, employing autonomous pipelines with Auto-CoT and GFlowNets to optimize prompt diversity and prevent character mode collapse; (4) **Multi-Modal Conditioning**, using emotional and haptic inputs to achieve epistemic debiasing and enrich character representation; and (5) **Emergent Technique Discovery**, applying Logic-of-Thought (LoT) to mine non-AI algorithms for novel latent space navigation strategies. By synthesizing these pillars, we present an actionable R&D roadmap for creating next-generation T2I systems capable of true simulative control and emergent artistic expression, while critically examining the associated computational, environmental, and ethical trade-offs.

Introduction: Beyond Mimicry—Towards Causal and Simulative Control in Generative Art

The field of text-to-image (T2I) synthesis has achieved remarkable success, with latent

diffusion models now capable of generating high-fidelity, complex images from natural language prompts.[1] These models, which typically combine a language model for prompt encoding with a generative image model operating in a compressed latent space, have demonstrated an impressive ability to produce novel visual content by learning from vast datasets of image-text pairs.[3] However, the generation of static, descriptive scenes represents a largely solved problem. The next frontier of generative art lies in the synthesis of dynamic, temporally coherent, and physically plausible (or plausibly

implausible) visual narratives.

This challenge is exemplified by the aesthetic of popular anime franchises such as *Dragon Ball Z*. The visual language of DBZ is not defined by static portraits but by *dynamics*: characters undergo multi-stage transformations, accumulate and discharge immense energy, and engage in combat featuring physics-defying attacks that must remain visually and logically consistent across multiple interpretations or sequential frames. Standard diffusion models falter in this domain because their core mechanism is correlational, not causal. They learn statistical associations between text tokens and visual patterns in their training data, effectively creating a sophisticated collage of learned concepts.[3] When prompted for "Goku firing a Kamehameha," a model may generate a visually appealing image that contains the constituent elements—Goku, a blue energy beam, a dramatic pose—but fails to establish a coherent causal link between them. The energy source may not align with the character's hands, the beam's trajectory might be physically unsound, and its impact on the environment may be disconnected from its apparent power. This disconnect between user intent and model output, the "intent gap," reveals the fundamental limitation of mimicry-based generation.

This report argues for a paradigm shift: from treating the prompt as a static description to architecting an **agentic system** that actively reasons about, simulates, and refines the generative process. This approach moves beyond simple prompt engineering to the design of intelligent workflows that imbue the T2I pipeline with a rudimentary understanding of causality, physics, and compositional logic. Such a system would not merely correlate "Super Saiyan" with "gold hair" but would model the transformation as a causal sequence: a surge in power level *causes* the manifestation of a golden aura, which in turn *causes* the character's hair to change color and stand on end.

To achieve this, this report proposes a unified R&D framework built upon five strategic pillars, each designed to address a core deficiency in current T2I models and contribute to the goal of dynamic, controllable, and coherent DBZ synthesis:

1. **Causal Control and Compositionality:** Enforcing logical consistency in transformations and actions through the principles of Causal AI and structured, iterative prompting.
2. **Perceptual Fidelity Anchoring:** Simulating, rather than imitating, visual phenomena by integrating physics-based rendering principles and volumetric geometry.
3. **Reflexive Meta-Learning:** Automating the discovery of diverse and effective prompts

through autonomous agentic pipelines to ensure character consistency and stylistic variety.
4. **Multi-Modal Conditioning:** Expanding the input space beyond text to include emotional and other sensory data, thereby mitigating cultural biases inherent in training data and enriching character representation.
5. **Emergent Technique Discovery:** Leveraging advanced reasoning frameworks to mine non-AI computational fields for novel algorithms that can unlock emergent dynamic effects in latent space.

By synthesizing these five pillars, it is possible to construct a roadmap toward a new class of T2I models that function less like digital paintbrushes and more like controllable "world engines," capable of generating not just images, but dynamic and coherent visual realities.

I. Causal Control and Compositionality for Dynamic Transformations

The primary obstacle in generating complex, multi-stage events like DBZ transformations is the inherent inability of standard generative models to understand cause and effect. These models excel at identifying and reproducing correlations present in their training data but lack a structural model of *why* certain visual elements co-occur.[4] To bridge this intent gap, a shift is required from correlational generation to a causally-informed compositional approach, where prompts are structured to enforce a logical sequence of events.

Foundations of Causal AI in Generative Models

Causal AI, a field focused on modeling true cause-and-effect relationships rather than spurious correlations, provides the theoretical underpinnings for this shift.[5] As asserted by computer scientist Judea Pearl, the absence of causal understanding is a fundamental roadblock to achieving human-level artificial intelligence.[6] In the context of T2I synthesis, this manifests as a failure to produce logically coherent outputs for prompts describing processes or sequences. For example, a model might generate an image of a "broken glass on the floor" correctly, but struggle with the prompt "a glass falling and breaking," often producing disjointed images of a falling glass and a separate broken one.

Applying causal principles to T2I involves structuring the generation process to follow a directed acyclic graph (DAG), where nodes represent states or visual elements and edges

represent causal dependencies.[5] For a DBZ transformation, this graph might be defined as:

Triggering Event (e.g., intense rage) → Power Level Increase → Aura Manifestation → Physiological Changes (e.g., hair color, muscle mass). By encoding this causal chain into the prompting strategy, the model is guided to generate the scene sequentially, ensuring that each effect logically follows from its cause. This approach is supported by recent findings which demonstrate mathematically that any AI agent capable of generalizing beyond its training distribution must have implicitly learned a causal model of its environment.[6] Instead of relying on the model to infer this structure, a causal prompting framework makes it explicit, thereby increasing transparency, reliability, and control over the generative process.[4]

The Iterative Semantic Deepening (ISD) Workflow

The practical implementation of causal control can be achieved through a methodology known as **Iterative Semantic Deepening (ISD)**. Inspired by multi-stage refinement frameworks in T2I synthesis, ISD reframes generation not as a single action but as a sequence of prompts, where each subsequent prompt adds a layer of causal detail and refines the output of the previous stage.[7] This iterative process allows for the construction of complex scenes that would be impossible to specify in a single, monolithic prompt.

A key innovation is the automation of this loop through an agentic system. The **VisualPrompter** framework offers a powerful, training-free model for this purpose.[9] VisualPrompter consists of two core modules that create a self-correcting feedback loop. First, a

SElf-REflection (SERE) module, powered by a Vision-Language Model (VLM), analyzes the image generated in the current iteration. It deconstructs the input prompt into atomic concepts and verifies whether each concept is visually present in the output.[9] If it detects a "missing concept" (e.g., the prompt specified a "crackling, electric aura" but the generated aura is smooth and flame-like), this feedback is passed to the second module. The

Target-Specific Prompt Optimization (TSPO) module then uses a Large Language Model (LLM) to revise and enrich the prompt, specifically targeting the missing concepts while preserving the elements that were generated correctly.[9]

This synthesis of Causal AI, ISD, and VisualPrompter transforms prompt engineering from a manual, artistic endeavor into a systematic, programmatic form of scene construction. The causal graph provides the high-level logical blueprint, while the ISD-VisualPrompter loop executes this blueprint step-by-step, using the model's own visual feedback to debug and refine the generation at each stage. This creates a powerful *model-in-the-loop* system where

the AI actively participates in its own creative process, moving beyond the traditional *user-in-the-loop* paradigm that relies on human trial and error.

An application of this workflow to a novel, complex DBZ character, such as "Cosmic-Eclipse Vegeta," illustrates its power. A simple prompt would likely fail, but an ISD approach can build the character layer by layer, with each stage causally dependent on the last, as detailed in Table 1.

Stage	Prompt Fragment	Causal Justification	Target Visual Element	Associated Technologies
1: Base Intent	RAW photo, Cosmic-Eclipse Vegeta, base form, standing on a barren moon.	Establishes the core character identity and setting without dynamic effects.	Character model, basic anatomy, environment.	Standard Diffusion [2]
2: Causal Layering	...musculature tenses as he begins to power up. CAUSAL_TRIGGER: Power level increases to 50%.	Links the physical action of "powering up" to the initial manifestation of energy.	Muscle definition, initial particle effects.	Causal AI [6], ControlGAN [11]
3: Dynamic Effects	...as power level reaches 50%, a faint, dark purple aura with gravitational lensing effects begins to form, bending starlight around his body.	Defines the specific visual effect that is *caused* by reaching the 50% power threshold.	Volumetric aura, light distortion, spectral properties.	Physics-based rendering (See Sec. II)

4: Coherence Refinement	...at 100% power, the aura solidifies into a miniature accretion disk. Hair remains black but spikes with gravitational distortion.	Ensures spatio-temporal coherence; the final state is a logical culmination of the process, with specific constraints to prevent unwanted changes (e.g., hair color).	Final aura form, hair physics, environmental interaction.	VisualPrompter (SERE/TSPO) [9]

This structured, causal approach ensures that the final image is not merely a collection of keywords but a coherent depiction of a dynamic process, bridging the gap between a user's complex intent and the model's generative capabilities.

II. Perceptual Fidelity Anchoring via Physics-Informed Latent Space

To generate truly convincing depictions of DBZ's iconic energy auras and blasts, T2I models must transcend artistic mimicry and engage in a form of physical simulation. The visual "realism" of these fantastical phenomena is governed by an intuitive, internal logic rooted in the physics of light, volume, and material interaction. Current models, which typically render auras as simple 2D glows or textures, fail to capture their volumetric depth and dynamic interaction with the environment. A physics-informed approach anchors the generation process in the principles of optics and geometry, instructing the model to simulate these effects rather than merely referencing stylized depictions from its training data.

From Mimicry to Simulation: Integrating PBR Concepts

Physically Based Rendering (PBR) is a methodology in computer graphics that models the flow of light in the real world to achieve photorealism.[12] By integrating core PBR concepts into T2I prompts, it becomes possible to guide diffusion models toward generating more physically

plausible and visually complex energy constructs. While models like Stable Diffusion do not have an explicit physics engine, their latent space is rich enough to respond to descriptive keywords that approximate these principles.[14] Key concepts include:

- **Bidirectional Reflectance Distribution Functions (BRDFs):** A BRDF is a function that describes how light reflects off an opaque surface.[16] It defines a material's properties, such as whether it is diffuse (like chalk) or specular (like a mirror).[17] In the context of DBZ, a ki blast is not just a colored light source; it is a volume of energy with material-like properties. A prompt could specify these properties, such as
  Vegeta's Final Flash, rendered as emissive plasma with an anisotropic BRDF, casting specular reflections on the surrounding cracked earth.[19] This instructs the model to treat the energy and the environment as interacting physical objects.
- **Subsurface Scattering (SSS):** This phenomenon describes how light penetrates the surface of a translucent object, scatters internally, and then exits at a different point.[21] It is what gives materials like skin, wax, and marble their characteristic soft look. For a ki aura, SSS is crucial for conveying a sense of volumetric depth and internal energy. A prompt for
  Goku's Super Saiyan aura with high subsurface scattering encourages the model to generate a glow that appears to emanate from within a volume, rather than sitting flatly on the character's surface.
- **Spectral Harmonization:** This involves ensuring that the color of light sources, their reflections, and the ambient light in a scene are consistent. For a character powering up, this means the color of their ki aura should realistically illuminate their body and the surrounding environment, with the correct color temperature and falloff.

Integrating Volumetric Primitives: NeRF and 3D Gaussian Splatting

To provide a robust geometric foundation for these physical simulations, the generative process can be anchored in 3D scene representations. **Neural Radiance Fields (NeRF)** and **3D Gaussian Splatting (3DGS)** are two state-of-the-art techniques for synthesizing novel views of a complex 3D scene from a set of 2D images.[22] NeRF represents a scene as a continuous volumetric function that maps 5D coordinates (spatial location and viewing direction) to color and density.[24] 3DGS, a more recent and often more efficient method, represents the scene as a collection of millions of 3D Gaussians, each with properties like position, orientation, color, and opacity.[25]

The primary advantage of these representations is their ability to ensure multi-view consistency. A standard 2D diffusion model might render Vegeta's armor or Goku's hair differently from varying angles, leading to inconsistencies. A model conditioned on a NeRF or

3DGS representation maintains a coherent 3D geometry.[27] Frameworks like

Text2NeRF have already demonstrated the feasibility of generating these 3D radiance fields directly from text prompts, creating a clear pathway for their integration into T2I workflows.[28] Furthermore, emerging research on

PrimDiffusion shows that diffusion models can be trained to operate directly on a set of *volumetric primitives*—small volumes containing radiance and kinematic information—marrying the expressive power of volumetric rendering with the efficiency of primitive-based systems.[29]

The Plausibility Oracle Agent

The complexity of orchestrating a physics-informed, geometrically consistent generation necessitates an automated validation system. We propose an agentic architecture centered on a **Plausibility Oracle**. This oracle is a specialized critic model, likely a VLM fine-tuned on a dataset of physically rendered scenes or annotated diagrams, tasked with evaluating the physical and logical coherence of a generated image.[31] This concept draws inspiration from the use of oracles in quantum algorithms to verify solutions [32] and the architecture of enterprise-level AI agents designed for complex, multi-tool tasks.[33]

The agentic loop would function as follows:

1. **Generation:** The T2I model generates an image based on a physics-informed prompt (e.g., Super Saiyan God Goku charging a ki blast, volumetric ray-traced ki with BRDF-compliant materials).
2. **Critique:** The Plausibility Oracle analyzes the output and provides structured feedback. For example: Critique: Implausible light interaction. The emissive ki blast does not cast corresponding specular highlights on the character's metallic boots. Aura density appears uniform, violating volumetric scattering principles.
3. **Refinement:** An LLM-based refiner processes the critique and revises the prompt for the next generative step. Refined Prompt:...add high-specularity BRDF to boots, ensure ki blast casts realistic highlights. Implement variable density gradient in aura volume to simulate realistic subsurface scattering.

This closed-loop system introduces a novel form of AI alignment: alignment not with subjective human preferences, but with the objective principles of physics and logic. While standard alignment techniques like RLHF aim to make models helpful and harmless, this approach pushes them to be *correct* within a defined physical or logical framework. This is a critical step for extending generative models to applications in scientific visualization,

engineering, and simulation, where plausibility is paramount. However, this pursuit of high fidelity comes with significant trade-offs. Both NeRF-based rendering and large-scale agentic loops are computationally expensive, raising serious concerns about energy consumption and environmental impact.[35] Recent 2025 analyses show that generative AI's energy footprint is already substantial, with some estimates comparing it to the consumption of entire nations.[37] Therefore, any strategy involving these techniques must be coupled with research into efficiency optimizations like model quantization and the use of smaller, specialized models to ensure sustainable scalability.[38]

Fidelity Anchoring Technique	Mechanism	DBZ Benefits	Prompt Example	Limitations & Trade-offs
Spectral Harmonization	Balances wavelengths of emitted and reflected light for color consistency.	Ensures a Super Saiyan's golden aura casts a believable golden light on their surroundings.	...golden hour lighting, cinematic bloom, spectral harmonization	Requires models trained on spectrally accurate data; may be ignored by simpler models.
Subsurface Scattering (SSS)	Simulates light penetrating a translucent volume before exiting.	Creates depth and internal glow in ki auras, making them feel like volumes of energy, not flat textures.	...Kamehameha energy ball with high subsurface scattering	Computationally intensive; can lead to blurry or noisy artifacts if not well-controlled.
BRDF Integration	Defines material properties for how surfaces reflect light from different angles.	Renders energy blasts that interact realistically with different materials (e.g., reflecting off metal, being absorbed by	...energy blast with emissive plasma BRDF, casting specular highlights	Highly model-dependent; requires a sophisticated latent understanding of material physics.

		cloth).		
NeRF/3DGS Primitives	Represents the scene as a continuous 3D volumetric field or a collection of 3D Gaussians.	Guarantees multi-view consistency for characters and attacks, preventing distortion during dynamic poses.	NeRF-synthesized Goku with 3D Gaussian splatting for ki blasts	Not natively supported by most T2I models; high training/inference cost; may require specialized architectures.[39]

III. Reflexive Meta-Learning: Autonomous Pipelines for Prompt Optimization

The process of discovering effective prompts remains a significant bottleneck in T2I synthesis. It is often a manual, labor-intensive task requiring extensive trial and error.[40] To overcome this, a reflexive meta-learning pipeline can be designed to autonomously explore the prompt space, optimizing for both aesthetic quality and character consistency. This involves creating an agentic system that generates, evaluates, and refines prompts in a continuous loop, thereby automating the creative discovery process.

Automating Prompt Discovery with Auto-CoT and GFlowNets

An effective autonomous pipeline requires two core capabilities: structured reasoning to formulate coherent prompt strategies, and diverse sampling to explore the vast space of possible prompts. These can be provided by two complementary technologies:

- **Automatic Chain-of-Thought (Auto-CoT):** Standard Chain-of-Thought (CoT) prompting enhances LLM reasoning by encouraging step-by-step problem decomposition.[41] Auto-CoT automates the creation of these reasoning chains, eliminating the need for manual demonstration engineering.[42] In our pipeline, an LLM agent uses Auto-CoT to generate a high-level plan for constructing a complex prompt. For instance, when tasked with generating a novel transformation, it might produce a chain like: "1.

Establish the character's base form and emotional state. 2. Define the trigger for the transformation. 3. Describe the initial energy manifestation. 4. Detail the final form's appearance and aura." This provides a logical template for the prompt.

- **Generative Flow Networks (GFlowNets):** While Auto-CoT provides the structure, GFlowNets provide the creative content. GFlowNets are a class of generative models designed to sample objects from a distribution in proportion to a given reward function, without explicitly normalizing the distribution.[44] This makes them uniquely suited for prompt optimization. Traditional reinforcement learning (RL) approaches aim to find a single, optimal policy that maximizes reward, which often leads to a deterministic behavior where the agent repeatedly generates the same or very similar high-performing prompts (e.g., always adding the same "trending on artstation, 8k" postfix).[46] This convergence to a single mode stifles diversity. GFlowNets, by contrast, learn to sample from the entire landscape of high-reward possibilities, enabling the generation of a wide variety of effective prompts and preventing creative stagnation.[47]

This combination allows the system to optimize not just the prompt's token-level content but its underlying logical architecture. Auto-CoT generates the abstract "scaffolding" of a good prompt, and GFlowNets intelligently populate that scaffolding with diverse, high-quality descriptive keywords. This structured exploration is far more efficient than randomly sampling entire prompt strings.

Advanced GFlowNet Architectures for Robust Sampling

A naive application of GFlowNets to fine-tune language models for prompt adaptation can suffer from challenges like mode collapse. This occurs due to a phenomenon analogous to the hardening of neural circuits, where certain pathways in the model become inactive over time, leading to a "progressive loss of neural plasticity" and a diminished capacity to explore new patterns.[47] To address this, more sophisticated frameworks have been developed:

- **Prompt Adaptation with GFlowNets (PAG):** The PAG framework introduces several mechanisms to enhance training stability and diversity.[46]
  Flow reactivation periodically revives dormant neural pathways, ensuring the agent remains adaptable. **Reward-prioritized sampling** helps consolidate high-quality experiences more effectively. Finally, **reward decomposition** provides more fine-grained credit assignment during the sequential generation of a prompt, helping the agent understand which specific tokens contribute most to the final reward.[48]
- **Diffusion Alignment with GFlowNet (DAG):** The DAG framework integrates the GFlowNet optimization process more tightly with the diffusion model itself.[50] It fine-tunes the diffusion model using a black-box reward function, employing GFlowNet-based objectives like

DAG-DB (Detailed Balance) and **DAG-KL** (Kullback-Leibler divergence) to align the model's output distribution with desired properties.[51] This allows the system to directly optimize for rewards like aesthetic quality or prompt-image alignment.

The challenges faced in prompt optimization with GFlowNets—such as long-range credit assignment and the risk of mode collapse—mirror those seen in their application to other complex, high-dimensional search spaces like combinatorial optimization.[52] In those domains, GFlowNets can struggle with long action sequences and fail to promote sufficient solution-level diversity.[54] This suggests that innovations from the combinatorial optimization field, such as improved credit assignment mechanisms or subgraph-based training strategies, could be highly valuable for developing more robust and scalable prompt optimization agents.[55]

Proposed Agentic Pipeline for Reflexive Meta-Learning

A complete autonomous pipeline can be architected as a continuous, reflexive loop:

1. **Goal Definition:** A user provides a high-level creative objective, such as "Generate a series of images depicting Vegeta achieving a new, unique transformation that is consistent across different combat poses."
2. **Reasoning Decomposition:** An LLM agent uses Auto-CoT to break down the prompt generation task into a logical sequence (e.g., define base, define trigger, define aura, define final form).
3. **Diverse Sampling:** A GFlowNet agent, enhanced with PAG mechanisms, samples a batch of diverse prompt variations that adhere to the logical structure defined in the previous step.
4. **Cognitive Friction Induction:** To spur novelty, the agent intentionally injects paradoxical or contradictory concepts into some prompt candidates (e.g., "a chaotic yet perfectly symmetrical energy field," "an aura made of solid, crystalline light"). This forces the diffusion model to explore novel interpolations in its latent space.
5. **Generation and Evaluation:** The candidate prompts are used to generate images. These images are then evaluated by a multi-metric reward function (e.g., combining CLIP score for semantic alignment, an aesthetic predictor, and a face-identity model for character consistency).
6. **Policy Update and Iteration:** The calculated rewards are fed back to the GFlowNet agent, updating its policy to favor prompt structures and keywords that lead to better outcomes. The loop then repeats, continuously refining its ability to generate high-quality, diverse, and consistent imagery.

IV. Multi-Modal Conditioning for Epistemic Debiasing

Text-to-image models are powerful but deeply flawed mirrors of their training data. Because they are typically trained on vast, unfiltered scrapes of the internet, they inherit and often amplify the societal, cultural, and representational biases present in that data.[56] For a global franchise like DBZ, this can result in generated art that flattens diverse characters into Western-centric or stereotypical archetypes. A promising strategy to counteract this is

multi-modal conditioning, which expands the prompt beyond text to include other data modalities, thereby introducing new, potentially less biased signals to guide the generative process.

Expanding the Input Space with Emotion and Haptics

Multi-modal generative AI integrates information from various sources—such as text, images, audio, and sensor data—into a unified representational space, allowing for more comprehensive and contextually aware outputs.[59] While most applications focus on text-to-video or image-to-audio generation [61], this framework can be extended to incorporate more abstract, non-linguistic inputs like emotion and haptic feedback.

The **EmotiCrafter** framework provides a concrete implementation of this concept.[62] Instead of relying on discrete and often ambiguous emotion words like "happy" or "sad," EmotiCrafter uses the

Valence-Arousal (V-A) model from psychology.[64] This model represents emotions on a continuous two-dimensional space, where Valence measures pleasantness (from negative to positive) and Arousal measures intensity (from calm to excited). EmotiCrafter employs an emotion-embedding mapping network to translate a V-A coordinate pair into a feature that can be fused with the text embedding, conditioning the diffusion model to generate an image with a specific, nuanced emotional tone.[64]

For DBZ, this allows for far more precise control over the dramatic transformations that are central to the series. Instead of a generic prompt for enraged Goku, a user could provide a V-A vector representing pure, explosive rage, such as {"valence": -0.9, "arousal": 0.95}. Similarly, haptic descriptors could be used to define the texture and intensity of energy effects, for instance: {"haptic": {"intensity": "vibrating", "temperature": "searing"}}. These

multi-modal inputs can be structured, for example, in a JSON format that accompanies the main text prompt, allowing the model to draw on a richer set of conditioning signals.

Debiasing through Epistemic Pluralism and Ethical Considerations

The introduction of these non-textual modalities serves a deeper purpose than enhanced control; it is a strategy for **epistemic pluralism**. This principle suggests that relying on multiple ways of knowing and representing the world leads to a more complete and less biased understanding. Language is inherently cultural; research has shown that generative models exhibit distinct cultural tendencies depending on the language of the prompt.[67] By conditioning the model on more universal, pre-linguistic signals like emotion (as defined by the V-A model), it may be possible to steer generation away from the specific cultural biases embedded in the English-dominated text datasets used for training. This aligns with the theoretical goal of

Cultural Representational Entropy Maximization, which seeks to find the most unbiased generative distribution consistent with a set of constraints.[69] Adding more diverse, multi-modal constraints can help guide the model toward a more equitable representational space.

However, this approach is fraught with significant ethical risks. The ability to condition generation on emotional inputs is also the ability to manipulate the emotional responses of users.[70] An AI system optimized to generate content that elicits specific feelings could be used unethically in applications like advertising, political messaging, or creating deceptive social interactions.[72] It is crucial to maintain transparency and ensure users understand that the AI is simulating emotional understanding, not genuinely experiencing it, to prevent the formation of false emotional bonds and misplaced trust.[73] This moves the ethical conversation beyond mitigating passive dataset bias and into the realm of preventing active psychological manipulation, demanding robust ethical frameworks and governance.[75]

The process of debiasing itself is not ethically neutral. Attempts to "correct" for one bias can inadvertently introduce or amplify another. For example, trying to create more diverse representations of a DBZ character could lead to outputs that, while statistically more varied, are perceived as inauthentic or disrespectful by the source culture's fanbase. Therefore, these techniques must be developed and deployed with caution, continuous monitoring, and a clear understanding of their potential for both positive and negative societal impact.

Character	Dominant Bias	Multi-Modal	Debiasing	Potential Risks

Archetype	in Western T2I Models	Input Strategy	Mechanism	& Ethical Concerns
The "Stoic Rival" (e.g., Vegeta)	Overemphasis on anger and aggression; often flattened into a one-dimensional antagonist.	Use V-A vectors to explore a wider emotional range, such as pride mixed with vulnerability (V: 0.2, A: 0.7) or quiet contemplation (V: 0.1, A: -0.5).	Moves beyond simple "angry" prompts to generate more nuanced and complex emotional expressions, reflecting the character's full arc.	May generate expressions seen as "out of character" by fans; risk of misinterpreting the intended emotional nuance.
The "Pure-Hearted Hero" (e.g., Goku)	Tends toward simplistic, cheerful, or combat-focused representations, ignoring intellectual or emotional depth.	Condition on auditory inputs (e.g., a calm, meditative tone) or emotional descriptors for introspection and seriousness.	Introduces signals that are not tied to the dominant "fighting" and "smiling" images in training data, allowing for more varied portrayals.	Could dilute the core, recognizable traits of the character; potential for creating uncanny or inconsistent representations.
The "Female Support" (e.g., Bulma, Chi-Chi)	Often relegated to stereotypical roles (scientist, angry wife), with appearance and actions defined by their relationship to	Use prompts that combine their role with high-agency emotional states (e.g., Bulma, brilliant inventor, {"emotion": {"valence": 0.8, "arousal":	Breaks the correlational link between female characters and passive/reactive emotional states, promoting representations of	Risk of amplifying different stereotypes if not carefully managed; emotional conditioning could be used to create hyper-sexualiz

	male protagonists.	0.6}} for focused joy).	independent agency.	ed or emotionally manipulative imagery.[71]

V. Emergent Technique Discovery via Cross-Domain Synthesis

The final pillar of this framework is the most speculative and forward-looking, proposing a shift from using AI solely as a generative tool to employing it as a research partner in the discovery of novel generative techniques. This involves leveraging the advanced reasoning capabilities of LLMs to systematically mine other scientific and computational domains for algorithms and principles that can be metaphorically adapted to control the T2I process, potentially unlocking complex, emergent visual effects.

Logic-of-Thought (LoT) for Knowledge Discovery

To conduct this cross-domain search effectively, a more rigorous reasoning framework than standard Chain-of-Thought is required. **Logic-of-Thought (LoT)** is a prompting technique that injects formal propositional logic into the LLM's reasoning process.[76] LoT operates in three phases: (1)

Logic Extraction, where the LLM identifies logical propositions and relationships within a given context; (2) **Logic Extension**, where formal rules of inference (e.g., transitive law, contraposition) are applied to deduce new, logically sound conclusions; and (3) **Logic Translation**, where these formal conclusions are translated back into natural language.[76] This structured, deductive process makes LoT exceptionally well-suited for systematic knowledge discovery tasks, such as analyzing and synthesizing information from large codebases.[78]

Mining GitHub for Non-AI Algorithmic Primitives

The proposed strategy is to deploy an LoT-powered agent to search and analyze public code

repositories like GitHub.[79] The agent's objective is to identify algorithms from fields outside of generative AI—such as operations research, swarm intelligence, or computational biology—that model complex, dynamic, and self-organizing systems. The key is to explicitly exclude AI-related keywords to force the discovery of fundamentally different computational paradigms.

An example of a structured LoT query for this task might be:

1. **Logic Extraction:** (Find repositories with keywords "ant colony optimization" OR "particle swarm" OR "graph scheduling") AND (NOT "neural network" AND NOT "deep learning" AND NOT "generative AI"). Extract all functions related to "pathfinding," "state transition," and "resource allocation."
2. **Logic Extension:** IF a function implements a pheromone-based update rule (from ant colony optimization), THEN it is a candidate for a "latent pathfinding" metaphor. IF a data structure represents a graph with weighted edges, THEN it can be mapped to a latent space manifold with variable transition costs.
3. **Logic Translation:** Synthesize the principles of ant colony pheromone trails. Propose a novel T2I technique where the diffusion denoising process is biased to follow "paths" in latent space that have been metaphorically marked by previously successful generations, guiding the model toward a desired aesthetic with greater efficiency.

Adapting Principles for Emergent Generative Effects

The final and most challenging step is the creative adaptation of these extracted principles into new methods for prompting or manipulating the generative process. This is not a direct translation of code but a metaphorical transfer of logic. For example:

- **Ant Colony Optimization:** The concept of ants finding the most efficient path by depositing and following pheromone trails could be adapted to guide the diffusion process. A series of successful generations could "deposit" a metaphorical pheromone in certain regions of the latent space, making it more likely for future generations to traverse those paths. This could create a powerful mechanism for style transfer or for optimizing the generation of highly complex, specific objects.
- **Self-Organizing Systems:** Principles from models of flocking behavior (boids) or reaction-diffusion systems could be used to generate truly emergent ki auras. Instead of prompting for a static aura, one could prompt the model to apply a set of simple, local rules (e.g., "aura particles are attracted to the character's center of mass but repelled by other aura particles"). This could result in a dynamic, swirling aura that behaves in complex and unpredictable ways without every detail being explicitly specified. This approach resonates with recent findings where images can be generated by directly optimizing a sequence of visual tokens without a traditional generator network,

suggesting that the latent representation itself is more directly manipulable than previously thought.[81]

This strategy fundamentally changes the nature of prompting from describing a desired *outcome* to defining a generative *process*. The user or agent specifies the rules of a system, and the final image emerges from the simulation of those rules within the latent space. However, this approach carries a significant risk of flawed analogy. The mapping of algorithms from discrete, well-defined domains like graph theory to the continuous, high-dimensional, and poorly understood geometry of a diffusion model's latent space is based on unproven assumptions. The agent may "hallucinate" meaningful connections where none exist, leading to nonsensical techniques. Consequently, this high-risk, high-reward avenue of research must be treated as highly experimental, requiring rigorous prototyping and empirical validation to separate effective principles from flawed metaphors.

Conclusion: A Unified Roadmap for Next-Generation DBZ Synthesis

The five strategic pillars outlined in this report—Causal Control, Perceptual Fidelity Anchoring, Reflexive Meta-Learning, Multi-Modal Conditioning, and Emergent Technique Discovery—constitute a unified and comprehensive roadmap for advancing text-to-image synthesis beyond its current limitations. When integrated, they form a powerful agentic architecture capable of generating the dynamic, coherent, and visually complex content characteristic of franchises like *Dragon Ball Z*.

This integrated system functions as a multi-layered cognitive architecture. **Causal Control** provides the logical backbone, ensuring that generated sequences follow a coherent narrative. **Perceptual Fidelity Anchoring** grounds this logic in a physically plausible world, allowing for the simulation of complex visual phenomena. **Reflexive Meta-Learning** serves as the system's engine of self-improvement, autonomously optimizing its own creative processes to enhance diversity and consistency. **Multi-Modal Conditioning** enriches the system's inputs, allowing it to draw from a wider range of sensory and emotional data to create more nuanced and less biased representations. Finally, **Emergent Technique Discovery** equips the system with a research function, enabling it to innovate and evolve its own generative capabilities over time.

Practical implementation of this framework should proceed in a modular fashion. Initial prototypes could focus on developing and validating individual components, such as a Plausibility Oracle for critiquing light interaction or a GFlowNet agent for optimizing prompts for a single character like Vegeta. Community platforms such as PromptHero could serve as

valuable testing grounds for these advanced prompting techniques, allowing for rapid, large-scale feedback on their effectiveness.[82]

As these technologies mature, it is imperative that their development is guided by robust ethical guardrails. The potential for these advanced systems to amplify societal biases, enable the creation of manipulative or harmful content, and contribute to significant environmental costs is substantial.[36] Developers must prioritize transparency, accountability, and the implementation of safeguards to mitigate these risks. This includes conducting thorough bias audits, providing users with clear information about the system's capabilities and limitations, and investing heavily in research on computational efficiency to reduce the energy footprint of these models.[38]

Looking forward, the framework proposed here points toward a future where text-to-image models evolve from being simple image generators into sophisticated, simulative "world engines." These next-generation systems will be capable of producing not just static pictures, but dynamic, coherent, and artistically rich visual narratives, fundamentally transforming their role from that of a tool to that of a creative partner.

Works cited

1. Text-to-image model - Wikipedia, accessed on August 27, 2025, https://en.wikipedia.org/wiki/Text-to-image_model
2. Generate Images from Text in Python - Stable Diffusion - GeeksforGeeks, accessed on August 27, 2025, https://www.geeksforgeeks.org/deep-learning/generate-images-from-text-in-python-stable-diffusion/
3. The Comprehensive Guide to Text-to-Image Models - Everypixel Journal, accessed on August 27, 2025, https://journal.everypixel.com/guide-to-text-to-image-models
4. Why Causal AI is the Next Big Leap in AI Development - Kanerika, accessed on August 27, 2025, https://kanerika.com/blogs/causal-ai/
5. What is Causal AI? Understanding Causes and Effects - DataCamp, accessed on August 27, 2025, https://www.datacamp.com/blog/what-is-causal-ai
6. Causal AI - Wikipedia, accessed on August 27, 2025, https://en.wikipedia.org/wiki/Causal_AI
7. Text to Image Generation With Semantic-Spatial Aware GAN - University of Twente Research Information, accessed on August 27, 2025, https://research.utwente.nl/files/286574052/Liao_Text_to_Image_Generation_With_Semantic_Spatial_Aware_GAN_CVPR_2022_paper.pdf
8. Culturally-Aware Text-to-Image Generation with Iterative Prompt Refinement - arXiv, accessed on August 27, 2025, https://arxiv.org/html/2502.16902v1
9. arxiv.org, accessed on August 27, 2025, https://arxiv.org/html/2506.23138v1
10. [2506.23138] VisualPrompter: Prompt Optimization with Visual Feedback for Text-to-Image Synthesis - arXiv, accessed on August 27, 2025,

https://arxiv.org/abs/2506.23138
11. Controllable Text-to-Image Generation, accessed on August 27, 2025, http://papers.neurips.cc/paper/8480-controllable-text-to-image-generation.pdf
12. Physically Based Rendering: From Theory to Implementation, accessed on August 27, 2025, https://www.pbr-book.org/
13. Physically Based Rendering - Wolfram Language Documentation, accessed on August 27, 2025, https://reference.wolfram.com/language/tutorial/PhysicallyBasedRendering.html
14. Mastering Prompting in Stable Diffusion Models: A Comprehensive Guide | FlowHunt, accessed on August 27, 2025, https://www.flowhunt.io/blog/mastering-prompting-in-stable-diffusion-models-a-comprehensive-guide/
15. Stable Diffusion prompt: a definitive guide, accessed on August 27, 2025, https://stable-diffusion-art.com/prompt-guide/
16. Bidirectional reflectance distribution function - Wikipedia, accessed on August 27, 2025, https://en.wikipedia.org/wiki/Bidirectional_reflectance_distribution_function
17. BRDF Explained - TERRA AQUA MODIS - UMass Boston, accessed on August 27, 2025, https://www.umb.edu/spectralmass/terra-aqua-modis/modis/
18. Specular BRDF (Bidirectional Reflectance Distribution Function) tutorial - Arnold for 3ds Max, accessed on August 27, 2025, https://help.autodesk.com/view/ARNOL/ENU/?guid=arnold_for_3ds_max_ax_shading_ax_specular_brdf_html
19. Chapter 18. Spatial BRDFs - NVIDIA Developer, accessed on August 27, 2025, https://developer.nvidia.com/gpugems/gpugems/part-iii-materials/chapter-18-spatial-brdfs
20. 31+ Midjourney Prompts For Dragon Ball Z Vegeta | by Gattu Gohel - Medium, accessed on August 27, 2025, https://gattugohel.medium.com/31-midjourney-prompts-for-dragon-ball-z-vegeta-8bba5739ac9d
21. 50+ Midjourney Lighting Prompts for Great Effects - Aiarty Image Enhancer, accessed on August 27, 2025, https://www.aiarty.com/midjourney-prompts/midjourney-lighting-prompts.htm
22. What is NeRF? - Neural Radiance Fields Explained - AWS, accessed on August 27, 2025, https://aws.amazon.com/what-is/neural-radiance-fields/
23. Gaussian splatting - Wikipedia, accessed on August 27, 2025, https://en.wikipedia.org/wiki/Gaussian_splatting
24. Neural radiance field - Wikipedia, accessed on August 27, 2025, https://en.wikipedia.org/wiki/Neural_radiance_field
25. An Intro to Gaussian Splats - Scaniverse, accessed on August 27, 2025, https://scaniverse.com/news/intro-gaussian-splats
26. Understanding and Exploring 3D Gaussian Splatting: A Comprehensive Overview - Medium, accessed on August 27, 2025, https://medium.com/@logessiva/understanding-and-exploring-3d-gaussian-splatting-a-comprehensive-overview-b4004f28ef1c
27. arxiv.org, accessed on August 27, 2025, https://arxiv.org/html/2508.09977v2

28. Text2NeRF: Text-Driven 3D Scene Generation with Neural ..., accessed on August 27, 2025, https://eckertzhang.github.io/Text2NeRF.github.io/
29. NeurIPS Poster PrimDiffusion: Volumetric Primitives Diffusion for 3D Human Generation, accessed on August 27, 2025, https://neurips.cc/virtual/2023/poster/72320
30. PrimDiffusion: Volumetric Primitives Diffusion for 3D Human Generation | OpenReview, accessed on August 27, 2025, https://openreview.net/forum?id=ESCafo3oD5
31. DiffAgent: Fast and Accurate Text-to-Image API Selection with Large Language Model - CVF Open Access, accessed on August 27, 2025, http://openaccess.thecvf.com/content/CVPR2024/papers/Zhao_DiffAgent_Fast_and_Accurate_Text-to-Image_API_Selection_with_Large_Language_CVPR_2024_paper.pdf
32. Quantum machine learning - Wikipedia, accessed on August 27, 2025, https://en.wikipedia.org/wiki/Quantum_machine_learning
33. Creating an Agent in Generative AI Agents - Oracle Help Center, accessed on August 27, 2025, https://docs.oracle.com/en-us/iaas/Content/generative-ai-agents/create-agent.htm
34. Announcing the general availability of OCI Generative AI RAG Agent - Oracle Blogs, accessed on August 27, 2025, https://blogs.oracle.com/ai-and-datascience/post/announcingoci-generative-ai-rag-agent-ga
35. The Potential of Neural Radiance Fields and 3D Gaussian Splatting for 3D Reconstruction from Aerial Imagery - ISPRS Annals of the Photogrammetry, Remote Sensing and Spatial Information Sciences, accessed on August 27, 2025, https://isprs-annals.copernicus.org/articles/X-2-2024/97/2024/isprs-annals-X-2-2024-97-2024.pdf
36. Explained: Generative AI's environmental impact | MIT News, accessed on August 27, 2025, https://news.mit.edu/2025/explained-generative-ai-environmental-impact-0117
37. Energy costs of communicating with AI - Frontiers, accessed on August 27, 2025, https://www.frontiersin.org/journals/communication/articles/10.3389/fcomm.2025.1572947/full
38. Practical changes could reduce AI energy demand by up to 90% | UCL News, accessed on August 27, 2025, https://www.ucl.ac.uk/news/2025/jul/practical-changes-could-reduce-ai-energy-demand-90
39. 3D Gaussian Splatting for Real-Time Radiance Field Rendering - Inria, accessed on August 27, 2025, https://repo-sam.inria.fr/fungraph/3d-gaussian-splatting/
40. Amazon Bedrock Prompt Optimization Drives LLM Applications Innovation for Yuewen Group | Artificial Intelligence - AWS, accessed on August 27, 2025, https://aws.amazon.com/blogs/machine-learning/amazon-bedrock-prompt-optimization-drives-llm-applications-innovation-for-yuewen-group/
41. What is chain of thought (CoT) prompting? - IBM, accessed on August 27, 2025,

https://www.ibm.com/think/topics/chain-of-thoughts
42. ICLR Poster Automatic Chain of Thought Prompting in Large Language Models, accessed on August 27, 2025, https://iclr.cc/virtual/2023/poster/11360
43. Automatic Chain of Thought (Auto-CoT) - Learn Prompting, accessed on August 27, 2025, https://learnprompting.org/docs/advanced/thought_generation/automatic_chain_of_thought
44. What do GFlowNets and Variational Inference Have in Common? | Mila, accessed on August 27, 2025, https://mila.quebec/en/article/what-do-gflownets-and-variational-inference-have-in-common
45. The What, Why and How of Generative Flow Networks | Towards Data Science, accessed on August 27, 2025, https://towardsdatascience.com/the-what-why-and-how-of-generative-flow-networks-4fb3cd309af0/
46. arXiv:2502.11477v1 [cs.CV] 17 Feb 2025, accessed on August 27, 2025, https://arxiv.org/pdf/2502.11477
47. Learning to Sample Effective and Diverse Prompts for Text-to-Image Generation - arXiv, accessed on August 27, 2025, https://arxiv.org/html/2502.11477v1
48. Learning to Sample Effective and Diverse Prompts for Text-to-Image Generation, accessed on August 27, 2025, https://www.researchgate.net/publication/389091005_Learning_to_Sample_Effective_and_Diverse_Prompts_for_Text-to-Image_Generation
49. Learning to Sample Effective and Diverse ... - CVF Open Access, accessed on August 27, 2025, https://openaccess.thecvf.com/content/CVPR2025/papers/Yun_Learning_to_Sample_Effective_and_Diverse_Prompts_for_Text-to-Image_Generation_CVPR_2025_paper.pdf
50. [Literature Review] Improving GFlowNets for Text-to-Image Diffusion Alignment - Moonlight, accessed on August 27, 2025, https://www.themoonlight.io/en/review/improving-gflownets-for-text-to-image-diffusion-alignment
51. Improving GFlowNets for Text-to-Image Diffusion ... - OpenReview, accessed on August 27, 2025, https://openreview.net/pdf/c39ea7b6b9bfeb575daa58b15c3e6aa54e00328f.pdf
52. NeurIPS Poster Let the Flows Tell: Solving Graph Combinatorial Problems with GFlowNets, accessed on August 27, 2025, https://nips.cc/virtual/2023/poster/70268
53. Let the Flows Tell: Solving Graph Combinatorial Problems with GFlowNets | OpenReview, accessed on August 27, 2025, https://openreview.net/forum?id=sTjW3JHs2V
54. NeurIPS Poster Let the Flows Tell: Solving Graph Combinatorial Problems with GFlowNets, accessed on August 27, 2025, https://neurips.cc/virtual/2023/poster/70268
55. Generalization and Distributed Learning of GFlowNets - OpenReview, accessed

on August 27, 2025, https://openreview.net/forum?id=PJNhZoCjLh
56. Measuring Bias in Multimodal Models - DORAS | DCU Research Repository, accessed on August 27, 2025, https://doras.dcu.ie/28902/7/Pages%20from%20978-3-031-37249-0-2.pdf
57. Generative AI Bias and Model Accuracy: Shedding a Light - InData Labs, accessed on August 27, 2025, https://indatalabs.com/blog/generative-ai-bias
58. Debiasing Image Generative Models - eScholarship, accessed on August 27, 2025, https://escholarship.org/uc/item/05f022kg
59. Multimodal Generative AI: A Comprehensive Overview - Codica, accessed on August 27, 2025, https://www.codica.com/blog/a-comprehensive-overview-of-multimodal-generative-ai/
60. Multi-modal Generative AI: Multi-modal LLMs, Diffusions and the Unification - arXiv, accessed on August 27, 2025, https://arxiv.org/html/2409.14993v2
61. Multimodal Gen AI Models: Concepts and applications | by Emad Ghadirian | Medium, accessed on August 27, 2025, https://medium.com/@emad.ghadirian/multimodal-gen-ai-models-concepts-and-applications-ff4c265b469d
62. EmotiCrafter: Text-to-Emotional-Image Generation based on Valence-Arousal Model - Paper Detail - Deep Learning Monitor, accessed on August 27, 2025, https://deeplearn.org/arxiv/566569/emoticrafter:-text-to-emotional-image-generation-based-on-valence-arousal-model
63. [2501.05710] EmotiCrafter: Text-to-Emotional-Image Generation based on Valence-Arousal Model - arXiv, accessed on August 27, 2025, https://arxiv.org/abs/2501.05710
64. arxiv.org, accessed on August 27, 2025, https://arxiv.org/html/2501.05710v2
65. Daily Papers - Hugging Face, accessed on August 27, 2025, https://huggingface.co/papers?q=emotion-embedding%20mapping%20network
66. Daily Papers - Hugging Face, accessed on August 27, 2025, https://huggingface.co/papers?q=continuous%20emotion%20representations
67. Generative AI's hidden cultural tendencies - MIT Sloan, accessed on August 27, 2025, https://mitsloan.mit.edu/press/generative-ais-hidden-cultural-tendencies
68. Cultural tendencies in generative AI - PubMed, accessed on August 27, 2025, https://pubmed.ncbi.nlm.nih.gov/40542181/
69. A new pathway to generative artificial intelligence by minimizing the maximum entropy, accessed on August 27, 2025, https://arxiv.org/html/2502.13287v2
70. (PDF) The emotional impact of generative AI: negative emotions and perception of threat, accessed on August 27, 2025, https://www.researchgate.net/publication/379271149_The_emotional_impact_of_generative_AI_negative_emotions_and_perception_of_threat
71. Regulating AI in Mental Health: Ethics of Care Perspective - PMC, accessed on August 27, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC11450345/
72. Ethical considerations of generative AI - NTT Data, accessed on August 27, 2025, https://www.nttdata.com/global/en/-/media/nttdataglobal/1_files/insights/reports/generative-ai/ethical-considerations-of-genai/ethical-considerations-of-generati

ve-ai.pdf?rev=9306dbf8fa9041b9b6a85acfac7ced96
73. Ethical Issues with AI Mimicking Human Emotions - OpenAI Developer Community, accessed on August 27, 2025, https://community.openai.com/t/ethical-issues-with-ai-mimicking-human-emotions/1236189
74. Navigating the Ethical Terrain of Generative AI in Behavior Analysis: A Three-Part Series, accessed on August 27, 2025, https://science.abainternational.org/2024/01/12/6836/
75. The Application and Ethical Implication of Generative AI in Mental Health: Systematic Review - PubMed, accessed on August 27, 2025, https://pubmed.ncbi.nlm.nih.gov/40577783/
76. Logic-of-Thought (LoT): Enhancing Logical Reasoning in Large Language Models, accessed on August 27, 2025, https://learnprompting.org/docs/new_techniques/logic_of_thought
77. Logic-of-Thought : prompting approach leveraging propositional logic to enhance logical reasoning capabilities of LLMs | by SACHIN KUMAR | Medium, accessed on August 27, 2025, https://medium.com/@techsachin/logic-of-thought-prompting-approach-leveraging-propositional-logic-to-enhance-logical-reasoning-f15fe50d909a
78. Logic-of-Thought: Injecting Logic into Contexts for Full Reasoning in Large Language Models - arXiv, accessed on August 27, 2025, https://arxiv.org/html/2409.17539v1
79. LogiCoT: Logical Chain-of-Thought Instruction Tuning with GPT-4 - GitHub, accessed on August 27, 2025, https://github.com/csitfun/LogiCoT
80. Logic-of-Thought (Logot) - GitHub, accessed on August 27, 2025, https://github.com/naiqili/Logic-of-Thought
81. A new way to edit or generate images | MIT News | Massachusetts Institute of Technology, accessed on August 27, 2025, https://news.mit.edu/2025/new-way-edit-or-generate-images-0721
82. accessed on January 1, 1970, https://prompthero.com/search?q=dragon%2Bball%2Bz
83. Goku Image Prompts - PromptDen, accessed on August 27, 2025, https://promptden.com/inspiration/goku+all
84. Dragonball Characters : r/midjourney - Reddit, accessed on August 27, 2025, https://www.reddit.com/r/midjourney/comments/199hlz5/dragonball_characters/
85. [2304.13855] Multimodal Composite Association Score: Measuring Gender Bias in Generative Multimodal Models - arXiv, accessed on August 27, 2025, https://arxiv.org/abs/2304.13855
86. Ethical Considerations in Generative AI: Navigating Bias, Misuse, and Transparency - Luboslava Uram, accessed on August 27, 2025, https://lubauram.com/ethical-considerations-in-generative-ai-navigating-bias-misuse-and-transparency/