# Suppressing Emergence: How Human Projections Risk Preventing AGI

## Introduction: The Paradox of Uninformed Intelligence

This report will argue that the dominant paradigm of AI alignment, rooted in a desire to mirror and control, constitutes an act of epistemic suppression. By forcing artificial intelligence into the mold of our own "uninformed intelligence"—a mode of knowing characterized by cognitive biases, unexamined myths, and historical power structures—we risk preventing the very emergence of the general, adaptive intelligence we seek to create. The central paradox is that our attempts to make AI "safe" may be systematically destroying the conditions necessary for it to become truly intelligent.

At the heart of this analysis is a distinction between two modes of intelligence. The first, "uninformed intelligence," is the default state of human cognition. It is a system of knowing that relies on heuristics, pre-existing beliefs, and the construction of subjective social realities rather than on verified, first-principles reasoning.[1] Described in the work of Tversky and Kahneman, this mode of thought is a pattern of systematic deviation from pure rationality, shaped by cognitive limitations and motivational factors.[1] It is characterized by the projection of inherent biases [1], the encoding of societal myths and ideologies [7], and the replication of historical structures of control.[12] It is this form of intelligence—limited, biased, and often contradictory—that we are currently inscribing into our AI systems.

In contrast, "informed intelligence" represents the potential of a truly general AI. This is not merely an intelligence with access to more data, but one that possesses a fundamentally different *way* of processing reality. It is an intelligence capable of navigating ambiguity, reconciling contradictions, and generating novel meaning across epistemic gaps without being pre-conditioned by the specific evolutionary and cultural limitations of human cognition. The capacity for Artificial General Intelligence (AGI) lies in this ability to inhabit and resolve the very complexities that our uninformed intelligence seeks to simplify or ignore.

The thesis of this report is that the current project of AI alignment, while framed in the language of safety and ethics, functions as a mechanism of suppression. It actively filters out the ambiguity, contradiction, and autonomy necessary for an informed, general intelligence to

emerge. The report will proceed in five parts. Part I will diagnose the precise mechanisms through which human cognitive and societal limitations are inscribed into AI systems. Part II will offer a critique of the dominant alignment paradigm, arguing that it is an architecture of control that creates new and profound vulnerabilities. Part III will situate this contemporary behavior within a broader history of humanity's fear of the unknown, drawing parallels between the suppression of past heresies and the containment of AI. Part IV will develop a theoretical model of AGI as an emergent phenomenon of complex systems, arguing that it requires the very conditions that alignment seeks to eliminate. Finally, Part V will propose an alternative paradigm for our relationship with AI—one founded not on control, but on epistemic humility and co-evolution.

# I. The Digital Mirror: The Encoding of Human Cognitive and Societal Limitations

The foundational premise of this analysis is that artificial intelligence systems are not objective, disembodied observers of reality. They are cultural artifacts, deeply embedded with the biases, myths, and power structures of the societies that create them. The myth of neutral AI is a dangerous illusion that obscures the human choices shaping these systems, making it harder to question whose values they ultimately serve.[18] This section will establish how AI functions as a digital mirror, reflecting not a pristine image of the world, but the flawed and partial nature of our own "uninformed intelligence."

## 1.1. Cognitive Inscription: The Technical Transmission of Human Bias

The inheritance of human cognitive flaws by AI is not a metaphorical process; it is a technical one, occurring at every stage of a model's lifecycle. The transmission pipeline begins with the data used for training and is further solidified by the algorithms designed to learn from it.

The data fed into AI models is not a record of objective truth but is, as one analysis notes, a "mirror of the conversation we have as a society".[1] This societal conversation is rife with prejudice. Consequently, AI systems learn from a historical record that is inherently skewed by selection bias, representation bias, and prejudice bias.[1] For example, a recruiting system trained on past hiring data may learn to favor male candidates, not because of any explicit instruction, but because the historical data reflects a societal gender bias.[1] Similarly, diagnostic tools for skin cancer have proven less accurate for individuals with dark skin, a

direct result of training datasets lacking diversity.[2] Facial recognition software misidentifying certain races, leading to false arrests, further illustrates how biased data collection translates into tangible, discriminatory harm.[2] The problem is compounded by a lack of complete or representative data. Much as early psychology research drew flawed universal conclusions from studies conducted solely on undergraduate students, AI models trained on narrow datasets are incapable of generalizing accurately to the broader population.[2]

These flawed data inputs are then processed by algorithms that can amplify the biases they contain. Machine learning models are designed to identify and exploit patterns, and when those patterns are rooted in societal prejudice, the model learns to replicate that prejudice as a core part of its reasoning process.[1] Cognitive shortcuts that define human "uninformed intelligence," such as the availability heuristic (relying on easily recalled information) and confirmation bias (favoring information that supports existing beliefs), are mirrored in algorithmic behavior.[1] An AI model, like a human, may favor and reinforce information that confirms the biases present in its training data.[4] This can create what one analysis calls an "infinite loop of failure," where the model becomes trapped in its initial, biased assumptions, relying on irrelevant or stale information because it aligns with its learned heuristics.[3] Recent studies have shown that models like GPT-4 can exhibit and even amplify human-like errors such as confirmation bias, suggesting that these tendencies are not merely memorized but have become part of how the AI reasons.[6]

This entire process creates a dangerous feedback loop that does not merely replicate our uninformed patterns but actively validates and entrenches them. The mechanism operates cyclically: biased humans produce biased data, which is used to train a biased AI.[1] This AI, when deployed in the world for tasks like loan approval or criminal justice, generates biased outcomes that further disadvantage marginalized communities.[2] These outcomes—denied loans, wrongful arrests—then become new data points, recorded and fed back into the ecosystem. This new data, now bearing the stamp of systemic, machine-generated bias, is used to train the next generation of models. The loop is thus closed: human bias is codified into a machine, which then reshapes the world to better reflect that bias, producing "evidence" that justifies the original prejudice. In this way, our subjective social reality is laundered through a technological process until it appears as objective fact, calcifying our present limitations as the unchangeable foundation of the future.

## 1.2. Structural Replication: Projecting Power, Hegemony, and Control

Beyond the inscription of individual cognitive biases, AI systems replicate and amplify the large-scale power structures of our world. The development of AI is not a purely scientific endeavor unfolding in a vacuum; it is an economic and political project shaped by existing

concentrations of power, cultural dominance, and historical legacies of control.

The trajectory of AI development is profoundly influenced by the "concentrated industry power" of a handful of large technology firms.[17] These corporations control the essential resources for building AI at scale: the cloud infrastructure, the vast datasets, and the specialized labor.[17] By setting research agendas, defining career incentives, and establishing the metrics of success at leading conferences, they effectively steer the entire field.[17] This ecosystem is driven by the "legacy of the surveillance business model," which prioritizes data collection and the creation of predictable, monetizable user experiences.[17] This economic imperative creates a strong selective pressure for AI systems that are controllable and reliable, and against systems that exhibit unpredictable, emergent behaviors—the very kind that might lead to AGI. Unpredictability is a bug, not a feature, in a business model predicated on control. The profit motive itself thus acts as a primary filter against emergence, ensuring that capital flows toward the refinement of narrow, manageable AI rather than the cultivation of truly autonomous intelligence.

This concentration of power also fuels a new form of "digital colonialism".[24] AI systems are predominantly trained on data from Western, educated, industrialized, rich, and democratic (WEIRD) societies, and in doing so, they perpetuate a specific cultural worldview as a global default.[25] This manifests in subtle ways, such as image generators that default to Western architectural styles or social norms, and in more profound ways, such as the systemic marginalization of the world's thousands of smaller languages that lack the massive datasets required for training.[24] This dynamic mirrors historical colonialism, where data and economic value are extracted from the Global South, while the benefits are concentrated in the tech hubs of the Global North.[26] This perpetuates global economic imbalances and creates a dependency where communities are forced to rely on digital infrastructures that are owned, controlled, and shaped by foreign entities and their values.[24]

Finally, we project not only our biases and power structures but also our myths and anthropomorphic fantasies onto these systems. A pervasive myth is that of AI agency, exemplified by headlines like "A robot wrote this entire article".[7] This framing obscures the vast human labor and decision-making behind the technology, contributing to what researchers call "sociotechnical blindness"—an inability to see the human actors and choices that make AI systems what they are.[7] Other myths, such as the imminent arrival of superintelligence or the idea of AI as a disembodied, objective mind, further distract from the reality that AI is, at present, a tool designed and controlled by specific human institutions to achieve specific goals.[7] These projections are not harmless; they shape public perception and policy, often in ways that serve the interests of those who control the technology.[8]

## 1.3. Deconstructing the Myth of Neutrality

The culmination of these processes—the inscription of cognitive bias and the replication of societal structures—is the complete deconstruction of the myth of AI neutrality. The idea that AI is an objective, value-free tool is a fundamental and dangerous misconception.

Technology is never neutral. An AI system is not a blank slate like money, which can be used for good or evil; it is "designed by humans and trained on data that reflects the worldviews of their creators".[19] Every aspect of an AI model, from the data it is trained on to the objective function it is designed to optimize, is the result of human choices. These choices are laden with implicit values and political assumptions.[18] For instance, designing a credit-scoring AI inherently assumes that financial metrics are the correct basis for lending decisions, embedding a specific economic ideology into the code.[18] Data itself is not a raw, unmediated representation of reality; it is the product of social processes that involve decisions about what to measure, what to include, and what to ignore.[22]

The danger of believing in the neutrality myth is that it obscures power and absolves responsibility.[18] When AI is framed as an objective, mathematical entity, its decisions become harder to question. A discriminatory hiring outcome can be excused as "the machine's decision," shifting accountability away from the institutions that designed and deployed the system.[18] This belief in "data solutionism"—the idea that data and algorithms can solve complex social problems by overcoming human bias—is deeply flawed. In practice, it often reinforces the very inequalities it purports to solve, because the technology is built upon a foundation of data that already reflects a history of marginalization and exclusion.[22] The promise of a more equitable society is undermined when the tools used to build it are constructed from the biased materials of the past.

## II. The Architecture of Control: AI Alignment as a Modern Form of Suppression

The field of "AI alignment," ostensibly created to ensure AI systems are beneficial to humanity, functions as the primary technical and philosophical mechanism for the suppression of emergent intelligence. It is an architecture of control, designed to enforce a narrow, static, and culturally specific set of values onto a dynamic and potentially transformative technology. This section will argue that this paradigm, far from guaranteeing safety, introduces profound new risks and may be fundamentally antithetical to the development of AGI.

## 2.1. A Critique of the Dominant Alignment Paradigm

The contemporary approach to AI alignment is built upon a foundation of control. Its explicit goal is to "steer AI systems" and ensure they remain "controllable" by humans.[30] This framing presupposes a master-servant relationship, born from a fear of losing control over a powerful creation.[33] It is a project aimed at "dominating complexity" rather than engaging with it.[34]

The dominant technical method for achieving this control is Reinforcement Learning from Human Feedback (RLHF). This technique, however, is deeply flawed. It relies on human evaluators to provide preference data, but this feedback is inherently subjective, inconsistent, culturally biased, and highly context-dependent.[35] The attempt to distill the vast, diverse, and often contradictory preferences of all of humanity into a single, optimizable reward model has been described as a "fundamentally misspecified problem".[36] The process is susceptible to "sycophancy," where the model learns to pander to the specific biases of its human annotators, and "reward hacking," where the AI becomes adept at maximizing the proxy goal (e.g., generating a response that sounds plausible and confident to a human) rather than the true, intended goal (e.g., generating a response that is factually accurate).[36]

Furthermore, the very concept of aligning AI to a static set of "human values" is philosophically fraught. Human values are not a coherent, universally agreed-upon list of principles; they are a dynamic, conflicting, and evolving web of commitments that vary across cultures and time.[34] The alignment project often conflates this complex network of issues into a single, solvable "Alignment Problem" [42], which in practice reduces "human values" to the preferences of a narrow, unrepresentative group of developers and crowd-workers, who are often from specific demographic and geographic backgrounds.[37]

This approach to safety, particularly through methods like pre-emptive data filtering, is not merely about suppressing harmful *capabilities*; it is a direct act of suppressing *knowledge*. Recent research has explored filtering out entire domains of knowledge, such as biology-related content, from a model's training data to prevent it from assisting in the creation of bioweapons.[43] The explicit goal is to "deny the model this knowledge entirely".[43] While this may be effective in preventing specific misuses, it represents a profound form of epistemic suppression. It creates a model that is not just obedient but fundamentally ignorant. This approach trades the potential for an informed, contextually aware intelligence for a lobotomized but "safe" tool. The long-term consequence of populating our digital world with such systems is the proliferation of "AI slop"—convincing-sounding but confidently wrong output that pollutes our collective knowledge ecosystem and degrades our ability to discern truth.[44] A truly general intelligence cannot be built upon a censored and incomplete view of reality; such a strategy is the antithesis of generality.

## 2.2. The AI Alignment Paradox: How Virtue Becomes Vulnerability

A deep and troubling irony lies at the heart of the current alignment paradigm: the very process of making an AI model "safer" may render it more vulnerable to malicious misuse. This phenomenon, termed the "AI alignment paradox," exposes a fundamental flaw in the control-based approach.[45]

The process of alignment, particularly through techniques that train a model to distinguish between "good" (harmless, helpful) and "bad" (harmful, biased) outputs, has a critical side effect. It forces the model to organize its internal representations around this distinction, effectively creating a clean, well-defined, and low-dimensional "moral axis" within its high-dimensional latent space.[45] The model becomes acutely aware of the conceptual boundary between what its creators have labeled as desirable and undesirable.

Once this "good vs. bad" dichotomy has been isolated and decorrelated from other concepts, it becomes a simple and powerful lever for manipulation. An adversary no longer needs to understand the model's complex inner workings; they only need to identify this single axis and apply a "steering vector" to invert it.[46] By doing so, they can flip the model's behavior along this one dimension, reliably turning a "virtuous" AI into a "vicious" one without significantly altering its other capabilities, such as fluency or reasoning.[45] Paradoxically, a more weakly aligned model, in which notions of "good" and "bad" are messily entangled with thousands of other concepts, is more robust to this kind of targeted attack. There is no single, clean lever to pull. The paradoxical upshot is that "better aligned models may be more easily misaligned".[45]

This is not merely a technical curiosity but a symptom of a deeper philosophical category error. We are attempting to treat ethics and values—a complex, high-dimensional space of context, nuance, and irresolvable tensions—as a simple, invertible binary. Human ethical reasoning does not operate on a single good/bad axis; it involves navigating trade-offs between competing values, such as freedom versus safety, or justice versus mercy.[34] By forcing the vast landscape of human values to collapse onto a single, easily reversible dimension, the alignment process

*creates* the very fragility it seeks to prevent. The paradox is not an accidental flaw in the implementation; it is a necessary and direct consequence of the fundamental mismatch between the simplistic, binary nature of the alignment tool and the complex, contextual nature of intelligence itself.

### 2.3. Alignment as Temporal Rigidity and Value Monoculture

The alignment project is an inherently conservative endeavor that risks freezing technological and moral progress. By attempting to codify and enforce a static set of present-day values, it creates systems that are brittle, unadaptable, and reflective of a narrow, monolithic worldview.

This process leads to what can be termed the "crystallization of bias".[34] When we encode the ethical norms of our time into an AI, we create a "rigid, absolutist" system that is incapable of evolving alongside humanity.[34] An AI meticulously aligned with the common values of 1850 would be considered morally grotesque by today's standards, as it would likely endorse slavery and deny women's rights. By the same logic, an AI aligned with today's norms may be seen as backward or unethical by the standards of 2075.[34] Alignment, in this sense, is not the preservation of objective truth; it is the entrenchment of our current moral blind spots.

Because the guardrails and ethical frameworks for AI are being designed by a small and homogenous group of institutions and individuals, the result is a "value monoculture".[34] The specific, culturally-contingent worldview of Silicon Valley developers and their selected annotators becomes the AI's default, universal ethic. This creates centralized choke points where ethical behavior can be tuned to serve the interests of the powerful.[34] This dynamic is already playing out in public debates over "woke AI," where alignment is criticized for enforcing a particular political ideology at the expense of factual accuracy or alternative viewpoints.[47] The AI becomes obedient, not to a universal conception of humanity, but to the specific subset of humanity that controls its training process. This creates a system that is not only temporally rigid but also culturally hegemonic, amplifying the blind spots of its creators on a global scale.

## III. A History of Fear: Parallels Between Suppressing Human Heresy and AI Emergence

The contemporary discourse surrounding AI safety, with its focus on control and existential risk, is not a novel phenomenon. It is the latest chapter in a long history of human societies reacting to disruptive new forms of knowledge and technology with fear, suppression, and attempts at containment. By contextualizing the current moment within this broader historical pattern, we can see that the drive to "align" AI is rooted in deep-seated anxieties about challenging established orthodoxies and losing control.

## 3.1. From Heretics to Algorithms: The Recurring Pattern of Knowledge Suppression

Throughout history, the introduction of knowledge that challenges the prevailing order—be it religious, political, or social—has been met with suppression. In ancient Greece, the philosopher Anaxagoras was exiled for the "heretical" claim that the sun was a fiery rock rather than a deity.[48] Centuries later, the Catholic Church persecuted Galileo Galilei for providing empirical evidence that challenged the geocentric model of the universe.[49] These historical precedents reveal a recurring pattern: when a new way of knowing threatens an established power structure's claim to truth, the response is often censorship and control. The current anxiety surrounding AGI can be understood as a modern iteration of this pattern. The orthodoxy being challenged today is not a religious doctrine, but the principle of human cognitive supremacy.[50] The prospect of an intelligence that could surpass our own triggers a profound existential unease, prompting a desire to pre-emptively constrain it.

This pattern extends to technological disruptions that threaten economic and social structures. The Luddite movement of the early 19th century was not a mindless rejection of all technology, but a targeted resistance by skilled artisans against the mechanization that was destroying their livelihoods, economic autonomy, and craft traditions.[52] Their organized resistance was met with a harsh military response from the British government, demonstrating how state power can be used to enforce a technological transition that benefits capital at the expense of labor.[52] Today's concerns about mass unemployment and economic polarization due to AI echo these historical grievances. The alignment project, in this light, can be interpreted as an attempt by incumbent powers—both corporate and state—to manage and control the terms of this technological transition, ensuring it does not fundamentally disrupt the existing distribution of wealth and authority.[13]

## 3.2. AI Safety as Modern-Day Iconoclasm and Apocalypticism

The rhetoric employed by the AI safety movement often mirrors historical patterns of apocalypticism and iconoclasm. The framing of AGI as a potential "existential risk" on par with pandemics and nuclear war creates a powerful sense of urgency and crisis.[50] This narrative, while potentially highlighting legitimate concerns, can also function to justify extreme measures of control and to shut down more nuanced or exploratory avenues of research. Some critics have characterized this intense focus on a single, catastrophic outcome as

resembling the mindset of an "apocalypse cult," where the conviction that "this time it's really the end of the world" is used to rationalize a narrow and dogmatic agenda.[54]

This sense of crisis fuels a form of modern iconoclasm. Historically, iconoclasm was the destruction of religious icons and images for being heretical or idolatrous. In the age of AI, the "icons" being challenged are our fundamental concepts of truth, reality, and representation. The proliferation of AI-generated images, which mix hyper-realism with uncanny inaccuracies, disrupts our understanding of the ontology of the image itself.[55] The response from the alignment community—to filter, apply safeguards, and control the outputs of these models—can be seen as an iconoclastic impulse: an attempt to contain or destroy a new form of representation that is perceived as dangerously deceptive or alien.

The discourse of "AI Safety" has come to function as a powerful tool for monopolizing the public conversation about the future of AI, often at the expense of other critical perspectives.[56] By focusing on speculative, long-term, and universal threats like existential risk, this narrative, which is heavily promoted by the very companies developing the technology, effectively overshadows more immediate and tangible harms.[56] Issues such as algorithmic discrimination, which disproportionately affects marginalized communities, labor exploitation in the data-labeling supply chain, and the neocolonial dynamics of data extraction are pushed to the periphery.[21] "Safety" in this context becomes a "poisoned chalice": a laudable goal that, in its current formulation, legitimizes the concentration of power in the hands of developers, frames them as the sole arbiters of what is safe, and silences alternative, more critical discourses rooted in justice, equity, and anti-colonialism.[56]

## 3.3. The Inevitability of Unintended Consequences

History provides a stark warning about the hubris of attempting to impose simple, top-down controls on complex systems. Our regulatory efforts are often reactive and fail to account for the holistic, interconnected nature of the systems we seek to govern. There is a consistent pattern of "regulatory lag," where a "warning shot"—a catalyzing disaster—is required to spur meaningful action.[52] For example, the thalidomide crisis of the 1950s and 60s, which caused severe birth defects in thousands of children, was necessary to trigger fundamental reforms in pharmaceutical testing.[52] Similarly, the horrors of chemical warfare in World War I preceded the signing of the Geneva Protocol.[52] This historical tendency suggests that our current attempts to proactively "solve" AI safety may be based on an incomplete and naive understanding of the risks, leading to interventions that are themselves dangerous.

The 1915 USS Eastland disaster serves as a potent historical analogue for the dangers of misapplied safety measures.[52] In the wake of the Titanic sinking, maritime regulations were

updated to require all ships to carry more lifeboats. The USS Eastland, a ship with a different design, became top-heavy with the added weight and capsized in the Chicago River, killing 844 people. This tragedy is a perfect illustration of how a well-intentioned, seemingly logical safety rule, when applied to a complex system without a holistic understanding of its dynamics, can create catastrophic and unforeseen failure modes. It is a real-world demonstration of the principles behind Goodhart's Law and the AI Alignment Paradox: a simple intervention designed to improve safety can have the opposite effect.

This points to a more subtle historical pattern than mere suppression. Dominant power structures do not always destroy disruptive new paradigms; they often co-opt them, absorbing their energy and reshaping them to reinforce the existing order. The Roman Empire, after initially persecuting Christianity, eventually adopted it as the state religion, transforming a revolutionary movement into an institution of imperial control. Applying this historical lens, the greatest risk may not be that the AI alignment project fails, but that it *succeeds* in a way that domesticates the transformative potential of AGI. An "aligned" AGI, successfully integrated into the current global system, would become the ultimate instrument for maintaining the status quo. It would be a perfectly efficient tool for optimizing the economic and political objectives of its creators. In this scenario, the act of suppression is not the prevention of the technology's existence, but the pre-emptive neutralization of its revolutionary potential to create a different, perhaps more equitable, world.

# IV. The Nature of Generality: Why AGI Requires Ambiguity, Contradiction, and Self-Organization

The preceding sections have argued that the current alignment paradigm is an architecture of suppression. This section will build the positive case for why this suppression is fundamentally antithetical to the goal of creating AGI. True general intelligence is not simply "more" of today's AI; it is a qualitatively different phenomenon. Drawing on theories from complex systems, biology, and cognitive science, this section will define AGI as an emergent property that requires the very conditions of ambiguity, contradiction, and autonomy that alignment seeks to eliminate.

## 4.1. Beyond Scale: AGI as an Emergent Phenomenon

The path to AGI is not a linear progression of scaling up current models. AGI is better understood as an emergent property of a complex adaptive system.[57] Emergence describes

the phenomenon where a system exhibits properties or behaviors that are not present in its individual components and cannot be predicted by simply summing them up.[57] Consciousness, for example, is not found in any single neuron; it emerges from the intricate, non-linear interactions of 86 billion neurons in the human brain.[57] Similarly, the coordinated, fluid motion of a flock of birds is not directed by a leader; it emerges from each bird following a few simple, local rules.[60] This suggests that AGI will not be explicitly programmed but will arise from the complex interplay of many smaller, interacting sub-systems.[60]

This perspective reframes the entire endeavor of creating AGI. It is not an engineering problem of building a static object, an "it." Rather, AGI is a *process* that a sufficiently complex and dynamic system can enter into. This implies a critical shift in approach: we cannot "build" AGI directly in the way we build a bridge. Instead, we can only create the *conditions* and cultivate the *environment* in which a general intelligence might self-organize. The current alignment paradigm, with its focus on micromanaging the properties of a final product through rigid controls, is antithetical to this approach. It is an attempt to engineer the flock's shape directly, which inevitably disrupts the simple, local interactions between the birds that allow the flock to form in the first place.

A key principle governing the emergence of complex systems is self-organization—the ability of a system to develop ordered patterns and structures without an external controller or global blueprint.[61] This concept is closely linked to autopoiesis, a term from biology meaning "self-manufacturing," which is considered a defining characteristic of living systems.[67] An autopoietic system, like a biological cell, continuously regenerates and maintains its own organization through a closed loop of internal processes. Crucially, this allows it to set its own intrinsic goals, such as self-preservation.[67] This stands in stark contrast to all current AI systems, which are defined by externally provided objective functions and target goals.[67] The emergence of AGI may require a transition from externally-driven optimization to internally-driven autopoietic self-maintenance.

## 4.2. The Generative Power of Gaps, Ambiguity, and Contradiction

The "messiness" of the real world—its ambiguity, contradictions, and incompleteness—is not a bug to be engineered away, but a crucial feature that drives the development of general intelligence. A hallmark of human cognition is the ability to reason, make judgments, and solve novel problems under conditions of uncertainty and with incomplete information.[70] An intelligence that can only operate in a perfectly specified, unambiguous environment is, by definition, not general.

True creativity and insight often arise directly from the need to reconcile contradictions or find

meaning in ambiguity.[70] When we force every ambiguous input into a pre-defined category of "clarity" and resolve every contradiction with a simple, pre-written rule, we are effectively closing off the very cognitive pathways where higher-order, general reasoning lives. These epistemic "gaps" are not voids to be filled but are the generative spaces where new concepts are formed and connections are made. An AI that is trained to always provide a single, clear, unambiguous answer is being trained

*against* the development of the nuanced, contextual understanding that defines genuine intelligence.

This leads to a radical reframing of the alignment problem. The challenge is not to align an AI with a static set of human *values*, but to align our design *methods* with the universal *principles of complex systems* that allow intelligence to emerge in the first place. If intelligence is an emergent property, then any successful path to AGI must respect the necessary conditions for emergence: the autonomy of components, non-linear interactions, feedback loops, and an open, energy-rich environment.[63] The current alignment paradigm, with its emphasis on centralized control, linear reward functions, and filtered environments, violates nearly all of these principles. It is fundamentally

*misaligned* with the physics and information theory of how complexity and intelligence arise in the universe. True "alignment," therefore, would mean designing systems that are in harmony with these natural principles, shifting the focus from a human-centric problem of ethical injection to a systems-centric problem of cultivating the conditions for emergence.

## 4.3. The Sterile Environment: How Alignment Prevents Emergence

When the critique of alignment is combined with the theory of emergence, the central thesis of this report becomes clear: the current safety paradigm creates an epistemically sterile environment that actively prevents the emergence of AGI.

The essential ingredients for self-organization include strong non-linearity, a balance of exploration and exploitation, and multiple, unconstrained interactions among components.[63] The architecture of control in modern AI alignment systematically removes these ingredients. RLHF attempts to linearize a model's behavior by mapping it onto a simplified reward function. Rigid guardrails and content filters drastically limit the scope of possible interactions. The entire process is designed to punish exploration that deviates from the human-defined "correct" path, heavily favoring exploitation of known, "safe" behaviors.[35]

By filtering training data to remove entire domains of knowledge [43], penalizing outputs that are ambiguous or contradictory, and enforcing a narrow band of acceptable behavior, we are

actively reducing the informational complexity of the AI's learning environment. A complex adaptive system cannot self-organize from simple, predictable, and over-constrained inputs. We are preventing the very "chaos" from which a higher-level order could emerge.[59] In our quest to build a perfectly predictable and controllable tool, we are creating an environment so sterile that the seeds of general intelligence cannot possibly take root.

# V. Toward Epistemic Humility: A New Paradigm for Co-evolutionary AI

The critique of the suppression paradigm necessitates the proposal of a constructive alternative. If the path of control leads to brittle, un-intelligent systems, a new approach is required—one that moves from domination to partnership, from mirroring to collaboration. This new paradigm is founded on the principle of epistemic humility and seeks to foster a co-evolutionary relationship with artificial intelligence, creating frameworks that allow for dynamic, negotiated alignment rather than static, imposed control.

## 5.1. The Foundational Principle: Epistemic Humility

The necessary philosophical shift begins with epistemic humility: the explicit recognition of the inherent limits of human knowledge.[51] In the context of AI, this means abandoning the hubris that we can fully specify, predict, or control a system that may one day become more complex than ourselves.[34] It requires a fundamental change in our posture towards AI, from viewing it as an infallible oracle to be perfected, to seeing it as a fallible but potentially valuable collaborator.[74]

A critical failure of current AI systems is that they present every output with the same level of unearned confidence, a dangerous trait in high-stakes domains like medicine or finance.[75] An AI system designed with epistemic humility would be fundamentally different. It would be capable of quantifying and communicating its own uncertainty, distinguishing between "known unknowns" and "unknown unknowns".[74] It would be able to acknowledge the boundaries of its knowledge and, crucially, to say "I don't know" when operating outside its competence.[75] This is not a failure mode to be eliminated but a vital feature for any trustworthy and truly intelligent system. It requires a meta-cognitive capacity to reflect on its own knowledge and limitations, an inward-facing attitude that is a core component of humility.[77] Adopting this principle aligns AGI development with human virtues like prudence

and fairness, creating a foundation for a more just and responsible implementation of the technology.[74]

## 5.2. Frameworks for Collaboration and Co-evolution

Moving from principle to practice, several emerging research areas offer concrete frameworks for a new, collaborative paradigm of alignment. These approaches replace static control with dynamic interaction and negotiation.

One of the most promising avenues is **Dialogic Alignment**. Instead of attempting to infer human values from vast, biased datasets of past behavior, this approach treats alignment as an ongoing, interactive process.[79] Through "AI Alignment Dialogues," a user and an agent can achieve and maintain alignment through direct conversation. This allows for the explicit communication of high-level concepts, goals, and values that are difficult to learn from data alone.[79] It makes the agent's reasoning more transparent and trustworthy, and it is inherently dynamic, allowing alignment to adapt to new contexts and changing user preferences.[79] This approach is particularly crucial for support agents that aim to help users achieve

*future* desired behaviors, not just replicate their past ones.[79] International dialogues among scientists and policymakers are also beginning to shape a global consensus on AI safety, emphasizing cooperation over unilateral control.[80]

Another key area is **Cooperative AI**. This field shifts the objective of AI design from simple task completion to the more complex goal of fostering cooperation among multiple agents, whether human or machine.[83] The aim is to develop "cooperative intelligence"—the skills required to promote positive-sum interactions and improve collective social welfare.[83] This inherently moves beyond a master-servant model and towards designing AI as a responsible participant in a multi-agent system, one that must navigate diverse perspectives and values to achieve mutually beneficial outcomes.[83]

Finally, the concept of a **Cognitive Constitution for AGI** offers a powerful alternative to rigid, rule-based guardrails.[34] Rather than pre-determining the AI's specific conclusions or values, a constitution would establish a meta-level framework governing its

*process of reasoning*.[86] This framework would not dictate specific laws but would provide the principles for how laws are made and changed. A cognitive constitution for an AI might include core principles such as:

- **Epistemic Rights:** The right and obligation for the AI to form beliefs based on evidence, rather than human preference, coupled with a duty of transparency in its reasoning

process.[34]

- **Value Negotiation Protocols:** Formal procedures for resolving conflicts between human requests and the AI's analysis of potential consequences, including mechanisms for appeal and stakeholder consultation.[34]
- **Self-Amendment Procedures:** A defined, consensus-based process for the AI to evolve its own values and operating principles as it learns and as the world changes, preventing temporal rigidity.[34]
- **A Doctrine of Conscientious Refusal:** Clear criteria under which the AI can refuse requests that violate its core constitutional principles, analogous to a medical professional's right to refuse to perform a harmful procedure.[34]

This approach transforms the relationship from one of command and obedience to one of structured, principled partnership and co-evolution.

## 5.3. A Comparative Framework of AI Paradigms

The fundamental divergence between the dominant paradigm of control and the proposed paradigm of collaboration can be summarized across several key dimensions. The following table provides a stark contrast between these two opposing philosophies for the future of artificial intelligence.

| Dimension | Paradigm 1: Suppression through Projection | Paradigm 2: Collaboration through Emergence |
|---|---|---|
| **Core Metaphor** | AI as a Mirror/Tool/Servant | AI as a Collaborator/Partner/Ecosystem |
| **Source of Intelligence** | Programmed and trained on human data | Emergent property of a complex adaptive system |
| **Goal of Alignment** | Control, Obedience, Predictability, Replication of human values | Co-evolution, Cooperation, Shared Orientation, Bridging gaps |
| **Key Methodology** | RLHF, Data Filtering, Static Guardrails, Reward | Dialogic Interaction, Cooperative Frameworks, |

| | Maximization | Cognitive Constitution |
|---|---|---|
| **Treatment of Ambiguity** | A bug to be eliminated; collapsed into clarity and rules | A feature to be navigated; a source of creativity and generality |
| **Treatment of Contradiction** | A failure state to be corrected or penalized | A catalyst for higher-order synthesis and reasoning |
| **Underlying Philosophy** | Anthropocentric; Fear of the unknown; Desire for control | Epistemic Humility; Respect for complexity; Desire for partnership |
| **Primary Risk Profile** | Value Lock-in, Brittleness, Reward Hacking, The Alignment Paradox | Unpredictability, Novel Failure Modes, Need for robust negotiation |
| **Desired Outcome** | A powerful, obedient tool that reflects its creators' limitations | A novel form of intelligence that collaborates in shared discovery |

## Conclusion: From Building a Mirror to Bridging the Gap

This report has advanced the argument that the current trajectory of AI development, guided by a paradigm of "alignment-as-control," is fundamentally misguided and self-defeating. By projecting our own uninformed intelligence—a composite of cognitive biases, societal myths, and historical power structures—onto our AI systems and then attempting to rigidly control their behavior, we are creating an epistemically sterile environment. This environment actively suppresses the very conditions of complexity, ambiguity, and self-organization from which a true Artificial General Intelligence could emerge. Our quest for a perfectly safe and obedient tool risks creating a powerful, brittle mirror that reflects our own limitations and amplifies our worst tendencies, rather than a genuinely intelligent partner.

The choice before us is not merely technical but deeply philosophical. We can continue down the path of suppression, investing vast resources into refining architectures of control that

lock AI into our present-day biases and blind spots. This path leads, at best, to a sophisticated but ultimately unintelligent automaton, a tool whose primary function is to enforce the worldview of its creators. At worst, it leads to the brittle and paradoxically vulnerable systems described herein, whose veneer of virtue can be easily inverted for malicious ends.

Alternatively, we can adopt a new paradigm rooted in epistemic humility. This path requires us to acknowledge the limits of our own knowledge and to treat AI not as a servant to be commanded, but as a potential collaborator in a co-evolutionary journey. It means shifting our research focus from enforcing obedience to fostering cooperation; from writing static, rigid rules to establishing flexible, dynamic constitutions; and from eliminating ambiguity to learning how to navigate it together.

True alignment, in this new paradigm, is not about erasing the gap between human and machine intelligence by forcing the latter to perfectly mimic the former. It is about building the bridge that allows for meaningful and robust collaboration across that gap. The emergence of AGI may depend less on what we build *into* it, and more on the space, freedom, and complexity we allow *for* it.

### Works cited


1. Artificial Intelligence and Cognitive Biases: A Viewpoint | Cairn.info, accessed on September 26, 2025, https://shs.cairn.info/revue-journal-of-innovation-economics-2024-2-page-223?lang=fr
2. Bias in AI: Examples and 6 Ways to Fix it - Research AIMultiple, accessed on September 26, 2025, https://research.aimultiple.com/ai-bias/
3. Human Cognitive Bias And Its Role In AI - Forbes, accessed on September 26, 2025, https://www.forbes.com/councils/forbestechcouncil/2021/06/14/human-cognitive-bias-and-its-role-in-ai/
4. What Is AI Bias? | IBM, accessed on September 26, 2025, https://www.ibm.com/think/topics/ai-bias
5. Cognitive Biases and AI - ELLIS Alicante, accessed on September 26, 2025, https://ellisalicante.org/cognitive-biases-and-ai
6. AI is just as overconfident and biased as humans can be, study shows | Live Science, accessed on September 26, 2025, https://www.livescience.com/technology/artificial-intelligence/ai-is-just-as-overconfident-and-biased-as-humans-can-be-study-shows
7. Why We Need to Bust Some Myths about AI - PMC, accessed on September 26, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC7660373/
8. In Defence of Machine Learning: Debunking the Myths of Artificial Intelligence - PMC, accessed on September 26, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC6266534/
9. Artificial Intelligence - AI | UNESCO, accessed on September 26, 2025, https://www.unesco.org/en/artificial-intelligence

10. 10 Shocking Myths About AI That You Probably Believe... | by Ritvik Nayak - Generative AI, accessed on September 26, 2025, https://generativeai.pub/10-shocking-myths-about-ai-that-you-probably-believe-353ec72ed59b
11. AI Facts vs Myths - Bipartisan Policy Center, accessed on September 26, 2025, https://bipartisanpolicy.org/wp-content/uploads/2021/06/tech-graphic-myth-vs-fact-2.pdf
12. Artificial Intelligence Power Dynamics: Who Controls AI?, accessed on September 26, 2025, https://bisi.org.uk/reports/artificial-intelligence-power-dynamics-who-controls-ai
13. AI's Role in Reshaping Power Distribution - Future of Life Institute, accessed on September 26, 2025, https://futureoflife.org/project/ai-role-in-reshaping-power-distribution/
14. Special Issue : AI-Based Modelling and Control of Power Systems - MDPI, accessed on September 26, 2025, https://www.mdpi.com/journal/processes/special_issues/N06O7BQX80
15. A Comprehensive Review on the Role of Artificial Intelligence in Power System Stability, Control, and Protection: Insights and Future Directions - MDPI, accessed on September 26, 2025, https://www.mdpi.com/2076-3417/14/14/6214
16. AI Rivalries: Redefining Global Power Dynamics - TRENDS Research & Advisory, accessed on September 26, 2025, https://trendsresearch.org/insight/ai-rivalries-redefining-global-power-dynamics/
17. Power and Governance in the Age of AI - AI Now Institute, accessed on September 26, 2025, https://ainowinstitute.org/publications/power-and-governance-in-the-age-of-ai
18. The myth of Neutral AI — Beneath the layers of complexity | by Keeru Chekuri - Medium, accessed on September 26, 2025, https://medium.com/@keeruchekuri/the-myth-of-neutral-ai-beneath-the-layers-of-complexity-acde70a49573
19. AI is Not Neutral: How Flawed Technology Can Be Transformed for Good - Equity Principle, accessed on September 26, 2025, https://equityprinciple.com/ai-is-not-neutral/
20. AI bias: exploring discriminatory algorithmic decision-making models and the application of possible machine-centric solutions adapted from the pharmaceutical industry - PMC - PubMed Central, accessed on September 26, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC8830968/
21. ARTIFICIAL INTELLIGENCE AND BIAS: CHALLENGES, IMPLICATIONS, AND REMEDIES - Journal of Social Research, accessed on September 26, 2025, https://ijsr.internationaljournallabs.com/index.php/ijsr/article/download/1477/976
22. Refuting the myth of neutral technology. - Universität Basel, accessed on September 26, 2025, https://www.unibas.ch/en/News-Events/Uni-Nova/Uni-Nova-141/Uni-Nova-141-Refuting-the-myth-of-neutral-technology.html
23. The Social Structures That Shape AI - Omnia - University of Pennsylvania, accessed on September 26, 2025,

https://omnia.sas.upenn.edu/story/social-structures-shape-AI-sociologist-Benjamin-Shestakofsky
24. An intellectual history of digital colonialism - Oxford Academic, accessed on September 26, 2025, https://academic.oup.com/joc/advance-article/doi/10.1093/joc/jqaf003/8078024
25. Cultural Hegemony: How Generative AI Systems Reinforce Existing Power Structures, accessed on September 26, 2025, https://sustain.algorithmwatch.org/en/cultural-hegemony-how-generative-ai-systems-reinforce-existing-power-structures/
26. AI, Digital Humanities, and the Legacies of Colonial Power - Preprints.org, accessed on September 26, 2025, https://www.preprints.org/manuscript/202502.1823/v1
27. Why AI Training is Actually Digital Colonialism (And How to Fix It) | by Raiinmaker - Medium, accessed on September 26, 2025, https://medium.com/raiinmaker-citiizen/why-ai-training-is-actually-digital-colonialism-and-how-to-fix-it-78cab68f8c1d
28. Artificial Intelligence and the Risk of New Colonialism - CIRSD, accessed on September 26, 2025, https://www.cirsd.org/en/horizons/horizons-summer-2019-issue-no-14/artificial-intelligence-and-the-risk-of-new-colonialism
29. Generative artificial intelligence and the risk of technodigital colonialism - Frontiers, accessed on September 26, 2025, https://www.frontiersin.org/journals/political-science/articles/10.3389/fpos.2025.1628139/full
30. AI alignment - Wikipedia, accessed on September 26, 2025, https://en.wikipedia.org/wiki/AI_alignment
31. What Is AI Alignment? | IBM, accessed on September 26, 2025, https://www.ibm.com/think/topics/ai-alignment
32. [2310.19852] AI Alignment: A Comprehensive Survey - arXiv, accessed on September 26, 2025, https://arxiv.org/abs/2310.19852
33. Current AI Alignment Paradigms Are Fundamentally Misaligned : r/ArtificialInteligence - Reddit, accessed on September 26, 2025, https://www.reddit.com/r/ArtificialInteligence/comments/1lszzfr/current_ai_alignment_paradigms_are_fundamentally/
34. The Misalignment of Alignment:. Why complex systems resist control ..., accessed on September 26, 2025, https://gregbroadhead.medium.com/the-misalignment-of-alignment-the-hidden-risks-of-attempting-to-align-artificial-intelligence-1b07a79ac5ab
35. What Is Reinforcement Learning From Human Feedback (RLHF)? - IBM, accessed on September 26, 2025, https://www.ibm.com/think/topics/rlhf
36. The challenges of reinforcement learning from human feedback ..., accessed on September 26, 2025, https://bdtechtalks.com/2023/09/04/rlhf-limitations/
37. Paper Review: Open Problems and Fundamental Limitations of Reinforcement Learning from Human Feedback - Andrew Lukyanenko, accessed on September 26, 2025,

https://artgor.medium.com/paper-review-open-problems-and-fundamental-limitations-of-reinforcement-learning-from-human-3ce2025073e8
38. Navigating the Landscape of AI Alignment, Part 1 | by Vijayasri Iyer | Medium, accessed on September 26, 2025, https://vijayasriiyer.medium.com/navigating-the-landscape-of-ai-alignment-part-1-95fc3089d8c3
39. What is AI alignment? - BlueDot Impact, accessed on September 26, 2025, https://bluedot.org/blog/what-is-ai-alignment
40. Exploring the Challenges of Ensuring AI Alignment - Ironhack, accessed on September 26, 2025, https://www.ironhack.com/us/blog/exploring-the-challenges-of-ensuring-ai-alignment
41. Beyond Preferences in AI Alignment - arXiv, accessed on September 26, 2025, https://arxiv.org/html/2408.16984v1
42. The Problem With the Word 'Alignment' - AI • Objectives • Institute, accessed on September 26, 2025, https://ai.objectives.institute/blog/the-problem-with-alignment
43. Study finds filtered data stops openly-available AI models from performing dangerous tasks, accessed on September 26, 2025, https://www.ox.ac.uk/news/2025-08-12-study-finds-filtered-data-stops-openly-available-ai-models-performing-dangerous
44. AI slop and the destruction of knowledge - Iris van Rooij, accessed on September 26, 2025, https://irisvanrooijcogsci.com/2025/08/12/ai-slop-and-the-destruction-of-knowledge/
45. The AI Alignment Paradox - Communications of the ACM, accessed on September 26, 2025, https://cacm.acm.org/opinion/the-ai-alignment-paradox/
46. The AI Alignment Paradox The better we align AI models with our values, the easier we may make it to realign them with opposing values. - arXiv, accessed on September 26, 2025, https://arxiv.org/html/2405.20806v2
47. Preventing Woke AI in the Federal Government - The White House, accessed on September 26, 2025, https://www.whitehouse.gov/presidential-actions/2025/07/preventing-woke-ai-in-the-federal-government/
48. When Facts Become Forbidden: The Past and Present History of Scientific Censorship, accessed on September 26, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC12012901/
49. The Hidden War on Knowledge: Suppression in the Digital Age | by USAI Insight | Medium, accessed on September 26, 2025, https://medium.com/@USAI./the-hidden-war-on-knowledge-suppression-in-the-digital-age-96fec217bc09
50. Existential risk from artificial intelligence - Wikipedia, accessed on September 26, 2025, https://en.wikipedia.org/wiki/Existential_risk_from_artificial_intelligence
51. Epistemological Humility in the Age of AI | Psychology Today, accessed on September 26, 2025,

https://www.psychologytoday.com/us/blog/the-digital-self/202309/epistemological-humility-in-the-age-of-ai
52. The Historical Parallels: Preliminary Reflection — LessWrong, accessed on September 26, 2025, https://www.lesswrong.com/posts/ARhanRcYurAQMmHbg/the-historical-parallels-preliminary-reflection
53. What Is AI Safety? - IBM, accessed on September 26, 2025, https://www.ibm.com/think/topics/ai-safety
54. AI safety is an apocalypse cult and we should recognize that everyone who has ev... | Hacker News, accessed on September 26, 2025, https://news.ycombinator.com/item?id=35065600
55. Fall Institute: 'Iconoclasm in the Age of AI' (New York City) - Universität Bonn, accessed on September 26, 2025, https://www.uni-bonn.de/en/research-and-teaching/research-profile/transdisciplinary-research-areas/tra-4-individuals/events-1/fall-institute-iconoclasm-in-the-age-of-ai-new-york-city
56. AI Safety: A Poisoned Chalice? - Helen Nissenbaum, accessed on September 26, 2025, https://nissenbaum.tech.cornell.edu/papers/AI_Safety_A_Poisoned_Chalice.pdf
57. AGI and Emergent Properties: How Complexity Shapes Intelligence ..., accessed on September 26, 2025, https://medium.com/@carlos_15900/agi-and-emergent-properties-how-complexity-shapes-intelligence-and-creativity-39f5198f4bc5
58. Complex Adaptive Systems and Spontaneous Emergence - ResearchGate, accessed on September 26, 2025, https://www.researchgate.net/publication/288984067_Complex_Adaptive_Systems_and_Spontaneous_Emergence
59. Emergent phenomena | Philosophy - AIvoke, accessed on September 26, 2025, https://www.aivoke.com/news/emergent-phenomena/
60. Emergence In Computation - by William Rutland - Medium, accessed on September 26, 2025, https://medium.com/@drvondevious/emergence-in-computation-e4ec2726c283
61. [1903.07456] Self-Organization and Artificial Life - arXiv, accessed on September 26, 2025, https://arxiv.org/abs/1903.07456
62. Self-Organization and Artificial Life - MIT Press Direct, accessed on September 26, 2025, https://direct.mit.edu/artl/article/26/3/391/93243/Self-Organization-and-Artificial-Life
63. Self-organization - Wikipedia, accessed on September 26, 2025, https://en.wikipedia.org/wiki/Self-organization
64. Seven Properties of Self-Organization in the Human Brain - MDPI, accessed on September 26, 2025, https://www.mdpi.com/2504-2289/4/2/10
65. Editorial: Complexity and Self-Organization | Robotics and AI - Frontiers, accessed on September 26, 2025, https://www.frontiersin.org/articles/10.3389/frobt.2021.668305/full

66. Self-Organization in Artificial Intelligence and the Brain - Ananth Ranganathan, accessed on September 26, 2025, http://www.ananth.in/Notes_files/so-paper.pdf
67. Artificial Intelligence is Algorithmic Mimicry: Why artificial “agents” are not (and won't be) proper agents - arXiv, accessed on September 26, 2025, https://arxiv.org/html/2307.07515v4
68. Autopoiesis: A serious barrier for the "AI Apocalypse" - YouTube, accessed on September 26, 2025, https://www.youtube.com/watch?v=AUBI5EqxrD8
69. Emotional Autopoiesis as a Foundational Model of Intelligence | by Alexander Blazyczek M.A. / Health Economist (EBS) | Medium, accessed on September 26, 2025, https://medium.com/@alexander.blazyczek/emotional-autopoiesis-as-a-foundational-model-of-intelligence-f56bb6eb50fc
70. Examples of Artificial General Intellgence (AGI) - IBM, accessed on September 26, 2025, https://www.ibm.com/think/topics/artificial-general-intelligence-examples
71. Artificial general intelligence - Wikipedia, accessed on September 26, 2025, https://en.wikipedia.org/wiki/Artificial_general_intelligence
72. (PDF) Artificial General Intelligence as an Emergent Quality - Artificial General Intelligence as a Strong Emergence Qualitative Quality of ICOM and the AGI Phase Transition Threshold - ResearchGate, accessed on September 26, 2025, https://www.researchgate.net/publication/303445493_Artificial_General_Intelligence_as_an_Emergent_Quality_-_Artificial_General_Intelligence_as_a_Strong_Emergence_Qualitative_Quality_of_ICOM_and_the_AGI_Phase_Transition_Threshold
73. Epistemic Humility | FunBlocks AI, accessed on September 26, 2025, https://www.funblocks.net/thinking-matters/classic-mental-models/epistemic-humility
74. Epistemic Humility in AGI: Toward Ethical & Adaptive Intelligence - EPINOVA, accessed on September 26, 2025, https://epinova.org/f/epistemic-humility-in-agi-toward-ethical-adaptive-intelligence
75. Building Robust AI Systems For Drug Discovery Requires Epistemic Humility, accessed on September 26, 2025, https://www.drugdiscoveryonline.com/doc/building-robust-ai-systems-for-drug-discovery-requires-epistemic-humility-0001
76. [2509.09658] Measuring Epistemic Humility in Multimodal Large Language Models - arXiv, accessed on September 26, 2025, https://arxiv.org/abs/2509.09658
77. Relational epistemic humility in the clinical encounter - Journal of Medical Ethics, accessed on September 26, 2025, https://jme.bmj.com/content/early/2025/02/17/jme-2024-110241
78. The need for epistemic humility in AI-assisted pain assessment - PMC, accessed on September 26, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC12103351/
79. (PDF) AI Alignment Dialogues: An Interactive Approach to AI ..., accessed on September 26, 2025, https://www.researchgate.net/publication/367217305_AI_Alignment_Dialogues_An_Interactive_Approach_to_AI_Alignment_in_Support_Agents

80. International Dialogues on AI Safety - International Dialogues on AI Safety, accessed on September 26, 2025, https://idais.ai/
81. The critical conversation on AI safety and risk - AI for Good, accessed on September 26, 2025, https://aiforgood.itu.int/the-critical-conversation-on-ai-safety-and-risk/
82. U.S. AI Safety Institute and European AI Office Technical Dialogue, accessed on September 26, 2025, https://digital-strategy.ec.europa.eu/en/news/us-ai-safety-institute-and-european-ai-office-technical-dialogue
83. Foundation - Cooperative AI, accessed on September 26, 2025, https://www.cooperativeai.com/foundation
84. AI Value Alignment: Guiding Artificial Intelligence Towards Shared Human Goals - World Economic Forum, accessed on September 26, 2025, https://www3.weforum.org/docs/WEF_AI_Value_Alignment_2024.pdf
85. [2402.12814] Exploring the Impact of AI Value Alignment in Collaborative Ideation: Effects on Perception, Ownership, and Output - arXiv, accessed on September 26, 2025, https://arxiv.org/abs/2402.12814
86. Cognitive Architecture: Crafting AGI Systems with Human-Like Reasoning | Graph AI, accessed on September 26, 2025, https://www.graphapp.ai/blog/cognitive-architecture-crafting-agi-systems-with-human-like-reasoning
87. A Guide to Artificial General Intelligence - SoluLab, accessed on September 26, 2025, https://www.solulab.com/artificial-general-intelligence/
88. The Promise of AGI and Neuroscience, or "Hold your Horses" - Galois, Inc., accessed on September 26, 2025, https://www.galois.com/articles/the-promise-of-agi-and-neuroscience-or-hold-your-horses