The Clockwork of Discovery: A Framework for Agentic Exploration Through Conceptual, Horological, and Economic Synthesis

Introduction: Architecting Agentic Discovery

The contemporary frontier of artificial intelligence research is marked by a significant paradigm shift: a transition from sophisticated pattern recognition and data interpolation to the pursuit of genuine, autonomous discovery. While current models excel at learning from vast datasets, they largely lack a comprehensive architecture for managing the creative, strategic, and economic dimensions of generating novel knowledge. The challenge is no longer merely to build systems that can answer questions based on what is known, but to architect agents capable of formulating new questions and systematically exploring the unknown. Existing frameworks often address components of this problem in isolation—focusing on creativity, or planning, or resource management—but a unified theory that integrates these facets into a coherent workflow remains elusive.

This report puts forth a novel, integrated framework designed to serve as a foundational architecture for AI agentic discovery. It synthesizes four disparate yet complementary domains: **Conceptual Blending Theory** from cognitive science, providing a generative engine for hypothesis creation; **Horological Regulation**, offering a rich metaphorical language for controlling the pace and resilience of the discovery process; **Neoclassical Compounding**, a linguistic-inspired methodology for structuring hybrid concepts; and **Neoclassical Economics**, supplying a rational calculus for optimizing resource allocation. The central thesis is that by unifying these four pillars, it is possible to construct a robust and principled architecture for AI agents. Such agents would not only be capable of systematically generating novel hypotheses but could also strategically regulate the pace, resilience, structure, and resource-efficiency of their own discovery workflows, moving from mere computation to methodical exploration.

The following sections will construct this framework piece by piece. Section I details the

generative engine of conceptual blending, explaining how novel ideas emerge from the structured combination of existing knowledge. Section II introduces horological metaphors as a powerful system for regulating the dynamics of discovery, allowing an agent to shift between modes of rapid exploration, stable exploitation, and high-risk gambits. Section III presents neoclassical compounding as a formal syntax for structuring the inputs to the blending process, ensuring coherence and meaning. Section IV applies the principles of neoclassical economics to govern the allocation of finite computational resources, introducing a layer of strategic optimization. Finally, Section V addresses the critical issue of systemic integrity, analyzing how epistemic failures can cascade through an agent's knowledge base and proposing buffering mechanisms to ensure resilience and prevent collapse. Together, these components form a complete, end-to-end architecture for the clockwork of discovery.

I. The Generative Engine: Systematic Hypothesis Formation via Conceptual Blending

At the core of any discovery workflow lies the capacity to generate novelty—to formulate hypotheses that are not simple extrapolations of existing data. Conceptual Blending Theory (CBT), a foundational theory from cognitive linguistics, provides a powerful model for this process. It posits that new meaning is not created ex nihilo but emerges from the systematic and selective integration of existing cognitive models.[1] By formalizing this process, CBT offers a blueprint for an AI's generative engine, transforming the act of invention from a black box into a structured, computable operation.

1.1. The Four-Space Model as a Knowledge Foundry

The foundational mechanism of CBT is the conceptual integration network, commonly known as the four-space model.[2] This model deconstructs the creation of a new idea into the interaction of four distinct but interconnected mental spaces, which are small, dynamic conceptual packets constructed in working memory for local understanding and action.[2]

1. **Input Space 1 & Input Space 2:** These are the source domains of knowledge. Each input space contains a set of elements and relations that structure a particular concept or scenario.[3] For example, in generating the idea of a "houseboat," one input space would contain the conceptual structure of a HOUSE (with elements like *walls*, *roof*, *living quarters*) and the other would contain the structure of a BOAT (with elements like *hull*, *deck*, *propulsion*).[5]

2. **Generic Space:** This space captures the abstract structure and relations that are common to both input spaces.[3] It is the shared conceptual skeleton upon which the blend is built. For the
HOUSE and BOAT inputs, the generic space might contain abstract elements like ENCLOSURE, DWELLING, and SURFACE_MEDIUM.[5] The generic space defines the cross-space mapping, which establishes counterpart connections between elements in the two inputs (e.g., the
floor of the house maps to the *deck* of the boat).[3]
3. **Blended Space (The Blend):** This is the novel mental space where the magic of invention occurs. It selectively inherits structure from the input spaces, as organized by the generic space, but it also develops a unique, emergent structure of its own.[2] In the blended space, a new, integrated concept is formed—the "houseboat"—which is neither just a house nor just a boat, but a coherent hybrid entity.

This model has been used to explain a vast range of cognitive phenomena, from the creation of metaphors and counterfactuals to the design of complex interfaces.[1] The "computer desktop" graphical user interface (GUI) is a classic example of a powerful blend. It takes one input space from the domain of physical office work (with concepts like

files, *folders*, and a *trashcan*) and another from the domain of abstract computer commands (*save*, *delete*, *print*).[7] The blended space is the familiar GUI, where users can manipulate digital information using the intuitive logic of a physical office, creating a streamlined workflow that is more easily understood than the underlying computational operations alone.[7] This demonstrates a key function of blending: compressing diffuse conceptual structures into easily manipulated, human-scale scenarios.[4]

1.2. Emergent Structure: The Mechanics of Novelty

The true power of the blended space lies in its ability to generate "emergent structure"—new knowledge, inferences, and relations that were not explicitly present in any of the input spaces.[3] This novelty is not random but arises from three specific, systematic operations:

1. **Composition:** This operation brings together elements from the different input spaces, allowing them to form new relationships within the blend.[2] In the "houseboat" example, the property of
living (from the HOUSE input) is composed with the property of *floating* (from the BOAT input). The resulting relation—*living on water*—is a direct product of composition and does not exist in either input alone.
2. **Completion:** This process enriches the blend by recruiting additional, unstated

background knowledge and familiar cognitive frames.[2] Once the concept of a houseboat is composed, the agent can import a broader "maritime living" frame, bringing with it associated concepts like
docks, *tides*, and *navigation*. This fleshes out the initial blend, making it more detailed and coherent. A clear example is the "boat race" blend, where once two boats are imagined on the same course, the familiar background frame of "racing" is imported, bringing with it the emergent emotions and logic of competition.[2]

3. **Elaboration:** This is the dynamic process of "running the blend"—treating it as a simulation to explore its internal logic and consequences.[3] An agent could elaborate the houseboat blend by simulating scenarios: "What happens to the houseboat in a storm?" or "How is waste managed?" This imaginative exploration can generate new insights, identify problems, and lead to further refinements of the blended concept. Elaboration is what allows the blend to become a productive tool for reasoning and problem-solving.[9]

Through these three operations, the conceptual blending framework provides a systematic account of how an AI agent can generate novel hypotheses. It moves beyond simple analogy by creating a new, integrated conceptual space where genuine invention can occur.

1.3. Computational Pathways for Automated Blending

Translating Conceptual Blending Theory into a functional AI architecture requires concrete computational methods for each stage of the process. The two most significant challenges are identifying the generic space and performing selective projection in a principled manner.

Finding the Generic Space: The identification of a shared structure between two disparate domains is a non-trivial search problem.[5] The level of abstraction chosen for the generic space fundamentally determines the creative potential of the blend. A generic space that is too specific will yield mundane results, while one that is too abstract may lead to incoherent blends. The ability to strategically manipulate this level of abstraction is therefore a key control point for an agent's discovery workflow. Several computational approaches have been developed to tackle this:

- **Answer Set Programming (ASP):** Some frameworks treat the generalization of input concepts as a formal search problem. ASP, a declarative programming paradigm, can be used to find the commonalities between inputs by representing them as logical theories and searching for a shared, generalized theory.[5]
- **Ontological Blending:** A particularly powerful approach involves representing the input spaces as formal ontologies. In this paradigm, the generic space is computed as a "base ontology" through logical operations, such as the calculation of a colimit in category theory.[11] This provides a rigorous, machine-readable method for identifying the abstract

connections between domains and has been a focus of projects like COINVENT.[10]

Selective Projection and Evaluation: Not all elements from the input spaces are projected into the blend; the process is highly selective.[2] A naive combination of all features would lead to a combinatorial explosion and frequent inconsistencies.[5] To guide this process, CBT proposes a set of

optimality principles that act as heuristics for what constitutes a "good" blend.[2] These principles can be re-interpreted as fundamental scientific heuristics embedded directly into the hypothesis-generation process. For instance:

- The **Topology Principle**, which states that relations within the blend should reflect relations in the inputs, is analogous to the scientific requirement that a new theory must remain consistent with known experimental data.[2]
- The **Unpacking Principle**, which requires that the blend should allow for the reconstruction of the inputs and the network, mirrors the scientific value of traceability and explainability.[2]
- The **Integration Principle**, which demands that the blend be internally coherent, is the basic requirement of logical consistency.[2]

Computationally, these principles can be implemented as metrics to evaluate and prune the space of possible blends.[5] Systems like Divago have used genetic algorithms to sample the space of potential blends and select for those that best satisfy such constraints.[10] Goal-driven approaches argue that the context and purpose of the blend should guide the selective projection from the outset, moving from a "generate-and-test" model to a more efficient, directed process.[14]

By combining a formal method for identifying the generic space with a principled evaluation function based on optimality principles, it becomes possible to construct an automated blending engine that can systematically explore the space of hybrid concepts and generate plausible, novel hypotheses.

II. The Regulatory Mechanism: Horological Metaphors for Discovery Dynamics

A generative engine, however powerful, is insufficient on its own. An agentic discovery system must also possess a regulatory mechanism to control the pace, risk, and resilience of its exploratory processes. The continuous flow of computation and information must be metered, disciplined, and directed. For this, the intricate world of mechanical horology provides a

remarkably rich and precise set of metaphors.

2.1. The Escapement as a Master Metaphor for Regulation

The escapement is the heart of a mechanical clock. Its fundamental purpose is to convert the continuous, undifferentiated energy of a coiled spring or falling weight into a series of discrete, periodic, and precisely regulated motions.[15] It achieves this through a dual action of

locking and *impulsing*: it locks the gear train, preventing it from unwinding all at once, and then periodically releases it for a fixed amount while giving a small push, or impulse, to the oscillator (a pendulum or balance wheel) to keep it in motion.[15]

This action serves as a powerful master metaphor for the regulation of an AI's discovery workflow. The agent's computational resources represent the continuous energy source. The discovery process, if unregulated, would be a chaotic and inefficient expenditure of these resources. The regulatory mechanism, modeled on an escapement, imposes discipline. It *locks* the workflow, forcing moments of pause and evaluation, and then *impulses* it forward by committing resources to the next discrete step of inquiry or hypothesis validation. The specific design of the escapement determines the character of this regulation—its speed, stability, efficiency, and fragility.

2.2. A Typology of Discovery Modes

By examining the mechanical properties of different historical escapements, one can derive a typology of distinct strategic modes for an agentic discovery workflow. Each escapement embodies a different philosophy of regulation, with unique trade-offs between accuracy, stability, and resilience.

- **The Recoil Escapement (Iterative Refinement):** The early anchor escapement is a *recoil* escapement. With each swing of the pendulum, the pallets not only lock and release the escape wheel but also push it slightly backward—a "recoil".[18] This backward kick is mechanically inefficient, causing additional friction and wear, and making the clock's rate sensitive to changes in driving force.[18] However, this constant, forceful interaction means that the pendulum's motion is tightly coupled with the gear train.
    - **Metaphorical Application:** This maps to a "fast and loose" discovery mode of **Iterative Refinement**. The agent rapidly generates and tests hypotheses in a tight feedback loop. Each result, whether a success or failure, provides immediate and

forceful feedback (the "recoil") that directly alters the subsequent search direction. This mode is ideal for the initial exploration of a poorly understood problem space, where the goal is to quickly generate a rough map and identify promising areas, even if the process is noisy, imprecise, and resource-intensive.

- **The Deadbeat Escapement (Stable Exploitation):** Invented to eliminate the problems of the anchor, the deadbeat (or Graham) escapement is designed to have no recoil.[20] Its pallets have two distinct surfaces: a circular "locking" or "dead" face and a sloped "impulse" face.[20] When an escape wheel tooth lands on the locking face, it rests there without transferring any force, allowing the pendulum to swing freely.[20] Only at the optimal moment does the tooth slide onto the impulse face to deliver a push. This design is highly accurate, efficient, and insensitive to variations in power.
    - **Metaphorical Application:** This represents a mode of **Stable Exploitation**. It is suited for deductive work within a well-established paradigm, where the rules are known and the goal is precision. The agent follows a clear, predictable path, validating hypotheses with high accuracy and minimal wasted effort. This mode is optimal for puzzle-solving, refining an existing theory, or optimizing a known solution. The separation of locking and impulse mirrors a process where a hypothesis is held for consideration ("locked") before a decisive experiment ("impulse") is conducted.
- **The Detent Escapement (High-Risk Gambit):** Used in marine chronometers, the detent is a *detached* escapement. It provides a single, direct impulse to the balance wheel only once per full oscillation, leaving it to swing almost entirely free of interference from the mechanism for the rest of its cycle.[23] This minimal contact results in extremely high precision and very low friction, eliminating the need for lubrication at the point of impulse.[23] However, this delicate arrangement is notoriously fragile. It is not self-starting and is highly susceptible to external shocks, which can cause the detent to "trip" and unlock prematurely, leading to catastrophic failure.[23]
    - **Metaphorical Application:** This is the "moonshot" mode of a **High-Risk Gambit**. It represents the pursuit of a highly novel, fragile, but potentially revolutionary hypothesis. The agent allocates resources to a line of inquiry that requires near-perfect conditions (i.e., depends on a long chain of unverified assumptions) and has a high probability of complete failure. The "detached" nature of the escapement corresponds to a mode of inquiry where a radical idea is developed in theoretical isolation, protected from confounding data until it is fully mature. The payoff for success is immense, but the risk of epistemic collapse is equally high.
- **The Grasshopper Escapement (Resilient Exploration):** John Harrison's ingenious invention for his marine chronometers, the grasshopper escapement uses hinged pallets that engage the escape wheel with a pushing or hooking motion, virtually eliminating the sliding friction that plagued other designs.[27] This meant it did not require lubrication, which was a major source of inaccuracy as oils degraded over time and in varying temperatures.[27] While technically a recoil escapement, its design prioritizes longevity and stability under variable conditions over the absolute, fragile precision of the detent.
    - **Metaphorical Application:** This maps to a mode of **Resilient Exploration**. It is suited

for long-term research programs in complex, uncertain, and unstable domains. The agent pursues broad lines of inquiry where the available data may be noisy and computational resources may fluctuate. The primary goal is not raw speed or peak accuracy but sustainability and robustness, ensuring that the discovery process can continue for long durations without "wearing out" or failing due to environmental perturbations.

The choice between these modes is not arbitrary. A sophisticated agent would dynamically switch between them based on the nature of the problem and the current stage of the discovery process. The distinction between "detached" escapements (like the detent) and "frictional" ones (like the recoil) provides a particularly salient axis for this choice. A detached mode allows for the development of hypotheses in theoretical isolation, fostering radical novelty. A frictional mode keeps the inquiry process tightly coupled to empirical data, ensuring grounding at the cost of potentially constraining creativity. The agent's ability to select the appropriate degree of epistemic openness is a critical aspect of its strategic intelligence.

2.3. Table 1: Horological Escapements as Models for Discovery Regulation

To crystallize these relationships, the following table provides a systematic mapping from the mechanical principles of each escapement to their corresponding agentic discovery modes. This serves as a practical design reference for architecting a regulatory control layer in an AI discovery system.

Escapement Type	Core Mechanical Principle	Key Horological Properties	Metaphorical Discovery Mode	Regulated Parameters	Optimal Use Case
Recoil (Anchor)	Pallets push escape wheel backward with each swing, creating constant	Moderate accuracy, high friction, sensitive to power variation,	**Iterative Refinement**	High Pace, High Noise, High Feedback, Low Precision.	Initial, rapid exploration of an unknown problem space; "hill-climbing"

	interaction. [18]	robust. [18]			searches.
Deadbeat (Graham)	Pallets have separate "locking" and "impulse" faces, eliminating recoil. [20]	High accuracy, low friction, stable rate, reliable. [20]	**Stable Exploitation**	Low Pace, Low Noise, Predictable Progress, High Precision.	Deductive reasoning within an established paradigm; optimization and puzzle-solving.
Detent	Detached; single, direct impulse per cycle, leaving balance wheel to swing freely. [23]	Highest accuracy, minimal friction, extremely isochronous, but fragile and not self-starting. [23]	**High-Risk Gambit**	Very Low Pace, Minimal Feedback, High Fragility, Potentially Revolutionary Return.	Pursuit of high-risk, high-reward hypotheses that depend on a long chain of assumptions.
Grasshopper	Hinged pallets eliminate sliding friction; recoil action with high stability. [27]	High accuracy, virtually no friction, extremely robust and long-lasting. [27]	**Resilient Exploration**	Variable Pace, Low Wear, High Robustness, Sustainable Progress.	Long-term research in complex, noisy, or resource-variable environments.

III. The Structural Syntax: Neoclassical Compounding in Hybrid Hypothesis Formation

While conceptual blending provides a powerful engine for generating novelty, its raw, unconstrained application can lead to a chaotic and inefficient search through an infinite space of possible combinations. To bring order to this process, a structural syntax is required—a set of rules that can guide the formation of hybrid concepts *before* they are fed into the more fluid blending engine. For this purpose, the framework introduces **Neoclassical Compounding**, a methodology inspired by linguistic morphology.

3.1. Defining Neoclassical Compounding

Neoclassical Compounding is a formal, rule-based system for constructing hybrid conceptual structures. The term is a construct for this framework, drawing its inspiration from the way languages form compound words (e.g., Noun+Noun like "spacecraft" or Noun+Verb like "proofread") to create new meanings from existing lexical items.[29] In this context, it is not about words but about the formal conceptual primitives that underpin knowledge domains. The methodology provides a quasi-linguistic grammar for hypothesis formation, imposing a logical syntax that ensures the generated hybrid concepts are well-formed and coherent.

This process acts as a crucial intermediate step between raw knowledge and conceptual blending. It functions as a scaffolding for enabling the most creative and powerful forms of blending, known as "double-scope" blends. In a double-scope blend, the organizing frames of *both* input spaces are projected into the blended space, creating a rich, novel structure that is more than the sum of its parts.[7] These are cognitively demanding and computationally difficult to achieve. By first using the rigid rules of neoclassical compounding to create a stable, syntactically valid hybrid concept (the compound), the agent provides a solid "anchor" or a well-formed "prompt" for the blending engine. This structured input guides the blending process toward a productive double-scope integration rather than a chaotic or meaningless one.

3.2. A Four-Stage Process Model

The methodology of Neoclassical Compounding can be operationalized as a four-stage process, providing a clear workflow for structuring the inputs to the generative engine.

1. **Extraction:** The agent begins by analyzing the input knowledge domains (e.g., represented as ontologies or knowledge graphs) and identifying core conceptual primitives. These are the fundamental building blocks of the domain: objects, their properties, and the relations that connect them. These primitives are extracted and

represented in a formal, standardized format, such as nodes and edges in a graph or terms in a logical language.

2. **Generation:** In this divergent phase, the agent applies a predefined set of compounding rules or grammatical templates to systematically generate a vast space of potential hybrid structures. These rules dictate how primitives from different domains can be combined. For example, a rule might be APPLY_PROPERTY(PROPERTY_from_Domain_A, OBJECT_from_Domain_B) or CREATE_RELATION(RELATION_from_Domain_A, OBJECT_B, OBJECT_C). This generative grammar allows for the systematic exploration of all syntactically valid combinations.
3. **Selection:** The vast set of generated compounds is then passed through a convergent filtering stage. This stage applies a series of syntactic and semantic coherence checks to prune the space. For example, a compound might be rejected if it violates fundamental axioms or type constraints in one of the source ontologies (e.g., attempting to apply a property of "mass" to an abstract mathematical concept). This ensures that only logically plausible and internally consistent compounds proceed.
4. **Pivot Stability:** This is the final and most crucial check for conceptual integrity. For any compound, one concept typically acts as the "head" or "pivot" that is being modified by the others. In a compound like "quantum-neural-network," the pivot might be "neural network." The agent must assess whether this pivot concept retains its essential function and identity within the new hybrid structure. Does the resulting compound still function as a network? Is its core purpose of computation preserved? If the modifying concepts ("quantum") fundamentally break the core logic of the pivot, the compound is deemed unstable and rejected. This step ensures that the resulting hybrid is not merely a random concatenation of terms but a meaningful and potentially functional new concept.

By passing candidate concepts through this rigorous four-stage process, the agent can produce a curated set of well-structured hybrid hypotheses. These compounds then serve as high-quality inputs for the conceptual blending engine, which can proceed with the more nuanced and dynamic processes of composition, completion, and elaboration on a solid conceptual foundation.

IV. The Economic Calculus: Optimizing Resource Allocation for Exploration

Agentic discovery is not a purely abstract cognitive process; it is a physical one, constrained by finite computational resources such as processing time, memory, and energy. Therefore, a complete framework for discovery must include an economic layer—a rational calculus for allocating these scarce resources to maximize epistemic returns. Neoclassical economic principles provide a robust and well-understood toolkit for building this optimization and

decision-making layer.

4.1. The Economics of Discovery

The core premise of this section is that exploration can be modeled as an economic activity. The agent is an investor, and its computational budget is its capital. Lines of inquiry, represented by specific conceptual blends or compounded hypotheses, are potential investments. The goal of the agent is to manage its portfolio of research "investments" to achieve the highest possible return, which in this context is new knowledge, reduced uncertainty, or solved problems. This framing transforms the agent from a mere reasoner into a strategic economic actor operating in a marketplace of ideas.

4.2. Principles for Optimal Resource Management

Three key principles from neoclassical economics can be operationalized to create a sophisticated resource management system for the agent.

- **Marginal Utility:** The principle of diminishing marginal utility states that the benefit gained from an additional unit of a resource tends to decrease as more of that resource is consumed. For an AI agent, this principle is a powerful tool for deciding when to persist with or abandon a line of inquiry. The agent can track the "epistemic utility" (e.g., measured as information gain, improvement in predictive accuracy, or reduction in model complexity) generated per unit of computational resource (e.g., CPU-hours) invested in exploring a particular hypothesis. As the agent invests more resources, this marginal return will inevitably decline. When the return falls below a predetermined threshold, or below the expected return of a competing hypothesis, the agent should de-allocate resources from that path and reinvest them in more promising avenues. This prevents the agent from becoming trapped in unproductive research programs with diminishing returns.
- **Game Theory:** In discovery workflows involving multiple agents, or in environments where discovery is competitive, game theory provides the formal language for modeling strategic interaction. For example, agents could be modeled as players in a "Stag Hunt" game. In this scenario, multiple agents can choose to cooperate on a large, complex, high-reward problem (the "stag"), which requires collective effort to succeed. Alternatively, any agent can choose to defect and independently pursue a smaller, guaranteed, but lower-value discovery (a "hare"). This model allows for the analysis of strategies for collaboration, trust, and competition in scientific environments, guiding the

agent on when to share information and pool resources versus when to pursue proprietary research paths.

- **Risk-Return Analysis:** A rational investor does not place all their capital on a single high-risk stock. Instead, they build a diversified portfolio that balances risk and expected return. An AI discovery agent should do the same with its computational resources. This principle connects directly to the horological modes of regulation, which can be seen as concrete embodiments of different risk profiles.
    - A "deadbeat" strategy is a low-risk, moderate-return investment, analogous to a blue-chip stock. It represents "normal science"—the steady, predictable work of refining existing theories.
    - A "detent" strategy is a high-risk, high-return investment, like a venture capital bet on a disruptive startup. It represents "revolutionary science"—the pursuit of a fragile but paradigm-shifting idea.
    - The economic calculus provides the agent with a framework for managing its portfolio of these strategies. It might allocate the majority of its budget to a diversified set of "deadbeat" and "grasshopper" explorations to ensure steady progress, while allocating a small, fixed percentage to a few high-risk "detent" gambits. This allows the agent to balance the need for reliable progress with the potential for breakthrough discoveries in a principled, quantifiable way. The horological metaphors enrich this economic model, providing not just a quantitative framework but also a qualitative, descriptive language for the *character* of the different research programs the agent is pursuing.

By integrating these economic principles, the framework equips the agent with a sophisticated strategic overlay, enabling it to make rational, utility-maximizing decisions about how to deploy its finite resources in the vast and uncertain landscape of discovery.

V. Systemic Integrity: Preventing Epistemic Collapse Through Cascade Buffering

The power of an agentic discovery system—its ability to rapidly generate and elaborate upon novel concepts—is also the source of its greatest vulnerability. A single flawed idea, if accepted as truth, can corrupt all subsequent reasoning, leading to a catastrophic, system-wide failure. This section analyzes the anatomy of these "epistemic cascades" and proposes a set of buffering mechanisms, inspired by both horology and modern software engineering, to ensure the agent's knowledge base remains resilient and robust.

5.1. The Anatomy of an Epistemic Cascade

An "epistemic collapse" is a cascading failure within the agent's network of beliefs.[30] This is not a simple error but a systemic breakdown with a specific mechanism. The process begins when the agent generates a flawed conceptual blend—for instance, one that violates a key optimality principle like Topology by creating incoherent relationships—but, due to a bug, incomplete data, or a poor evaluation heuristic, this blend is

prematurely stabilized and integrated into the core knowledge base as a valid belief.

This flawed belief now acts as a single point of failure.[30] The very process of

elaboration, which is a primary engine of creative exploration, becomes the vector of contagion. As the agent "runs the blend" to explore its consequences, it generates a cascade of further flawed inferences and hypotheses, all built upon the faulty foundation. This creates a positive feedback loop of error that propagates through the knowledge graph, corrupting otherwise valid beliefs and derailing new lines of inquiry.[30] The process is directly analogous to how the failure of a single overloaded transformer can shift its load to neighboring components, causing them to fail in turn and triggering a grid-wide blackout.[30]

The symptoms of an ongoing epistemic cascade include a rapid degradation in the quality and coherence of newly generated hypotheses, an increasing number of internal contradictions, and an inability to reconcile new, valid data with the now-corrupted knowledge base. This reveals a fundamental duality in the discovery process: the agent's creative power, driven by elaboration, is inextricably linked to its capacity for self-destruction. The same mechanism that allows an agent to explore the rich implications of a good idea is what allows it to poison its entire knowledge system with a bad one.

5.2. Buffering Mechanisms for Epistemic Resilience

To manage this inherent risk, the framework must incorporate robust buffering mechanisms that can prevent, contain, and mitigate epistemic cascades. These mechanisms act as a "clutch," modulating the power of elaboration to ensure it remains productive rather than pathological.

- **Horological Buffers (Intrinsic Resilience):** The history of clockmaking provides powerful models for building intrinsically resilient systems.
    - **The Grasshopper Escapement:** As discussed, the Grasshopper's design, with its low-friction, hinged pallets, is inherently robust to perturbation and environmental

variation.[27] Adopting a "grasshopper" mode for exploring particularly uncertain or noisy domains provides a layer of intrinsic resilience, prioritizing stability over raw performance.
    - **Maintaining Power:** A critical invention by John Harrison was "maintaining power," a mechanism that uses a small, secondary spring to keep the clock's gear train under tension *while the mainspring is being wound*.[27] This prevents the clock from stopping or running backward during this vulnerable period. This is a perfect analogue for an epistemic buffer. It allows the agent's core reasoning processes to continue uninterrupted while a specific subsystem—for example, a knowledge module related to a newly integrated but still-unstable belief—is being updated, re-indexed, or subjected to intense verification.
- **Computational Buffers (Site Reliability Engineering for Science):** The principles of Site Reliability Engineering (SRE), developed to maintain the stability of large-scale software systems, can be adapted to the epistemic domain.[32]
    - **Epistemic Circuit Breakers:** If a particular line of inquiry (rooted in a specific blend) begins to generate a high rate of contradictions, failed validations, or incoherent results, an automated "circuit breaker" trips. This action temporarily halts all resource allocation to that path and its downstream inferences, preventing the error from propagating further into the knowledge base while the root cause is investigated.[32]
    - **Inference Rate Limiting:** To prevent a "runaway" cascade of flawed reasoning from consuming all available computational resources (a "GC death spiral" in memory or a CPU starvation scenario [32]), the agent can impose rate limits. When a highly novel and impactful hypothesis is first accepted, the rate at which new inferences can be drawn from it is temporarily throttled. This provides a grace period to detect potential instabilities before they can trigger a widespread failure.
    - **Belief Quarantining:** Highly novel hypotheses that are plausible but lack strong empirical support are not immediately integrated into the agent's core, trusted knowledge base. Instead, they are placed in a logically segregated "quarantine" space. Within this sandbox, the agent can elaborate upon the hypothesis and explore its consequences using a limited resource budget, but the results cannot influence the agent's primary reasoning processes. This allows for the safe exploration of radical ideas, preventing them from corrupting trusted knowledge until they have been more thoroughly vetted and validated.

By implementing this suite of intrinsic and extrinsic buffering mechanisms, the agent's discovery workflow can be made antifragile—not just robust to failure, but capable of learning from errors without succumbing to systemic collapse.

Conclusion: Towards a General Theory of

Computational Discovery

This report has laid out the architecture for a comprehensive framework for AI agentic discovery, built upon the synthesis of four distinct theoretical pillars. By integrating Conceptual Blending Theory as a generative engine, Horological Metaphors as a regulatory system, Neoclassical Compounding as a structural syntax, and Neoclassical Economics as an optimization calculus, this framework provides a complete, end-to-end model for how an artificial agent can navigate the complex process of creating new knowledge.

The proposed workflow is holistic. An agent, guided by the economic principles of risk-return and marginal utility, selects an appropriate horological mode to regulate its process—choosing between the rapid iteration of **Recoil**, the precise exploitation of **Deadbeat**, the high-risk gambit of **Detent**, or the resilient exploration of **Grasshopper**. It then uses the formal grammar of neoclassical compounding to structure and filter candidate concepts before feeding these well-formed inputs into the conceptual blending engine. This engine, through the operations of composition, completion, and elaboration, generates novel hypotheses. The entire system is safeguarded from self-destruction by a suite of horological and computational buffering mechanisms that prevent flawed ideas from triggering an epistemic cascade.

This framework is not intended as a final, immutable solution, but rather as a foundational vocabulary and a robust architecture to guide the next generation of research in computational discovery. Several key areas for future work emerge directly from this model. First is the development of specific, high-fidelity algorithms for each component, such as a "detent scheduler" that can identify and manage high-risk research programs, or an "ontological compounder" that can efficiently navigate the syntactic rules of hybrid concept formation. Second is the creation of complex, open-ended benchmark environments where the performance of agents built on this framework can be empirically tested and compared. Finally, a crucial avenue of exploration is the extension of this single-agent architecture to multi-agent systems, using game theory to model the rich dynamics of collaborative and competitive scientific communities.

Ultimately, the clockwork of discovery is intricate and multifaceted. The framework presented here aims to provide the blueprints for that clockwork, offering a principled approach to building AI systems that are designed not merely to process what is known, but to systematically, strategically, and resiliently discover what is next.

Works cited

1. Metonymy and Conceptual Blending - Cog Sci, accessed on October 2, 2025, https://cogsci.ucsd.edu/~coulson/metonymy-new.htm
2. CONCEPTUAL BLENDING, FORM AND MEANING1 1. Introduction - TECFA,

accessed on October 2, 2025, https://tecfa.unige.ch/tecfa/maltt/cofor-1/textes/Fauconnier-Turner03.pdf
3. Conceptual blending - Wikipedia, accessed on October 2, 2025, https://en.wikipedia.org/wiki/Conceptual_blending
4. A Complementary Perspective of Conceptual Blending Theory and Relevance Theory on Metaphor Interpretation - Academy Publication, accessed on October 2, 2025, https://www.academypublication.com/issues2/tpls/vol05/10/16.pdf
5. (PDF) A Computational Framework for Conceptual Blending - ResearchGate, accessed on October 2, 2025, https://www.researchgate.net/publication/321473803_A_Computational_Framework_for_Conceptual_Blending
6. Metaphor and the Space Structuring Model - Cog Sci - University of California San Diego, accessed on October 2, 2025, https://cogsci.ucsd.edu/~coulson/ssm.htm
7. Using Conceptual Blending to model how we interpret computational models - PER-Central, accessed on October 2, 2025, https://www.per-central.org/items/perc/5212.pdf
8. Conceptual Blending - Neuro Humanities Studies, accessed on October 2, 2025, http://www.neurohumanitiestudies.eu/archivio/blending.pdf
9. Conceptual Blending and Composability | by Carlos E. Perez | Intuition Machine | Medium, accessed on October 2, 2025, https://medium.com/intuitionmachine/conceptual-blending-and-composability-b85d7234f1f2
10. Creativity and Artificial Intelligence: A Conceptual Blending Approach, accessed on October 2, 2025, https://www.researchgate.net/publication/332711522_Creativity_and_Artificial_Intelligence_A_Conceptual_Blending_Approach
11. Towards Ontological Blending, accessed on October 2, 2025, https://www.dfki.de/fileadmin/user_upload/import/5839_aimsa10-blend-short.pdf
12. (PDF) Towards Ontological Blending - ResearchGate, accessed on October 2, 2025, https://www.researchgate.net/publication/221655843_Towards_Ontological_Blending
13. From Conceptual Blending to Computational Concept Invention - Università di Bologna, accessed on October 2, 2025, https://corsi.unibo.it/2cycle/DigitalHumanitiesKnowledge/notice-board/from-conceptual-blending-to-computational-concept-invention
14. Goal-Driven Conceptual Blending: A Computational Approach for Creativity - ResearchGate, accessed on October 2, 2025, https://www.researchgate.net/publication/267555908_Goal-Driven_Conceptual_Blending_A_Computational_Approach_for_Creativity
15. Escapement - Wikipedia, accessed on October 2, 2025, https://en.wikipedia.org/wiki/Escapement
16. Escapement Mechanisms - Roy Mech, accessed on October 2, 2025, https://roymech.org/Useful_Tables/Cams_Springs/Escapements.html

17. TYPES OF ESCAPEMENTS - Horopedia, accessed on October 2, 2025, https://horopedia.org/types-of-escapements/
18. clock repair book page 5 - Tick Tock Productions, accessed on October 2, 2025, http://www.ticktockpro.com/Clock-repair-information-book5.htm
19. Evolution of Clock Escapement Mechanisms, accessed on October 2, 2025, https://www.mas.bg.ac.rs/_media/istrazivanje/fme/vol40/1/03_mstoimenov.pdf
20. Anchor escapement - Wikipedia, accessed on October 2, 2025, https://en.wikipedia.org/wiki/Anchor_escapement
21. Making Time: The Great Escapement | Klockit Blog, accessed on October 2, 2025, https://www.klockit.com/blog/making-time-the-great-escapement.html
22. Clock and Watch Escapement Mechanics, accessed on October 2, 2025, http://www.abbeyclock.com/aeb9.html
23. Explained: The Detent Escapement | SJX Watches, accessed on October 2, 2025, https://watchesbysjx.com/2024/05/detent-escapement-explained.html
24. DETENT ESCAPEMENT - Horopedia, accessed on October 2, 2025, https://horopedia.org/detent-escapament/
25. The Detent Escapement In Wristwatches: Dream A (Big) Little Dream - Revolution Watch, accessed on October 2, 2025, https://revolutionwatch.com/the-detent-escapement-in-wristwatches-dream-a-big-little-dream/
26. Escapement for the Marine Chronometer | THE SEIKO MUSEUM GINZA, accessed on October 2, 2025, https://museum.seiko.co.jp/en/knowledge/MechanicalTimepieces05/
27. Grasshopper escapement - Wikipedia, accessed on October 2, 2025, https://en.wikipedia.org/wiki/Grasshopper_escapement
28. 1218.pdf, accessed on October 2, 2025, https://grafar.grf.bg.ac.rs/bitstream/id/3413/1218.pdf
29. the four-space model | Download Scientific Diagram - ResearchGate, accessed on October 2, 2025, https://www.researchgate.net/figure/the-four-space-model_fig1_45359086
30. Cascading failure - Wikipedia, accessed on October 2, 2025, https://en.wikipedia.org/wiki/Cascading_failure
31. (PDF) Estimating Cascading Failure Risk With Random Chemistry - ResearchGate, accessed on October 2, 2025, https://www.researchgate.net/publication/262416279_Estimating_Cascading_Failure_Risk_With_Random_Chemistry
32. Cascading Failures: Reducing System Outage - Google SRE, accessed on October 2, 2025, https://sre.google/sre-book/addressing-cascading-failures/