Drift, Friction, and Co-Design: An Analysis of Free AI Tools as Emergent Socio-Technical Systems

Introduction: Beyond the Interface – AI Tools as Socio-Technical Systems

This report posits that free Artificial Intelligence (AI) tools are not static artifacts but complex, emergent systems co-constituted by their designers' intentions, their technical affordances, and the dynamic practices of their user communities. The behaviors, functions, and even the core identity of these tools evolve in ways that are often unpredictable and unintended by their creators. This phenomenon, known as emergent behavior, arises from the intricate interactions of simpler components within a system—be it algorithms, users, or environmental factors—leading to novel, system-level properties that were not explicitly programmed.[1] The unpredictability inherent in emergence creates a dual potential for both profound innovation and significant, unforeseen risks.[3]

To investigate this dynamic, this report undertakes a deep sociotechnical analysis of three prominent free AI tools, each representing a distinct creative and professional domain: **ChatGPT** (language generation), **Stable Diffusion** (image synthesis), and **Tabnine** (code assistance). These case studies serve as exemplars of different architectural philosophies, community structures, and co-design models. The analysis will focus on three core phenomena:

1. **Drift:** The process by which tools are adopted for purposes far removed from their original, intended functions.
2. **Friction:** The cognitive, practical, and ethical barriers users encounter when a tool's behavior mismatches their expectations or values.
3. **Co-design:** The implicit and explicit ways in which users and their communities actively shape a tool's ongoing development, meaning, and norms of use.

The investigation is structured by a four-lens analytical framework designed to systematically deconstruct these emergent dynamics: **Typological Drift**, which examines the shifting functional identity of tools; **Cognitive Friction**, which analyzes the sources of user struggle and the coping mechanisms that arise; **Community Epistemology**, which explores how communities socially construct knowledge and "correct" usage; and **Interoperability**, which assesses how technical architectures for integration enable or constrain co-design. By applying this framework, the report moves beyond a surface-level review of features to reveal the complex, recursive relationships between technology, users, and the social worlds they mutually create.

Table 1: Comparative Overview of Case Study Systems

Case Study	Declared Purpose	Primary Emergent Uses (Drift)	Key Sources of Cognitive Friction	Dominant Community Platform	Interoperability Model
ChatGPT	"AI Writing Partner" [5]; "AI research and deployment company" [6]	Full essay generation [7]; Emotional support [8]; Philosophical partner [9]	Hallucinations; Conversational degradation; Factual inaccuracy [10]	Reddit (r/ChatGPT, r/ChatGPTPro)	Controlled Plugin Ecosystem
Stable Diffusion	"Text-to-image synthesis" [12]; "Generate stunning images" [13]	NSFW content generation [14]; Specialized aesthetic creation; Watermark removal [16]	Representational harm; Technical complexity; Usability barriers [16]	Reddit (r/StableDiffusion), GitHub	Decentralized Open-Source
Tabnine	"AI-powered code completion tool" [18]; "AI coding assistant" [19]	Learning/educational aid [20]; Debugging partner; Organizational standards enforcement [18]	Intrusiveness; Inaccurate suggestions; High resource usage [20]	GitHub Discussions, JetBrains Marketplace, Reddit	Enterprise-Focused Integration

---

Part I: The Four-Lens Analysis of Emergent AI Systems

Section 1: Typological Drift – From Intended Function to Lived Reality

The functional identity of an AI tool is not a fixed property determined solely by its creators. It is a fluid construct that drifts over time, shaped by the evolving practices of its users. This section analyzes this typological drift, examining how the "declared purpose" marketed by developers diverges into the "lived reality" of community use, often giving rise to powerful "shadow functions" that were never intended.

1.1 The Pre-Emptive Drift of Declared Purpose

Typological drift is not a phenomenon that begins only after a tool's release; it is often initiated by developers and marketers themselves through a strategy of "pre-emptive drift." By framing their tools in the broadest possible terms and enumerating a vast array of potential applications, they create a permissive environment that implicitly encourages users to explore the tool's boundaries and adapt it to novel contexts. This marketing-driven expansion of a tool's identity lays the groundwork for more radical, user-driven drift to occur later.

This strategy is exemplified by the marketing for text generation tools. Grammarly, for instance, is not simply an "AI writing tool." Its website presents a sprawling menu of highly specific micro-tools, each targeting a distinct use case, from professional tasks like drafting a "Business Plan" or "Letter of Resignation" to academic needs like writing an "Abstract" or "Thesis Statement," and even creative pursuits like composing a "Poem".[5] Similarly, Lindy is marketed as a "no-code platform" for building a wide variety of "AI agents" capable of handling everything from email triage and meeting summaries to the automatic generation of sales proposals triggered by CRM updates.[23]

This approach contrasts sharply with the more focused "declared purpose" of early image generators like DALL-E 2, which was simply described as an "AI system that can generate stunning images and artwork from a simple text description".[13] The effect of Grammarly's and Lindy's strategy is to collapse the distinction between a single-purpose utility and a versatile platform. By pre-emptively defining dozens of "legitimate" uses, the developers are communicating to the user that the tool is not meant to be a hammer for a specific nail, but a malleable material to be shaped for any number of problems. This perception of inherent flexibility, cultivated by the marketing itself, becomes a direct antecedent to emergent user behavior, lowering the psychological barrier for users to push the tool into entirely unforeseen domains.

1.2 Emergent Use-Cases and "Shadow Functions"

Once in the hands of users, tools undergo a more profound drift, developing emergent use-cases that can diverge significantly from, or even subvert, their declared purpose. In some cases, these new uses become so dominant within certain communities that they constitute "shadow functions"—primary modes of interaction that exist largely outside the official narrative of the tool.

ChatGPT: The most significant drift observed with ChatGPT is its evolution from a "writing partner" [5] or assistant to a full-fledged content generator. Academic research has meticulously documented this shift, creating custom platforms to track student interactions and finding that common query patterns include not just "proofreading" or "asking for examples," but also "generating the entire essay".[7] This use, often framed as academic dishonesty, represents a clear drift from assistance to outsourcing. Beyond academia, community forums reveal an even wider functional drift. Users engage with ChatGPT for complex problem-solving that rivals professional consultation [10], for emotional support and companionship, with some users reporting that "ChatGPT Makes Me Feel Like a Human Being" [8], and as a quasi-sentient interlocutor for philosophical debates about its own consciousness and nature.[9] These uses transform the tool from a utility into a partner, therapist, and intellectual sparring partner.

Stable Diffusion: The drift of Stable Diffusion is perhaps the most dramatic, branching into two powerful streams. The first is a drift towards extreme specialization. While the declared purpose is general text-to-image generation, communities have become focused on mastering highly specific aesthetics, such as creating "action figure" images [16] or photorealistic styles that require complex, multi-step workflows.[26] The second, more prominent drift is the emergence of Not Safe For Work (NSFW) content generation as a major shadow function. Despite official

safety filters, a significant portion of community discussion and technical development is dedicated to circumventing these restrictions.[14] Users share guides on how to disable the

safety_checker function, fine-tune models on explicit material sourced from community sites like Civitai, and troubleshoot issues specific to generating pornographic or violent imagery.[27] This shadow function is not a fringe activity but a core driver of innovation and engagement within influential sub-communities, representing a fundamental divergence from the sanitized, publicly-facing image of the tool.

Tabnine: The drift of Tabnine is more subtle but equally significant, marking a shift from an individual productivity tool to an instrument of organizational control. While its declared purpose is to be an "AI-powered code completion tool" for individual developers [18], its emergent use extends into education and team management. Users describe it as a tool that helps them "learn new languages" [20], functioning as an educational aid. More consequentially, the enterprise version of the tool has drifted into a mechanism for enforcing institutional norms. Companies can create "team-trained models" that learn from a shared codebase, ensuring that code suggestions "reflect the coding standards, patterns, and frameworks specific to the team".[18] This represents a typological drift from a personal assistant that adapts to an individual's style to a managerial tool that imposes a collective style, transforming it from a tool of personal expression to one of organizational governance.

1.3 Reflexive Inquiry on Typological Drift

A critical self-assessment of this analysis reveals potential oversimplifications and areas of invisibility that warrant further investigation.

- **Collapsed Distinctions:** Does the common framing of student use of ChatGPT as "cheating" collapse important distinctions in user intent and practice? The data from empirical studies suggests a spectrum of engagement, from students using the tool for brainstorming and overcoming writer's block, to those seeking grammatical correction, to those who copy-paste entire essays.[7] A monolithic label like "cheating" obscures the nuanced reality of how students are integrating the tool into their writing process, potentially conflating legitimate assistance with academic misconduct. This collapsed distinction hinders the development of effective pedagogical strategies that could guide students toward productive and ethical use.
- **Invisible Data:** What are the underlying motivations and social contexts driving the "shadow function" of NSFW generation in Stable Diffusion? The available material is rich with the technical "how"—discussions of specific models, prompts, and software forks [15]—but is largely silent on the sociological "why." The user base for this shadow function is likely diverse, potentially including digital artists exploring expressive freedom, individuals creating commercial pornography, users engaging in personal fantasy, and those ideologically committed to the principle of uncensored AI. Without ethnographic data on these user groups, the analysis of this drift remains incomplete, risking a flattened understanding of a complex social phenomenon.
- **Lens Oversimplification:** Does the concept of "drift" adequately capture the evolution of Tabnine? For Tabnine, the shift towards an enterprise tool for enforcing coding standards is not a passive, user-led drift away from an original design. It is a deliberate, strategic pivot by the company towards a more lucrative B2B market, as evidenced by their marketing of enterprise-grade security and team-specific models.[18] Framing this as mere "drift" risks obscuring the agency of the developer in actively steering the tool's typology. This phenomenon is better understood as a "strategic pivot" driven by market incentives, a distinction that the more passive concept of drift may blur.

Section 2: Cognitive Friction – The User's Struggle for Control and Understanding

Cognitive friction arises from the gap between a user's mental model of how a system should work and its actual, often opaque, behavior. This section explores the primary sources of this friction across the three case studies—from interpretive incoherence and representational harm to workflow disruption—and analyzes the "folk theories" and coping mechanisms users develop in response. These user-generated strategies are a form of co-design, representing a collective effort to make unpredictable systems usable.

2.1 Sources of Friction: A Taxonomy of User Frustration

The nature of cognitive friction varies significantly depending on the tool's domain and architecture, revealing the unique challenges posed by different types of AI systems.

- **Interpretive Incoherence (ChatGPT):** The primary source of friction for ChatGPT users is its profound unpredictability and lack of reliability. Users frequently encounter what they describe as "hallucinations," where the model confidently fabricates information.[10] An academic study surveying user dissatisfaction found that while failures in understanding user intent were the most frequent complaint, issues of "information accuracy" caused the most *severe* dissatisfaction.[11] This friction is compounded by the phenomenon of conversational degradation, where users report that the model can "become stupid halfway through a conversation" [8], losing context and coherence over time. For developers using the API for coding tasks, this manifests as the generation of "outright incorrect" code, sometimes inventing "nonexistent language features".[10] This interpretive incoherence forces the user into a constant state of vigilance, unable to trust the system's output without rigorous verification.
- **Representational and Usability Friction (Stable Diffusion):** Users of Stable Diffusion confront a dual friction. The first is a significant usability barrier. The open-source nature of the tool, while empowering, creates a high technical threshold for entry. Community forums are filled with users "vent[ing] about Stable diffusion installation" [16] and seeking basic guidance on "How can I get better results?" [32], indicating that the path from installation to effective use is fraught with technical hurdles. The second, and more profound, source of friction is representational. Academic studies involving human subjects reveal a deep disconnect between the tool's outputs and users' sense of self and community. Participants from marginalized groups, when prompted to generate images of their own identities, expressed sentiments of anger and alienation, stating, "I don't see myself represented here at all".[17] They found the outputs subjected them to "incorrect and often dehumanizing stereotypes" [33], constituting a significant representational harm that goes far beyond simple usability issues.
- **Perceptual and Workflow Friction (Tabnine):** For Tabnine users, friction is primarily experienced as a direct disruption of the coding workflow. User reviews across multiple platforms are replete with complaints about the quality of suggestions, which are often described as "garbage code," "poor suggestions," and "irrelevant responses".[20] The tool's behavior is perceived as intrusive, with one user noting that it "constantly interrupts my import/export module syntax with nonsensical suggestions that don't resemble the project's structure at

all".[21] This leads to a negative productivity loop where developers spend more time correcting the AI's mistakes than they would have spent writing the code manually, with some reporting they became "more productive" after uninstalling the plugin.[21] This is exacerbated by technical issues like high memory usage, which can slow down the IDE, and frequent crashes or login problems that further fragment the development process.[20]

2.2 Folk Theories and Mental Models as Coping Mechanisms

In the face of opaque and unpredictable systems, users do not remain passive. They actively construct "folk theories"—informal, explanatory frameworks—to make sense of the AI's behavior and develop strategies to manage the friction. This process of creating and sharing mental models is a crucial form of community-driven co-design, effectively creating a user-generated manual where the official documentation fails.

When a system is a "black box," its internal logic is hidden from the user. This opacity creates a significant gap between the user's expectations and the system's performance, resulting in high cognitive friction. Official documentation often proves insufficient to bridge this gap, as it typically describes *what* a feature does, not *why* the system behaves in a particular, often counterintuitive, way. In response to this failure of developer transparency, users turn to their communities to collaboratively reverse-engineer the system's "psychology." They conduct informal experiments, share observations, and collectively build usable mental models that allow them to predict and control the unpredictable.

A sophisticated example of this is the folk theory developed by the ChatGPT community to manage its conversational degradation. One user on Reddit outlines a detailed model, explaining that the model's tendency to "slip into a roleplay" and generate misleading, person-like responses can be triggered by using anthropomorphic language like "Do *you* understand?".[9] They advise other users to avoid such prompts to maintain conversational stability. This is not just a simple tip; it is a coherent theory about the model's internal state ("character-mode") and a procedural guide for controlling it. Simpler, more anthropomorphic mental models also serve this purpose, such as the belief that ChatGPT "gets offended" or "gets lazy".[8] While less technically accurate, these narratives provide a predictable framework for a user to rationalize a failed interaction ("The AI is being difficult today") and adjust their strategy.

For Stable Diffusion, the entire practice of advanced prompt engineering can be seen as a highly proceduralized set of folk theories for controlling an otherwise chaotic system. The community-developed rituals of crafting detailed positive and negative prompts, selecting specific samplers, adjusting the CFG scale, and combining various LoRA models [35] are not features explicitly taught by the developers in a comprehensive way. They are a body of craft knowledge, built through collective trial and error and shared in forums, that allows users to impose a degree of order and intentionality onto the generative process. This collective knowledge-building compensates for the system's inherent unpredictability, turning friction into a manageable, albeit complex, craft.

2.3 Reflexive Inquiry on Cognitive Friction

Further reflection on the analysis of cognitive friction reveals areas where distinctions may be collapsed and where the analytical lens itself may be insufficient.

- **Collapsed Distinctions:** Does the category of "poor suggestions" from Tabnine users [20] group together fundamentally different experiences of friction? It is plausible that this single complaint masks a spectrum of user contexts. A junior developer might find a suggestion "poor" because it uses an advanced syntax they do not yet understand, creating a friction of comprehension. Conversely, a senior developer might find the same suggestion "poor" because it is simplistic, inefficient, or genuinely incorrect, creating a friction of expertise. The user's skill level is a critical but often invisible variable in the data, and collapsing these distinct experiences into one category risks misdiagnosing the root cause of the friction.
- **Invisible Data:** The analysis of friction is heavily based on the vocal complaints found in community forums and review sites. This creates a potential survivorship bias. For every user who posts a detailed complaint about Tabnine's intrusiveness [21], how many others encounter the same friction and simply uninstall the tool in silence? This "dark matter" of user churn—the quiet abandonment of a tool—is a crucial form of data that is largely invisible in the available material. Without large-scale telemetry on uninstallation rates and the specific "last straw" friction points that trigger them, the understanding of user frustration may be skewed towards the experiences of the most persistent or vocal minority.
- **Lens Oversimplification:** Does the term "cognitive friction," a concept rooted in the functional and usability-oriented field of Human-Computer Interaction, adequately capture the full impact of certain system failures? For a user of Stable Diffusion whose gender or nationality is misrepresented with "dehumanizing stereotypes" [17], the experience is not merely one of a difficult-to-use interface. It is an experience of social and emotional harm—of anger, disappointment, erasure, and alienation. The relatively neutral, technical lens of "friction" may fail to encompass the profound ethical and psychological weight of such representational violence, highlighting a potential limitation of applying a purely functionalist framework to socio-technical systems with deep social implications.

Section 3: Community Epistemology – The Social Construction of "Correct" Use

AI tools are not used in a vacuum. They are embedded in communities that actively generate, validate, and disseminate knowledge about how to use them. This section examines these "epistemic cultures," analyzing how different communities develop distinct ways of knowing and defining "correct" or "good" use. This process of collective knowledge creation is a fundamental form of co-design, where the community itself establishes the norms, standards, and "meta" that guide interaction with the tool.

3.1 The Three Epistemic Cultures

The communities surrounding ChatGPT, Stable Diffusion, and Tabnine have evolved into distinct epistemic cultures, each with its own values, practices, and forms of knowledge.

- **ChatGPT Communities: The Artisanal Prompt Crafters.** The epistemic culture of communities like r/ChatGPT and r/ChatGPTPro is centered on the linguistic art of prompt engineering. Knowledge is often narrative, heuristic, and focused on mastering the craft of conversation with the AI. The community values

cleverness, efficiency, and novelty in prompting. This is evident in subreddits dedicated to finding and sharing "the best AI prompts" [6] and in detailed, user-generated guides on how to "humanize" the AI's output by avoiding certain linguistic patterns and adopting others.[37] A significant part of this culture involves "jailbreaking"—designing prompts to bypass the model's safety guardrails—and exploring the model's philosophical and creative limits.[6] Furthermore, a sophisticated subculture has emerged around benchmarking, where users run detailed tests comparing the coding or reasoning capabilities of different model versions (e.g., GPT-4o vs. GPT-4.1-mini) and share their findings in detailed reports.[10]

- **Stable Diffusion Communities: The Technical Workflow Architects.** The culture of the r/StableDiffusion community is highly technical, procedural, and asset-based. The primary form of knowledge is the "workflow"—a complex, often multi-step process for achieving a specific visual aesthetic. This knowledge is shared through node graphs for interfaces like ComfyUI [26], detailed textual guides [16], and the curation of community-made assets. These assets include fine-tuned checkpoint models designed for specific styles (e.g., anime, photorealism), LoRAs (small patch files that modify a model's output), and textual inversions.[36] The community is in a constant state of flux, with users frequently asking about the "new meta" [16], indicating a shared understanding that the "best" way to use the tool is continuously evolving through collective experimentation.
- **Tabnine Communities: The Pragmatic Troubleshooters.** The epistemic culture surrounding Tabnine, found in forums like GitHub Discussions, the JetBrains Marketplace, and technical subreddits (e.g., r/webdev), is overwhelmingly utilitarian and problem-solution oriented. The community is not focused on creative exploration but on practical application and troubleshooting. The GitHub discussions space is dominated by Q&A threads, bug reports (e.g., "IDEA插件无法登录" - "IDEA plugin cannot log in"), and feature requests, such as asking for a lower price to "match github copilot".[38] Reddit threads are characterized by direct, pragmatic comparisons between Tabnine and its competitors (GitHub Copilot, Codeium, etc.) [39] and by users seeking help with specific IDE integration issues and errors.[21] The knowledge generated is not about artistic craft but about making a professional tool work as intended in a specific development environment.

3.2 The Co-Design of "Meta" and Norms

Within these epistemic cultures, communities collaboratively define the "meta"—the dominant strategies, techniques, and standards that constitute "good" or effective use of the tool. This socially constructed meta is a powerful, informal mode of co-design, establishing norms that can influence the behavior of thousands of users and eventually feed back into the tool's official development.

The emergence of a community meta is a direct consequence of a tool's complexity and opacity. When there are countless ways to interact with a system and the optimal path is not obvious, users engage in widespread experimentation. Successful strategies are shared, imitated, and refined, while unsuccessful ones are abandoned. Over time, these successful strategies coalesce into a dominant "meta" that functions as a de facto set of best practices. This user-generated meta becomes an informal design specification, signaling to developers what users value and how they are actually working with the tool. Observant developers can then formalize these popular community techniques, incorporating them as default settings or simplified features in future versions of the official product. This completes a powerful co-design loop, where emergent user practices are codified back into the technology itself.

In the ChatGPT community, a "good" prompt is not merely one that elicits a factually correct answer. The meta values prompts that are elegantly constructed, that "jailbreak" the AI in a novel way, or that produce uniquely creative or philosophically interesting output.[6] This shared aesthetic of "prompt crafting" guides the community's creative and exploratory efforts. In the Stable Diffusion community, the meta is even more pronounced. A "good" image is not just one that matches the prompt, but one that demonstrates technical mastery over the current meta, which might involve a specific combination of a base model, a LoRA, a particular sampler, and a complex negative prompt.[36] This shared understanding of what constitutes a high-quality generation drives a continuous, community-wide effort to discover and refine the next dominant workflow.

3.3 Reflexive Inquiry on Community Epistemology

A critical examination of this analysis of community epistemology highlights potential biases and oversimplifications.

- **Collapsed Distinctions:** Does an analysis focused on the dominant Reddit communities for a tool like Stable Diffusion risk collapsing the distinction between different user groups and their unique epistemic cultures? The r/StableDiffusion subreddit appears to be heavily populated by hobbyists and artists focused on character generation and technical experimentation. This focus may render invisible other significant user communities—such as professional architects, product designers, or medical illustrators—who may use the tool for entirely different purposes and congregate in different online spaces, such as private Discord servers or professional forums. Their unique knowledge, workflows, and definitions of "good" use would be absent from this analysis, leading to an incomplete picture of the tool's overall epistemic landscape.
- **Invisible Data:** What knowledge remains unshared within these communities? Public forums are primarily showcases for success. Users post their best images, their most clever prompts, and their most effective workflows. The vast "dark matter" of failed experiments, frustrating dead-ends, abandoned projects, and the sheer volume of mundane or mediocre output is largely invisible. This creates a skewed perception of the tool's ease of use and capabilities, a form of collective survivorship bias. The academic studies that systematically capture user dissatisfaction and failure [11] offer a rare but valuable glimpse into this invisible reservoir of negative experience, providing a necessary corrective to the triumphalism of community show-and-tell.
- **Lens Oversimplification:** Is "epistemology"—the formal philosophical term for the theory of knowledge—the most accurate lens for understanding these community dynamics, or does it risk over-intellectualizing what is often a deeply social and affective process? The existence of posts titled "A word of thanks to the Stable Diffusion community" [16] points to a strong emotional component that underpins these groups. They are not just sites of information exchange; they are communities built on mutual support, shared identity, and social bonding. The lens of epistemology, with its focus on knowledge, may under-emphasize these crucial affective dimensions that give these communities their coherence and resilience.

Section 4: Interoperability – The Technical Scaffolding for Co-Design

Interoperability—the capacity of a system to connect and exchange data with other systems—is not merely a

technical feature. It is the architectural foundation upon which co-design is built. The specific ways in which developers allow their AI tools to be integrated with other software, through APIs, plugins, or open-source code, fundamentally enable, constrain, and reflect different models of collaboration between the company and its user community. This section analyzes the distinct interoperability models of the three case studies and reveals how they shape the emergent culture of each tool.

4.1 A Spectrum of Interoperability Models

The three case studies occupy distinct points on a spectrum of interoperability, from a centrally controlled platform to a decentralized, emergent ecosystem to a privately contained garden.

- **ChatGPT: The Controlled Platform Ecosystem.** OpenAI's strategy is to position ChatGPT as a central platform, extended through a formal, curated, and controlled ecosystem. This is primarily achieved through a well-documented API and an official plugin store.[43] The API provides clear protocols for authentication, rate limiting, and interaction, allowing developers to build ChatGPT's capabilities into their own applications in a structured way.[43] The plugin store represents an even tighter form of control, where OpenAI vets and approves third-party partners like Zapier, Canva, and Expedia, giving them privileged access to the ChatGPT interface.[44] This top-down model ensures a degree of safety and consistency but centralizes power, making OpenAI the gatekeeper of its ecosystem.
- **Stable Diffusion: The Decentralized, Emergent Ecosystem.** In stark contrast, the interoperability of Stable Diffusion is characterized by bottom-up, emergent decentralization. While Stability AI provides an official API for commercial use [47], the true dynamism of the ecosystem stems from the model's open-source nature.[49] This allows the community to bypass official channels entirely and build new layers of functionality and integration from the ground up. The most prominent example is the development of alternative user interfaces like ComfyUI, which provides a sophisticated node-based workflow system that is far more powerful and flexible than any official tool.[26] This community-led development extends to integration services; platforms like Albato and Dify offer connections between Stable Diffusion and hundreds of other applications, driven by user demand rather than official corporate partnerships.[50]
- **Tabnine: The Walled Garden Ecosystem.** Tabnine's interoperability model is best described as a "walled garden," designed for deep integration within a private, secure enterprise environment. The focus is not on public-facing extensibility but on connecting securely to a company's internal assets. Marketing materials and documentation repeatedly highlight privacy, the ability to run on-premises or in a private cloud, and the feature of training personalized models on a company's own private code repositories hosted on Bitbucket, GitHub, or GitLab.[18] While Tabnine integrates with major IDEs as a plugin [54], its core interoperability promise is about securely ingesting a company's private data, not about providing a general-purpose API for broad, public innovation. This is a controlled, inwardly-focused, enterprise-centric model.

4.2 Interoperability as the Architecture of Co-Design

The choice of an interoperability model is not a neutral technical decision; it is a direct expression of a company's business strategy and its philosophy of user engagement. This architecture, in turn, becomes a primary causal

driver of the tool's emergent community culture, as it defines the very possibilities for co-design. A close examination of the three case studies reveals a clear causal chain: the underlying business model dictates the interoperability architecture, which then creates specific affordances for co-design, ultimately shaping the type of community that forms around the tool.

1. **ChatGPT's Platform Model:** OpenAI operates on a **Platform-as-a-Service** business model. Its controlled API and curated plugin store [44] are the architectural manifestations of this strategy. This architecture affords co-design primarily through formal partnerships and API usage. The community's role is largely defined as being consumers of these official integrations or developers building on the API within the rules set by OpenAI. This fosters a community culture focused on the *creative use* of the available tools—mastering prompt engineering, exploring plugin combinations—rather than on fundamentally altering the tool's core infrastructure.
2. **Stable Diffusion's Open-Source Model:** Stability AI's initial strategy was rooted in an **Open-Source/Community-Driven** model, releasing the core model weights publicly.[49] The architectural consequence is that co-design can happen at the most fundamental level. The community is not limited to using an API; they can fork the code, build entirely new interfaces (like ComfyUI [26]), and create and distribute their own fine-tuned models.[36] This open architecture fosters a highly technical community of *builders, tinkerers, and modifiers* who are deeply engaged in co-designing the tool's infrastructure itself.
3. **Tabnine's Enterprise Model:** Tabnine pursues an **Enterprise B2B** business model. Its interoperability architecture reflects this by prioritizing secure, private integration with a company's internal code repositories.[53] Here, co-design is afforded at the organizational level. The "community" that co-designs the tool's behavior is not a public forum but the development team within a single company, which fine-tunes a private model on its own codebase. This architecture cultivates a user base of *professional teams* whose engagement is focused on workflow efficiency, security, and enforcing internal standards, rather than public exploration or infrastructure building.

4.3 Reflexive Inquiry on Interoperability

A critical self-reflection on the interoperability lens reveals potential areas of collapsed distinction, invisible data, and oversimplification.

- **Collapsed Distinctions:** Does the use of a single term like "API" or "integration" risk collapsing the distinction between vastly different levels of functionality and control? The "Stable Diffusion API" offered by a third-party service like stablediffusionapi.com, which provides access to over 100 community-trained models, inpainting, and other advanced features [55], is a fundamentally different entity from a simple, official inference endpoint that only serves the base model. A nuanced analysis must be sensitive to the varying depth, power, and provenance of different integrations, as lumping them together under one label obscures the rich diversity of the ecosystem.
- **Invisible Data:** What is the invisible labor required to sustain these interoperability ecosystems? For community-driven projects like ComfyUI [26] or the hundreds of fine-tuned models on platforms like Hugging Face and Civitai [36], who is performing the continuous, often uncompensated, labor of maintenance, bug-fixing, and updating these tools to remain compatible with the latest official releases? This maintenance work is the invisible scaffolding that supports the entire decentralized ecosystem, yet it is rarely acknowledged in analyses that focus on the visible outputs of innovation.

- **Lens Oversimplification:** Does a purely technical focus on "interoperability" obscure the significant power dynamics at play? OpenAI's governance of its plugin store [44] is not just a technical act of managing APIs; it is a political and economic act of gatekeeping. The decision to approve one plugin and reject another can determine the success or failure of third-party businesses, shaping the entire market that builds upon its platform. This exercise of platform power is a critical socio-economic dimension that a purely technical lens of interoperability might understate or ignore.

Part II: Synthesis and Recommendations

Section 5: Synthesizing the Lenses – A Holistic Model of Emergent AI Co-Design

The preceding analysis, viewed through four distinct lenses, reveals that drift, friction, and co-design are not isolated phenomena but are deeply interconnected components of a continuous, cyclical process. This section synthesizes these findings into a holistic model—the Co-Design Flywheel—that illustrates the dynamic interplay between tool, user, and community.

The **Co-Design Flywheel** describes an iterative cycle that drives the evolution of these emergent AI systems:

1. **Typological Drift Creates Demand:** The cycle begins when users push a tool beyond its intended function, a process of **Typological Drift**. Students using ChatGPT to write entire essays [7], or Stable Diffusion users generating NSFW content [15], are creating new, often unarticulated, user requirements. This drift reveals what users
actually want to do, regardless of the developer's original vision.
2. **Friction Signals Misalignment:** This drift inevitably leads to **Cognitive Friction** when the tool's design or safeguards resist the new use case. The ChatGPT model begins to hallucinate under the strain of a long, complex task [10], or Stable Diffusion's safety filters block NSFW prompts and its training data produces biased, stereotypical representations of identity.[15] This friction is a critical signal of misalignment between user desires and system capabilities.
3. **Community Epistemology Develops Solutions:** This friction, in turn, galvanizes **Community Epistemology**. Faced with a system that does not work as desired, users congregate in communities to collectively develop workarounds and new knowledge. They create sophisticated prompt engineering techniques to "humanize" ChatGPT's output [37], build and share custom models to bypass Stable Diffusion's filters [36], and troubleshoot Tabnine's integration bugs.[21] This socially constructed knowledge is a direct response to the friction encountered.
4. **Interoperability Formalizes Solutions:** Finally, this community-generated knowledge is often codified and scaled through the system's **Interoperability** architecture. In the decentralized Stable Diffusion ecosystem, community-built tools like ComfyUI directly implement the advanced workflows developed by users.[26] In more controlled ecosystems, developers may observe the community "meta" and formalize it into official features in a future release—for example, adding a dedicated "negative prompt" field to a UI after seeing its widespread

community adoption. This act of formalization closes the loop, creating a more powerful and flexible tool that, in turn, enables new and more ambitious forms of drift, restarting the flywheel.

This model demonstrates that users are not passive recipients of AI technology. Through a continuous cycle of drifting, encountering friction, building knowledge, and leveraging technical affordances, they act as powerful, albeit often informal, co-designers of these emergent systems.

Table 2: Synthesis of Findings Across the Four-Lens Framework

Analytical Lens	ChatGPT	Stable Diffusion	Tabnine
Typological Drift	From "writing partner" to full "essay generator," "therapist," and "philosopher".[5] Drift is broad and conceptual.	From "image generator" to specialized aesthetic tool and a dominant "shadow function" of NSFW content creation.[13] Drift is functional and subversive.	From individual "code completion" tool to team-level "standards enforcement" mechanism and learning aid.[18] Drift is from individual productivity to organizational governance.
Cognitive Friction	**Interpretive Incoherence:** Severe user dissatisfaction stems from factual inaccuracy, hallucinations, and conversational degradation.[8]	**Representational & Usability Friction:** High technical barriers to entry and profound representational harms (stereotyping, erasure) cause user alienation.[16]	**Perceptual & Workflow Friction:** Intrusive, inaccurate suggestions, and high resource usage disrupt the coding flow and reduce productivity.[20]
Community Epistemology	**Artisanal Prompt Crafters:** Community knowledge centers on the linguistic art of prompt engineering, "jailbreaking," and model benchmarking.[6]	**Technical Workflow Architects:** Knowledge is procedural and asset-based, focused on sharing complex workflows, custom models, and LoRAs.[16]	**Pragmatic Troubleshooters:** Knowledge is utilitarian, focused on bug reporting, technical support, and direct competitor comparisons.[21]
Interoperability	**Controlled Platform Ecosystem:** A curated plugin store and formal API centralize control, fostering a community of creative *users* of official extensions.[43]	**Decentralized Open-Source Ecosystem:** The open-source model allows fundamental, bottom-up innovation, fostering a community of technical *builders* of new infrastructure.[26]	**Walled Garden Ecosystem:** Secure, private integrations with enterprise repositories foster a "community" of professional *teams* focused on internal workflows.[18]

Section 6: Implications and Recommendations for Human-Centered AI Development

The analysis of these emergent AI systems yields critical lessons for the future of human-centered AI. The findings

point toward a necessary shift in design philosophy, moving from a paradigm of control and prevention to one that embraces emergence, learns from friction, and engages deeply with the social and ethical dimensions of use. The following recommendations are offered to researchers, designers, and policymakers seeking to build more responsible, responsive, and empowering AI systems.

6.1 Designing for Emergence, Not Just Use

The analysis demonstrates the inevitability of typological drift across all three case studies.[7] Users will consistently push tools into unforeseen domains. A design philosophy focused solely on enforcing a narrow, intended use is destined to fail and will only serve to alienate creative users.

Recommendation: Developers should transition from a mindset of preventing unintended use to one of *designing for emergence*. This involves creating systems with built-in flexibility and observable "seams." Rather than attempting to rigidly lock down functionality, developers could create "sandbox" modes or experimental feature flags that allow users to safely explore the tool's boundaries. This would not only satisfy user creativity but also provide developers with invaluable, real-world data on emergent behaviors, allowing them to anticipate and adapt to drift rather than being caught by surprise.

6.2 Friction as a Design Resource

The rich data from community forums shows that user-reported friction is a powerful, if unstructured, source of design insight.[8] The folk theories and mental models that users develop to cope with this friction are not evidence of user error, but signals of a failure in system transparency and documentation.[9]

Recommendation: AI development organizations should establish formal channels to systematically capture and analyze community-reported friction as a primary source of usability data and user requirements. This involves treating platforms like Reddit, GitHub Issues, and user review sites as vital, continuous feedback mechanisms. Furthermore, the emergence of complex "folk theories" should be treated as a design bug, prompting developers to improve the system's explainability and documentation to bridge the gap between the user's mental model and the system's behavior.

6.3 From Algorithmic Audits to 'Thick Evaluations'

The case of Stable Diffusion demonstrates a critical failure of purely computational approaches to safety and ethics. While the model may pass internal, quantitative bias audits, user-centered studies reveal it perpetrates profound representational harms, causing anger and alienation among users from marginalized groups.[17] There is a

deep disconnect between the "thin" data of algorithmic audits and the "thick" reality of lived experience.[57]

Recommendation: For AI systems with significant social and cultural impact, developers must move beyond a sole reliance on computational bias audits. It is imperative to adopt **"thick evaluation"** methodologies grounded in qualitative, community-centered research. This means conducting studies with users from the specific communities the system claims to represent, listening to their experiences, and using their definitions of appropriate and harmful representation to guide model training, fine-tuning, and mitigation efforts. Without this deep, qualitative understanding, AI systems will continue to perpetuate and amplify societal biases, regardless of their technical sophistication.

6.4 Ethical and Empowering Interoperability

The analysis of interoperability reveals that API and plugin architectures are not neutral technical constructs; they are powerful mechanisms of governance that shape markets and community possibilities. A controlled, centralized model like ChatGPT's plugin store risks creating a "permission-based" innovation culture, while a fully decentralized model like Stable Diffusion's can facilitate harmful "shadow functions" alongside creative freedom.

Recommendation: The design of interoperability systems should be treated as an ethical and political act, not just a technical one. Governance models for API access and plugin stores must be transparent, with clearly articulated principles for moderation, access, and revenue sharing. The goal should be to create architectures that empower community innovation and ensure fair value distribution, rather than simply enabling value extraction by the platform owner. This involves fostering a collaborative ecosystem where the co-design relationship between the developer and the community is one of partnership, not just control.

Works cited

1. aiethicslab.rutgers.edu, accessed June 14, 2025, https://aiethicslab.rutgers.edu/e-floating-buttons/emergent-behavior/#:~:text=Emergent%20Behavior%20in%20the%20context,or%20intended%20by%20the%20designers.
2. Emergent Behavior - AI Ethics Lab, accessed June 14, 2025, https://aiethicslab.rutgers.edu/e-floating-buttons/emergent-behavior/
3. What is emergent behavior in AI? | TEDAI San Francisco, accessed June 14, 2025, https://tedai-sanfrancisco.ted.com/glossary/emergent-behavior/
4. Emergent Behavior | Deepgram, accessed June 14, 2025, https://deepgram.com/ai-glossary/emergent-behavior
5. Free AI Writing & Text Generation Tools - Grammarly, accessed June 14, 2025, https://www.grammarly.com/ai/ai-writing-tools
6. Best ChatGPT Posts - Reddit, accessed June 14, 2025, https://www.reddit.com/t/chatgpt/
7. An Empirical Study to Understand How Students Use ChatGPT for Writing Essays - arXiv, accessed June 14, 2025, https://arxiv.org/html/2501.10551v2
8. Is there a peer support forum for ChatGPT? - Reddit, accessed June 14, 2025, https://www.reddit.com/r/ChatGPT/comments/1d99vo9/is_there_a_peer_support_

forum_for_chatgpt/
9. ChatGPT - Reddit, accessed June 14, 2025, https://www.reddit.com/r/ChatGPT/new/
10. ChatGPTPro - Reddit, accessed June 14, 2025, https://www.reddit.com/r/ChatGPTPro/
11. Understanding Users' Dissatisfaction with ChatGPT Responses: Types, Resolving Tactics, and the Effect of Knowledge Level - arXiv, accessed June 14, 2025, https://arxiv.org/html/2311.07434v3
12. SDXL: Improving Latent Diffusion Models for High-Resolution Image Synthesis - arXiv, accessed June 14, 2025, http://arxiv.org/pdf/2307.01952
13. Top 7 Free AI Tools for 2025 You Don't Want to Miss - Retable, accessed June 14, 2025, https://www.retable.io/blog/7-free-ai-tools-you-should-know-in-2025
14. I need help using stable diffusion : r/StableDiffusion - Reddit, accessed June 14, 2025, https://www.reddit.com/r/StableDiffusion/comments/1gusjgu/i_need_help_using_stable_diffusion/
15. Does Stable Diffusion Allow NSFW - Viblo, accessed June 14, 2025, https://viblo.asia/p/does-stable-diffusion-allow-nsfw-E1XVO6nZLMz
16. Stable Diffusion Explained : r/StableDiffusion - Reddit, accessed June 14, 2025, https://www.reddit.com/r/StableDiffusion/comments/1kh3i2o/stable_diffusion_explained/
17. arxiv.org, accessed June 14, 2025, https://arxiv.org/html/2408.01594v1
18. 20 Best AI-Powered Coding Assistant Tools in 2025 - Spacelift, accessed June 14, 2025, https://spacelift.io/blog/ai-coding-assistant-tools
19. How can an AI coding assistant help reduce production bugs ..., accessed June 14, 2025, https://www.tabnine.com/blog/how-does-ai-assist-in-reducing-production-bugs/
20. Tabnine Pros and Cons | User Likes & Dislikes - G2, accessed June 14, 2025, https://www.g2.com/products/tabnine/reviews?qs=pros-and-cons
21. Tabnine: AI Chat & Autocomplete for JavaScript, Python & more ..., accessed June 14, 2025, https://plugins.jetbrains.com/plugin/12798-tabnine-ai-chat--autocomplete-for-javascript-python--more/reviews
22. Tabnine Reviews & Ratings 2025 - TrustRadius, accessed June 14, 2025, https://www.trustradius.com/products/tabnine/reviews
23. The 10 Best AI Text Generators for Writing: Free and Paid | 2025 - Lindy, accessed June 14, 2025, https://www.lindy.ai/blog/best-ai-text-generator
24. An Empirical Study to Understand How Students Use ChatGPT for Writing Essays - arXiv, accessed June 14, 2025, https://arxiv.org/abs/2501.10551
25. University students' reactions to ChatGPT: Enhancing academic ..., accessed June 14, 2025, https://www.ijirss.com/index.php/ijirss/article/download/5048/803/8148
26. Stable Diffusion 3.5 Workflow Tutorial in ComfyUI, accessed June 14, 2025, https://comfyui-wiki.com/en/tutorial/advanced/stable-diffusion-3-5-comfyui-workflow

27. Detailed Guide: How to use Stable Diffusion to Create NSFW Online - Dr.Fone, accessed June 14, 2025, https://drfone.wondershare.com/ai-tools/use-stable-diffusion-to-create-nsfw-online.html
28. Stable Diffusion NSFW and Its Alternatives - Analytics Insight, accessed June 14, 2025, https://www.analyticsinsight.net/artificial-intelligence/stable-diffusion-nsfw-and-its-alternatives
29. Hold on, how did I create NSFW content like this? : r/StableDiffusion - Reddit, accessed June 14, 2025, https://www.reddit.com/r/StableDiffusion/comments/16kj3et/hold_on_how_did_i_create_nsfw_content_like_this/
30. Version 1.8.0 has many errors in generating NSFW content · AUTOMATIC1111 stable-diffusion-webui · Discussion #15136 - GitHub, accessed June 14, 2025, https://github.com/AUTOMATIC1111/stable-diffusion-webui/discussions/15136
31. Accelerate your development process with Tabnine AI and Bitbucket - Atlassian Community, accessed June 14, 2025, https://community.atlassian.com/forums/Bitbucket-articles/Accelerate-your-development-process-with-Tabnine-AI-and/ba-p/2576062
32. What is best community to post stable diffusion art? : r/StableDiffusion - Reddit, accessed June 14, 2025, https://www.reddit.com/r/StableDiffusion/comments/1bh0dzl/what_is_best_community_to_post_stable_diffusion/
33. [2408.01594] "I don't see myself represented here at all": User Experiences of Stable Diffusion Outputs Containing Representational Harms across Gender Identities and Nationalities - arXiv, accessed June 14, 2025, https://arxiv.org/abs/2408.01594
34. Tabnine Reviews 2025: Details, Pricing, & Features | G2, accessed June 14, 2025, https://www.g2.com/products/tabnine/reviews
35. Stable Diffusion 3.5 Prompt Guide - Stability AI, accessed June 14, 2025, https://stability.ai/learning-hub/stable-diffusion-3-5-prompt-guide
36. Stable Diffusion Models: a beginner's guide, accessed June 14, 2025, https://stable-diffusion-art.com/models/
37. Ask for help- What is the ideal prompt to make the content more "humanized"?, accessed June 14, 2025, https://community.openai.com/t/ask-for-help-what-is-the-ideal-prompt-to-make-the-content-more-humanized/1124343
38. codota TabNine · Discussions · GitHub, accessed June 14, 2025, https://github.com/codota/TabNine/discussions
39. In your IDE, are you using tabnine or jetbrains ai ? : r/webdev - Reddit, accessed June 14, 2025, https://www.reddit.com/r/webdev/comments/1d7tly7/in_your_ide_are_you_using_tabnine_or_jetbrains_ai/
40. GitHub Copilot vs Tabnine : r/webdev - Reddit, accessed June 14, 2025, https://www.reddit.com/r/webdev/comments/10e8nht/github_copilot_vs_tabnine/

41. Blackbox vs GitHub Copilot vs TabNine – A brutally honest comparison - Reddit, accessed June 14, 2025, https://www.reddit.com/r/BlackboxAI_/comments/1konk1q/blackbox_vs_github_copilot_vs_tabnine_a_brutally/
42. TabNine vs Co-Pilot in VSCode: Can they run in parallel? Should they? - Reddit, accessed June 14, 2025, https://www.reddit.com/r/vscode/comments/148skqx/tabnine_vs_copilot_in_vscode_can_they_run_in/
43. API Reference - OpenAI Platform, accessed June 14, 2025, https://platform.openai.com/docs/api-reference/chat
44. ChatGPT plugins - OpenAI, accessed June 14, 2025, https://openai.com/index/chatgpt-plugins/
45. OpenAI API | Documentation | Postman API Network, accessed June 14, 2025, https://www.postman.com/devrel/openai/documentation/k25n3c8/openai-api
46. 12 Most Useful ChatGPT Plugins in 2025 - TTMS, accessed June 14, 2025, https://ttms.com/12-most-useful-chatgpt-plugins-in-2025/
47. Stability AI Stable Diffusion 3.5 API Node ComfyUI Official Example, accessed June 14, 2025, https://docs.comfy.org/tutorials/api-nodes/stability-ai/stable-diffusion-3-5-image
48. Stable Diffusion 3: Research Paper - Stability AI, accessed June 14, 2025, https://stability.ai/news/stable-diffusion-3-research-paper
49. The Stable Diffusion Guide - Hugging Face, accessed June 14, 2025, https://huggingface.co/docs/diffusers/v0.12.0/en/stable_diffusion
50. Stable Diffusion Integrations | Connect Stable Diffusion to your apps with Albato, accessed June 14, 2025, https://albato.com/apps/stable_diffusion
51. StableDiffusion - Dify, accessed June 14, 2025, https://docs.dify.ai/guides/tools/tool-configuration/stable-diffusion
52. Tabnine Integrations - SourceForge, accessed June 14, 2025, https://sourceforge.net/software/product/Tabnine/integrations/
53. Plans & Pricing | Tabnine: The AI code assistant that you control, accessed June 14, 2025, https://www.tabnine.com/pricing/
54. Supported IDEs - Tabnine Docs, accessed June 14, 2025, https://docs.tabnine.com/main/welcome/readme/supported-ides
55. Stable Diffusion API Docs, accessed June 14, 2025, https://stablediffusionapi.com/docs/
56. Text to Image Endpoint | Stable Diffusion API Documentation, accessed June 14, 2025, https://stablediffusionapi.com/docs/stable-diffusion-api/text2img/
57. arXiv:2503.19075v1 [cs.CY] 24 Mar 2025, accessed June 14, 2025, https://www.arxiv.org/pdf/2503.19075