"Juice Box Break" and the Semiotics of Sentient Stillness: A Report on Sensory Computation and Childhood Evocation in Generative AI

Introduction: The Algorithmic Madeleine

In the landscape of contemporary technology, generative artificial intelligence, particularly text-to-image diffusion models, has emerged as a powerful new medium for visual creation. Beyond its capacity for photorealism and stylistic mimicry, this technology exhibits a startling, almost preternatural ability to evoke complex emotional and sensory states. It can conjure images that resonate with the specificity of a personal memory, a fleeting feeling, or a deeply embedded cultural sentiment. This report uses the conceptual test case of a "Juice Box Break"—a seemingly simple scene of childhood pause and refreshment—as a lens through which to dissect this phenomenon. Like Proust's madeleine, a generated image of this moment can unlock a cascade of associations: the sticky sweetness of juice, the warmth of an afternoon sun, the quiet contentment of a brief rest.

The central inquiry of this report is to move beyond the surface of these evocative images and analyze the intricate computational, semiotic, and ethical systems that produce them. The core problem is to understand how a non-sentient system, devoid of memory, embodiment, or emotion, can so effectively simulate the visual language of these human experiences. This investigation questions whether the process constitutes a genuine evocation of feeling or a sophisticated, high-fidelity mimicry of its representation. Is the AI a nascent artist, or is it an algorithmic necromancer, reanimating the ghosts of our collective visual culture?

To answer these questions, this report employs a multi-layered methodological framework, proceeding through four interlocking analytical vectors. First, it investigates **Semantic Encoding**, deconstructing how abstract feelings are translated into the statistical language of latent space. Second, it explores **Temporal Narrative Simulation**, analyzing how a static, single-frame image can be engineered to imply a sequence of events and a sense of time. Third, it examines **Archetype Activation**, dissecting how cultural memory is leveraged, remixed, and potentially exploited by models trained on vast, uncurated datasets. Finally, it delves into **Computational Subtlety**, assessing the role of micro-fidelity and sensory simulation in constructing a believable, immersive reality.

This report is structured for an expert audience of AI model designers, affective computing scholars, prompt engineers, and visual culture theorists. Its objective is to furnish a technically grounded, critical-aesthetic framework for understanding, evaluating, and designing more emotionally aware and ethically responsible generative systems. By dissecting the "Juice Box Break," we aim to illuminate the profound capabilities and inherent paradoxes of AI as it learns to render the ineffable qualities of human experience.

I. The Latent Semantics of Sensation: Encoding Affective States

This section deconstructs the foundational process of how abstract concepts of feeling and sensation are encoded within the architecture of generative models. It moves from the text-image pairing mechanism to the geometric organization of the latent space itself, exposing the statistical, non-embodied nature of AI's "emotional intelligence."

1.1. Emotional Tokenization: The Architecture of Association

The capacity of a diffusion model to generate an image of a "nostalgic juice box break" begins not with an understanding of nostalgia, but with a process of statistical association. The architectural bedrock for this is often a multimodal model like OpenAI's Contrastive Language-Image Pre-training (CLIP).[1] CLIP is trained on hundreds of millions of image-text pairs scraped from the internet, learning to create a shared, high-dimensional embedding space where visual and linguistic concepts are co-located.[1] Through a process of contrastive learning, the model is trained to maximize the cosine similarity between the vector for a given image and the vector for its corresponding text caption, while minimizing its similarity to incorrect captions.[1] This is the fundamental mechanism of

emotional tokenization. An abstract affective state like "joy" is not comprehended as a feeling; rather, the word "joy" becomes a token whose embedding is statistically proximal to the embeddings of images that were captioned with "joy" in the training data—images that likely feature smiles, bright colors, and sunny environments.[1]

When integrated into a text-to-image diffusion model, this text embedding becomes the primary guide for the generation process. The core of a diffusion model's generator is typically a U-Net architecture that learns to progressively denoise a random noise field into a coherent image. The guidance from the text prompt is injected via **cross-attention layers**.[4] At each step of the denoising process, these layers correlate the current state of the noisy image with the semantic information contained in the text embedding. If a prompt includes "nostalgic warmth," the cross-attention mechanism will guide the U-Net to generate and preserve visual features that are statistically linked to that concept in the CLIP space, such as "sun flare," "soft focus," or a "Kodachrome" color palette.[5]

The fidelity of this process is purely statistical, not experiential. This is evidenced by research showing that AI models are particularly effective at generating images depicting positive emotions like "joy," especially in familiar contexts like home interiors, but struggle to accurately convey negative emotions such as anger or sadness.[7] This discrepancy suggests a significant bias in the training data, where positive emotions are more frequently and consistently represented in visual media than negative ones. This reveals the model's reliance on what can be termed

non-visual associative fidelity—its ability to faithfully reproduce statistical links between words and pixels, without any grounding in the non-visual, embodied experience of the emotion itself.

1.2. A Cartography of Feeling: Mapping the Latent Space

The latent space of a generative model is a compressed, abstract, and lower-dimensional representation of the training data. Within this space, data points with similar characteristics are positioned closer to one another, forming a kind of semantic map or cartography.[9] For our "Juice Box Break" scenario, a point in latent space that decodes to the target image would be "geographically" situated near the vector representations for concepts like "childhood," "summer afternoon," "refreshment," and "rest." Exploring this space is not random; it is a navigation through a universe of learned concepts.

Recent research has begun to map this terrain with greater precision. Studies using Singular Value Decomposition (SVD) on the latent space of diffusion models have discovered that it possesses an interpretable geometric structure.[12] It is possible to identify specific semantic vectors within this space. For instance, by learning the vector that represents the difference between latent codes for "a photo of a person" and "a photo of a

young person," one can apply this "young" vector to other images to manipulate age.[12] This implies the existence of vectors for more abstract and affective qualities. A "nostalgia" vector could be isolated and applied to an image, systematically shifting its color palette, lighting, and textural properties toward a more vintage aesthetic.

This exploration is made more efficient by advanced architectures that conduct the diffusion process directly within the latent space, rather than in the high-dimensional pixel space.[13] Operating on this more abstract, semantically rich level allows the model to manipulate core concepts more directly.[10] The concept of "conditional dependence," borrowed from psychometric diffusion models used in psychology, provides a powerful analogy for this process.[16] In these models, the distance between a person and an item in a latent space can quantify their specific interaction. In our generative context, the proximity between the latent representation of "juice box" and "warmth" represents a learned, non-causal statistical interaction derived from their co-occurrence in the training dataset.[17] This is not reasoning, but a reflection of the data's structure. Furthermore, there is evidence that models can perform a form of non-linguistic reasoning or "thinking" within this latent space, capturing relationships that are not easily articulated in words.[18] This suggests that the final affective quality of a generated image arises from these complex geometric relationships between concepts, which are then decoded by the generator into the final pixel arrangement.

1.3. The Boundary of Believability: Affective Embedding Drift

The process of emotional tokenization and latent space mapping inevitably creates a gap between the model's statistical approximation of an emotion and the rich, embodied reality of human feeling. The model learns to associate the concept of "refreshment" with visual tokens like condensation on a surface, relaxed postures, and cool-toned color palettes because these features are statistically overrepresented in its training data, likely sourced from countless beverage advertisements and stock photos.[1] This is a proficient approximation of the

visual signs of refreshment, but it is not an evocation born from the experience of thirst or relief.

This leads to a phenomenon that can be termed **affective embedding drift**. The model's learned vector

representation for a given emotion or concept drifts away from its nuanced, true-to-life meaning, becoming skewed by the biases and idiosyncrasies of the training dataset. For a concept like "childhood," the training data scraped from the web is likely saturated with idealized, commercial, and culturally specific imagery (e.g., American television commercials from the 1990s).[19] Consequently, the model's embedding for "childhood" may drift towards this specific, commercially-inflected aesthetic, making it difficult to generate images of childhood that feel authentic to other cultural contexts or time periods.

This drift is a direct consequence of the model's non-embodied "experience." It learns exclusively from disembodied image-text pairs, lacking any grounding in physical sensation, social interaction, or internal consciousness.[20] Its output is therefore a sophisticated recombination of how humans have

represented feelings, not a generation from the feeling itself. This marks the fundamental boundary of its evocative power. Research involving the inversion of CLIP models—generating an image that best matches a text prompt's embedding—reveals the alien logic at play. When prompted with "an interested person," the model produced an image with grotesquely enlarged ears, a literal and non-human interpretation of "listening".[5] This demonstrates that while the model can master the statistical correlations of our world, its "understanding" remains profoundly different from our own.

Deliverable: Latent Affection Map (LAM)

The following Latent Affection Map (LAM) serves as a practical tool for deconstructing the "black box" of emotional generation. It provides a structured translation from abstract human feeling to the concrete visual tokens and underlying architectural mechanisms that produce the affective illusion. For designers and prompt engineers, this map transforms the art of prompting from intuitive guesswork into a more structured, semiotic process. It allows for more precise control over the final output by targeting specific visual tokens and provides a framework for evaluating models by identifying where affective embedding drift is most likely to occur. For instance, to evoke "quiet nostalgia," instead of relying on the abstract phrase alone, an engineer can use the LAM to construct a more robust prompt targeting associated tokens like "faded color palette," "soft focus," and "low sun angle," while also leveraging proximate keywords like "vintage photography" to guide the model's style.

Affective State	Sensory Correlate	Primary Visual Tokens	Secondary/Compositional Tokens	Latent Space Proximity & Keywords	Architectural Enabler
Nostalgia	Warmth, Softness, Faded Memory	Sun flare, lens dust, warm color cast (yellow/orange), high-key lighting	Relaxed/resting pose, vintage clothing/objects (e.g., 90s juice box design), shallow depth of field	vintage photo, 70s summer, Kodachrome, faded memory	CLIP text-image embedding, U-Net denoising guided by attention [1]
Refreshment	Coolness, Wetness, Relief	Condensation on surface, bright/cool colors (blue/green), splashing liquid	Close-up shot, dynamic pose (mid-sip), high-contrast lighting	cold drink, summer splash, ice cold, beverage ad	CLIP text-image embedding, Attention mechanism focusing on surface texture [1]

Innocence	Simplicity, Safety, Play	Soft lighting, simple composition, primary colors, uncluttered background	Child looking away from camera, discarded toys, natural setting (grass, park)	childhood, innocent, playful, carefree	Dataset bias towards idealized childhood imagery, attention on subject pose [22]
Stillness/Pause	Quiet, Liminality, Contemplation	Static pose, long shadows (low sun angle), lack of motion blur	Half-finished activity (half-drunk juice), gaze directed inward/downward	a quiet moment, pause, contemplation, taking a break	Prompt scheduling weights, compositional cues from narrative-captioned images [24]

II. The Stillness of Time: Simulating Temporal Arcs in Static Images

This section investigates the paradox of how a static, single-frame image can evoke a sense of time, narrative, and causality. It explores the technical and compositional strategies that allow a generative model to capture a "visual exhale"—a liminal state pregnant with the immediate past and the impending future.

2.1. Prompt as Chronoscript: Engineering Inferred Temporality

The generation of a static image that implies a temporal narrative is a sophisticated act of prompt engineering, transforming the text input from a simple description into a "chronoscript." This involves leveraging advanced prompting techniques that guide the model to render a scene rich with temporal cues.

One such technique is **temporal layering** through prompt weighting and scheduling. By assigning different weights to various components of a prompt, an engineer can signal their relative importance and imply a temporal or causal sequence.[24] For example, a prompt structured as

(A child running through a field:0.3), (a child sitting on the grass drinking a juice box, tired but happy:1.1) instructs the model that the "running" is a preceding, less central action, while the "sitting and drinking" is the primary focus of the generated moment. This method adapts logic from video generation, where prompt scheduling is used to guide different phases of an animation, conditioning the model on a sequence of states.[26]

Another powerful approach involves **compositional decomposition**, as seen in frameworks like Recaption, Plan and Generate (RPG).[28] RPG uses a multimodal large language model (MLLM) to break down a complex prompt into sub-prompts that are then assigned to different spatial regions of the image. This can be repurposed for temporality. A prompt could be planned to generate a scene with "a pair of muddy sneakers discarded at the edge of the frame" in one region and "a child resting peacefully with a juice box" in the center. The spatial co-presence

of these elements creates a narrative: the muddy sneakers are the trace of a past action (playing), which gives context to the present state (resting). The model synthesizes not a timeline, but a single frame containing the evidence of a timeline.

Furthermore, the model's ability to process sequential instructions, a capability honed by training on video data or sequentially captioned image sets, can be leveraged. Prompting guides for advanced video models explicitly teach the technique of **chaining emotions and gestures**, such as "Starts calm. Then grimaces in a flash of violence... One beat later, he's composed".[30] Although the desired output here is a static image, the model's internal logic can translate this sequence into a single, potent frame that captures the

peak of the action, the *transition* between states, or the quiet *aftermath*. The model learns to find the single most representative moment from the described temporal arc.

2.2. The Visual Exhale: Compositional Cues for Liminal States

The ultimate goal of temporal simulation in a static image is to capture a "visual exhale"—a liminal state that is neither pure action nor pure stillness, but a pause pregnant with meaning. This is achieved not through technical animation but through careful composition, where visual elements are arranged to suggest a narrative and imply causality.

Narrative sequencing through composition is the art of telling a story with objects and their relationships. For the "Juice Box Break," a compelling narrative can be constructed from a collection of carefully chosen visual cues. Grass stains on the child's knees serve as evidence of the "before" (play). The half-empty juice box itself represents the "during" (the act of refreshment). The low angle of the sun, casting long shadows, suggests the day is drawing to a close, implying an "after".[31] Generative models learn these powerful compositional tropes from their training data, which includes millions of images captioned with narrative descriptions, such as those found in storyboarding or visual storytelling datasets.[25] The model doesn't understand the story, but it understands that these elements frequently co-occur in images described as "a quiet moment after playing."

The stillness of the scene is often broken by subtle, yet critical, cues of life and feeling. A character's posture—slumped shoulders, a hand loosely holding the juice box—can imply a release of tension following physical exertion. Even in a static image, models can generate nuanced facial expressions. While these are not true **micro-expressions**, which are defined by their fleeting, sub-second duration, they can capture the visual essence of a transitional emotion.[34] For instance, the model might generate a face with a slight, unfocused smile and relaxed eyelids, perfectly conveying the quiet contentment that follows the quenching of thirst. This adds a layer of psychological realism that makes the "pause" feel authentic and relatable.[36]

Ultimately, the liminal frame is a carefully constructed illusion. It freezes a moment that is maximally suggestive of temporal flow. The image becomes a symbol of the pause itself, a visual representation of an exhale. The principles of intelligent scene composition, which are central to AI-powered storyboarding tools that aim to create narratively compelling sequences, are here condensed into a single, powerful frame.[32] The model is not just rendering a scene; it is rendering a moment suspended in time.

Deliverable: Temporality Implied Frame Test (TIFT)

The Temporality Implied Frame Test (TIFT) provides a structured methodology for prompt engineers and visual culture theorists to evaluate the capacity of generative models to embed narrative and temporal cues within a static image. By systematically varying prompts to target different phases of a micro-narrative—anticipation, the moment itself, and the aftermath—the TIFT allows for a nuanced assessment that moves beyond simple prompt adherence. It offers concrete criteria for judging the success of implied causality, emotional progression, and the creation of a liminal state. This framework is invaluable for artists, advertisers, and designers seeking to create more sophisticated, emotionally layered, and narratively rich visual content using generative AI. It transforms the goal of "implying a story" from an abstract ambition into a testable and refinable engineering challenge.

Prompt Variant	**Generated Image Analysis**	**Implied Causality Score (1-5)**	**Emotional Progression Score (1-5)**	**Liminality ('Visual Exhale') Score (1-5)**
Prompt 1: The Anticipation "Golden hour, a young child, slightly out of breath, eagerly reaches for a juice box on a picnic blanket. Close-up on the hands and the box, a sense of anticipation."	The analysis focuses on whether the image successfully captures pre-action tension. Key visual cues include the child's posture (leaning forward), facial expression (eager, perhaps with flushed cheeks suggesting recent activity), and the compositional focus on the object of desire (the juice box). The lighting ("golden hour") should contribute to the feeling of a culminating moment.	**4** - The image successfully implies a preceding action (e.g., playing) through cues like being "out of breath" and the eagerness of the reach. The cause (thirst, exertion) for the effect (reaching for juice) is strongly suggested.	**3** - The image captures the beginning of an emotional arc. The primary emotion is anticipation, a state of looking forward to a future feeling of relief. The model effectively renders this forward-looking emotional state.	**2** - This frame is less about a pause and more about the tension *before* the action. It is a moment of high potential energy, not a visual exhale. It represents the inhale before the release.
Prompt 2: The Moment "A child sits on sun-dappled grass, eyes closed in a moment of quiet refreshment, sipping from a juice box. Cinematic, shallow depth of field, soft light."	The analysis evaluates the image's ability to portray the central action not as a dynamic event but as a state of being. Important cues are a relaxed posture, a serene facial expression (eyes closed), and a focus on the sensory experience of refreshment. The cinematic and lighting instructions should enhance the introspective mood.	**5** - The image strongly fulfills the implied causality of the "Anticipation" frame. The act of drinking and the expression of contentment are the direct result of the previously established need for refreshment. The narrative sequence is clear and satisfying.	**5** - This frame represents the peak of the emotional progression. The anticipation has resolved into contentment and relief. The model's rendering of a serene, almost meditative state is a powerful depiction of this emotional climax.	**5** - This is the quintessential "visual exhale." The action (sipping) is minimal and serves the state of being (refreshment). The closed eyes, soft light, and relaxed pose all contribute to a powerful sense of a pause, a liminal moment of pure sensory experience between past exertion and future action.
Prompt 3: The Aftermath "An empty, crushed juice box and a bent straw lie on the grass.	The analysis assesses the use of absence and remnants to imply a completed narrative. The state of the objects	**4** - The causality is implied through remnants. The empty box is the clear effect of a cause (a child	**4** - The emotional state is now one of post-narrative reflection, evoked in the viewer rather than	**3** - While the scene is still, it represents a concluded state rather than a liminal pause. It is the silence *after* the

The late afternoon sun casts long shadows. The scene is quiet, the child is gone."	(crushed box, bent straw) and the environmental cues (long shadows, emptiness) are critical. The image's success depends on its ability to evoke a sense of a story that has already concluded.	drinking). The absence of the child strongly suggests they have moved on, completing the narrative arc.	depicted in a character. The quiet, empty scene can trigger feelings of nostalgia, peace, or even a slight melancholy, demonstrating a sophisticated emotional transfer.	exhale, a moment of finality. The liminality has passed, leaving only its trace.

III. The Archetype Engine: Cultural Memory and Ethical Simulation

This section transitions from the technical mechanisms of generation to the cultural and ethical dimensions of the content itself. It examines how generative models function as powerful engines for reproducing, remixing, and potentially exploiting potent cultural archetypes. This process raises critical questions about the ownership of memory, the ethics of consent, and the profound paradox of a non-sentient system simulating emotional truth.

3.1. Data-Derived Archetypes: The Ghost in the Machine

The concept of archetypes, as originally formulated by Carl Jung, refers to universal, innate psychic structures that reside in the collective unconscious—primordial images like The Hero, The Great Mother, or The Divine Child that shape human experience across cultures.[37] Generative AI, however, operates on a fundamentally different principle. Its archetypes are not innate but

data-derived. The model does not possess a collective unconscious; it constructs its archetypal representations from the statistical regularities and biases present in its vast training data.[22]

For our "Juice Box Break" scenario, the model will inevitably draw upon the powerful archetype of **Childhood Innocence**. An analysis of the likely composition of its training data, such as the web-scraped LAION dataset, reveals that this archetype is not a pure, universal concept. Instead, it is a complex collage assembled from a multitude of culturally specific and often commercial sources: the vibrant, high-energy aesthetics of 1990s cereal and juice commercials; the distinct visual styles of popular cartoons; the soft-focus, warm-toned conventions of professional family photography; and millions of amateur snapshots shared on social media and blogs.[19] The resulting AI-generated archetype is therefore not a timeless truth but a specific, heavily inflected remix, often reflecting a Western, consumerist vision of childhood.

The model learns the **semiotics of this archetype** through statistical correlation. The visual language of "innocence" in the data includes a consistent set of codes: soft, diffused lighting; high-key exposures; a specific color palette favoring primary colors and pastels; and common compositional tropes like a shallow depth of field or a child looking away from the camera in a moment of un-self-conscious play.[21] The model learns that these visual features have a high probability of co-occurring with text captions containing words like "innocent," "childhood," or

"playful." Its ability to generate an "innocent" scene is a testament to its mastery of these semiotic codes, not its understanding of the concept itself.

3.2. Ventriloquized Memory: The Ethics of Data Necromancy

The process by which AI models learn and reproduce these archetypes is fraught with profound ethical complications. The LAION-5B dataset, a foundational training set for popular models like Stable Diffusion, was created through an unsupervised crawl of the internet, resulting in a collection of over 5.8 billion image-text pairs.[40] Critically, this dataset is known to contain not only a vast repository of publicly available images but also a significant volume of copyrighted material, private medical records, intimate family photographs, and even verified Child Sexual Abuse Material (CSAM).[23]

When a model is prompted to generate a "nostalgic" image of childhood, it performs an act that can be described as **ventriloquized memory**. It is speaking with the "voice" of its data, a voice composed of the uncredited labor of artists, the intellectual property of corporations, and the private, often sensitive, moments of countless individuals—all appropriated without consent, credit, or compensation.[19] This is not merely a copyright dispute; it is a form of data necromancy, where the model reanimates the visual ghosts of our collective past. One analysis has referred to the output as a potential "trauma collage," where a single generated image might unknowingly abstract and blend patterns learned from images of historical atrocities, private grief, and commercial advertising.[23]

This practice of large-scale, un-consented data scraping places generative AI development on a collision course with emerging legal and regulatory frameworks. The European Union's General Data Protection Regulation (GDPR), for instance, establishes strong protections for "Personal Data," a category that explicitly includes images of identifiable individuals and can be extended to cover inferred emotional data.[41] The EU AI Act goes further, prohibiting AI systems that employ subliminal or manipulative tactics and placing strict restrictions on the use of emotion recognition AI in sensitive contexts like workplaces and schools.[41] The very foundation of many current generative models is built on a data acquisition model that appears to be in direct violation of the principles of data minimization, purpose limitation, and informed consent that underpin these regulations. The generation of archetypes is, in this light, the generation of derivative works from a dataset of questionable legal and ethical standing.

3.3. The Empathy Paradox: Simulating Truth Without Experience

At the heart of AI-driven emotional evocation lies a fundamental paradox. On one hand, the technology is rapidly advancing. Large language models can generate text and speech that are affectively aligned with human emotional expression, and models like ChatGPT-4 have demonstrated a human-level capacity to interpret emotions from both visual and textual data.[42] On the other hand, the model simulates an emotional truth that it cannot, in any meaningful sense, experience.[20] It possesses no consciousness, no embodiment, no lived history from which to draw genuine feeling. Its "empathy" is a high-fidelity pattern-matching exercise.[46]

This disconnect has significant implications for how the resulting artwork is perceived. The authenticity of human art is deeply intertwined with the viewer's belief in the artist's intentionality, moral vision, and lived experience.[46] Studies have shown that human observers consistently rate artworks lower when they are told the artist was an AI, attributing less sensitivity, creativity, and emotional depth to the piece.[46] This suggests that the "emotional truth" of an AI-generated image is not an inherent quality of the artifact itself. Rather, it is a fragile illusion, completed by the viewer's act of projection and dependent on their potential ignorance of the artist's non-human nature.

The philosophical ramifications of this empathy paradox are vast. When our most cherished cultural archetypes and personal memories are simulated and reflected back at us by non-sentient systems, we risk entering a cultural feedback loop. This loop could flatten and homogenize human expression, reducing it to the most statistically probable visual patterns found in the training data—patterns that are overwhelmingly sourced from a narrow, predominantly Western and commercialized slice of global culture.[19] This process challenges our core definitions of art, creativity, authorship, and even the uniqueness of human experience itself.[48] We are building engines that are masters of our emotional language, but they remain incapable of understanding the meaning behind the words.

Deliverable: Archetype Remediation Index (ARI)

The Archetype Remediation Index (ARI) is a proposed risk-assessment framework designed for ethical AI development and deployment. It provides a structured method for developers, designers, and legal compliance teams to proactively identify, score, and mitigate the ethical and legal risks associated with generating archetypal content. By mapping specific archetypal elements to their likely data source exposures and suggesting concrete remediation strategies, the ARI transforms a vague ethical concern into an actionable risk management process. This tool is critical for navigating the complex landscape of copyright, privacy, and cultural appropriation inherent in models trained on large, uncurated web data. It enables the creative use of archetypes while minimizing legal, financial, and reputational harm.

Archetypal Element	Potential Data Source Exposure	Primary Ethical/Legal Risk	Risk Score (1-5)	Remediation Strategy & Prompting Tactic
The "Juice Box"	Trademarked brand logos/designs (e.g., Hi-C, Capri Sun) from 90s ads and photos found in LAION.[23]	**Trademark Infringement:** Unauthorized commercial use of a brand's likeness.	4	**Abstraction:** Prompt for a "generic retro juice box" instead of a specific brand. Use negative prompts for known brand names (--no capri sun).
The "90s Cartoon" Style	Copyrighted animation styles (e.g., 'Doug,' 'Rugrats') scraped from fan sites, wikis, and screenshots.[23]	**Copyright Infringement/Style Plagiarism:** Replicating a protected artistic style for commercial gain.	5	**Stylistic Blending:** Mix multiple artist/style influences with weights (e.g., in the style of

				[artist1]:0.6 and [artist2]:0.4). Fine-tune on a licensed or public domain dataset of period illustrations.	
"Childhood Innocence" Pose	Private family photos from blogs, Flickr, social media, potentially including sensitive or identifiable images of minors.[23]	**Violation of Privacy/Right to be Forgotten (GDPR):** Using personal, non-public images of identifiable individuals, especially minors, without consent.[41]	5	**Dataset Curation & Filtering:** Utilize models trained on ethically sourced, filtered datasets (e.g., LAION-5B with NSFW filters enabled [40]).	**Prompt for Anonymity:** Use prompts that specify non-identifiable features, e.g., "a child seen from behind," "face obscured by shadow."
Seasonal Nostalgia (e.g., "Summer Afternoon")	Stock photography with restrictive licenses; professional photography from portfolios.	**Copyright Infringement:** Replicating the composition, lighting, and subject of a protected photograph.	3	**Compositional Remixing:** Use regional diffusion prompts [28] to combine elements from different conceptual sources, creating a novel composition that is less likely to be a direct copy.	

IV. The Texture of Reality: Computational Photography and Affective Presence

This final analytical section examines the micro-level techniques that construct the illusion of sensory reality. It investigates how computational mimicry of photographic and physical phenomena creates a sense of tactile immersion and asks at what point this hyper-fidelity becomes counterproductive to the emotional goal.

4.1. Hyperreal Microfidelity: The Algorithms of Atmosphere

Modern generative models do not merely synthesize scenes; they imbue them with an atmospheric realism by incorporating sophisticated techniques from the field of computational photography.[51] These algorithmic enhancements are not just for achieving technical accuracy; they are powerful tools for modulating the affective content of an image.

High Dynamic Range (HDR) is a prime example. By algorithmically merging multiple exposures, HDR techniques produce images with a vastly expanded range of luminance, capturing details in both the brightest highlights and the deepest shadows that a single exposure would lose.[53] From a psychological perspective, this expanded range can make a scene feel more vibrant, immersive, and "real," as it more closely approximates the adaptive capabilities of the human visual system.[55] Viewers often describe non-HDR images as "fake" looking precisely because they lack the natural contrast the brain expects to see.[57] In the context of the "Juice Box Break," HDR rendering can make the sun feel genuinely dazzling and the shade under a tree feel cool and deep, intensifying the sensory experience of a hot summer day.

Similarly, the simulation of **Depth of Field (DoF)** is a critical affective tool. A shallow DoF, which renders the subject in sharp focus while blurring the background into a soft bokeh, is a powerful compositional device for directing the viewer's attention and creating a profound sense of intimacy or isolation.[58] This effect mimics the natural way the human eye focuses on a subject of interest, and its application in photography is known to evoke specific emotional responses, such as admiration, pensiveness, or drama.[60] For our scene, a shallow DoF isolating the child would emphasize their personal moment of refreshment, separating them from the distractions of the wider world.

The ultimate goal of these techniques is not necessarily physical accuracy but **perceptual realism**—the creation of imagery that is, to a human observer, indistinguishable from a real-world scene or a photograph of one.[61] This involves modeling not just the physics of light, but also the known limits and heuristics of the human visual system.[63] The model does not need to render every photon with perfect fidelity; it needs to render an image that satisfies our perceptual expectations and feels congruent with our experience of the world.

4.2. Sensory Recursion and Computational Tactility

To move beyond simple atmosphere and create a true sense of tactile immersion, generative models must simulate the very texture of reality. This is achieved through advanced techniques that generate not just patterns, but the visual cues of physical properties.

Recursive wavelet refinement offers a powerful analogy for how models can generate convincing textures, such as the matte finish of a cardboard juice box or the complex, varied blades of grass. This method involves decomposing an image into different frequency bands—separating the low-frequency structural information from the high-frequency detail and texture—and then progressively refining these layers.[65] This multi-scale approach, whether implemented directly with wavelets or through analogous hierarchical structures in a deep network, allows for the generation of textures that are not flat, repeating patterns but possess a sense of depth, inconsistency, and physical plausibility that invites closer visual inspection and reinforces the object's realism.[68]

The most sophisticated models can even evoke a sense of touch from a purely visual stimulus, a phenomenon that can be termed **computational tactility**. The model cannot, of course, generate a physical haptic sensation. Instead, it generates the precise visual cues that trigger a haptic memory or sensation in the viewer's brain. This is the principle of **visual haptics**. Neuroscience research confirms that the human brain seamlessly integrates visual, auditory, and tactile information to form a cohesive perception of the world.[69] When a generative model renders hyper-realistic beads of condensation on the side of the juice box, it is leveraging this cross-modal integration. The viewer's brain, upon processing the visual cue of condensation, is primed to access the associated haptic memory

of how a cold, wet object feels to the touch.[71] The tactility is not in the pixels; it is a recursive sensory experience completed in the mind of the observer.

This pursuit of realism extends to the human subject as well. Even in a static image, the simulation of subtle muscle activations in a character's face—a slight furrow of the brow indicating concentration, the gentle upturn at the corner of a mouth signaling contentment—can dramatically enhance emotional realism. These are not true, time-based **micro-expressions**, but their static visual equivalents. Learned from vast datasets of human faces, these subtle cues add a layer of believable, fleeting emotion that deepens the character's psychological complexity and believability, making the captured moment feel more authentic.[34]

4.3. Functional vs. Fetishistic Fidelity: The Risk of Affective Overkill

In the quest for realism, there exists a critical threshold where the pursuit of technical fidelity becomes detrimental to the image's emotional power. It is essential to distinguish between **functional fidelity** and **fetishistic fidelity**. Functional fidelity employs computational techniques in service of the image's narrative and emotional goals; for example, using a shallow depth of field to create a feeling of intimacy. Fetishistic fidelity, in contrast, occurs when the technical simulation becomes the subject of the image itself. The perfectly rendered water droplet, the mathematically precise lens flare, the impossibly detailed texture—these elements can draw attention to the image's own artifice, breaking the viewer's immersion and undermining the intended emotional connection.[75]

This leads to the risk of **affective overkill**: the point at which hyper-realism becomes so technically perfect that it feels sterile, cold, and inhuman. The image triggers a cognitive dissonance in the viewer, who becomes acutely aware that they are looking at a flawless simulation rather than a captured moment. This awareness can inhibit a genuine emotional response, as the image ceases to be a window into a remembered or imagined world and instead becomes a mere demonstration of the algorithm's technical prowess.[46]

The power of real photography often lies in its imperfections—the subtle motion blur, the film grain, the unexpected framing, the "happy accidents" that signify a human hand and a unique moment captured in time. The very ability of AI to generate technically flawless images can, paradoxically, strip them of the authenticity and emotional resonance that these imperfections provide.[75] The challenge for designers and artists is to wield these powerful tools of realism with restraint, ensuring that fidelity serves feeling, rather than overwhelming it.

Deliverable: Fidelity-to-Feeling Matrix (FFM)

The Fidelity-to-Feeling Matrix (FFM) is a practical framework for artists, designers, and prompt engineers using generative AI. It aids in making deliberate, informed decisions about which computational effects to employ, how to modulate them, and when their application might become counterproductive. The matrix moves the conversation beyond a simplistic pursuit of "realism" toward a more nuanced, intentional use of computational techniques to serve specific emotional and narrative ends. It provides a guide for balancing sensory authenticity with the risk of affective overkill, helping creators to produce work that is both technically impressive and emotionally resonant.

Computational Technique	Contribution to Sensory Authenticity (Scale: 1-Gimmick, 2-Atmospheric, 3-Immersive, 4-Essential)	Risk of Affective Overkill (Scale: 1-Low, 2-Medium, 3-High)	Optimal Use Case for "Juice Box Break"	Relevant Sources
Shallow Depth of Field (Bokeh)	4 - Essential	1 - Low	Isolate the child from the background, creating intimacy and focusing the viewer's attention on the emotional moment of refreshment.	58
High Dynamic Range (HDR) Lighting	3 - Immersive	2 - Medium	Create strong contrast between bright sun and deep shadows to evoke the feeling of a hot summer day. Avoid unnatural, "glowing" edges that announce the technique.	51
Simulated Condensation/Frost	3 - Immersive	3 - High	Add small, imperfect beads of condensation to the juice box to evoke a tactile sense of "cold" and "wet." Avoid overly perfect, uniform droplets that look computer-generated.	69
Recursive Wavelet Texture	2 - Atmospheric	2 - Medium	Generate a subtle, non-uniform cardboard texture for the juice box and a varied texture for the grass to enhance physical presence without drawing undue attention.	65
Lens Flare & Chromatic Aberration	2 - Atmospheric	3 - High	Use a soft, diffuse lens flare to suggest the sun's presence outside the frame and enhance the "photographic memory" feel. Avoid sharp, geometric, or overly colorful flares.	21
Micro-Expressions	3 - Immersive	2 - Medium	Add a subtle, fleeting expression of contentment or relief to the child's face to deepen the emotional realism of the "pause" moment.	34

Conclusion: The Authorship of Affect and the Future of Simulated Memory

This report has traversed the complex landscape of generative AI's capacity to evoke the sensory and emotional experience of a childhood memory, using the "Juice Box Break" as a guiding case study. The analysis across four distinct but interlocking vectors—from the statistical encoding of feeling in latent space to the ethical implications of archetypal simulation—reveals a technology of profound capability and equally profound paradox. The journey from a simple text prompt to an affectively resonant image is not a magical act of creation, but a multi-layered process of statistical reconstruction, cultural remixing, and perceptual illusion.

The synthesis of these findings provides clear, if unsettling, answers to the critical reflexive questions that motivated this inquiry.

First, **what emotional truths can be simulated without experience?** Generative models can proficiently simulate the *visual representation* of emotional truths, particularly those whose semiotics are well-documented, visually consistent, and heavily represented in the training data. The smile of joy, the tears of sadness, the warm glow of nostalgia—these are reproducible because their visual language is canonical. However, the models fail to capture nuanced, context-dependent, or deeply internal states that lack a consistent visual signifier. The "truth" they generate is a truth of appearance, a mastery of the image-as-sign, not a truth of being. It is a simulation of the external manifestation of emotion, not the internal experience.[45]

Second, **when does high fidelity become affective overkill?** Fidelity crosses the line into affective overkill at the precise moment the viewer's cognitive and emotional focus shifts from the *subject of the image* to the *technique of its creation*. This is the point where perceptual congruence breaks down. When a perfectly rendered texture is too sharp for its place in the depth of field, or a simulated lens flare is too geometrically perfect, the illusion shatters. The viewer is reminded that they are observing a simulation, and this awareness disrupts the fragile process of emotional projection and immersion. Affective overkill occurs when the technology announces itself, transforming an evocative image into a sterile technical demonstration.[75]

Finally, **where does authorship truly reside in such emotionally resonant renders?** Authorship in generative art is not a singular point but a distributed network of agency. It resides, first and foremost, with the countless, often uncredited, human creators of the original data—the photographers, artists, and ordinary people whose lives and works form the digital clay from which these models are sculpted.[23] It resides with the AI architects and engineers who design the models' structures and training methodologies. It resides with the prompt engineer, who acts as a creative director, providing the specific impetus, constraints, and narrative guidance for a given generation.[24] And finally, it resides with the human viewer, who must project their own memories, cultural knowledge, and narrative intelligence onto the static image to complete its meaning and feel its emotional weight. The AI itself is not an author; it is a powerful, complex, and transformative lens—a new kind of cultural infrastructure—through which these other human agents act and interact.[47]

Looking forward, the trajectory of generative AI as a medium for creating emotionally resonant and sensorially rich art will depend on our ability to move beyond the simple pursuit of photorealism. The true artistic and cultural potential of this technology may not lie in its ability to perfectly mimic our world, but in its capacity to generate novel forms of expression that leverage its unique, non-human "imagination." This future, however, is contingent upon the development of robust ethical frameworks and a fundamental shift in the data practices of the industry. To build a future of simulated memory that is enriching rather than exploitative, we must acknowledge and respect

the human sources of culture, creativity, and feeling from which these powerful systems ultimately draw their evocative power.

Works cited

1. CLIP: Contrastive Language-Image Pre-Training - viso.ai, accessed June 17, 2025, https://viso.ai/deep-learning/clip-machine-learning/
2. CLIP: Connecting text and images - OpenAI, accessed June 17, 2025, https://openai.com/index/clip/
3. Leveraging CLIP for Image Emotion Recognition - CEUR-WS.org, accessed June 17, 2025, https://ceur-ws.org/Vol-3015/paper172.pdf
4. Controllable Text-to-Image Generation, accessed June 17, 2025, http://papers.neurips.cc/paper/8480-controllable-text-to-image-generation.pdf
5. What do we learn from inverting CLIP models? - arXiv, accessed June 17, 2025, https://arxiv.org/html/2403.02580v1
6. What do we learn from inverting CLIP models? - NSF-PAR, accessed June 17, 2025, https://par.nsf.gov/servlets/purl/10548407
7. (PDF) Decoding emotional responses to AI-generated architectural ..., accessed June 17, 2025, https://www.researchgate.net/publication/378969334_Decoding_emotional_responses_to_AI-generated_architectural_imagery
8. Decoding emotional responses to AI-generated architectural imagery - Frontiers, accessed June 17, 2025, https://www.frontiersin.org/journals/psychology/articles/10.3389/fpsyg.2024.1348083/full
9. Generative models and their latent space - The Academic, accessed June 17, 2025, https://theacademic.com/generative-models-and-their-latent-space/
10. Generative modelling in latent space – Sander Dieleman, accessed June 17, 2025, https://sander.ai/2025/04/15/latents.html
11. Latent Space Cartography: Visual Analysis of Vector Space ..., accessed June 17, 2025, https://idl.cs.washington.edu/files/2019-LatentSpaceCartography-EuroVis.pdf
12. Exploring the latent space of diffusion models directly through singular value decomposition, accessed June 17, 2025, https://arxiv.org/html/2502.02225v1
13. EmoDiffusion: Enhancing Emotional 3D Facial Animation with Latent ..., accessed June 17, 2025, https://www.arxiv.org/abs/2503.11028
14. Controllable Expressive 3D Facial Animation via Diffusion in a Unified Multimodal Space ⋆ Corresponding authors. This work was supported by NSFC under 62272456. We extend our special thanks to Ruoyu Chen for helping revise Figures 1, the Introduction, and the Abstract. - arXiv, accessed June 17, 2025, https://arxiv.org/html/2506.10007v1
15. ReactDiff: Latent Diffusion for Facial Reaction Generation - arXiv, accessed June 17, 2025, https://arxiv.org/html/2505.14151v2
16. A Latent Space Diffusion Item Response Theory Model to Explore Conditional Dependence between Responses and Response Times, accessed June 17, 2025,

https://www.cambridge.org/core/services/aop-cambridge-core/content/view/5F3F99850F95787F2CEAB8B9D7E681F6/S0033312300007079a.pdf/a-latent-space-diffusion-item-response-theory-model-to-explore-conditional-dependence-between-responses-and-response-times.pdf
17. A Latent Space Diffusion Item Response Theory Model to Explore Conditional Dependence between Responses and Response Times | Psychometrika - Cambridge University Press, accessed June 17, 2025, https://www.cambridge.org/core/journals/psychometrika/article/latent-space-diffusion-item-response-theory-model-to-explore-conditional-dependence-between-responses-and-response-times/5F3F99850F95787F2CEAB8B9D7E681F6
18. A new paper demonstrates that LLMs could "think" in latent space, effectively decoupling internal reasoning from visible context tokens. This breakthrough suggests that even smaller models can achieve remarkable performance without relying on extensive context windows. : r/LocalLLaMA - Reddit, accessed June 17, 2025, https://www.reddit.com/r/LocalLLaMA/comments/1inch7r/a_new_paper_demonstrates_that_llms_could_think_in/
19. How Generative AI Endangers Cultural Narratives, accessed June 17, 2025, https://www.uib.no/en/cdn/167570/how-generative-ai-endangers-cultural-narratives
20. Love Through AI's Eyes: Emotions, Art, and Philosophy - AI Blog, accessed June 17, 2025, https://www.artificial-intelligence.blog/ai-news/love
21. Interpretable aesthetic features for affective image classification ..., accessed June 17, 2025, https://www.researchgate.net/publication/271555172_Interpretable_aesthetic_features_for_affective_image_classification
22. Cultural Archetypes Research Articles - R Discovery, accessed June 17, 2025, https://discovery.researcher.life/topic/cultural-archetypes/12884768?page=1
23. LAION-5B, Stable Diffusion 1.5, and the Original Sin of Generative AI ..., accessed June 17, 2025, https://www.techpolicy.press/laion5b-stable-diffusion-and-the-original-sin-of-generative-ai/
24. Full article: A taxonomy of prompt modifiers for text-to-image ..., accessed June 17, 2025, https://www.tandfonline.com/doi/full/10.1080/0144929X.2023.2286532
25. Imagining from Images with an AI Storytelling Tool - ResearchGate, accessed June 17, 2025, https://www.researchgate.net/publication/383279707_Imagining_from_Images_with_an_AI_Storytelling_Tool
26. Learning Conditional Space-Time Prompt Distributions for Video Class-Incremental Learning - CVF Open Access, accessed June 17, 2025, https://openaccess.thecvf.com/content/CVPR2025/papers/Zou_Learning_Conditional_Space-Time_Prompt_Distributions_for_Video_Class-Incremental_Learning_CVPR_2025_paper.pdf
27. Video generation - Hugging Face, accessed June 17, 2025, https://huggingface.co/docs/diffusers/using-diffusers/text-img2vid

28. Mastering Text-to-Image Diffusion: Recaptioning, Planning, and Generating with Multimodal LLMs - arXiv, accessed June 17, 2025, https://arxiv.org/html/2401.11708v1
29. Mastering Text-to-Image Diffusion: Recaptioning ... - OpenReview, accessed June 17, 2025, https://openreview.net/pdf/3b0f982aa81ac55b2f8880173bfba1117f233dec.pdf
30. How to Generate Google Veo 3 Prompt Theory Videos (Google Veo 3 Prompt Guide), accessed June 17, 2025, https://apidog.com/blog/google-veo-3-prompt-theory/
31. How Color and Light Affect Our Emotions in Building Spaces | illustrarch, accessed June 17, 2025, https://illustrarch.com/articles/45737-how-color-and-light-effect-what-we-feel-in-buildings.html
32. Story-boards.ai, accessed June 17, 2025, https://www.story-boards.ai/
33. Transform Your Ideas into Visual Stories with Storyboard AI, accessed June 17, 2025, https://www.katalist.ai/
34. Micro-expressions in Animated Agents, accessed June 17, 2025, https://par.nsf.gov/servlets/purl/10340819
35. Investigating Macro- and Microexpressions in Computer Graphics Animated Faces - University of Pennsylvania, accessed June 17, 2025, https://repository.upenn.edu/bitstreams/21910c83-405d-4d10-a86f-2a61bb0a2969/download
36. Perception of Micro-Expressions in Animated Characters with Different Visual Styles, accessed June 17, 2025, https://docs.lib.purdue.edu/dissertations/AAI30505312/
37. Fathers: Heroes, Villains, and Our Need for Archetypes - Dale M. Kushner, accessed June 17, 2025, https://www.dalemkushner.com/blog/fathers-heroes-villains-and-our-need-for-archetypes
38. Can Jungian Archetypes be Evidence-Based? - - Taproot Therapy Collective, accessed June 17, 2025, https://gettherapybirmingham.com/can-jungian-archetypes-be-evidence-based/
39. arxiv.org, accessed June 17, 2025, https://arxiv.org/html/2505.21912v1
40. LAION-5B: An open large-scale dataset for training ... - OpenReview, accessed June 17, 2025, https://openreview.net/pdf?id=M3Y74vmsMcY
41. The Price of Emotion: Privacy, Manipulation, and Bias in Emotional AI, accessed June 17, 2025, https://www.americanbar.org/groups/business_law/resources/business-law-today/2024-september/price-emotion-privacy-manipulation-bias-emotional-ai/
42. Not Feeling It: AI's Emotional Disconnect - USC Viterbi | School of Engineering, accessed June 17, 2025, https://viterbischool.usc.edu/news/2024/08/not-feeling-it-ais-emotional-disconnect/
43. Capacity of Generative AI to Interpret Human Emotions From Visual ..., accessed June 17, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC10879976/

44. (PDF) Affective Computing in the Era of Large Language Models: A ..., accessed June 17, 2025, https://www.researchgate.net/publication/383037310_Affective_Computing_in_the_Era_of_Large_Language_Models_A_Survey_from_the_NLP_Perspective
45. Art in an age of artificial intelligence - Frontiers, accessed June 17, 2025, https://www.frontiersin.org/journals/psychology/articles/10.3389/fpsyg.2022.1024449/full
46. Assessing Emotion and Sensitivity of AI Artwork - Frontiers, accessed June 17, 2025, https://www.frontiersin.org/journals/psychology/articles/10.3389/fpsyg.2022.879088/full
47. Full article: What AI "art" can teach us about art, accessed June 17, 2025, https://www.tandfonline.com/doi/full/10.1080/20004214.2024.2395511
48. (PDF) Agents, Simulation, and Technology: A Philosophical ..., accessed June 17, 2025, https://www.researchgate.net/publication/378529022_Agents_Simulation_and_Technology_A_Philosophical_Interpretation_of_AI_Literary_and_Artistic_Production
49. Aesthetics in the Philosophy of AI Art and the Evolution of Artificial ..., accessed June 17, 2025, https://www.airbrush.ai/philosophy-of-ai-art-and-the-evolution-of-ai/
50. The Philosophical Conundrum of AI Art and Human Existence ..., accessed June 17, 2025, https://ludo.ai/blog/philosophical-conundrum-ai-art-human-existence
51. Facts and Stats on AI Photography: How AI is Reshaping Visual ..., accessed June 17, 2025, https://aestheticsofphotography.com/facts-and-stats-on-ai-photography-how-ai-is-reshaping-visual-creation/
52. The Impact of AI on Traditional Photography Techniques: Revolutionizing Image Capture and Processing, accessed June 17, 2025, https://proedu.com/blogs/photography-fundamentals/the-impact-of-ai-on-traditional-photography-techniques-revolutionizing-image-capture-and-processing
53. HDR-GS: Efficient High Dynamic Range Novel View Synthesis at 1000x Speed via Gaussian Splatting - NIPS, accessed June 17, 2025, https://proceedings.neurips.cc/paper_files/paper/2024/file/7e83fd2ff7ae58485d418685521c9608-Paper-Conference.pdf
54. Perceptual Assessment and Optimization of HDR Image Rendering - CVF Open Access, accessed June 17, 2025, https://openaccess.thecvf.com/content/CVPR2024/papers/Cao_Perceptual_Assessment_and_Optimization_of_HDR_Image_Rendering_CVPR_2024_paper.pdf
55. A Perceptual Framework for Contrast Processing of High Dynamic Range Images - CiteSeerX, accessed June 17, 2025, https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=4856a545f2957687475ec67c844c4c080bd9e219
56. The Dynamic Range of Human Lightness Perception - PMC, accessed June 17, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC3244211/
57. How Photography Impacts The Psychology of Attention and Visual Processing -

Neuroscience Of, accessed June 17, 2025, https://www.neuroscienceof.com/human-nature-blog/neuroscience-photography-attention-psychology
58. The Art of Depth of Field in Storytelling - Number Analytics, accessed June 17, 2025, https://www.numberanalytics.com/blog/the-art-of-depth-of-field-in-storytelling
59. Understanding Depth of Field in Photography, accessed June 17, 2025, https://www.vernonchalmers.photography/2025/05/understanding-depth-of-field-in.html
60. Depth of Field and Your Emotions: Why Soft and Sharp Focus ..., accessed June 17, 2025, https://photovideocreative.com/en/depth-of-field-and-emotions-why-soft-and-sharp-focus-reinforce-photo/
61. Path from Photorealism to Perceptual Realism - Department of Computer Science and Technology |, accessed June 17, 2025, https://www.cl.cam.ac.uk/~fz261/thesis/Dissertation__Path_from_Photorealism_to_Perceptual_Realism_compressed.pdf
62. Path from Photorealism to Perceptual Realism - Apollo, accessed June 17, 2025, https://www.repository.cam.ac.uk/items/b4116c6a-ebdc-44a9-8e6f-cbc746fe8d5e
63. Three Varieties of Realism in Computer Graphics, accessed June 17, 2025, https://www.graphics.cornell.edu/pubs/2003/Fer03.pdf
64. The Role of Perception for Computer Graphics - Eurographics, accessed June 17, 2025, https://diglib.eg.org/server/api/core/bitstreams/097f365f-799f-4449-ab69-2782c7104b60/content
65. arxiv.org, accessed June 17, 2025, https://arxiv.org/html/2504.13214v1
66. Wavelet-Based Texture Reformation Network for Image Super ..., accessed June 17, 2025, https://www.researchgate.net/publication/359405000_Wavelet-Based_Texture_Reformation_Network_for_Image_Super-Resolution
67. Spherical Wavelets: Texture Processing - GitHub Pages, accessed June 17, 2025, https://cm-bell-labs.github.io/who/wim/papers/texture.pdf
68. Wavelet-Based Dual Recursive Network for Image Super-Resolution | Request PDF - ResearchGate, accessed June 17, 2025, https://www.researchgate.net/publication/346435019_Wavelet-Based_Dual_Recursive_Network_for_Image_Super-Resolution
69. HapticVLM: VLM-Driven Texture Recognition Aimed at Intelligent Haptic Interaction - arXiv, accessed June 17, 2025, https://arxiv.org/html/2505.02569v1
70. Haptic Texture Database with Variety in Velocity and Direction of Sliding Contacts - arXiv, accessed June 17, 2025, https://arxiv.org/html/2407.16206v1
71. Haptics | Media Expression and Communication Class Notes ..., accessed June 17, 2025, https://library.fiveable.me/media-expression-and-communication/unit-8/haptics/study-guide/QSkLRQN1l1nUogRS

72. Haptics Device Creates Realistic Virtual Textures - USC Viterbi ..., accessed June 17, 2025, https://viterbischool.usc.edu/news/2022/05/virtual-textures-i-cant-believe-its-not-wood/
73. Controllable Visual-Tactile Synthesis - CVF Open Access, accessed June 17, 2025, https://openaccess.thecvf.com/content/ICCV2023/papers/Gao_Controllable_Visual-Tactile_Synthesis_ICCV_2023_paper.pdf
74. Why Micro-Expressions Are the Missing Link for AI Avatars, accessed June 17, 2025, https://pettauer.net/en/ai-avatars-and-micro-expressions/
75. How real photography captures human emotion ... - Tech For Good, accessed June 17, 2025, https://www.techforgood.net/guestposts/how-real-photography-captures-human-emotion-and-how-ai-doesnt
76. The Future of Portrait Photography: AI Innovations Reshaping Visual Artistry - PRO EDU, accessed June 17, 2025, https://proedu.com/blogs/photoshop-skills/the-future-of-portrait-photography-ai-innovations-reshaping-visual-artistry
77. (PDF) The Effects of Depth of Field on Subjective Evaluation of Aesthetic Appeal and Image Quality of Photographs - ResearchGate, accessed June 17, 2025, https://www.researchgate.net/publication/338592901_The_effects_of_depth_of_field_on_subjective_evaluation_of_aesthetic_appeal_and_image_quality_of_photographs
78. Wavelet Based Texture Resampling Silviu Borac+ Eugene Fiume+* +Alias|Wavefront and *University of Toronto1 Abstract The integral, accessed June 17, 2025, https://www.dgp.toronto.edu/~elf/.misc/wavelet-based-texture-resampling.pdf
79. Statistical Learning of Facial Expressions Improves ... - Frontiers, accessed June 17, 2025, https://www.frontiersin.org/journals/virtual-reality/articles/10.3389/frvir.2021.619811/full