AI for Social Good: Architecting a Just and Sustainable Future

Part I: The Promise and Pillars of Algorithmic Benevolence

The discourse surrounding artificial intelligence (AI) is increasingly framed by a dual narrative of unprecedented opportunity and existential risk. Within this landscape, a powerful and influential movement known as "AI for Good" (AI4Good) has emerged, positing that AI, if intentionally directed, can serve as a transformative force for positive societal change.[1] This movement seeks to harness the immense computational power of AI not merely for commercial gain but to address the world's most intractable problems, from climate change and disease to poverty and inequality. It represents a global effort to align the trajectory of technological development with the collective aspirations of humanity, most notably encapsulated in the United Nations' 17 Sustainable Development Goals (SDGs).[2] This initial part of the report will define the AI4Good imperative, map its institutional ecosystem, and explore its foundational pillars before examining concrete case studies where this promise is being put into practice.

Section 1: Defining the AI for Good (AI4Good) Imperative

The AI for Good movement is predicated on the intentional application of artificial intelligence technologies to solve pressing social, environmental, and economic challenges, with a deliberate focus on promoting positive outcomes for humanity and improving societal well-being.[1] This endeavor is not a fringe activity but a rapidly institutionalizing field, driven by a global coalition of governments, international organizations, technology corporations, academic institutions, and non-profits.[1] The central premise is that AI's capacity to analyze vast datasets, identify complex patterns, and automate decision-making can be a powerful lever for accelerating

progress on global development targets.

This imperative has gained significant traction against a backdrop of global urgency. A 2023 special edition of the UN's Sustainable Development Goals Report delivered a stark assessment: at the midpoint of the 2030 Agenda, progress has stalled or regressed on 80% of the SDG targets, with only 15% on track.[5] This implementation gap, exacerbated by a confluence of crises including the climate emergency, geopolitical conflict, and the lingering effects of the COVID-19 pandemic, has created a powerful narrative impetus for new solutions.[6] In this context, AI is positioned not as a panacea, but as a critical catalyst with the potential to enable progress on as many as 79% of the 169 SDG targets, breathing new life into a flagging global agenda.[3]

The operationalization of this vision rests on three key pillars that are consistently cited by proponents of the movement: inclusivity, sustainability, and ethical governance.[1]

The Three Pillars of AI4Good

Inclusivity: At its core, the inclusivity pillar seeks to ensure that the benefits of AI are distributed equitably and that the technology serves to bridge, rather than widen, the digital divide. The goal is to democratize access to technology and information, thereby reducing systemic inequalities in critical areas such as education, healthcare, and economic opportunity.[1] This involves designing AI tools that are accessible to all, including individuals with disabilities or those with limited resources, and localizing them to specific cultural and socioeconomic contexts.[1] A key component of inclusivity is capacity building—proactively developing AI skills and literacy in marginalized and underserved communities that are often at the greatest risk of being left behind by technological disruption.[1]

Sustainability: The sustainability pillar explicitly aligns AI with the global effort to create an environmentally sustainable future. This involves leveraging AI to address a spectrum of ecological challenges, including climate change mitigation, deforestation, and wildlife conservation.[1] Applications in this domain include environmental monitoring, where AI is used to track and predict environmental changes like air pollution and biodiversity loss, enabling more proactive policymaking. It also encompasses smart agriculture, where AI optimizes farming practices to enhance crop yields while minimizing water and energy consumption, and disaster prediction,

where AI models analyze vast datasets to forecast natural disasters like floods and hurricanes, improving response times and saving lives.[1]

Ethical Governance: Recognizing the profound ethical questions raised by the rapid proliferation of AI, this pillar prioritizes the responsible and ethical deployment of the technology. The primary objective is to ensure that AI benefits society without exacerbating existing inequalities or introducing new, unforeseen risks.[1] This entails a focus on several key areas. First is mitigating bias and promoting fairness, ensuring that algorithms do not perpetuate discrimination based on race, gender, or socioeconomic status. Second is transparency, which demands that the decision-making processes of AI systems are understandable to stakeholders. Third is accountability, which holds the developers and deployers of AI systems responsible for their societal impacts, including any potential harms.[1]

The Global Ecosystem

The AI4Good movement is championed by a powerful and interconnected global ecosystem of actors. At the forefront are international bodies like the United Nations, which provides the overarching ethical and developmental framework through the SDGs. The UN's specialized agency for information and communication technologies, the International Telecommunication Union (ITU), has become a central convener, organizing the annual AI for Good Global Summit in Geneva.[4] This summit brings together over 40 UN agencies, governments, and leaders from technology and civil society to identify innovative AI applications and advance partnerships to solve global challenges.[4]

Major technology corporations are also key players, launching significant initiatives that provide the funding, computational resources, and technical expertise to power the movement. Microsoft's AI for Good Lab, established in 2018, is a prominent example, focusing on using AI to tackle challenges in sustainability, humanitarian action, and health, often in alignment with the SDGs.[8] Similarly, Google.org launched a $25 million open call for organizations using AI to accelerate progress on the Global Goals, building on years of support for AI-powered social impact projects.[10] These corporate initiatives are complemented by dedicated non-profits like the AI for Good Foundation, which deploys AI-driven solutions in areas like humanitarian aid and refugee assistance.[11]

This convergence of actors—international organizations providing the platform and normative framework, and technology companies providing the capital and technical solutions—creates a powerful synergy. However, it also reveals a more complex dynamic at play. The "AI for Good" movement is not just a philanthropic endeavor; it functions as a potent geopolitical and corporate narrative. For international bodies like the ITU, it is a means to assert relevance and a guiding role in the governance of a disruptive new technology. For technology giants, it provides a "social license to operate," allowing them to position themselves as benevolent and indispensable partners in solving global problems. This framing helps shape global governance norms in a manner that is favorable to their interests, often emphasizing voluntary, multi-stakeholder approaches over binding regulation. The consistent positioning of AI as the primary solution to the lagging progress on the SDGs creates a sense of technological inevitability, forging a symbiotic relationship where tech companies gain legitimacy and influence over global policy, and international bodies reinforce their role as conveners. The "good" in AI4Good thus becomes a vehicle for soft power, market positioning, and the shaping of future governance landscapes.

Section 2: AI in Action: Case Studies in Global Problem-Solving

Moving from the conceptual framework to tangible impact, a growing portfolio of case studies demonstrates how AI is being applied to address specific global challenges. These initiatives, often born from the collaborative ecosystem described above, provide a grounded view of AI's potential across various domains aligned with the Sustainable Development Goals. They illustrate not only the power of the technology but also a recurring pattern in its most effective applications.

Climate & Environment (SDG 13)

The fight against climate change and environmental degradation has become a flagship area for AI4Good applications. One of the most prominent examples is **Google's Flood Forecasting Initiative**. Launched as a pilot in India in 2018, the program has since expanded to provide flood warnings in over 100 countries, covering more than 700 million people globally by 2024.[2] A significant challenge in flood prediction is the lack of high-quality, long-term hydrological data in many

developing nations, which are often the most vulnerable. Traditional models require extensive calibration and are often ineffective in these data-scarce regions. Google's approach circumvents this limitation by using advanced machine learning to generate accurate predictions even in unmeasured locations. This is particularly critical given that 90% of the 1.8 billion people at high risk of flood damage live in low- and middle-income countries, where early warning systems can reduce flood-related fatalities by as much as 43%.[2]

Similarly, **Microsoft's AI for Earth** program and its AI for Good Lab have spawned numerous projects targeting environmental sustainability. These initiatives leverage satellite imagery and machine learning to tackle complex problems at scale. In one project, AI is used to monitor and mitigate deforestation in the Amazon by analyzing satellite data to detect illegal logging and mining activities in near real-time.[8] Other projects focus on biodiversity, using AI to map and analyze species distribution to aid conservation efforts, and on renewable energy, where AI models are developed to forecast solar energy production, helping to optimize grid management and advance the transition to clean energy.[8]

Health & Well-being (SDG 3)

In global health, AI is being deployed to augment the capabilities of healthcare systems, especially in areas with limited resources. **RAD-AID**, a recipient of a Google.org grant, is using AI to address the critical shortage of radiologists in low- and middle-income countries.[10] Millions of patients undergo X-rays and other scans but lack access to qualified professionals to interpret them. RAD-AID supports low-resource hospitals by using AI to triage and provide preliminary interpretations of these scans, with a focus on high-impact areas like respiratory disease and breast cancer. The initiative aims to improve diagnostic technology and patient outcomes for over 30 million people across 20 countries.[10]

Beyond specialized diagnostics, AI is also enhancing primary care access. AI-powered diagnostic tools are being developed to detect conditions like cancer and cardiovascular diseases at earlier, more treatable stages.[1] Furthermore, AI-driven telemedicine platforms are helping to bridge the healthcare gap in underserved communities. These platforms can provide remote medical consultations, monitor chronic conditions, and deliver health information, connecting patients with care

regardless of their geographical location.[1]

Agriculture & Food Security (SDG 2)

To combat food insecurity and support rural livelihoods, AI is being used to deliver timely and localized agricultural knowledge. A compelling case study is the **Ulangizi chatbot**, developed by Opportunity International for farmers in rural Malawi.[2] Ulangizi, which translates to 'Advice,' is a generative AI tool that communicates in Chichewa, the local language. It provides farmers with tailored guidance on critical topics like crop management, irrigation techniques, and climate adaptation strategies. By making expert knowledge accessible through a simple, language-appropriate interface, the chatbot empowers farmers to make more informed decisions, leading to improved crop yields and enhanced food security. This model is significantly more scalable and cost-effective than traditional agricultural extension services, which often struggle to reach remote communities.[2]

Humanitarian Aid & Justice (SDG 16)

In the realm of humanitarian response and justice, AI is being employed to streamline aid delivery and empower crisis-affected populations. The AI for Good Foundation's **LifeForce Ukraine** platform serves as a powerful example. It functions as Ukraine's largest humanitarian aid matching system, providing civilians with real-time access to essential resources such as food, shelter, and medicine across a network of over 37,000 service locations.[11]

Another innovative application is the International Rescue Committee's (IRC) **Signpost Project**. Recognizing that information is a critical form of aid for displaced populations, Signpost uses AI-powered chatbots to provide reliable, 24/7 answers to common questions about services, legal rights, and safety.[2] A crucial element of this project is the human-in-the-loop design. While the AI handles the initial interaction, local human moderators review the responses to ensure they are accurate, culturally sensitive, and contextually appropriate. This hybrid model combines the scalability of AI with the nuance and trust of human oversight.[2]

A close examination of these successful case studies reveals a consistent underlying pattern. The most impactful AI4Good applications are not autonomous, all-knowing systems that parachute in to solve problems independently. Rather, they function as **"last-mile" augmentation tools**. Their success is not derived from replacing human expertise but from amplifying it and overcoming specific, well-defined bottlenecks. The Ulangizi chatbot succeeds because it solves a communication and accessibility bottleneck, translating expert agricultural knowledge into a local language and delivering it via a common technology. Google's flood model is valuable precisely because it addresses a data bottleneck, using machine learning to generate forecasts where historical data is absent. RAD-AID and AI-powered telemedicine tools tackle a human resource bottleneck, allowing a limited number of medical professionals to extend their reach and serve a larger population. This pattern challenges the techno-solutionist fantasy of a purely automated benevolence. It underscores that the true value of AI in a social good context lies in its ability to be deeply integrated into existing human and social infrastructures, acting as a connector, a translator, and a data interpolator within a collaborative, human-centric system.

Part II: Confronting the Algorithmic Shadow

While the promise of AI for Good is compelling and its applications are genuinely impactful, an exclusive focus on this optimistic narrative would be incomplete and dangerously naive. Beneath the surface of algorithmic benevolence lies a darker, often-hidden "algorithmic shadow"—a set of systemic costs and harms that are not merely unfortunate side effects but are deeply embedded in the current architecture of AI development and deployment. To architect a truly just and sustainable future with AI, it is imperative to confront this shadow directly. This part of the report will dissect the material shadow of AI—its staggering energy consumption and e-waste generation—and its socio-ethical shadow, characterized by the perpetuation of digital colonialism and epistemic injustice.

Section 3: The Material Shadow: Energy, E-Waste, and the AI Supply Chain

The ethereal nature of "the cloud" and the abstract quality of algorithms mask a

profoundly material reality. The AI revolution is built upon a physical infrastructure of unprecedented scale, with a voracious appetite for energy and a rapid cycle of hardware production and disposal that casts a long environmental shadow over its purported green credentials.

The Voracious Energy Appetite

The computational demands of training and running large-scale AI models are immense and growing at an exponential rate. The physical hubs of AI, massive data centers, are on a trajectory to more than double their global electricity demand by 2030, reaching an estimated 945 terawatt-hours. This figure is slightly more than the entire current electricity consumption of Japan.[12] The scale of this energy use is difficult to comprehend without concrete examples. The training of a single large language model, GPT-3, was estimated to consume 1,287 megawatt-hours of electricity, enough to power an average U.S. household for 120 years.[13] Another analysis highlighted in the 2023 Stanford AI Index Report found that the training run for the BLOOM model emitted 25 times more carbon than a single passenger on a one-way flight from New York to San Francisco.[14] This massive energy consumption directly contributes to greenhouse gas emissions, particularly when data centers are powered by fossil fuel-based energy grids.

The Rising Tide of E-Waste

The relentless pace of innovation in AI hardware, particularly in specialized chips like GPUs (Graphical Processing Units), creates a powerful dynamic of accelerated obsolescence. As newer, more powerful models are developed, the hardware required to train and run them efficiently changes, rendering previous generations of equipment less viable. This rapid turnover contributes significantly to the global electronic waste (e-waste) crisis, which is already one of the fastest-growing waste streams worldwide.[15] E-waste is not benign; it contains a cocktail of valuable precious metals like gold and copper alongside toxic substances such as lead and mercury, which pose severe environmental and public health risks if not managed properly.[15] The material footprint of AI, therefore, extends from the energy consumed during its

operation to the physical waste generated at the end of its hardware's lifecycle.

Architectural Solutions for Material Sustainability

Addressing this material shadow requires a multi-pronged approach that re-architects AI systems from the level of fundamental physics to the level of economic policy.

Hardware and Materials Science: A promising frontier of innovation lies in creating new materials and hardware architectures that are fundamentally more energy-efficient. Researchers at Clemson University, for instance, have developed a new polymer material called pTPADTP. This material can be used to build a "memristor," a novel electronic component that can both store and process information in the same physical location.[12] In classical computer architecture, data must be constantly shuttled back and forth between separate memory and processing units, a process that consumes significant time and energy. By eliminating this data-shuttling bottleneck, memristor-based chips could dramatically reduce the energy required for AI computation.[12]

Software and Optimization: Alongside hardware innovation, sophisticated software optimization can yield substantial energy savings. A team at the University of Michigan developed an open-source framework called **Zeus**, which can reduce the energy footprint of AI model training by up to 75% with only minor impacts on training time.[13] Zeus works by monitoring the training process in real-time and dynamically tuning two key software parameters: the GPU's power limit and the model's "batch size" (how many data samples it processes at once). It finds the optimal trade-off point where energy use is minimized for a given training speed.[13] The same team developed a complementary software called

Chase, which introduces carbon-aware computing. Chase layers on top of Zeus and makes decisions based on the real-time carbon intensity of the energy grid. When low-carbon energy (like solar or wind) is abundant, it prioritizes training speed. During peak demand times, when the grid is more likely to rely on carbon-intensive sources like coal, it prioritizes energy efficiency at the expense of speed. This intelligent scheduling further reduces the actual carbon footprint of AI training.[13]

Circular Economy Models: Beyond optimizing individual systems, a broader shift

towards a circular economy for AI hardware is necessary. This begins with principles of sustainable design, including maximizing energy efficiency, selecting recyclable or lower-impact materials, and embracing modularity, which allows for easier upgrades and repairs, extending the lifespan of devices and reducing the need for complete replacements.[16] This design philosophy must be coupled with new business models. Companies like Dell have implemented closed-loop recycling programs, where old products are collected, refurbished, and resold, or their materials are recycled into new products. HP's "hardware-as-a-service" model, where customers lease devices and return them at the end of their life, further incentivizes reuse and responsible material management.[16] Paradoxically, AI itself can be a powerful tool in this process. AI-driven vision systems can be used to automate the sorting of complex e-waste streams, improving the efficiency and purity of material recovery and making recycling more economically viable.[15]

Policy Levers: A Carbon Tax for Computation

While technological and business model innovations are crucial, their voluntary adoption may be insufficient to counteract the exponential growth in AI's energy demand. A more systemic intervention is required to internalize the environmental costs of computation. A **carbon tax for computation** is a powerful policy lever that could achieve this. Under such a policy, a price is set that emitters must pay for each ton of greenhouse gas emissions they produce.[17]

In the context of AI, this tax could be levied "upstream" at the point of energy consumption—on the data centers or the electric utilities that power them—which is the most administratively simple approach.[17] The tax rate could be tied to the social cost of carbon and designed to escalate over time, providing a clear and predictable signal to the market that investments in energy-efficient hardware, optimization software like Zeus, and renewable energy sources are not just ethically desirable but economically rational.[17] The revenue generated from such a tax could be used to fund further research and development in green AI, to invest in renewable energy infrastructure for data centers, or to be returned to the public as a "carbon dividend" to offset any potential increases in the cost of digital services.[17]

The push to develop and deploy AI for sustainability solutions, such as complex climate models or AI-driven carbon credit validation systems, creates a fundamental **"Sustainability Paradox."** These beneficial applications are themselves

computationally intensive, and their development and operation contribute to the very environmental problems they aim to solve. The existence of solutions like Zeus and new energy-efficient materials is an implicit admission that the industry's default trajectory is unsustainable. Without a mandatory, systemic framework like a carbon tax for computation, the adoption of these more sustainable practices will remain voluntary, patchy, and likely insufficient to offset the explosive growth in AI's energy and material footprint. This creates a serious risk that the "AI for Sustainability" sector could become a net-negative for the environment—a form of greenwashing that promotes its beneficial applications while offloading its massive and growing environmental debt. A comprehensive, lifecycle-based accounting, enforced by robust policy, is essential to ensure that the promise of AI for environmental good does not become a mirage.

Section 4: The Socio-Ethical Shadow: Digital Colonialism and Epistemic Injustice

The material shadow of AI is paralleled by an equally pervasive and damaging socio-ethical shadow. The current architecture of AI development, dominated by a small number of actors from the Global North, is perpetuating historical patterns of power and control, leading to new forms of colonialism and injustice. These are not isolated instances of "bias" to be patched, but systemic harms that challenge the very foundation of the AI4Good narrative.

Digital Colonialism

A growing body of scholarship critiques the role of major technology companies as new colonizing forces.[19] The concept of

digital colonialism describes how a handful of (primarily U.S.-based) corporations design and deploy digital infrastructures to fit their own economic interests, establishing a global dominance that creates deep societal dependency on their products.[19] This manifests in several concrete and harmful ways:

- **Linguistic and Cultural Imperialism:** The digital world is overwhelmingly Anglophone and Western-centric. This is starkly illustrated by the fact that as of May 2024, the world's leading AI assistants—Apple's Siri, Google Assistant, and

Amazon's Alexa—collectively provided native support for zero of the approximately 2,000 languages spoken on the African continent, which is home to roughly one-third of the world's languages.[19] This is not a neutral technical oversight; it is a form of cultural erasure that reinforces the dominance of one worldview and marginalizes countless others, reflecting the biases of the small, predominantly white, male, and American group of software engineers who design these systems.[19]

- **Surveillance and Racialized Control:** AI-powered surveillance technologies are being deployed in ways that replicate and amplify historical injustices. In Johannesburg, South Africa, the private security company Vumacam has built a network of over 5,000 CCTV cameras that use AI analytics, including facial recognition technology known to be racially biased, to trigger security alerts.[19] Local activists argue that this system reinforces the profound racialized inequalities of the country's apartheid past, creating a modern apparatus of racialized surveillance.[19]
- **Extractive Benevolence:** A key tactic of digital colonialism is to mask extractive economic practices with a discourse of benevolence and philanthropy.[19] A prime example is Facebook's Internet.org initiative, later rebranded as Free Basics. It was promoted as a charitable project to "connect the world" and bring internet access to the Global South. However, leaked internal documents revealed that users in developing countries who were promised free data access were, in fact, being secretly charged for it, paying nearly $100 million in 2021 alone.[19] This practice exemplifies how the language of "doing good" can be used to obscure the extraction of wealth and data from vulnerable populations.

Aesthetic Injustice

The colonial gaze is also being automated and scaled through AI, leading to a phenomenon that can be termed **aesthetic injustice**.[21] This refers to the way AI systems perpetuate and globalize a narrow and biased set of standards, particularly around human beauty. When AI models are trained on vast, uncurated datasets scraped from the internet, they inevitably learn the biases present in that data. In the context of aesthetics, this often means absorbing and reinforcing a dominant, Eurocentric ideal of beauty.[22]

AI applications used for facial analysis, photo editing, and beauty scoring quantify and standardize beauty based on these biased datasets. They learn to associate "beauty"

with specific features, often marginalizing and devaluing the diverse forms of beauty found across different cultures and ethnicities.[22] This process is not neutral; it actively promotes a hegemonic standard, which can lead to what researchers have described as "appearance-based stress" and can put the preservation of shared cultural differences at risk.[21] The subjective, culturally rich, and historically diverse human understanding of beauty is replaced by a rigid, quantified, and biased algorithmic judgment.

The Need for Epistemic Reparations

The harms of digital colonialism and aesthetic injustice are not merely technical or economic; they are profoundly epistemic. They involve the silencing of voices, the erasure of histories, and the invalidation of entire ways of knowing and being. In response to such deep-seated harms, a new framework of justice is required: **epistemic reparations**.[23]

Epistemic reparations move beyond traditional forms of redress, like monetary compensation, to focus on the "right to be known".[23] This concept, developed through work with survivors of carceral injustice, sexual violence, and other gross human rights violations, recognizes that a critical part of healing and justice is the opportunity for victims to tell their own stories, in their own words, and to have those stories heard and acknowledged by the wider public.[23] It is a process of countering the dominant, often dehumanizing, narratives promoted by the state or other powerful actors. For example, for a person wrongfully convicted and vilified by the justice system, epistemic reparations involve creating platforms to share their true story, thereby restoring their identity and humanity.[23]

This framework is a necessary remedy for the epistemic violence perpetrated by biased AI systems. When an AI system fails to recognize a language, devalues a cultural expression of beauty, or reinforces a harmful stereotype, it is committing an epistemic wrong. It is implicitly declaring that a particular group's knowledge, culture, and experience are not valid.

The seemingly disparate problems of digital colonialism, aesthetic injustice, and the call for epistemic reparations are, in fact, deeply interconnected. They are all manifestations of a single, foundational harm: **ontological erasure**. The issue is not simply that AI systems produce biased outputs, such as an unfair loan denial or a

skewed hiring recommendation. The deeper harm is the implicit message that a person's language, their culture, their appearance, their history, and their entire way of life are not valid enough to be included in the new digital reality being constructed. This is an assault on their ontology—their very world. This reframing of the problem demonstrates that the solution cannot be merely to "fix the bias" in existing systems. Such a corrective approach is insufficient because the systems themselves are built on a flawed, universalist foundation. Instead, what is required is a fundamental architectural shift towards resisting this ontological capture, demanding solutions that are not just corrective but are profoundly and structurally pluralistic.

Part III: Architecting a New Foundation for AI

The critique of AI's material and socio-ethical shadows necessitates a radical rethinking of its foundational architecture. Merely patching flaws or mitigating harms on a case-by-case basis is insufficient. What is required is a move from a reactive to a proactive stance, embedding principles of accountability, relationality, and human-centricity into the very design of AI systems. This part of the report proposes a multi-layered architectural framework for building this new foundation, encompassing technical, cognitive, and human-centric dimensions.

Section 5: Technical Architectures for Trust and Accountability

To build trust in AI systems, particularly in high-stakes domains, we must move beyond opaque "black box" models and construct technical frameworks that make accountability not just a principle, but a verifiable reality. This requires new architectures for auditing, governance, and impact assessment.

Verifiable AI and Attestable Audits

A critical step towards accountability is the development of **Verifiable AI**. This refers to AI systems designed from the ground up to be transparent, auditable, and

accountable, allowing users and regulators to trace, understand, and validate their decisions.[24] Verifiable AI is a broader concept that encompasses, but is not limited to, Explainable AI (XAI). While XAI focuses on making the reasoning behind a specific AI decision understandable, Verifiable AI provides a more robust guarantee by enabling the entire process to be audited and proven to be compliant with a given set of rules.[24]

A key technical implementation of this concept is the use of **Attestable Audits running in Trusted Execution Environments (TEEs)**.[26] A TEE is a secure, isolated area within a computer's processor that protects code and data from being accessed or tampered with, even by the system's owner or cloud provider. The Attestable Audit protocol leverages this technology to create a trustless verification system. In this protocol, a model provider (e.g., a tech company) and an auditor (e.g., a regulator or a civil society group) can interact without having to trust each other. The model provider's proprietary AI model and the auditor's confidential audit dataset are both sent in encrypted form to the TEE. Inside this secure enclave, the audit is executed, and the TEE produces a cryptographically signed "attestation"—an unforgeable proof that a specific, identified model was tested against a specific, identified audit dataset and achieved a specific result. This process protects the intellectual property of the model provider and the sensitive data of the auditor while producing a verifiable, trustworthy audit record that can be made public.[26] This architecture transforms compliance from a matter of corporate self-declaration to a mathematically verifiable fact.

Decentralized Governance with DAOs

While Verifiable AI provides a mechanism for technical accountability, it does not answer the question of who decides what rules the AI should be accountable *to*. To address the concentration of power that characterizes digital colonialism, new models of governance are needed. **Decentralized Autonomous Organizations (DAOs)** offer a promising architectural solution.[27] DAOs are organizations that operate on blockchain technology, with rules encoded in smart contracts and decisions made through voting by their members.[28]

In the context of AI governance, a DAO can create a framework for community-led oversight. For instance, a DAO could be formed by the stakeholders of a particular AI application—such as farmers using an agricultural AI or citizens subject to a municipal AI system. The DAO's members could vote on the ethical principles, fairness metrics,

and acceptable risk thresholds for the AI system. The AI itself can serve the DAO by providing data-driven analysis and recommendations to inform these decisions. However, the ultimate authority remains with the human members, who vote to approve or reject proposals, ensuring that the AI's operation remains aligned with the community's democratically determined values and preferences.[27] This model directly counters the top-down, opaque governance structures of today's tech giants, offering a blueprint for a more democratic and decentralized AI ecosystem.

Multi-Modal Impact Assessments

Before an AI system is even built or deployed, a structured process for anticipating its potential impacts is essential. This requires comprehensive **Ethical and Multi-Modal Impact Assessments**. These are not one-time checklists but iterative processes that guide the entire lifecycle of an AI system. Frameworks like **UNESCO's Ethical Impact Assessment (EIA)** provide a robust methodology for this.[29] The UNESCO EIA compels developers and procurers of AI systems to engage in a rigorous, multi-stage inquiry. It begins with scoping questions that force a justification for using AI in the first place, an exploration of non-algorithmic alternatives, and an analysis of the project's proportionality.[29] It then requires a detailed stakeholder engagement plan and a diversity assessment of the development team itself. The core of the assessment involves evaluating the proposed system against key ethical principles, including safety, fairness, non-discrimination, sustainability, privacy, human oversight, and transparency.[29]

Similarly, in high-stakes sectors like healthcare, the World Health Organization (WHO) has released specific guidance for the use of Large Multi-Modal Models (LMMs), which can process diverse data types like text, images, and video.[30] The WHO guidance calls for mandatory post-release auditing and impact assessments by independent third parties, with the results disaggregated by user demographics (age, race, etc.) to detect biases and inequitable outcomes.[30] These structured assessments [33] shift the burden from reacting to harms after they occur to proactively designing systems to be ethical, safe, and beneficial from their inception.

These advanced governance models—Verifiable AI, DAOs, and EIAs—are not mutually exclusive. In fact, their true power lies in their potential integration. The following table provides a comparative analysis of these frameworks, highlighting how they can be

combined into a cohesive system for accountability.

Governance Model	Primary Goal	Core Mechanism	Key Challenge Addressed	Primary Stakeholders	Potential Integration
Verifiable AI (via TEEs)	Technical Audibility & Confidentiality	Trusted Execution Environments (TEEs), Cryptographic Attestation	Black-Box Problem, IP Protection, Lack of Verifiable Compliance	Regulators, Auditors, Enterprise Customers	Provides the secure execution layer for audits mandated by other frameworks.
Decentralized Autonomous Orgs (DAOs)	Democratic Control & Decentralization	Blockchain, Smart Contracts, Community Voting	Centralized Power, Stakeholder Exclusion, Lack of Legitimacy	Community Members, Users, Distributed Stakeholders	Can serve as the governing body that defines the rules, parameters, and acceptable risk levels for an AI system, which are then enforced and verified by other models.
Ethical Impact Assessments (e.g., UNESCO EIA)	Procedural Rigor & Ethical Foresight	Structured Questionnaires, Stakeholder Consultation, "Ethics by Design"	Ad-hoc Ethical Considerations, Unforeseen Negative Impacts	Developers, Policymakers, Civil Society Organizations	Sets the substantive ethical criteria and questions that a Verifiable AI audit must check and that a DAO must consider in its governance

					decisions.
Pluriversal Alignment (via PKGs)	Ontological Respect & Cultural Resonance	Pluriversal Knowledge Graphs, Context Engineering	Epistemic Erasure, Digital Colonialism, Universalist Bias	Indigenous Groups, Culturally Diverse Communities, Linguists	Provides the underlying knowledge architecture that informs the data and models used in all other governance frameworks, ensuring they are not built on a single, biased worldview.

A synthesis of these models could create a powerful, closed-loop system of accountability. One can envision a future where a DAO, representing the community affected by an AI system, votes to adopt the UNESCO EIA as its guiding framework. The specific ethical parameters and audit questions defined by this process are then encoded into a smart contract. This contract triggers a periodic, attestable audit of the AI system, which is executed within a Verifiable AI framework using TEEs. The cryptographic proof of the audit's result is then published back to the DAO's blockchain, providing all members with a transparent and unforgeable record of the system's compliance. This integrated architecture combines democratic legitimacy (the DAO), procedural rigor (the EIA), and technical verifiability (the TEE), creating a far more robust foundation for trust than any single approach could achieve on its own.

Section 6: Cognitive Architectures for Relational Intelligence

Technical architectures for accountability are necessary but not sufficient. If the underlying cognitive architecture of the AI itself remains rooted in a narrow, universalist conception of intelligence, it will continue to perpetuate the harms of ontological erasure. To build AI that is truly for social good, we must move beyond the flawed project of "aligning" AI with a single set of human values and instead architect

AI that can fluently and respectfully operate within a multiplicity of worldviews.

Beyond Universalism: The Philosophy of Pluriversal AI

The dominant Western (or WEIRD—Western, Educated, Industrialized, Rich, and Democratic) view of intelligence conceives of it as an abstract, quantifiable, and disembodied attribute of an individual.[34] This view is deeply embedded in the design of current AI systems, which are optimized for abstract linguistic performance, decontextualized problem-solving, and scalability. However, this is a culturally specific and historically recent understanding of cognition.

A powerful philosophical alternative is found in the concept of **pluriversalism**. Drawing on Amerindian perspectivism and other non-Western cosmologies, pluriversalism challenges the very idea of a single, objective reality (a "universe") that is merely interpreted differently by various cultures. Instead, it posits a "pluriverse"—a world composed of many worlds, or many distinct natures and ontologies.[34] From this perspective, intelligence is not a static attribute but a relational, embodied, and culturally-grounded capacity for "becoming-with" others in a specific ecological and historical context. Language is not a neutral tool for describing a pre-existing reality; it is a world-making force, and each language and its associated cosmology constitute a different ontology.[34] The standard AI alignment project, which seeks to distill a universal set of "human values" and encode them into a single AI, is thus seen as a form of cognitive imperialism—an attempt to reduce the pluriverse to a monologue of dominant reason.[34]

From Theory to Practice: Decolonial Scaffolding and Context Engineering

Operationalizing this pluriversal philosophy in AI requires new methods for interacting with and guiding AI models. One such method is **Decolonial Prompt Scaffolding**. This technique adapts pedagogical principles of scaffolding—providing incremental, supportive structures to guide a learner through complex or uncomfortable territory—to the context of AI.[36] In a decolonial context, this means using a carefully designed series of prompts not just to get a better answer, but to guide the AI away from its default, often colonial patterns of thought. It involves prompting the AI to

reflect on its own biases, to consider alternative perspectives, to engage with "unsettling feelings," and to move beyond oppressive or simplistic patternings.[38]

This advanced form of interaction is part of a broader practice known as **Context Engineering 2.0**. While basic prompt engineering involves writing a simple instruction, context engineering is the discipline of architecting the entire context window provided to an AI model.[40] This includes not just the user's query but also the system prompt (high-level instructions), dynamic elements (like the current date or user location), few-shot examples (demonstrations of desired behavior), and, most importantly, retrieved knowledge from external sources. By precisely controlling this entire context, a developer can guide the model's behavior with a high degree of precision, steering it towards more nuanced, culturally aware, and self-reflective responses.[40]

The Pluriversal Knowledge Graph (PKG)

While context engineering provides the method to guide an AI, it needs a new kind of data structure to guide it *towards*. Standard knowledge graphs, which organize information about entities and their relationships, are a powerful tool for grounding AI in factual knowledge.[41] However, they are typically built on a single, universal ontology or schema, implicitly assuming one correct way to structure the world's knowledge.

To build a truly pluriversal AI, a new cognitive architecture is needed: the **Pluriversal Knowledge Graph (PKG)**. Unlike a standard KG, a PKG would be explicitly designed to hold multiple, distinct, and sometimes even contradictory ontologies and worldviews within a single structure. It would not attempt to merge or reconcile these different ways of knowing into a single, universal truth. Instead, it would preserve their integrity as separate, valid subgraphs.

The operation of an AI system built on a PKG would be fundamentally different from today's models. Guided by sophisticated context engineering, the AI's first step in any interaction would be to identify the relevant cultural or ontological context. Is the user operating from a Western scientific worldview? An Indigenous Amazonian cosmology? A specific legal framework? Once the context is identified, the AI would activate the corresponding subgraph of the PKG. All subsequent reasoning, information retrieval, and response generation would then occur *within* that specific ontological framework.

This architecture represents a paradigm shift. The AI is no longer a universal translator trying to map all knowledge onto a single plane of understanding—an act that inevitably involves distortion and erasure. Instead, it becomes an ontologically fluent navigator, capable of inhabiting and operating within different worlds as they are, on their own terms. This approach moves the goal of AI development from the impossible and colonial task of "universal value alignment" to the more challenging, but ultimately more respectful and just, goal of achieving "ontological fluency." It is only by building AI that can recognize and honor the pluriverse that we can hope to create technology that truly serves the good of all humanity, not just a dominant fraction of it.

Section 7: Human-Centric Architectures for Synergistic Collaboration

The ultimate purpose of architecting AI for social good is to enhance human flourishing. This has led to a vision of human-AI collaboration as a "synergistic zone" where the unique strengths of humans and machines combine to create a form of collaborative intelligence greater than the sum of its parts. However, this optimistic vision is being challenged by the rapid advancement of recursive AI, which threatens to automate the very meta-skills thought to be uniquely human. This necessitates a fundamental re-evaluation of the human role in the age of AI and a new framework for the skills that will define human value in a world of increasingly capable machines.

The Synergistic Zone of Human-AI Collaboration

The idealized model of human-AI collaboration posits a clear and complementary division of labor.[42] In this model, AI systems leverage their vast data processing capabilities, pattern recognition, and predictive analytics to handle repetitive, mundane, and computationally intensive tasks. This frees up human workers to focus on what are often described as uniquely human skills: providing deep contextual understanding, exercising ethical judgment, applying creativity and critical thinking, and demonstrating empathy.[42]

In a customer service setting, for example, an AI might analyze a customer's speech patterns and keywords to detect dissatisfaction, while a human agent uses empathy to understand the nuanced, underlying reasons for their frustration, leading to a more

effective resolution.[42] In a business context, an AI might analyze market data to identify trends, but human strategists are needed to interpret those trends and craft a compelling brand narrative.[44] This partnership, termed "collaborative intelligence," is meant to enhance decision-making, boost productivity, and drive innovation by creating a symbiotic relationship where each partner complements the other's strengths.[43]

The Recursive Replacement Loop

This neat division of labor, however, may be built on a foundation of sand. A revolutionary development in AI, known as **Recursive AI**, is beginning to blur the lines between machine computation and human cognition.[45] Recursive AI refers to systems that can learn and evolve iteratively by analyzing, critiquing, and improving upon their own outputs in a continuous feedback loop.[45] This is not simply automation of a static task; it is the automation of the process of improvement itself.

Several techniques demonstrate this capability. The **Reflexion** method, for instance, showed that a model like GPT-4 could be prompted to engage in verbal self-criticism. After generating a piece of code, the model would test it, identify its own errors, verbalize the mistakes, and then use that critique to generate a corrected version, improving its success rate on a programming benchmark from 80% to 91%.[45] This is a profound development. The system is not just writing code; it is learning to debug, reflect, and refine—skills previously designated as part of the human domain of "critical thinking" and "problem-solving."

This creates the risk of a **Recursive Replacement Loop**. As AI models become adept at not just performing tasks but also at the meta-tasks of self-improvement and optimization, the "synergistic zone" for human contribution begins to shrink. The tasks that are left for humans become the next target for automation. This is not a theoretical risk. The economic logic of this loop is already playing out. In 2024, Microsoft, a leader in AI development, reported saving over $500 million in its call centers through the use of AI, while also cutting over 15,000 jobs in the same year.[46] The company's CEO noted that AI was already writing 20-30% of the code in its repositories, directly encroaching on the work of software engineers.[46] This dynamic suggests that any definition of human value based on a specific set of functional skills is built on a moving target.

Redefining Human Skill: The EPOCH Paradigm

The threat of the Recursive Replacement Loop reveals the flaw in the current "future of work" debate. The focus on identifying tasks that AI cannot *yet* perform is a losing game. The only stable and defensible ground for long-term human value is not functional but normative. It is the ability to ask *why* a task should be performed, to define the ethical boundaries of its execution, and to provide the deep, pluriversal context that gives the task meaning. The ultimate human role must shift from execution or even collaboration to normative direction and ethical governance.

In response to this challenge, a new educational and skills framework is required. This report proposes re-appropriating the machine learning term **'epoch'**—which refers to one complete pass of a training dataset through an algorithm, during which the model learns and updates itself [47]—to define a new paradigm of uniquely human competencies.

EPOCH Skills are the core capabilities needed to guide and govern AI in a just and sustainable manner:

- **E - Ethical Scrutiny:** This is the ability to critically analyze and establish the ethical guardrails for AI systems. It involves not just applying pre-defined ethical frameworks but engaging in deep philosophical inquiry, anticipating second- and third-order consequences, and facilitating democratic deliberation about the values an AI system should embody. It is the skill of conducting the UNESCO EIA or participating in the governance of an AI-focused DAO.
- **P - Purpose-Driven Direction:** This is the capacity to define the ultimate purpose and goals that AI systems should pursue. In a world where AI can optimize for any given objective, the human role is to choose the right objective. This requires strategic foresight, a deep understanding of human needs and aspirations, and the wisdom to direct AI's power towards genuinely beneficial ends, such as the Sustainable Development Goals.
- **O - Ontological Awareness:** This is the skill of recognizing and respecting the pluriverse. It is the ability to understand that different cultures and communities have different ways of knowing and being, and to ensure that these diverse ontologies are not erased by universalist technologies. It is the human-centric skill required to build and curate a Pluriversal Knowledge Graph and to practice decolonial context engineering.

- **C - Contextual Intelligence:** This is the deep, embodied, and often tacit understanding of specific social, cultural, and historical contexts that AI systems, trained on decontextualized data, inherently lack. It is the skill of the Signpost Project's local moderator who knows the specific needs of their community, or the agricultural expert who can interpret the Ulangizi chatbot's advice in light of local weather patterns and soil conditions. It is the "last-mile" human wisdom that makes AI useful and safe in the real world.
- **H - Human-Centric Judgment:** This is the irreducible capacity for empathy, compassion, and nuanced judgment in situations that cannot be reduced to an algorithm. It is the legal professional's ability to apply the spirit of the law, not just the letter; the doctor's ability to weigh a data-driven recommendation against a patient's personal values and quality of life; and the social worker's ability to build a relationship of trust.

The EPOCH paradigm proposes a new curriculum for the age of AI. It argues that we should focus less on training people for jobs that will inevitably be automated and more on cultivating these higher-order, normative competencies. The future of human work and value lies not in competing with AI on functional tasks, but in mastering the uniquely human art and science of guiding it with wisdom, purpose, and a profound commitment to justice.

Part IV: From Architecture to Action: A Framework for Generative Justice

The architectural principles outlined in the preceding sections—encompassing technical, cognitive, and human-centric dimensions—provide a robust foundation for a more responsible AI. However, principles are inert without action. The final part of this report translates this architectural vision into a concrete and actionable framework for implementation. It moves beyond mitigating harm to proactively using AI to advance equity and redress historical injustices, a concept termed **Generative Justice**. It concludes with a multi-layered blueprint for global governance, designed to operationalize this new paradigm.

Section 8: Generative Justice: Proactive Applications for Equity and Heritage

Generative Justice involves harnessing the power of AI not just to avoid negative outcomes, but to actively create positive ones, particularly for marginalized communities and in the service of preserving the diverse tapestry of human culture. Two key domains for this proactive application are access to legal justice and the preservation of cultural heritage.

AI for Access to Justice

Access to justice is a fundamental human right, yet for billions of people, it remains out of reach due to the complexity, cost, and slow pace of legal systems.[49] This "justice gap" disproportionately affects low-income and marginalized communities. Generative AI offers a generational opportunity to dismantle these barriers and democratize access to legal support.[49]

Applications in this space are rapidly emerging. AI-powered chatbots and virtual assistants can provide immediate, accessible legal information to the public, helping individuals understand their rights and navigate complex procedures.[51] For example, Legal Aid of North Carolina developed an AI chatbot called LIA to provide actionable resources for common legal matters like domestic violence and landlord-tenant disputes.[50] For under-resourced legal aid organizations and pro bono lawyers, AI tools can dramatically improve efficiency by automating repetitive tasks like document review, legal research, and drafting initial briefs.[50] This allows legal professionals to delegate time-consuming work to the AI and focus their limited resources on high-value activities that require human judgment, such as client advising and legal strategy, ultimately enabling them to serve more people more effectively.[50] AI can also enhance legal referral and triage systems, guiding users through a series of structured questions to identify their core legal issue and refer them to the appropriate service, optimizing the allocation of scarce resources.[49]

The deployment of AI in this sensitive domain requires stringent ethical oversight. Legal professional bodies like the American Bar Association (ABA) and state bar associations are developing guidelines to address these concerns.[52] Key considerations include ensuring the accuracy and reliability of AI-generated legal information, maintaining strict client confidentiality, preventing the unauthorized use of data to train models, and ensuring human lawyers retain ultimate responsibility and

oversight of all legal work.[51] When implemented responsibly, these tools can be a powerful force for closing the justice gap and making the legal system fairer and more inclusive.[49]

AI for Cultural Heritage Preservation

A second, crucial application of Generative Justice is the use of AI to preserve and revitalize humanity's diverse cultural heritage—a direct and powerful countermeasure to the threat of ontological erasure. AI technologies are providing revolutionary new ways to restore, understand, and share cultural artifacts, texts, and historical environments.[53]

Generative AI models can be trained on photographs, 3D scans, and historical documents to digitally restore damaged artifacts or virtually reconstruct lost or destroyed heritage sites.[53] This allows for the creation of immersive virtual reality (VR) and augmented reality (AR) experiences that can transport people to ancient sites, breaking down geographical and physical barriers to cultural access.[53] For instance, a digital twin of St. Peter's Basilica was created using photogrammetry and AI, allowing for detailed virtual exploration and helping preservationists identify and monitor cracks and other forms of deterioration.[9]

AI is also unlocking lost knowledge from the past. A notable example is the **Vesuvius Challenge**, where researchers used CT scans and AI algorithms to digitally "unroll" and read ancient papyrus scrolls that were carbonized by the eruption of Mount Vesuvius in 79 AD, revealing texts that had been unreadable for nearly two millennia.[54]

Furthermore, AI can make history more interactive and engaging. At the University of Miami, a course on Sephardic history experimented with generative AI to create personalized narratives of historical events and developed a chatbot avatar of the medieval philosopher Maimonides, trained on his books and letters, allowing students and visitors to "converse" with him.[55] These applications demonstrate that AI, when thoughtfully applied, can be a catalyst for recovering overlooked histories and re-engaging communities with their cultural legacies, transforming preservation from a static act of archiving into a dynamic process of dialogue between tradition and technology.[55]

Section 9: A Blueprint for Global Implementation and Governance

Architecting a future where AI serves the global good requires a coordinated, multi-layered approach to governance that translates the principles of accountability, pluriversalism, and human-centricity into enforceable global standards and actionable policies. This blueprint outlines an integrated framework for action across five critical layers.

The Material Layer: Internalizing Environmental Costs

To address AI's material shadow, the global community must move beyond voluntary corporate commitments and establish binding international policies.

- **A Global Carbon Tax for Computation:** International bodies like the OECD and the UN should work towards establishing a harmonized framework for a carbon tax applied to the energy consumption of large-scale data centers.[17] This would internalize the environmental cost of AI, creating a powerful economic incentive for investment in energy-efficient hardware, optimization software, and renewable energy sources. Revenue could seed a global fund for green AI research and equitable climate adaptation.
- **Mandatory Circular Economy Standards:** Building on policies like the EU's WEEE Directive, international standards for AI hardware should be established based on circular economy principles. This would include mandating modular design for easy repair and upgrades, restricting the use of hazardous materials, and implementing robust Extended Producer Responsibility (EPR) schemes that hold manufacturers financially and logistically responsible for the entire lifecycle of their products, from production to end-of-life collection and recycling.[9]

The Technical-Governance Layer: Building Verifiable and Democratic Systems

To build trust, AI governance must be both technically verifiable and democratically legitimate.

- **Promotion of Verifiable DAOs:** Governments and foundations should fund and

support the development of hybrid governance models that combine the democratic oversight of Decentralized Autonomous Organizations (DAOs) with the technical assurance of Verifiable AI.[26] This would involve creating open-source toolkits and legal frameworks that empower communities to establish their own DAOs to govern local AI systems, with the power to mandate and receive verifiable, cryptographically secure audits of those systems' compliance with community-defined rules.

The Cognitive-Ethical Layer: Fostering Ontological Fluency

To counter digital colonialism and ontological erasure, a massive, concerted effort is needed to build AI that respects the pluriverse.

- **Investment in Pluriversal Knowledge Graphs (PKGs):** A global, public-private consortium should be established to fund the creation of Pluriversal Knowledge Graphs. This initiative must go beyond computer science and involve deep collaboration with interdisciplinary teams of linguists, anthropologists, historians, and, most importantly, Indigenous knowledge keepers and representatives from diverse cultural communities. The goal is to build a foundational knowledge architecture for AI that is not universalist but fundamentally pluralistic.[34]
- **Development of Decolonial AI Methodologies:** Research funding should be directed towards developing and standardizing decolonial AI methodologies, including advanced context engineering and prompt scaffolding techniques designed to guide AI away from biased, colonial patterns of thought and towards more nuanced, culturally sensitive, and self-reflective reasoning.[38]

The Human-Social Layer: Redefining Skill and Ensuring Control

To prepare humanity for a future of synergistic collaboration with AI, a paradigm shift in education and regulation is required.

- **A Global Educational Shift to EPOCH Skills:** National education policies and workforce development programs should be reoriented around the cultivation of **E**thical, **P**urpose-Oriented, **O**ntological, **C**ontextual, and **H**uman-centric (EPOCH) skills. This means prioritizing critical thinking, ethical reasoning, cross-cultural

understanding, and systems-level thinking over rote memorization and task-specific training that is vulnerable to automation.[44]

- **International Red Lines on Recursive Autonomy:** The international community, through bodies like the UN, must engage in urgent dialogue to establish clear red lines on the development and deployment of fully autonomous, recursively self-improving AI systems. While recursive techniques can be beneficial in contained environments, the creation of systems that could enter an uncontrolled replacement loop, automating human normative and ethical judgment, poses a profound risk. Ensuring "meaningful human control" must remain a non-negotiable principle of global AI governance.

The Justice Layer: Funding Reparation and Equity

Finally, a just AI future requires dedicated resources to redress past harms and proactively build a more equitable world.

- **A Global Fund for Generative Justice:** A portion of the revenue from a global computation tax should be allocated to a dedicated fund for Generative Justice initiatives. This fund would provide grants to projects using AI to close the access to justice gap, support cultural heritage preservation, and develop other applications that directly serve marginalized communities.[49]
- **Permanent Platforms for Epistemic Reparations:** This fund should also be used to establish and sustain permanent, well-resourced, and independent platforms for epistemic reparations. These platforms would empower communities harmed by historical injustices—and the new injustices of the digital age—to record, preserve, and share their own stories, countering dominant narratives and reclaiming their "right to be known" in the historical and digital record.[23]

In conclusion, architecting AI for social good is not a technical problem with a singular solution. It is a continuous and complex political, ethical, and cultural process. It demands that we move beyond a naive techno-optimism and confront the profound material and social shadows of our current trajectory. By building a new foundation based on the integrated principles of material sustainability, verifiable and democratic governance, pluriversal cognitive architectures, and a renewed focus on human-centric skills and generative justice, we can begin to steer the awesome power of artificial intelligence away from a future of extraction and control, and towards a

future that is truly more just, sustainable, and worthy of all humanity.

Works cited

1. What Is AI4Good? Harnessing AI for Positive Social Impact - ICTworks, accessed July 14, 2025, https://www.ictworks.org/define-ai4good-harnessing-ai-for-social-impact/
2. AI for Social Good: Positive Use Cases - Japeto AI, accessed July 14, 2025, https://www.japeto.ai/ai-for-social-good-positive-use-cases/
3. Artificial Intelligence and the Sustainable Development Goals: Operationalizing technology for a sustainable future | UN Global Compact, accessed July 14, 2025, https://unglobalcompact.org/compactjournal/artificial-intelligence-and-sustainable-development-goals-operationalizing
4. AI for Good - Identifying innovative AI applications to solve global ..., accessed July 14, 2025, https://aiforgood.itu.int/
5. AI in Action: Accelerating Progress Towards the Sustainable Development Goals - Google Public Policy, accessed July 14, 2025, https://publicpolicy.google/resources/research-brief-ai-and-SDG.pdf
6. The Sustainable Development Goals Report 2023: Special Edition - UN Statistics Division, accessed July 14, 2025, https://unstats.un.org/sdgs/report/2023/
7. UN summit confronts AI's dawn of wonders and warnings | UN News, accessed July 14, 2025, https://news.un.org/en/story/2025/07/1165346
8. AI for Good: Why Microsoft is Using AI for Positive Change - AI Magazine, accessed July 14, 2025, https://aimagazine.com/articles/ai-for-good-why-microsoft-is-using-ai-for-positive-change
9. AI For Good Lab - Microsoft Research, accessed July 14, 2025, https://www.microsoft.com/en-us/research/group/ai-for-good-research-lab/
10. AI for the Global Goals 2023 - Google, accessed July 14, 2025, https://globalgoals.withgoogle.com/globalgoals
11. Home - AI for Good Foundation - AI for social impact, accessed July 14, 2025, https://ai4good.org/
12. New material created at Clemson University could curb AI's energy ..., accessed July 14, 2025, https://news.clemson.edu/new-material-created-at-clemson-university-could-curb-ais-energy-appetite/
13. Optimization could cut the carbon footprint of AI training by up to 75 ..., accessed July 14, 2025, https://news.engin.umich.edu/2023/04/optimization-could-cut-the-carbon-footprint-of-ai-training-by-up-to-75/
14. The 2023 AI Index Report | Stanford HAI, accessed July 14, 2025, https://hai.stanford.edu/ai-index/2023-ai-index-report
15. Advancing Circular Economy Through AI-Driven E-Waste ..., accessed July 14, 2025, https://www.preprints.org/manuscript/202503.1989/v1
16. Sustainable AI: Reducing E-Waste - Number Analytics, accessed July 14, 2025,

https://www.numberanalytics.com/blog/sustainable-ai-e-waste-reduction
17. Carbon Tax Basics - Center for Climate and Energy SolutionsCenter ..., accessed July 14, 2025, https://www.c2es.org/content/carbon-tax-basics/
18. Carbon Pricing Policies | ACEEE, accessed July 14, 2025, https://database.aceee.org/state/carbon-pricing
19. intellectual history of digital colonialism | Journal of Communication ..., accessed July 14, 2025, https://academic.oup.com/joc/advance-article/doi/10.1093/joc/jqaf003/8078024
20. www.cambridge.org, accessed July 14, 2025, https://www.cambridge.org/core/journals/data-and-policy/article/artificial-intelligence-digital-colonialism-and-the-implications-for-africas-future-development/4BD73E9129A9CD9E9301C61CB2401450#:~:text=Specifically%2C%20it%20explores%20how%20the,oppression%2C%20dispossession%2C%20and%20exploitation.
21. Aesthetic Justice → Area → Resource 1, accessed July 14, 2025, https://lifestyle.sustainability-directory.com/area/aesthetic-justice/resource/1/
22. A Study of Artificial Intelligence's Impact on Aesthetic Standards and Its Potential Social Dilemmas - Clausius Scientific Press, accessed July 14, 2025, https://www.clausiuspress.com/assets/default/article/2024/09/22/article_1727002363.pdf
23. Epistemic Reparations: Buffett Institute for Global Affairs ..., accessed July 14, 2025, https://buffett.northwestern.edu/research/global-working-groups/epistemic-reparations.html
24. What Is Verifiable AI? A Guide to Transparency and Trust in AI, accessed July 14, 2025, https://www.identity.com/what-is-verifiable-ai-a-guide-to-transparency-and-trust-in-ai/
25. AI transparency vs. AI explainability: Where does the difference lie? - TrustPath, accessed July 14, 2025, https://www.trustpath.ai/blog/ai-transparency-vs-ai-explainability-where-does-the-difference-lie
26. Attestable Audits: Verifiable AI Safety Benchmarks Using ... - arXiv, accessed July 14, 2025, https://arxiv.org/abs/2506.23706
27. Transforming Business Through DAO and AI Integration - TokenMinds, accessed July 14, 2025, https://tokenminds.co/blog/ai/decentralized-ai
28. AI-Enhanced Decentralized Autonomous Organizations (DAOs) - Kava.io, accessed July 14, 2025, https://www.kava.io/news/ai-enhanced-decentralized-autonomous-organizations-daos
29. Ethical impact assessment: a tool of the Recommendation on the ..., accessed July 14, 2025, https://unesdoc.unesco.org/ark:/48223/pf0000386276
30. WHO releases AI ethics and governance guidance for large multi-modal models, accessed July 14, 2025, https://www.who.int/news/item/18-01-2024-who-releases-ai-ethics-and-governa

nce-guidance-for-large-multi-modal-models
31. Large multi-modal models – the present or future of artificial intelligence in medicine? - PMC, accessed July 14, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC10915764/
32. WHO releases AI ethics and governance guidance for large multi-modal models, accessed July 14, 2025, https://formiventos.com/2024/01/20/who-releases-ai-ethics-and-governance-guidance-for-large-multi-modal-models/
33. The Ethical and Societal Considerations of an AI Impact Analysis ..., accessed July 14, 2025, https://www.schellman.com/blog/ai-services/ethical-and-societal-considerations-of-ai-impact-analysis
34. Intelligence in the Pluriverse: Critical Ethology and the ... - PhilArchive, accessed July 14, 2025, https://philarchive.org/archive/PRAIIT
35. Design principles of the pluriversal design paradigm - DRS Digital ..., accessed July 14, 2025, https://dl.designresearchsociety.org/cgi/viewcontent.cgi?article=3180&context=drs-conference-papers
36. What is Prompt scaffolding? | PromptLayer, accessed July 14, 2025, https://www.promptlayer.com/glossary/prompt-scaffolding
37. Scaffolding for AI: Building Competence, One Prompt at a Time - TxDLA, accessed July 14, 2025, https://www.txdla.org/scaffolding-for-ai/
38. (PDF) Unsettling feelings in the classroom: scaffolding pedagogies ..., accessed July 14, 2025, https://www.researchgate.net/publication/356566099_Unsettling_feelings_in_the_classroom_scaffolding_pedagogies_of_discomfort_as_part_of_decolonising_human_geography_in_higher_education
39. (PDF) Decolonizing Methods of Educational Research - ResearchGate, accessed July 14, 2025, https://www.researchgate.net/publication/374260548_Decolonizing_Methods_of_Educational_Research
40. Context Engineering Guide | Prompt Engineering Guide, accessed July 14, 2025, https://www.promptingguide.ai/guides/context-engineering-guide
41. Knowledge graphs | The Alan Turing Institute, accessed July 14, 2025, https://www.turing.ac.uk/research/interest-groups/knowledge-graphs
42. Human-AI Collaboration: Enhancing Human-Machine Synergy, accessed July 14, 2025, https://convin.ai/blog/human-ai-collaboration-insights
43. What is Human AI Collaboration? - Aisera, accessed July 14, 2025, https://aisera.com/blog/human-ai-collaboration/
44. 6 Crucial Human Skills in the Age of AI - Intersog Israel, accessed July 14, 2025, https://intersog.co.il/blog/6-crucial-human-skills-in-the-age-of-ai/
45. Recursive AI: How Models Are Learning from Their Own Outputs in ..., accessed July 14, 2025, https://www.careerera.com/blog/recursive-ai-how-models-are-learning-from-their-own-outputs-in-continuous-improvement-loops

46. Microsoft saves $500 million with AI, but with 15,000 job cuts, accessed July 14, 2025, https://economictimes.indiatimes.com/tech/artificial-intelligence/microsoft-saves-500-million-with-ai-but-with-15000-job-cuts/articleshow/122362978.cms
47. What is Epoch in Machine Learning? - Simplilearn.com, accessed July 14, 2025, https://www.simplilearn.com/tutorials/machine-learning-tutorial/what-is-epoch-in-machine-learning
48. What Is an Epoch in Machine Learning? - Coursera, accessed July 14, 2025, https://www.coursera.org/articles/epoch-in-machine-learning
49. How Generative AI Can Improve Access To Justice | Artifact - Voltaire Legal, accessed July 14, 2025, https://www.voltairelegal.com/blog-posts/how-generative-ai-can-improve-access-to-justice
50. AI and legal aid: A generational opportunity for access to justice - Thomson Reuters Institute, accessed July 14, 2025, https://www.thomsonreuters.com/en-us/posts/ai-in-courts/ai-legal-aid-generational-opportunity/
51. Access to Justice 2.0: How AI-powered software can bridge the gap - ABA Journal, accessed July 14, 2025, https://www.abajournal.com/columns/article/access-to-justice-20-how-ai-powered-software-can-bridge-the-gap
52. The Implications of Generative AI: From the Delivery of Legal ..., accessed July 14, 2025, https://iaals.du.edu/blog/implications-generative-ai-delivery-legal-services-delivery-justice
53. (PDF) Generative AI for Cultural Heritage Preservation - ResearchGate, accessed July 14, 2025, https://www.researchgate.net/publication/387668842_Generative_AI_for_Cultural_Heritage_Preservation
54. AI in Historic Preservation | Anderson Archival, accessed July 14, 2025, https://andersonarchival.com/blog/ai-in-historic-preservation/
55. Using AI to preserve and share cultural heritage - University of Miami News, accessed July 14, 2025, https://news.miami.edu/as/stories/2025/06/using-ai-to-preserve-and-share-cultural-heritage.html