The Architecture of Ambiguity: From Philosophical Inquiry to Computational Resolution and Conceptual Blending

Section I: The Philosophical Landscape of Interpretive Uncertainty

The study of ambiguity is, in essence, the study of the inherent multiplicity of meaning in language and thought. It is a phenomenon that has long captivated and troubled philosophers, linguists, and computer scientists alike. Historically, within the Western philosophical tradition, ambiguity was often viewed as a pathology of language—a defect to be excised in the pursuit of logical precision. However, a more nuanced understanding has emerged, one that recognizes ambiguity not merely as a source of error but as a fundamental feature of human communication, cognition, and even existence itself. This section establishes the foundational concepts necessary to navigate this complex landscape, beginning with the classical view of ambiguity as a logical impediment, drawing a critical distinction between ambiguity and the related phenomenon of vagueness, and exploring how different philosophical traditions have framed the problem and its implications.

1.1 Defining Ambiguity: Beyond Equivocation and Fallacy

The philosophical interest in ambiguity has historically been rooted in a deep concern for the integrity of logical argumentation.[1] In the project of formalizing reasoning, natural language was often seen as an imperfect medium, rife with imprecision and equivocation that could invalidate an otherwise sound argument. An ambiguous word, phrase, or sentence is one that possesses two or more distinct meanings.[3] The danger, for logicians, is that an argument's inferential power depends on the consistent use of terms; if a term's meaning shifts mid-argument, the conclusion may not follow from the premises, even if the argument's form

appears valid.[3]

This concern gave rise to the identification of several "fallacies of ambiguity," patterns of incorrect reasoning that exploit the imprecise use of language.[3] Understanding these fallacies is crucial as it reveals the classical motivation for disambiguation: the preservation of logical truth.

- **Equivocation:** This fallacy occurs when an ambiguous word or phrase is used with one meaning in one part of an argument and a different meaning in another. A classic example illustrates the point: "Really exciting novels are rare. But rare books are expensive. Therefore, really exciting novels are expensive".[3] Here, the word "rare" is used equivocally. In the first premise, it means "uncommon" or "infrequently encountered." In the second, it means "valuable due to scarcity." The argument's apparent logical chain is broken by this semantic shift.[3]
- **Amphiboly:** This fallacy arises not from an ambiguous word but from an ambiguous grammatical construction (syntactic ambiguity). The structure of the sentence itself allows for multiple interpretations. For example, a radio broadcast reporting, "A reckless motorist Thursday struck and injured a student who was jogging through the campus in his pickup truck," could lead to the fallacious conclusion that "it is unsafe to jog in your pickup truck".[3] The ambiguity lies in whether the student or the motorist was in the truck.
- **Accent:** This fallacy stems from an ambiguity created by a shift in spoken or written emphasis. The statement "Jorge turned in his assignment on time today" is a simple statement of fact if read without inflection. However, with heavy stress on the word "today," it can be used to imply the conclusion, "Therefore, Jorge usually turns in his assignments late".[3]
- **Composition:** This fallacy involves incorrectly inferring that a whole has a certain property because all of its individual parts have that property. For instance, arguing that because "every course I took in college was well-organized," the college as a whole must be well-organized, is a fallacy of composition. The properties of the parts do not necessarily transfer to the whole.[3]
- **Division:** The inverse of composition, the fallacy of division incorrectly infers that the parts must have a property because the whole does. The argument "Ocelots are now dying out. Sparky is an ocelot. Therefore, Sparky is now dying out" is fallacious. The property of the class ("dying out") does not necessarily apply to each individual member of that class.[3]

This traditional focus on ambiguity as a source of logical error reveals a particular philosophical project. Figures like Gottlob Frege, a foundational figure in analytic philosophy, worried about such phenomena intensely, going so far as to advocate for a "perfect language" for science and logic that would be entirely free from any multiplicity of sense.[1] For Frege and others in this tradition, the primary purpose of language under analysis was the precise and unwavering representation of logical truths. From this perspective, ambiguity is not just an inconvenience but a fundamental flaw, a "bug" in the system of natural language

that must be eradicated for "serious philosophical business" to proceed.[1]

However, this view of ambiguity as a purely negative phenomenon is contingent on the perceived purpose of language. A significant shift occurs when the philosophical lens moves from prescriptive logic to the descriptive reality of human communication and existence. Later philosophers, most notably Ludwig Wittgenstein, reframe language's purpose not as a tool for perfect logical representation but as a set of practices embedded within human activities—his "language games".[5] From this viewpoint, ambiguity is not a flaw to be eliminated but a natural feature of language that is routinely and effectively managed by context. Similarly, existentialist thinkers like Simone de Beauvoir take this reframing a step further, positing ambiguity not as a feature of language but as the constitutive feature of the human condition itself—the very source of freedom and the impetus for creating meaning.[6] This represents a profound teleological shift in the philosophical treatment of ambiguity: from a bug in a logical system to a feature of a communicative practice, and finally to the fundamental operating system of human existence. The "problem" of ambiguity is thus not a static concept; it is defined by the purpose for which language and meaning are being analyzed.

1.2 The Ambiguity-Vagueness Distinction: Drawing a Sharp Line Around Fuzziness

To navigate the landscape of interpretive uncertainty, no conceptual clarification is more critical than the distinction between ambiguity and vagueness. While both phenomena involve a one-to-many mapping from an utterance to a meaning, they are fundamentally different in their structure and implications.[7] Conflating the two leads to significant analytical confusion, as they give rise to different philosophical problems and demand different computational solutions.

Ambiguity refers to a linguistic expression that has two or more *discrete, distinct, and specifiable meanings*.[8] The uncertainty lies in choosing which of these established meanings is intended in a given context.[4] A word like "bank" is ambiguous because it can mean a financial institution or a riverside, two entirely separate concepts.[4] Similarly, the sentence "We saw her duck" is ambiguous because "her duck" can refer to a bird she owns (a noun phrase) or to the action of her lowering her head (a verb phrase).[10] The key characteristic of ambiguity is that it involves a choice between a finite set of clear interpretations; once the context makes the choice clear, the uncertainty vanishes.[7] This property allows for a linguistic test: an ambiguous expression can be used to say something that is simultaneously true and false of the same situation. For example, a comedian who is peculiar but not funny could be described as "funny, but not funny".[7]

Vagueness, by contrast, refers to a single meaning whose boundaries of application are *imprecise, fuzzy, or lack sharp cut-off points*.[8] The uncertainty does not concern a choice between different meanings, but rather whether a single meaning applies in a particular "borderline" case.[7] Predicates like "tall," "bald," "heap," or "old" are vague because there is no precise height at which one becomes tall, no exact number of hairs one must lose to be bald, and no specific number of grains of sand that constitute a heap.[12] The instruction "Go down the road a ways" is vague because "a ways" does not specify a precise distance.[8] Unlike ambiguity, vagueness is characterized by a continuum of cases and the lack of a clear dividing line.[9] It is important to note that many, if not most, words exhibit both properties; for example, "child" is ambiguous (offspring vs. young person) and both meanings are vague (at what age does one cease to be a "child"?).[9]

The following table crystallizes this crucial distinction, linking each phenomenon to its core features, canonical philosophical problem, and the class of computational strategies required to manage it. This framework serves as an essential reference for the remainder of this report.

Feature	Ambiguity	Vagueness
Definition	One word/phrase with multiple discrete, distinct meanings.	One meaning with unclear, fuzzy, or non-existent boundaries of application.
Nature of Uncertainty	Uncertainty about *which meaning* is intended. (A problem of decidability between options)	Uncertainty about *whether the meaning applies*. (A problem of boundary definition)
Classic Example	"The stool is in the garden." (chair vs. poop) [4]	"Go down the road a ways." (how far is "a ways"?) [8]
Linguistic Test	Can be true and false of the same situation under different interpretations (e.g., "She is funny [ha-ha] but not funny [strange]").	Leads to borderline cases where truth value is indeterminate.
Canonical Problem	Fallacies of Equivocation & Amphiboly.[3]	The Sorites Paradox (Paradox of the Heap).[12]

Resolution Strategy	Contextual clarification; choosing one meaning from a finite set.	Establishing a (potentially arbitrary) cutoff; degree theories; fuzzy logic.
Computational Analogy	Classification task (e.g., Word Sense Disambiguation).	Thresholding, regression, or fuzzy logic task.

1.3 The Sorites Paradox: The Anatomy of Vagueness

If fallacies of equivocation represent the core logical problem of ambiguity, the Sorites paradox is the quintessential philosophical problem of vagueness.[14] First formulated by the Greek philosopher Eubulides of Miletus in the 4th century BCE, the paradox (whose name derives from

soros, the Greek word for "heap") brilliantly exposes the tension between gradual change in the world and the discrete categories of language.[12]

The paradox is constructed from a "sorites series"—a sequence of objects where adjacent members are only incrementally different, but the first and last members are clearly different with respect to a vague predicate.[12] The classic formulation involves a heap of sand:

1. **Premise 1:** A collection of 1,000,000 grains of sand is a heap.
2. **Premise 2 (The Tolerance Principle):** If a collection of *n* grains of sand is a heap, then a collection of *n-1* grains of sand is also a heap. (Removing a single grain cannot make the difference between a heap and a non-heap).
3. **Conclusion:** By repeated application of Premise 2, a collection of 1 grain of sand is a heap.

The paradox can also be run in the opposite direction, starting from the premise that one grain is *not* a heap, and concluding that a million grains are also not a heap.[12] In either form, the paradox arises because seemingly impeccable premises and uncontroversial reasoning (modus ponens) lead to an absurdly false conclusion.[12] This structure applies to any vague predicate: 'bald', 'tall', 'old', 'blue', and even complex concepts like personal identity ("You weren't you as a zygote... so you don't exist").[12]

Proposed resolutions to the paradox are numerous and have profound philosophical implications:

- **Deny the Premise (Nihilism):** One radical solution is to deny the first premise by

claiming that there are, in fact, no such things as heaps. Philosopher Peter Unger defends this view, arguing that our vague concepts are incoherent and do not correspond to anything in reality.[16]

- **Set a Sharp Boundary:** A common-sense response is to define a precise cutoff—for example, a heap is any collection of 10,000 or more grains. However, this solution is widely seen as unsatisfactory because the boundary is arbitrary. There seems to be no significant difference between 9,999 grains and 10,000 grains, making the precision misleading and contrary to how natural language is used.[16]
- **The Epistemic Theory:** Epistemicists like Timothy Williamson argue that there *is* a sharp, non-arbitrary boundary for every vague predicate, but it is impossible for humans to know where it lies.[12] Vagueness, on this view, is a form of ignorance. There is a specific number of hairs that marks the line between bald and not-bald, but our knowledge is too coarse-grained to locate it. This preserves classical logic (every statement is either true or false) but at the cost of positing unknowable semantic facts.[12]
- **Semantic Approaches (Non-Classical Logic):** Most theorists treat vagueness as a semantic phenomenon, requiring a departure from classical logic. These approaches accept that for borderline cases, a statement like "x is a heap" is neither determinately true nor determinately false.
    - **Supervaluationism:** This approach admits truth-value gaps. A sentence is "super-true" if it is true on all admissible ways of drawing a sharp boundary (precisifications). In the Sorites argument, the universal premise ("For all *n*, if *n* grains is a heap, then *n-1* is a heap") is false, but no single instance of it is false, resolving the paradox in a complex manner.[12]
    - **Degree Theories / Fuzzy Logic:** These theories replace the binary true/false logic with a continuum of truth values, often represented by the interval $$. A statement like "This is a heap" can be partially true (e.g., true to degree 0.6). The Tolerance Principle is then seen as not perfectly true, as each step slightly diminishes the degree of truth, eventually leading to a falsehood.[12]

The Sorites paradox is far more than a clever logical puzzle. It functions as a powerful diagnostic tool that compels a philosophical commitment regarding the relationship between language, human concepts, and the world. The type of solution one favors reveals one's underlying theory of meaning. For instance, to posit a sharp, fixed boundary suggests a form of realism where concepts have precise, mind-independent definitions. The epistemic view decouples semantics (what words mean) from epistemology (what we can know). Contextualist solutions, which argue the boundary shifts with context, imply that meaning is radically dependent on speaker intent and situation. The paradox is the sharpest formulation of a fundamental question: Does our language carve reality at its natural joints, or does it impose artificial, discrete boundaries on a continuous and seamless world? The application of Sorites-style reasoning to complex social and scientific concepts like "population" and "race" shows that this is not an abstract debate; it demonstrates that vagueness is a central

battleground for competing theories of how our concepts hook onto reality.[17]

1.4 Wittgenstein's Language Games: Meaning as Contextual Practice

Ludwig Wittgenstein's later work, particularly his *Philosophical Investigations*, offers a radical departure from the classical view of ambiguity as a logical flaw. Instead of trying to "fix" language, Wittgenstein sought to dissolve philosophical problems by showing they arise from a misunderstanding of how language actually works.[18] His central concept is the "language game," which brings into prominence "the fact that the speaking of language is part of an activity, or of a life-form".[5]

For Wittgenstein, the meaning of a word is not a fixed, dictionary-like definition that it carries around. Instead, a word's meaning is its *use* within a particular context or activity.[18] He provides numerous examples of language games: "Giving orders, and obeying them"; "Reporting an event"; "Constructing an object from a description".[5] The famous example is the word "game" itself. We use this single word for board games, card games, sports, and war games. Wittgenstein argues there is no single essential property that all these activities share. Instead, they are connected by a network of overlapping similarities, which he calls a "family resemblance".[20]

This framework provides a powerful way to understand ambiguity. Ambiguity arises when we pull a word out of its context—its language game—and treat it as an isolated object with a fixed meaning. Within a specific, functioning language game, ambiguity often vanishes. The word "Water!" in isolation is highly ambiguous. But as part of the language game "requesting a drink at a restaurant," it is an unambiguous request. As part of the game "warning people about a poisoned well," it is an unambiguous warning.[5] The context of the activity clarifies the meaning. Similarly, a word like "socialism" may seem hopelessly ambiguous, but when we understand the specific language game being played—whether it's a historical debate, a political rally, or an academic discussion—the intended meaning often becomes much clearer.[13]

Wittgenstein's philosophy can be seen as a blueprint for the entire field of computational pragmatics. His core claim—that meaning is use in context—is the foundational assumption of modern Natural Language Processing (NLP) systems that aim to resolve ambiguity. The NLP concept of "context" maps directly to Wittgenstein's "situation" or "life-form." The NLP notion of "speaker intent" corresponds to the "goal" of the language game. Tasks like intent detection, dialogue management, and coreference resolution are, in essence, computational attempts to model the rules and context of a specific language game. Wittgenstein's philosophical work, decades before modern NLP, provides the precise conceptual apparatus

to understand why ambiguity cannot be solved by syntax or dictionary definitions (semantics) alone. It requires a pragmatic level of analysis that models the "game" in which the language is being used. The historical progression of NLP, from rigid syntactic parsers to context-aware large language models, can be viewed as a technological journey toward implementing this profound philosophical insight.

1.5 The Existential Dimension: De Beauvoir and the Ambiguity of Being

While analytic philosophy and linguistics treat ambiguity as a property of language to be analyzed or resolved, the existentialist tradition, particularly in the work of Simone de Beauvoir, elevates ambiguity to a central feature of the human condition itself.[6] In her seminal work,

The Ethics of Ambiguity, Beauvoir argues that the human experience is fundamentally ambiguous. We are a paradoxical synthesis of opposites: we are both a consciousness that is radically free (an internal subject) and a physical body that is an object in a deterministic world; we are caught between the states of life and the inevitability of death.[6]

For Beauvoir, this ambiguity is not a problem to be solved but the very source of our freedom and the necessary precondition for ethics. Many philosophical and religious systems, she argues, are attempts to escape this ambiguity through "bad faith"—by denying one side of the duality. For example, proclaiming religious immortality is a way to deny the finality of death; pure determinism is a way to deny our freedom and responsibility.[6] Beauvoir contends that ethics only becomes meaningful for an ambiguous creature. An angel or a stone has no need for ethics. It is only because of the tension between what we

are (our factual situation) and what we *could be* (our freedom to project ourselves into a future) that the question of what we *ought to do* arises.[6]

To be free, one must first recognize and embrace this fundamental ambiguity. It is in the space of uncertainty and doubt that we are forced to create our own meaning and values. Life does not come with a pre-existing meaning; we must *make* it meaningful through our choices and actions, such as standing up for the oppressed or protesting injustice.[6] Embracing ambiguity means opening up the future, keeping alive the ability to act in multiple ways rather than collapsing into a single, predetermined path.[6]

This existentialist perspective offers a profound counterpoint to the purely logical or computational views of ambiguity. It suggests a deep connection between a system's tolerance for ambiguity and its capacity for higher-order functions like creativity and ethical reasoning. Most AI systems are designed with a Fregean goal: to reduce or eliminate

ambiguity in order to provide a single, optimal, and predictable output. Beauvoir's philosophy serves as a powerful cautionary framework for AI development. If ambiguity is the space where humans forge new values and create new meanings, an AI system designed to eliminate all ambiguity might also inadvertently eliminate the potential for genuine creativity and nuanced ethical judgment. An AI that can only apply pre-programmed ethical rules to well-defined situations lacks the ambiguous space needed to navigate novel ethical dilemmas where new "oughts" must be forged from the tension between "what is" and "what could be." This suggests that building an AI that is "ethically ambiguous" in the Beauvoirian sense—one that can recognize, tolerate, and grapple with the inherent openness of a situation rather than simply applying a rule—may be a necessary, if daunting, step toward more robust and responsible artificial intelligence.

Section II: The Cognitive Architecture of Meaning-Making

Moving from the philosophical foundations of ambiguity to its processing within the human mind, this section explores the cognitive mechanisms that enable us to navigate and create meaning in a world of interpretive uncertainty. Humans are remarkably adept at resolving ambiguity, often doing so unconsciously and in milliseconds. This proficiency is not the result of a single, simple module but a complex interplay of predictive processing, cognitive control, and embodied experience. Furthermore, the same cognitive architecture that resolves ambiguity also appears to be the engine of creativity, capable of generating novel meanings through the powerful operation of conceptual blending.

2.1 Resolving the Unresolved: Cognitive Processes in Disambiguation

The human brain does not process language like a static tape being fed into a machine. Instead, ambiguity resolution is a dynamic, real-time process that unfolds as an utterance is heard or read. One of the clearest examples of this is the processing of "garden-path" sentences, which contain a temporary syntactic ambiguity that leads the listener down an incorrect interpretive path before a reanalysis is required.[21] For instance, in the sentence "The baker poured the egg in the bowl...," the phrase "in the bowl" is temporarily ambiguous. It could be a modifier specifying

which egg is being poured (if there are multiple eggs, one of which is in a bowl), or it could specify the *destination* of the pouring action.[22] The human sentence processor constructs an

initial syntactic structure on the fly, and if subsequent information proves it wrong (e.g., the sentence continues "...was cracked"), it must perform a rapid reanalysis.[21]

This process of managing and resolving competing interpretations is not passive; it actively engages high-level executive functions, collectively known as **cognitive control**. Cognitive control is the ability to regulate thoughts and actions in service of an internal goal, primarily by biasing attention toward task-relevant information and inhibiting irrelevant information.[23] When an ambiguous sentence is encountered, multiple interpretations are activated and compete for selection. Cognitive control is hypothesized to be the mechanism that allows the listener to inhibit the initially preferred but ultimately incorrect interpretation and shift to the correct one.[23] This link is not merely correlational; research has shown a causal relationship. For example, dynamically manipulating cognitive control using a Stroop task (where one must name the color of a word while ignoring the word itself) has been shown to facilitate the processing of difficult garden-path sentences. The engagement of cognitive control on a preceding Stroop trial appears to "prime" the system, making it better able to resolve the conflict between competing syntactic interpretations on a subsequent sentence trial.[23]

Neuroscientific studies reveal the temporal dynamics of this process, showing that ambiguity resolution unfolds across distinct stages of brain activity. When presented with ambiguous stimuli, the brain first shows increased frontal θ-band power, indicating a greater reliance on top-down processes like memory and expectation to manage the uncertainty.[24] This is followed by activity in the occipito-parietal region in the β-band, thought to reflect the core disambiguation process itself. Finally, a large-scale fronto-parietal β-band network becomes active, associated with the final decision-making and response formation.[24]

This evidence collectively reframes ambiguity resolution from a simple, modular parsing task into a dynamic, embodied, and predictive activity. The brain does not passively receive an ambiguous sentence and then hand it off to a "disambiguation module." Instead, it is a continuous process where the cognitive system:

1. **Predicts:** It uses top-down knowledge and context to form expectations about the likely meaning even before the full utterance is processed.[24]
2. **Manages Competition:** It uses cognitive control to actively manage the conflict between competing interpretations as they arise in real-time.[23]
3. **Integrates Embodied and Environmental Information:** The process is not confined to the linguistic input alone. The referential context—the objects present in the environment—profoundly influences syntactic choices. For example, hearing "Pour the egg in the bowl" in a scene with one egg and two bowls leads to a destination interpretation, while hearing it in a scene with two eggs (one in a bowl, one on a napkin) leads to a modifier interpretation.[22] The brain continuously evaluates sentence information against the properties of the contextual environment, including the "affordances" of objects (what actions they make possible).[22]

This perspective suggests that ambiguity is not a property of a sentence in isolation, but an emergent property of the interaction between language, the cognitive system, and the environment. This "enactive" view posits that language is a tool that enables complex interactions in a complex world, a process that feels simple ("simplexity") because our entire embodied cognitive system is engaged in making sense of it.[25] A truly human-like computational model of disambiguation, therefore, cannot be a simple feed-forward classifier. It must incorporate feedback loops, predictive mechanisms, and multi-modal grounding to replicate the dynamic and deeply integrated nature of human comprehension.

2.2 Conceptual Blending Theory: The Engine of Creativity

The cognitive processes that resolve ambiguity by integrating information from different sources are closely related to a more general and powerful mental operation known as **conceptual blending** or **conceptual integration**. Developed by Gilles Fauconnier and Mark Turner, this theory proposes that a fundamental aspect of human thought is the ability to combine elements from diverse "mental spaces" to create novel meanings and emergent conceptual structures.[26] While often used to explain spectacular cases of creativity and metaphor, blending is argued to be a ubiquitous, workaday process that underpins much of everyday reasoning and language.[27]

The theory is built upon a **Network Model** that typically involves at least four mental spaces, which are small, partial conceptual packets constructed in working memory as we think and talk.[26]

- **Input Spaces:** These are two (or more) distinct conceptual structures that provide the initial ingredients for the blend. For example, to understand the phrase "In the end, VHS delivered a knock-out punch to Betamax," one input space would contain knowledge about the VCR format war of the 1970s and 80s, and the other would contain knowledge about boxing.[30]
- **Generic Space:** This space contains the abstract schematic structure that the input spaces share. In the VHS example, the generic space would capture the frame of "two rivals competing for dominance".[26] This shared structure provides the initial link that allows the inputs to be brought together.
- **Blended Space (The Blend):** This is a new, hybrid mental space where selected elements from the input spaces are projected and combined. The blend develops an "emergent structure" that is not present in any of the inputs alone.[26] In our example, the blend is a conceptual reality where the VCR formats are personified as boxers engaged in a physical fight, which VHS ultimately wins by a knockout.[30]

This emergent structure is created through three key dynamic operations within the blend [26]:

1. **Composition:** Elements from the different inputs are composed to create new relationships in the blend. The "competitor" role from the VCR space is composed with the "punching" action from the boxing space, creating the novel event of one VCR format punching another.
2. **Completion:** The blend is fleshed out by bringing in additional background knowledge associated with one of the input frames. The boxing frame brings with it concepts like rounds, referees, and a ring, which can be used to structure the blend, even though they have no direct counterpart in the VCR format war.
3. **Elaboration:** The blend is "run" as a dynamic mental simulation to explore its consequences and generate new inferences. One can imagine the "fight" lasting for years, with various blows exchanged (e.g., price cuts, new features), culminating in the final "knock-out punch" (Betamax being pulled from the market). This elaboration is what gives the blend its inferential and rhetorical power.[26]

A famous example used by Fauconnier and Turner is a magazine report about a modern catamaran in 1993 attempting to beat the sailing record set by a clipper in 1853 over the same course. The report reads: "Rich Wilson and Bill Biewenga were barely maintaining a 4.5 day lead over the ghost of the clipper Northern Light".[28] Here, two distinct events from different centuries (Input Spaces) are blended into a single, imaginary event: a boat race (the Blend). This blend allows for a highly compressed and intuitive way to understand the complex relationship between the two voyages, creating the emergent idea of a "lead" and the drama of a race that could not exist otherwise.[28]

Conceptual blending theory offers a powerful framework that may unify several disparate linguistic and cognitive phenomena. It can be generalized to serve as a high-level cognitive model for ambiguity resolution itself. When a listener encounters an ambiguous phrase like "the bat," the cognitive task is to integrate two inputs: the ambiguous utterance (Input 1) and the surrounding context (Input 2). The process of disambiguation involves creating a blended space where elements from the context (e.g., a baseball game frame) are projected onto one of the potential meanings of the utterance to create a single, coherent interpretation (e.g., a wooden bat). This reframes ambiguity resolution from a simple selection task to a creative, constructive act.

Furthermore, this provides a plausible mechanism for semantic change. A novel blend, initially created for a one-off communicative purpose, can become conventionalized if it proves useful. Over time, the emergent structure of the blend can become a new, stored meaning for a word. For example, the word "virus" was once confined to a biological input space. Through blending with the domain of computers, a new, conventionalized meaning emerged. Conceptual blending thus provides a powerful, generalizable architecture that can account for how we generate new meanings (creativity), how we select from existing meanings (ambiguity resolution), and how the set of available meanings evolves over time (semantic

change).

Section III: A Taxonomy of Ambiguity in Language and Computation

To bridge the gap between the philosophical and cognitive understanding of ambiguity and its practical management in technology, it is essential to develop a formal classification. In the field of Natural Language Processing (NLP), ambiguity is not a monolithic problem but a multi-layered challenge that manifests at every level of linguistic analysis, from individual words to entire discourses. This section provides a systematic taxonomy of ambiguity as it is understood and addressed in computational linguistics, linking each type of ambiguity to the specific NLP tasks designed to resolve it.

3.1 Lexical and Semantic Ambiguity: The Problem of the Word

The most fundamental and frequently encountered form of ambiguity occurs at the level of the word, known as **lexical ambiguity**.[33] This arises when a single word form (a string of letters or sounds) is associated with multiple meanings.[35] This category is further divided into two main subtypes:

- **Homonymy:** This occurs when a single word form corresponds to two or more historically and semantically *unrelated* meanings. The fact that they share a form is essentially a linguistic accident. The classic example is the word bank, which can refer to a financial institution or the land alongside a river.[33]
- **Polysemy:** This occurs when a single word has multiple meanings that are *related* to each other, often through metaphorical extension or specialization. For example, the word foot can refer to the appendage at the bottom of a leg, the bottom of a mountain, or the bottom of a page. These meanings are distinct but share a common conceptual basis of being the "base" of something.[33] Similarly,
  newspaper can refer to the physical paper object or the organization that publishes it.[40]

While the distinction between homonymy and polysemy is important for lexicographers, in many NLP applications they are treated together as the general problem of lexical ambiguity.

Semantic ambiguity is a broader term that encompasses lexical ambiguity but also includes ambiguity at the phrase and sentence level, where the combination of individually

unambiguous words creates multiple potential meanings.[10] The sentence "We saw her duck" is a classic example of semantic ambiguity that is also syntactically ambiguous. The phrase "her duck" could mean a bird she possesses or the action she is performing.[10]

The primary computational task designed to resolve lexical and some forms of semantic ambiguity is **Word Sense Disambiguation (WSD)**. WSD is the process of computationally identifying and selecting the correct meaning (or "sense") of a word in a given context from a predefined set of possibilities, often provided by a sense inventory like WordNet or BabelNet.[36] Early WSD systems relied on knowledge-based rules or supervised machine learning with handcrafted features. However, the advent of modern contextualized language models, such as BERT, has revolutionized the field. These models generate different vector representations (embeddings) for a word depending on its surrounding context, thereby implicitly capturing its sense.[43] The task is now evolving further to suit the generative capabilities of Large Language Models (LLMs), for instance by prompting an LLM to generate the correct definition for a word in context or to select the correct definition from a list of options.[42]

The following table provides a structured overview of the major types of ambiguity addressed in NLP, serving as a roadmap for the technical landscape.

Ambiguity Type	Level	Definition	Classic Example	Primary NLP Resolution Task(s)
Lexical	Word	A single word form has multiple distinct meanings (homonymy/polysemy).	"The **bank** is closed." (financial vs. river) [4]	Word Sense Disambiguation (WSD) [42]
Syntactic	Sentence	A sentence's grammatical structure allows for multiple parse trees.	"I shot an elephant **in my pajamas**." (Groucho Marx) [46]	Syntactic Parsing (Constituency/ Dependency Parsing), Probabilistic Parsing [48]
Semantic	Phrase/Senten	A sentence's	"**Every** student	Semantic Role

	ce	meaning is unclear due to the combination of word meanings or scope.	read **two** poems." (different poems or the same two?) [51]	Labeling (SRL), Scope Disambiguation [53]
Pragmatic	Discourse	A sentence's meaning is unclear without understanding speaker intent, context, or world knowledge.	"The horse ran up the hill. **It** was steep." (It = horse or hill?) [54]	Coreference Resolution, Intent Detection, Dialogue Management [55]

3.2 Syntactic Ambiguity: The Problem of Structure

Syntactic ambiguity, also known as structural ambiguity or amphiboly, arises when a sequence of words can be grammatically structured in more than one way, leading to different meanings.[46] The ambiguity lies not in the words themselves, but in the relationships between them.

A common form is **attachment ambiguity**, where a modifying phrase, such as a prepositional phrase, could legally attach to more than one part of the sentence.[37] In the sentence, "He ate the cookies on the couch," the phrase "on the couch" could modify "the cookies" (specifying which cookies he ate) or it could modify the verb phrase "ate the cookies" (specifying where the eating took place).[10] Another key type is

scope ambiguity, which occurs when a sentence contains multiple quantifiers or logical operators, and their relative scope is unclear.[51] The sentence "Every student read two poems" could mean that for each student, there exist two poems that they read (potentially different poems for each student), or it could mean that there exist two specific poems that every student read.[51]

Linguists distinguish between **globally ambiguous** sentences, where the ambiguity remains even after the entire sentence is processed, and **locally ambiguous** sentences. Local

ambiguities often create a "garden-path" effect, where the reader is led towards an initial interpretation that turns out to be incorrect, forcing a re-evaluation of the sentence structure.[21]

The fundamental NLP task for resolving syntactic ambiguity is **syntactic parsing**, which is the process of analyzing a string of words to determine its grammatical structure, typically represented as a parse tree.[48] To handle ambiguity, parsers often employ probabilistic models. A

Probabilistic Context-Free Grammar (PCFG) is a CFG in which each production rule is assigned a probability.[58] For an ambiguous sentence with multiple possible parse trees, the PCFG can calculate the probability of each tree by multiplying the probabilities of all the rules used to generate it. The parser then selects the parse tree with the highest overall probability as the most likely intended structure, thus resolving the ambiguity based on statistical patterns learned from a large corpus of text.[50]

3.3 Pragmatic Ambiguity: The Problem of Intent

Pragmatic ambiguity occurs when the meaning of an utterance is unclear without appealing to factors beyond the sentence itself, such as the broader discourse context, speaker's intentions, and shared world knowledge.[56] This is often the most challenging type of ambiguity for NLP systems because it requires reasoning beyond the literal text.

Key forms of pragmatic ambiguity and their corresponding NLP solutions include:

- **Referential Ambiguity:** This arises when it is unclear what entity a pronoun or other referring expression is pointing to.[36] In the discourse "The horse ran up the hill. It was very steep," the pronoun "It" could refer to either "the horse" or "the hill".[54] The NLP task designed to solve this is
  Coreference Resolution. This task involves identifying all expressions (mentions) in a text that refer to the same real-world entity and clustering them together.[55] For example, in "Mary called. She left a message," the system would link "She" to "Mary."
- **Illocutionary Ambiguity:** This is uncertainty about the speaker's communicative goal or the "speech act" being performed. The sentence "Can you open the window?" could be a literal question about ability or, more commonly, an indirect request to open the window.[36] In conversational AI, this is handled by
  Intent Detection or **Dialogue Act Recognition**, which classifies an utterance into a category like "question," "request," or "statement" based on linguistic cues and conversational context.
- **Semantic Role Ambiguity:** Even in a syntactically unambiguous sentence, the underlying

event structure can be semantically ambiguous. To fully understand "John broke the window with a hammer," a system needs to identify "who did what to whom, with what, where, and when." **Semantic Role Labeling (SRL)** is the task that addresses this by identifying the main predicate (the verb, e.g., "broke") and assigning semantic roles to its arguments.[53] Common roles include AGENT (the instigator of the action, "John"), PATIENT or THEME (the entity affected, "the window"), and INSTRUMENT (the tool used, "a hammer").[64] By explicitly labeling these roles, SRL provides a shallow semantic representation that clarifies the event structure, resolving potential ambiguities about the relationships between participants.[67]

3.4 The Efficiency Principle: Ambiguity as a Feature, Not a Bug

Contrary to the classical philosophical view of ambiguity as a defect, a compelling theory from cognitive science and information theory argues that ambiguity is a vital feature that makes human language remarkably *efficient*.[51] This principle posits that ambiguity allows a language to achieve a higher rate of information transmission by reusing short, frequent, and phonetically simple words (e.g., "run," "get," "set," "mean") for many different purposes.[68]

The logic rests on a trade-off between the cognitive costs for the speaker and the listener. It is "cognitively cheaper" for a listener, who is immersed in a rich context, to disambiguate a simple, common word than it is for a speaker to always retrieve and produce a longer, more specific, and perfectly unambiguous term for every situation.[68] This creates a communication system that is optimized for efficiency by skewing toward ambiguity, offloading the work of interpretation to the listener's powerful cognitive machinery.[68] Human language, in this view, is a highly compressed communication protocol that presupposes a sophisticated, context-aware receiver capable of decompressing the signal.

This efficiency principle reveals a fundamental architectural gap between human and machine intelligence that helps explain the historical "brittleness" of AI systems. Humans possess cognitive hardware and software (e.g., world knowledge, theory of mind, cognitive control) that are exquisitely optimized for resolving ambiguity. This allows our communication protocol—natural language—to be efficient and compressed. Computers, conversely, have historically lacked this sophisticated disambiguation engine. Their native languages—programming languages—must therefore be perfectly unambiguous. A single misplaced semicolon in C++ is a fatal error because the compiler has zero tolerance for ambiguity. This creates a profound impedance mismatch: human language is a lossy, context-dependent compression algorithm, while traditional computer languages are lossless, context-independent protocols.

The history of NLP can be viewed as the slow and arduous process of building a computational disambiguation engine powerful enough to handle the compressed, context-reliant protocol of human language. Early rule-based systems were brittle because they tried to force human language into the rigid, unambiguous framework of computers. The success of modern LLMs stems from their unprecedented, though still imperfect, ability to approximate the human receiver. By learning statistical patterns from vast amounts of text, they have developed a powerful capacity to use context to disambiguate, finally beginning to close the human-computer ambiguity gap and enabling more natural and fluid interaction.[68]

Section IV: Ambiguity in Practice: Case Studies in Law and Business Process Modeling

The philosophical, cognitive, and computational dimensions of ambiguity are not merely academic abstractions. The successful management of ambiguity is a critical necessity in high-stakes, real-world domains where misinterpretation can lead to significant financial, legal, and operational consequences. This section examines two such domains—law and business process modeling—to illustrate how the theoretical concepts of ambiguity manifest in practice and the specialized doctrines and frameworks developed to control them.

4.1 The Letter of the Law: Ambiguity in Legal Contexts

In the legal field, language is the primary instrument for defining rights, obligations, and punishments. The interpretation of legal texts—statutes, contracts, wills, and treaties—is a central function of the judiciary, and ambiguity is a common and perilous challenge.[73] When the language of a legal document can be reasonably interpreted in more than one way, a dispute arises that the courts must resolve.[75] To manage this, legal systems have developed a crucial distinction between two types of ambiguity:

- **Patent Ambiguity:** This is an ambiguity that is apparent "on the face of the instrument" to any reasonable person reading it, even without knowledge of the external circumstances.[74] It arises from defective, obscure, or contradictory language within the document itself.[76] An example would be a bill of exchange drawn for "£245" in figures but "two hundred pounds" in words.[74] In cases of patent ambiguity, courts are generally reluctant to admit extrinsic evidence (or "parol evidence") about what the parties *intended* to write. The evidence is typically limited to explaining what *was* written.[74] In

government contracts, there is a duty on the contractor to seek clarification on a patent ambiguity before bidding; failure to do so may prevent them from later relying on their preferred interpretation.[77]

- **Latent Ambiguity:** This is a more subtle form of ambiguity that is not visible on the face of the document. The language itself is clear and intelligible, but an ambiguity appears when it is applied to the external subject matter.[74] The classic example is a will that leaves an estate "to my nephew, John," when the testator is discovered to have two nephews named John.[74] The text is facially clear, but its application to the world reveals the ambiguity. In these cases, because the ambiguity arises from extraneous circumstances, courts permit the use of parol evidence to ascertain the true intention of the parties as expressed in the agreement.[74]

To guide their interpretive task, courts employ several doctrines or "canons of construction":

- **The Rule of Lenity:** In criminal law, this doctrine holds that if a penal statute is ambiguous, it should be construed in the manner most favorable to the defendant, imposing the lesser penalty. This is founded on the principle that the power to define crimes and punishments rests with the legislature, not the judiciary, and on a "tenderness of the law for the rights of individuals".[74]
- ***Contra Proferentem* ("Against the Drafter"):** In contract law, this rule dictates that an ambiguous term should be construed against the party that drafted the contract.[74] The rationale is that the drafter, often the party with greater bargaining power and expertise, had the opportunity to make the language clear and should bear the risk of any resulting ambiguity.
- **Liberal Construction for Indigenous Treaties:** In jurisdictions like Canada, courts have established that treaties with Indigenous peoples should be liberally construed, with any "doubtful expressions resolved in favour of the Indians." This recognizes the historical power imbalances and differing cultural contexts in which these treaties were negotiated.[74]

Beyond these doctrines, courts use traditional rules of construction, such as reading the contract as a whole, giving words their plain and ordinary meaning, and striving to give purpose to every word used while avoiding interpretations that render parts of the text superfluous.[76]

4.2 Modeling the Flow: Ambiguity in Business Processes

In the field of Business Process Management (BPM), the goal is to design, model, execute, and monitor business processes to improve efficiency and effectiveness. Process-related artifacts—such as informal specifications, formal models, and event logs—are the blueprints

and records of these operations. Ambiguity within these artifacts is a significant quality issue that can lead to misunderstandings, modeling errors, non-conformance with regulations, and incorrect interpretations of performance data, ultimately undermining the success of the BPM initiative.[78]

A systematic taxonomy helps to identify where and how ambiguity can emerge in the BPM lifecycle [78]:

1. **Descriptive Ambiguity:** This occurs in the initial stages, within informal process specifications written in natural language. It can stem from:
    - **Linguistic Ambiguity:** The language itself is unclear. This includes *unclear relations* (e.g., "Treat any underlying cause... and/or review medication" does not specify if the actions are sequential, parallel, or exclusive) and *unclear references* (e.g., in a model of a sales process, a label simply says "Review request," which is ambiguous as to what request is being reviewed and by whom).[78]
    - **Epistemic Ambiguity:** The specification reflects a lack of knowledge about the process. This includes *underspecifications* (deliberately or unintentionally omitting details, creating conceptual gaps) and *inconsistent specifications* (presenting conflicting requirements that cannot all be satisfied).[78]
2. **Representational Ambiguity:** This arises within the formal process model itself (e.g., a diagram in Business Process Model and Notation, BPMN). It can be:
    - **Intrinsic:** Caused by the limitations of the modeling language itself, which may have a "representational bias" that cannot capture certain semantic nuances of the real-world process.[78]
    - **Extrinsic:** Caused by the modeling task, such as using inadequate software tools that don't fully support the modeling standard, or "modeling fuzziness," where the designer deviates from best practices (e.g., using the same name for different activities).[78]
3. **Observational Ambiguity:** This affects the event logs that record the actual execution of processes. It can relate to:
    - **Coverage:** The data is incomplete (*missing data*) or contains duplicates (*multiplied data*) or noise (*irrelevant data*).[78]
    - **Reliability:** The data is untrustworthy, containing *incorrect information* (misaligned with the known model), *imprecise information* (e.g., timestamps with too-coarse granularity, making the order of events uncertain), or *unexpected information* (values that deviate from expected norms).[78]

In the context of project management and business strategy, it is also vital to distinguish **uncertainty** from **ambiguity**. Uncertainty relates to a lack of information; it is what is not clearly known or determined. It can often be resolved by gathering more data and performing rational analysis. Ambiguity, in this context, is linked to complexity; it relates to having too many choices or possible interpretations, often because cause-and-effect relationships are unclear.[81] While uncertainty is generally something to be reduced, ambiguity can sometimes

be a deliberate strategic choice to maintain flexibility, avoid overly rigid models, or facilitate exploration in turbulent environments.[78]

Section V: A New Framework: Contextual Ambiguity Blending

The preceding sections have traversed the multifaceted landscape of ambiguity, from its philosophical roots and cognitive underpinnings to its formal classification in computation and its practical consequences in law and business. This analysis reveals a powerful convergence: philosophy underscores the primacy of context; cognitive science unveils the dynamic, constructive nature of comprehension; and computer science provides a toolkit for modeling these phenomena. A truly robust theory of ambiguity resolution must synthesize these perspectives, moving beyond a simplistic model of "selection" to a more powerful model of "construction." This final section proposes such a framework: Contextual Ambiguity Blending (CAB).

5.1 Synthesis of Preceding Analyses

The journey through the architecture of ambiguity has yielded several key principles that must inform any comprehensive model:

- **From Philosophy:** The critical distinction between ambiguity (discrete meanings) and vagueness (fuzzy boundaries) provides the foundational problem definition. Wittgenstein's concept of language games demonstrates that meaning is not inherent in words but is actualized in a context of use, making the resolution of ambiguity fundamentally a pragmatic, context-driven act.
- **From Cognitive Science:** The human mind does not resolve ambiguity passively. It is a dynamic, predictive process managed by cognitive control, integrating linguistic cues with environmental and embodied knowledge. The theory of conceptual blending provides a powerful, domain-general mechanism for how the mind creatively constructs new, coherent meanings by integrating information from disparate mental spaces.
- **From Computer Science:** NLP offers a formal taxonomy of ambiguity types (lexical, syntactic, semantic, pragmatic) and a corresponding set of specialized resolution tasks (WSD, parsing, coreference resolution, etc.). This demonstrates that ambiguity is a structured, multi-layered problem that can be systematically decomposed. The principle of efficiency also suggests that ambiguity is a feature of an optimized communication

system, not a flaw.

A unified model must therefore account for the discrete choices of ambiguity, the context-dependency of Wittgensteinian pragmatics, and the constructive, emergent nature of conceptual blending.

5.2 The Proposed Model: Contextual Ambiguity Blending (CAB)

The central thesis of the Contextual Ambiguity Blending (CAB) model is that **ambiguity resolution is not a selection process, but a creative construction process best modeled by conceptual blending.** When faced with an ambiguous utterance, the mind does not simply search a mental list of meanings and pick the one that best fits the context. Instead, it performs a dynamic integration, creating a new, singular, and coherent meaning in a blended mental space. This act is a form of "on-the-fly" meaning creation, homologous to the processes underlying metaphor and creativity.

The CAB model is structured using the network model of Fauconnier and Turner [26]:

- **Input Space 1: The Ambiguous Utterance.** This mental space contains the linguistic form itself, but crucially, its multiple potential meanings or structures are held in a state of cognitive superposition. For the utterance "There's a bat in the attic," this space contains the word form bat linked to both the ANIMAL concept and the SPORTING_IMPLEMENT concept.[33] For a syntactically ambiguous phrase like "the man with the telescope," this space contains the potential parse trees where the telescope is held by the man and where the telescope is the instrument of seeing.[82]
- **Input Space 2: The Context.** This is a rich, heterogeneous, and multi-layered mental space that contains all the non-linguistic information available to the comprehender. It is the computational instantiation of Wittgenstein's "language game" or "form of life".[5] This space includes:
  - **Discourse Context:** Information from previous sentences.
  - **Situational Context:** The physical environment, time, and place of the utterance.
  - **World Knowledge:** General knowledge about how the world works (e.g., flying mammals are often found in attics, baseball bats are not).
  - **Epistemic and Intentional Models:** Beliefs about the speaker's knowledge, goals, and intentions.
- **Generic Space:** This space captures the abstract structure of a successful communicative act, where a speaker produces a form intended to convey a single, coherent meaning to a hearer.
- **Blended Space: The Resolved Interpretation.** This is the output of the process. The blend is a new, integrated conceptualization that is more than the sum of its parts.

Elements from the Context space (Input 2) are selectively projected and composed with *one* of the potential meanings from the Ambiguous Utterance space (Input 1). The result is a single, coherent, and contextually grounded meaning. For "There's a bat in the attic," if the context is a conversation about pest control, the ANIMAL sense is projected into the blend, creating the specific, unambiguous meaning of a flying mammal in the attic.

5.3 Operations and Principles of CAB

The dynamic operations of conceptual blending are adapted in the CAB model to describe the specific cognitive work of disambiguation:

1. **Composition:** The context actively composes with a potential meaning to create a new, specific scenario. If the context (Input 2) is a baseball game, and the utterance (Input 1) is "He grabbed the bat," the SPORTING_IMPLEMENT sense is projected. This sense then composes with the "baseball game" frame from the context to form a coherent event: a player grabbing a wooden bat to go up to the plate. The meaning is not just "sporting implement"; it is this specific, situated action.
2. **Completion:** Once a potential meaning is projected into the blend, background knowledge associated with that meaning's frame is used to "complete" the picture and test for coherence. If the ANIMAL sense of "bat" is chosen in the context of a baseball game, the completion process would bring in associated knowledge (e.g., bats are nocturnal, they fly, they are not used for hitting balls). This creates an immediate clash with the "baseball game" frame, signaling that this interpretation is incoherent. This completion process is a powerful mechanism for filtering out contextually inappropriate meanings.
3. **Elaboration:** The constructed blend is "run" as a mental simulation to check its consistency and generate further inferences. This is analogous to the reanalysis process observed in cognitive studies of garden-path sentences.[21] The listener constructs an initial blend (the most likely interpretation). As more linguistic information arrives, it is fed into the blend. If the new information can be smoothly integrated (elaborated), the interpretation is maintained. If it creates a contradiction, the blend "breaks," forcing the cognitive system to backtrack and construct a new blend using an alternative projection from the ambiguous input space. This iterative process of blending, completion, and elaboration continues until a stable and coherent interpretation is achieved.

5.4 Applications and Future Directions

The Contextual Ambiguity Blending model offers a more cognitively plausible and powerful framework for understanding and modeling ambiguity resolution, with implications across several fields.

- **For Natural Language Processing:** CAB provides a new conceptual architecture for disambiguation in LLMs. Current models largely rely on massive statistical correlation learned during pre-training. CAB suggests a more structured approach where context and the ambiguous utterance are treated as distinct input spaces that are actively and dynamically integrated. This could inspire new prompting strategies (e.g., explicitly providing context as one input and the ambiguous sentence as another) or fine-tuning methods that train a model to perform the operations of composition, completion, and elaboration to generate a final, disambiguated output. It shifts the goal from merely predicting the right label to constructing a coherent, contextually-grounded representation.
- **For Philosophy and Law:** The model reframes the act of legal interpretation. Instead of a judge "discovering" the single, pre-existing "true" meaning of an ambiguous statute, CAB models the process as a constrained act of conceptual blending. The legal text with its latent or patent ambiguities is Input Space 1. The vast context of legal precedent, canons of construction, legislative history, and the specific facts of the case form Input Space 2. The judge's ruling is the Blended Space—a new, legally binding interpretation constructed by integrating these two inputs. This acknowledges the inherently creative and constructive, yet rule-governed, nature of legal reasoning.
- **For Cognitive Science:** The CAB framework serves to unify the process of ambiguity resolution with other high-level cognitive tasks like metaphor, analogy, and creativity under the single, powerful umbrella of conceptual blending. This lends significant support to the theory's claim to be a basic, domain-general mental operation.[27] It suggests that the same cognitive machinery that allows us to understand a poet's metaphor is at work when we effortlessly figure out whether "bat" means an animal or a piece of wood. It positions ambiguity not as a failure of language, but as an invitation for the mind to perform its most fundamental and powerful operation: the creation of new meaning.

Works cited

1. Ambiguity (Stanford Encyclopedia of Philosophy/Summer 2016 Edition), accessed on September 5, 2025, https://plato.stanford.edu/archives/sum2016/entries/ambiguity/
2. Ambiguity (Stanford Encyclopedia of Philosophy), accessed on September 5, 2025, https://plato.stanford.edu/entries/ambiguity/
3. Fallacies of Ambiguity - Philosophy Pages, accessed on September 5, 2025, https://www.philosophypages.com/lg/e06c.htm
4. Chapter 9: Vagueness, Ambiguity, and Philosophy, accessed on September 5, 2025, https://lucidphilosophy.com/chapter-9-vagueness-ambiguity-and-philosophy/

5. Philosophical Investigations - Wikipedia, accessed on September 5, 2025, https://en.wikipedia.org/wiki/Philosophical_Investigations
6. An Analysis of Ambiguity in Humans - SBU Brooklogue, accessed on September 5, 2025, https://sbubrooklogue.com/2021/11/12/an-analysis-of-ambiguity-in-humans/
7. 23. Ambiguity and Vagueness: An Overview - Chris Kennedy, accessed on September 5, 2025, https://semantics.uchicago.edu/kennedy/docs/ambivague.pdf
8. lucidphilosophy.com, accessed on September 5, 2025, https://lucidphilosophy.com/chapter-9-vagueness-ambiguity-and-philosophy/#:~:text=Vagueness%20refers%20to%20a%20lack,choose%20which%20meaning%20was%20intended.
9. Is there a difference between ambiguity and vagueness? - Philosophy Stack Exchange, accessed on September 5, 2025, https://philosophy.stackexchange.com/questions/93473/is-there-a-difference-between-ambiguity-and-vagueness
10. Ambiguity - Wikipedia, accessed on September 5, 2025, https://en.wikipedia.org/wiki/Ambiguity
11. Vagueness - 1000-Word Philosophy: An Introductory Anthology, accessed on September 5, 2025, https://1000wordphilosophy.com/2019/03/11/vagueness/
12. Sorites Paradox (Stanford Encyclopedia of Philosophy), accessed on September 5, 2025, https://plato.stanford.edu/entries/sorites-paradox/
13. What is the difference between vagueness and ambiguity and which do you think is a bigger issue in philosophical literature? : r/askphilosophy - Reddit, accessed on September 5, 2025, https://www.reddit.com/r/askphilosophy/comments/rrtkcr/what_is_the_difference_between_vagueness_and/
14. www.thecollector.com, accessed on September 5, 2025, https://www.thecollector.com/sorites-paradox-logic-not-omnipotent/#:~:text=The%20Sorites%20Paradox%20is%20a,spectrum%20of%20real%2Dworld%20phenomena.
15. Sorites paradox - Stanford Encyclopedia of Philosophy, accessed on September 5, 2025, https://plato.stanford.edu/archives/win2007/entries/sorites-paradox/
16. Sorites paradox - Wikipedia, accessed on September 5, 2025, https://en.wikipedia.org/wiki/Sorites_paradox
17. Populations, Race, and The Sorites Paradox - Anthrop 365, accessed on September 5, 2025, https://anthropology365.com/2018/06/26/science-race-and-the-sorites-paradox/
18. Language-games: Lessons from the later Wittgenstein - Triangle PR, accessed on September 5, 2025, https://www.trianglepr.co.uk/talking-shop/2023/01/language-games-lessons-from-the-later-wittgenstein
19. Wittgenstein's Language Games and the Critique of the Natural Abstraction Hypothesis, accessed on September 5, 2025, https://www.lesswrong.com/posts/7L7Mdmnoc7xAoTwdh/wittgenstein-s-language-games-and-the-critique-of-the

20. What does Wittgenstein mean when he says "Games can't be defined"? - Reddit, accessed on September 5, 2025, https://www.reddit.com/r/askphilosophy/comments/c7r62r/what_does_wittgenstein_mean_when_he_says_games/
21. ON TEMPORARY SYNTACTIC AMBIGUITY - Studia PsyPaed, accessed on September 5, 2025, https://studiapsypaed.com/wp-content/uploads/2021/09/2-2010-12.pdf
22. Actions and Affordances in Syntactic Ambiguity Resolution - Computational Cognitive Neuroscience of Language Lab, accessed on September 5, 2025, https://magnuson.psy.uconn.edu/wp-content/uploads/sites/1140/2015/01/chambers_tanenhaus_magnuson_JEPLMC2004.pdf
23. Full article: Linking cognitive control to language comprehension: proportion congruency effects in syntactic ambiguity resolution - Taylor & Francis Online, accessed on September 5, 2025, https://www.tandfonline.com/doi/full/10.1080/23273798.2024.2314027
24. Dissociating Cognitive Processes During Ambiguous Information Processing in Perceptual Decision-Making - PMC - PubMed Central, accessed on September 5, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC7370842/
25. Cognitive Mechanisms of Ambiguity Resolution, accessed on September 5, 2025, https://rsuh.ru/education/cognitive-center/files/publications/en/4_Vera_Zabotkina,_Didier_Bottineau,_Elena_Boyarskaya._Cognitive_Mechanisms_of_Ambiguity_Resolution.pdf
26. Conceptual blending - Wikipedia, accessed on September 5, 2025, https://en.wikipedia.org/wiki/Conceptual_blending
27. Conceptual Integration - Mark Turner, accessed on September 5, 2025, https://markturner.org/cinLEA.pdf
28. CONCEPTUAL BLENDING, FORM AND MEANING1 1. Introduction - TECFA, accessed on September 5, 2025, https://tecfa.unige.ch/tecfa/maltt/cofor-1/textes/Fauconnier-Turner03.pdf
29. Conceptual blending, creativity, and music - Lawrence Zbikowski, accessed on September 5, 2025, https://zbikowski.uchicago.edu/pdfs/Zbikowski_Conceptual_blending_creativity_and_music_2018.pdf
30. Definition and Examples of Conceptual Blending - ThoughtCo, accessed on September 5, 2025, https://www.thoughtco.com/what-is-conceptual-blending-cb-1689780
31. Blending and Metaphor - Cog Sci - University of California San Diego, accessed on September 5, 2025, https://cogsci.ucsd.edu/~coulson/joe1.html
32. Gilles Fauconnier and Mark Turner, accessed on September 5, 2025, https://muldisc.files.wordpress.com/2011/03/review-by-chf-of-fauconnie-turner-2002-in-ms-distibuted-rversion.pdf
33. What is lexical ambiguity? – Microsoft 365, accessed on September 5, 2025, https://www.microsoft.com/en-us/microsoft-365-life-hacks/writing/what-is-lexical-ambiguity-and-how-to-avoid-it
34. medium.com, accessed on September 5, 2025,

https://medium.com/womenintechnology/understanding-ambiguities-in-natural-language-processing-179212a23b55#:~:text=Lexical%20Ambiguity&text=During%20Lexical%20analysis%20given%20paragraphs,is%20known%20as%20Lexical%20Ambiguity.
35. Lexical Ambiguity Definition and Examples - ThoughtCo, accessed on September 5, 2025, https://www.thoughtco.com/what-is-lexical-ambiguity-1691226
36. Ambiguity in NLP and how to address them - GeeksforGeeks, accessed on September 5, 2025, https://www.geeksforgeeks.org/nlp/ambiguity-in-nlp-and-how-to-address-them/
37. Notes on Ambiguity - NYU Computer Science, accessed on September 5, 2025, https://cs.nyu.edu/~davise/ai/ambiguity.html
38. Lexical Ambiguity Jennifer Rodd Department of Experimental Psychology, University College London - OSF, accessed on September 5, 2025, https://osf.io/jwgnf/download/?version=1&displayName=Rodd%20Lexical%20Ambiguity%20FINAL%2030March2017-2017-05-26T14%3A44%3A04.904Z.pdf
39. Semantic ambiguity - Wikipedia, accessed on September 5, 2025, https://en.wikipedia.org/wiki/Semantic_ambiguity
40. Lexical Ambiguity - Linguistics - Oxford Bibliographies, accessed on September 5, 2025, https://www.oxfordbibliographies.com/abstract/document/obo-9780199772810/obo-9780199772810-0309.xml
41. Understanding Ambiguities in Natural Language Processing | by ..., accessed on September 5, 2025, https://medium.com/womenintechnology/understanding-ambiguities-in-natural-language-processing-179212a23b55
42. Exploring the Word Sense Disambiguation Capabilities of Large Language Models - arXiv, accessed on September 5, 2025, https://arxiv.org/html/2503.08662v1
43. [2109.13032] Patterns of Lexical Ambiguity in Contextualised Language Models - arXiv, accessed on September 5, 2025, https://arxiv.org/abs/2109.13032
44. [2008.11608] Analysis and Evaluation of Language Models for Word Sense Disambiguation, accessed on September 5, 2025, https://arxiv.org/abs/2008.11608
45. Exploring the Word Sense Disambiguation Capabilities of Large Language Models - arXiv, accessed on September 5, 2025, https://arxiv.org/abs/2503.08662
46. Syntactic ambiguity - Wikipedia, accessed on September 5, 2025, https://en.wikipedia.org/wiki/Syntactic_ambiguity
47. Definition and Examples of Syntactic Ambiguity - ThoughtCo, accessed on September 5, 2025, https://www.thoughtco.com/syntactic-ambiguity-grammar-1692179
48. A Truly Joint Neural Architecture for Segmentation and Parsing - arXiv, accessed on September 5, 2025, https://arxiv.org/html/2402.02564v1
49. A Survey of Unsupervised Dependency Parsing, accessed on September 5, 2025, https://arxiv.org/pdf/2010.01535
50. 9.1 Statistical Parsing - Alchemy, accessed on September 5, 2025, https://alchemy.cs.washington.edu/tutorial/9_1Statistical_Parsing.html
51. A Taxonomy of Ambiguity Types for NLP - arXiv, accessed on September 5, 2025,

https://arxiv.org/html/2403.14072v1
52. Scope Ambiguities in Large Language Models - arXiv, accessed on September 5, 2025, https://arxiv.org/html/2404.04332v1
53. arxiv.org, accessed on September 5, 2025, https://arxiv.org/html/2502.08660v1
54. Natural Language Processing - Introduction - Tutorials Point, accessed on September 5, 2025, https://www.tutorialspoint.com/natural_language_processing/natural_language_processing_introduction.htm
55. arXiv:2501.19316v2 [cs.CL] 21 Mar 2025, accessed on September 5, 2025, https://arxiv.org/pdf/2501.19316?
56. What is pragmatic ambiguity in NLP - Homework Guru, accessed on September 5, 2025, https://homeworkguru.com/What-is-pragmatic-ambiguity-in-NLP
57. A Survey of the Usages of Deep Learning for Natural Language Processing - arXiv, accessed on September 5, 2025, https://arxiv.org/pdf/1807.10854
58. Lexicalized and Probabilistic Parsing - Scaler Topics, accessed on September 5, 2025, https://www.scaler.com/topics/nlp/lexicalized-and-probabilistic-parsing/
59. Natural Language Processing Ch5. Lexicalized and Probabilistic Parsing, accessed on September 5, 2025, https://nlpr.ia.ac.cn/cip/TeachingMaterials/%E5%86%AF%E5%BF%97%E4%BC%9F/NLPR5-zwfeng.pdf
60. Lecture 10: Statistical Parsing with PCFGs, accessed on September 5, 2025, https://courses.grainger.illinois.edu/cs447/fa2018/Slides/Lecture10.pdf
61. VLC001: Pragmatic Ambiguity, accessed on September 5, 2025, https://oer-vlc.de/mod/page/view.php?id=17028
62. Coreference Resolution In NLP Simplified [8 Powerful Techniques], accessed on September 5, 2025, https://spotintelligence.com/2024/01/17/coreference-resolution-nlp/
63. Coreference Resolution without Span Representations, accessed on September 5, 2025, https://arxiv.org/abs/2101.00434
64. Semantic Role Labeling with BERT-Based Transformers: A Step-by-Step Guide - Medium, accessed on September 5, 2025, https://medium.com/@info.codetitan/semantic-role-labeling-with-bert-based-transformers-a-step-by-step-guide-7caf82ead8d1
65. Semantic Roles in NLP - GeeksforGeeks, accessed on September 5, 2025, https://www.geeksforgeeks.org/nlp/semantic-roles-in-nlp/
66. Semantic role labeling | Natural Language Processing Class Notes - Fiveable, accessed on September 5, 2025, https://library.fiveable.me/natural-language-processing/unit-4/semantic-role-labeling/study-guide/z4x5D7HUE4tgswhp
67. Semantic Role Labeling - Stanford University, accessed on September 5, 2025, https://web.stanford.edu/~jurafsky/slp3/21.pdf
68. The advantage of ambiguity | MIT News | Massachusetts Institute of ..., accessed on September 5, 2025, https://news.mit.edu/2012/ambiguity-in-language-0119
69. arXiv:2310.14610v1 [cs.CL] 23 Oct 2023, accessed on September 5, 2025, https://arxiv.org/pdf/2310.14610

70. Are some human languages significantly less ambiguous than others?, accessed on September 5, 2025, https://linguistics.stackexchange.com/questions/32423/are-some-human-languages-significantly-less-ambiguous-than-others
71. CLAMBER: A Benchmark of Identifying and Clarifying Ambiguous Information Needs in Large Language Models - arXiv, accessed on September 5, 2025, https://arxiv.org/html/2405.12063v2
72. arXiv:2405.12063v2 [cs.CL] 1 Jun 2024, accessed on September 5, 2025, https://arxiv.org/pdf/2405.12063
73. linguistics.ucsd.edu, accessed on September 5, 2025, https://linguistics.ucsd.edu/~schane/law/ambiguity.pdf
74. Ambiguity (law) - Wikipedia, accessed on September 5, 2025, https://en.wikipedia.org/wiki/Ambiguity_(law)
75. ambiguity | Wex | US Law | LII / Legal Information Institute, accessed on September 5, 2025, https://www.law.cornell.edu/wex/ambiguity
76. How Courts Resolve Ambiguous Contract Language: business law and litigation defense services - Plunkett Cooney, accessed on September 5, 2025, https://www.plunkettcooney.com/dontbetthebusinessblog/ambiguous-contracts
77. 73. Ambiguities | United States Department of Justice, accessed on September 5, 2025, https://www.justice.gov/archives/jm/civil-resource-manual-73-ambiguities
78. A Characterisation of Ambiguity in BPM - Andrea Burattin, accessed on September 5, 2025, https://andrea.burattin.net/public-files/publications/2023-er.pdf
79. A Characterisation of Ambiguity in BPM - DTU Orbit, accessed on September 5, 2025, https://orbit.dtu.dk/files/338332110/ER2023_Ambiguity_Taxonomy.pdf
80. Business process models: criteria for quality assessment - SAP Signavio, accessed on September 5, 2025, https://www.signavio.com/post/business-process-models-quality/
81. Ambiguity management, accessed on September 5, 2025, https://www.pmi.org/learning/library/ambiguity-management-new-frontier-6285
82. Syntactic ambiguity - (Intro to English Grammar) - Vocab, Definition, Explanations | Fiveable, accessed on September 5, 2025, https://library.fiveable.me/key-terms/fundamentals-of-the-grammar-of-standard-english/syntactic-ambiguity