The Axiom of Choice: Abstract Foundations and the Algorithmic Reality of AI

I. Introduction: The Axiom of Choice and the Computational Realm of AI

A. The Axiom of Choice: A Pillar of Modern Mathematics

The Axiom of Choice (AC) stands as a cornerstone of modern set theory and, by extension, much of contemporary mathematics. Introduced by Ernst Zermelo in 1904 [1], it addresses a fundamental question regarding collections of sets: it asserts that for any given collection of non-empty sets, it is possible to form a new set that contains exactly one element chosen from each set in the original collection.[2] This seemingly intuitive statement has profound consequences, particularly when dealing with infinite collections of sets, enabling mathematicians to extend reasoning and prove the existence of mathematical objects in contexts where finite constructive methods fall short.[3]

The Axiom of Choice is not derivable from the other standard axioms of set theory, namely the Zermelo-Fraenkel (ZF) axioms. Landmark work by Kurt Gödel in 1938 and Paul Cohen in 1963 established that AC is independent of the ZF axioms.[3] This means that both ZF set theory with the Axiom of Choice (ZFC) and ZF set theory with the negation of the Axiom of Choice (ZF¬C) are logically consistent systems (assuming ZF itself is consistent). Despite some of its counterintuitive consequences, ZFC has become the standard axiomatic foundation for the vast majority of working mathematicians, largely due to the powerful theoretical tools it provides.[3]

B. LLMs and AI Agents: The Algorithmic Frontier

In parallel to the abstract world of set theory, recent decades have witnessed the explosive growth of Artificial Intelligence (AI), particularly with the advent of Large Language Models (LLMs) and sophisticated AI agents. These systems are designed to perform a wide array of complex tasks, including understanding and generating human language, engaging in multi-step reasoning, planning, and interacting dynamically with complex environments.[6] At their core, LLMs and AI agents are computational entities. Their capabilities, however advanced, are fundamentally rooted in algorithms—precisely defined sequences of operations—executed on digital computers, processing data and relying on finite (though often vast) computational resources.

C. The Central Inquiry: Bridging Abstract Axiomatics and Concrete Computation

This report addresses the central inquiry: What are the effects, if any, of a profoundly non-constructive principle like the Axiom of Choice on the theory, development, and ultimate capabilities of LLMs and AI agents? There exists an apparent dichotomy: AC is a statement about the existence of objects, often involving an infinite number of arbitrary choices, where no explicit rule for making these choices can be given. In contrast, AI systems are inherently algorithmic, operating within the realm of finite representations and computable processes. This investigation seeks to explore potential connections, whether direct or indirect, and to illuminate the foundational implications and philosophical tensions that arise when considering the influence of such an abstract axiom on a field driven by concrete computation.

D. Navigating the Terrain

The primary influence of the Axiom of Choice on AI is anticipated to be indirect. It is unlikely that AC is directly applied within AI algorithms themselves, given their constructive and finite nature.[6] Instead, any impact is more

likely to be found in the foundational mathematical theories—such as functional analysis, abstract algebra, and general topology—that provide the theoretical underpinnings for various AI concepts and techniques. AC is a statement about existence, often non-constructive, particularly for infinite sets.[1] AI algorithms, by their very nature, are constructive and operate on finite data or finite approximations of infinite concepts. Thus, a direct operational use of AC in an AI algorithm appears improbable. However, AI leverages mathematical fields that, in their most general (often infinite-dimensional) forms, rely on AC for some of their foundational theorems. This suggests an indirect, foundational influence. This report will proceed by first defining the Axiom of Choice and its key implications, then outlining the computational basis of AI. Subsequently, it will delve into specific mathematical areas where AC's influence might be discerned, and finally, discuss the broader philosophical and theoretical considerations arising from this intersection.

II. Deciphering the Axiom of Choice

A. Formal Statements and Intuitive Meaning

The Axiom of Choice can be formulated in several equivalent ways, each shedding light on its fundamental nature.

Choice Functions: The most common statement posits the existence of a "choice function." For any given collection X of non-empty sets, AC asserts that there exists a function f (a choice function) such that for every set A in the collection X, f(A) is an element of A.2 In formal logical notation, this is expressed as:
$$ \forall X \left[\varnothing \notin X \implies \exists f: X \to \bigcup_{A \in X} A \quad \forall A \in X (f(A) \in A) \right] $$
.2 This means we can simultaneously "pick" one element from each set in the collection, even if the collection is infinite and there's no discernible rule for making these choices.
Cartesian Products: An equivalent formulation states that the Cartesian product of any indexed family of non-empty sets is itself non-empty.[3] If $\{S_i\}_{i \in I}$ is a collection of non-empty sets indexed by a set I, their Cartesian product $\prod_{i \in I} S_i$ is the set of all functions $f: I \to \bigcup_{i \in I} S_i$ such that $f(i) \in S_i$ for all $i \in I$. Each such function is precisely a choice function. Thus, AC guarantees that such a product will contain at least one element (i.e., at least one choice function exists). This formulation is particularly important for establishing the existence of mathematical objects that are constructed through a sequence of simultaneous choices.

The Axiom of Choice is primarily a powerful statement when dealing with *infinite* collections of sets. If the collection of sets is finite, one can prove the existence of a choice function by induction on the number of sets, without needing to invoke AC as a separate axiom.[3]

Bertrand Russell provided a famous analogy to illustrate the subtlety of AC.[4] For any collection of pairs of shoes (even an infinite one), one can define a choice function easily: for each pair, choose the left shoe. This works because "left shoe" is a definable property. However, for an infinite collection of pairs of socks (assuming socks have no distinguishing features like "left" or "right"), there is no obvious way to define a function that selects one sock from each pair without invoking the Axiom of Choice. AC asserts that such a selection is possible even in the absence of a defining rule. The challenge often arises from the lack of a *definable* way to make the choices. If a clear rule for selection exists for each set in an infinite collection (e.g., "pick the smallest number" from sets of positive integers), AC is not needed.[1] AC becomes indispensable when no such uniform defining property or selection rule is available across an infinite number of sets.

B. Crucial Equivalents and Their Significance

The Axiom of Choice is logically equivalent to several other fundamental mathematical principles. These equivalents often appear in different branches of mathematics and are sometimes more directly applicable than AC itself. Understanding these is crucial for recognizing AC's pervasive, if sometimes hidden, influence.

The Well-Ordering Theorem:

This theorem, also known as Zermelo's theorem, states that every set can be well-ordered.2 A set X is well-ordered by a strict total order if every non-empty subset of X has a least element under that ordering. Ernst Zermelo originally introduced AC in 1904 precisely to prove the well-ordering theorem.8

The significance of this theorem is profound; for instance, it implies that any set can be subjected to transfinite induction, a powerful proof technique extending induction beyond the natural numbers.8 However, the idea of a well-ordering for certain sets, like the set of real numbers R, can be highly counterintuitive. While AC guarantees such an ordering exists, it provides no method to construct it, leading to the famous quip that the Axiom of Choice is obviously true, the well-ordering principle obviously false, and who can tell about Zorn's lemma?.8

The equivalence can be seen by sketching the proofs: AC allows one to construct a well-ordering by repeatedly choosing an element from the set of not-yet-ordered elements. Conversely, if a set $U=\bigcup_{i\in I} S_i$ (the union of sets in a collection for which we want a choice function) can be well-ordered, then a choice function f for the collection $\{S_i\}_{i\in I}$ can be defined by setting $f(S_i)$ to be the least element of S_i with respect to the well-ordering on U.8

Zorn's Lemma:

This lemma states that if $(P,\leq)$ is a non-empty partially ordered set in which every chain (a totally ordered subset) has an upper bound in P, then P contains at least one maximal element.2 An element $m\in P$ is maximal if there is no $x\in P$ such that $m<x$.

Zorn's Lemma is an indispensable tool in many areas of abstract mathematics, particularly in algebra (e.g., proving that every vector space has a basis, or that every non-trivial ring with unity has a maximal ideal) and in functional analysis. It is often more convenient to apply in proofs than AC or the well-ordering theorem directly.3

Existence of a Hamel Basis for all Vector Spaces:

A direct consequence of Zorn's Lemma (and thus AC) is the theorem that every vector space, regardless of its dimension or underlying field, possesses a Hamel basis.2 A Hamel basis is a linearly independent set of vectors that spans the entire space, meaning every vector in the space can be expressed as a finite linear combination of basis vectors.

This is a cornerstone of linear algebra. While for finite-dimensional vector spaces, bases can be constructed algorithmically, the existence of a Hamel basis for arbitrary infinite-dimensional vector spaces (like the space of all real-valued functions on R, or R considered as a vector space over Q) is non-constructive and relies fundamentally on AC.2

Tychonoff's Theorem:

In general topology, Tychonoff's theorem states that the Cartesian product of any family of compact topological spaces is compact with respect to the product topology.4

Compactness is a critical topological property, generalizing the notion of closed and bounded sets in Euclidean spaces, and is used extensively in analysis to prove the existence of limits,

optima, and fixed points. The proof of Tychonoff's theorem for an arbitrary (potentially uncountable) product of compact spaces relies on the Axiom of Choice, often through Zorn's Lemma or an equivalent formulation like the Alexander Subbase Theorem (whose general proof also uses AC).12

These equivalences highlight the far-reaching impact of the Axiom of Choice. The pragmatic acceptance of AC by the mathematical community, despite its sometimes non-intuitive consequences [3], stems from the richness and utility of the mathematical theories it enables. This pragmatism may find a parallel in how AI researchers utilize powerful mathematical frameworks: the focus is often on the practical results and predictive power afforded by these tools, rather than a deep axiomatic scrutiny, as long as they are consistent and effective.

Table 1: Key Equivalents of the Axiom of Choice

Equivalent Statement	Brief Description	Typical Area of Application
Well-Ordering Theorem	Every set can be given a total order such that every non-empty subset has a least element.	Set Theory, Transfinite Mathematics
Zorn's Lemma	In a non-empty partially ordered set where every chain has an upper bound, at least one maximal element exists.	Algebra, Functional Analysis, Topology
Existence of a Hamel Basis (for all vector spaces)	Every vector space has a linearly independent spanning set (every vector is a finite linear combination of basis vectors).	Linear Algebra, Functional Analysis
Tychonoff's Theorem	The Cartesian product of any collection of compact topological spaces is compact.	General Topology, Functional Analysis
Every surjective function has a right inverse.	If a function $f:A\to B$ is onto, there exists a function $g:B\to A$ such that $f(g(b))=b$ for all $b\in B$.	Set Theory, Category Theory
Krull's Theorem	Every unital ring (other than the trivial ring) contains a maximal ideal.	Commutative Algebra
Every connected graph has a spanning tree.	For any connected graph, there is a subgraph that is a tree and includes all vertices.	Graph Theory

C. The Non-Constructive Essence of the Axiom of Choice

The most distinctive feature of the Axiom of Choice, and the primary reason for its controversial status in some mathematical circles, is its profoundly **non-constructive** nature. AC asserts the *existence* of a choice function or a set formed by choices, but it does not provide a *method*, *rule*, or *algorithm* to explicitly define or construct that function or set, particularly when dealing with infinite collections where no uniform selection criterion is apparent.[1]

This is in stark contrast to constructive proofs in mathematics, which demonstrate the existence of an object by providing an explicit procedure for building it.[10] For example, to prove that there exists an even number greater

than 10, one can simply construct it: 12. AC, however, allows for existence proofs where no such construction is forthcoming. If we have an infinite collection of non-empty sets of, say, real numbers, and these sets have no "special" elements (like a minimum, maximum, or an element definable by a simple property), AC guarantees we can still pick one element from each, but it doesn't tell us *how* to pick them.

This non-constructive power leads to some of the most famous and counterintuitive results in mathematics, most notably the **Banach-Tarski paradox**. This theorem, a consequence of AC, states that a solid 3-dimensional ball can be decomposed into a finite number of disjoint subsets (point sets), which can then be rearranged by rigid motions (rotations and translations) to form two complete, solid balls, each identical in size to the original.[3] The "pieces" in this decomposition are so bizarrely complex that they are non-measurable; they do not have a well-defined volume in the usual sense and cannot be physically realized. Such paradoxes underscore that AC allows for the existence of objects far removed from our everyday intuition and physical experience.

The philosophical debate around AC centers on this non-constructivity. Constructivist mathematicians, who believe that mathematical objects must be explicitly constructible to be considered existent, often reject AC (or accept only weaker, constructive versions of choice).[13] However, the majority of mathematicians working in diverse fields accept ZFC (ZF plus AC) because of its immense utility and the rich theoretical landscape it opens up.[3] The consensus is that the benefits of a more powerful and encompassing theory outweigh the discomfort of its non-constructive elements and paradoxical consequences, especially since these paradoxes typically arise in highly abstract, infinitary contexts rather than in more concrete applications.

D. Insights for Section II (Recap)

The Axiom of Choice is a fundamental assertion about the possibility of making infinitely many selections simultaneously, even without a rule. Its power is manifested through various equivalent principles like the Well-Ordering Theorem and Zorn's Lemma, which underpin vast areas of mathematics. However, its non-constructive nature, asserting existence without providing a method of construction, leads to both powerful theorems and counterintuitive paradoxes, placing it at the center of historical and ongoing foundational discussions. The mathematical community largely adopts AC for pragmatic reasons—the theoretical tools it enables are deemed indispensable for a comprehensive mathematical framework.

III. LLMs and AI Agents: A Computational Foundation

Understanding the potential influence of the Axiom of Choice on Large Language Models (LLMs) and AI agents necessitates a clear view of the computational principles upon which these AI systems are built. These systems, despite their sophisticated behaviors, are fundamentally algorithmic and operate within the constraints of computation.

A. Core Algorithmic Underpinnings

Transformer Architectures (LLMs):
The dominant architecture for modern LLMs is the Transformer, introduced in the paper "Attention Is All You Need".14 Its key components and processes include:

1. **Input Processing and Embeddings:** Text is first tokenized (broken into smaller units like words or subwords). Each token is then converted into a numerical vector through an embedding layer.[14] These embeddings are learned during training and aim to capture the semantic meaning of the tokens, such that similar words have similar vector representations in a high-dimensional space.[15] All encoders in a Transformer typically receive a list of vectors of a fixed size (e.g., 512 dimensions in the original Transformer paper).[15]

2. **Positional Encoding:** Since the Transformer architecture itself does not inherently process sequential order (unlike Recurrent Neural Networks, RNNs), positional encodings are added to the input embeddings.[15] These encodings provide information about the position of each token within the sequence, allowing the model to understand word order. The original paper proposed using sine and cosine functions of different frequencies for this purpose.[15]
3. **Attention Mechanism:** The core innovation of the Transformer is the self-attention mechanism (often multi-head attention).[14] For each token in a sequence, attention allows the model to weigh the importance of all other tokens in the sequence (within a context window) when computing its new representation. This enables the model to capture long-range dependencies and contextual relationships between words effectively.[14] The signal for key tokens is amplified, while less important tokens are diminished.[14]
4. **Encoder-Decoder Structure:** Many Transformer models, especially those used for tasks like machine translation, employ an encoder-decoder structure.[15] The encoder processes the input sequence and builds a rich contextual representation. The decoder then takes this representation and generates an output sequence, token by token.[15] Other models, like GPT-style LLMs, might use only a decoder stack for tasks like text generation.

Learning Paradigms:

The training of these complex models relies on established machine learning paradigms:

1. **Backpropagation:** This is the fundamental algorithm used to train most deep neural networks, including Transformers.[17] After a forward pass where the model makes predictions, a loss function quantifies the error between the predictions and the true targets. Backpropagation then efficiently calculates the gradient of this loss function with respect to all the model's weights and biases by applying the chain rule of calculus, working backward from the output layer to the input layer.[17] These gradients are then used by an optimization algorithm, typically a variant of stochastic gradient descent (SGD), to update the model parameters in a direction that minimizes the loss.[18] Key components in this process include the choice of loss function (e.g., cross-entropy for classification), activation functions within neurons (e.g., ReLU, GELU, SwiGLU [14]), and the network's weights and biases.[17]
2. **Pre-training and Fine-tuning:** LLMs are typically first pre-trained on massive, general datasets of text in a self-supervised manner (e.g., predicting the next word or masked words).[14] This allows the model to learn general language understanding capabilities. Subsequently, the pre-trained model can be fine-tuned on smaller, task-specific datasets to adapt it for particular applications.[14]

Agentic Behavior (AI Agents):

AI agents are systems that leverage models like LLMs to go beyond simple input-output processing and engage in more autonomous, goal-directed behavior.

1. **Definition and Capabilities:** An AI agent uses an LLM (or other AI models) to reason through a problem, create a plan, and interact with external tools or APIs to accomplish a task.[7] These agents can break down complex problems into smaller, manageable steps, evaluate alternative solutions, and make informed decisions.[6]
2. **Reinforcement Learning (RL):** RL is a common and powerful paradigm for training agents.[19] In RL, an agent learns through interaction with an environment. It takes actions, observes the resulting state of the environment, and receives a reward or penalty signal.[19] The agent's goal is to learn a policy—a strategy for choosing actions—that maximizes its cumulative reward over time.[20] Key concepts in RL include the agent,

environment, state, action, reward, policy, and value function (which estimates the expected future reward from a given state or state-action pair).[19] Many RL algorithms are formalized using the framework of Markov Decision Processes (MDPs), which model sequential decision-making under uncertainty.[19]

B. The Algorithmic and Finite Nature of Current AI Implementations

A crucial characteristic of all current LLMs and AI agents is their fundamentally algorithmic and finite nature. They operate based on precisely defined algorithms—step-by-step procedures that are executed by computers.6 Even when dealing with "vast" datasets, such as the large text corpora used for pre-training LLMs 14, these datasets are finite. The models themselves, while containing billions of parameters, are finite structures. Their computational resources (processing power, memory, training time) are also finite.

Complex behaviors like multi-step reasoning or planning are achieved by breaking down tasks into sequences of computable operations.6 The entire pipeline, from input processing to output generation or action execution, is a sequence of deterministic (or pseudo-random, in the case of sampling during generation) computational steps. This inherent finiteness and algorithmic basis is a key distinction from the realm where the Axiom of Choice often operates, which involves infinite sets and potentially unspecifiable choices. The process of embedding, for instance, maps potentially complex symbolic data like words into finite-dimensional vector spaces.14 This "finitization" is a prerequisite for algorithmic processing by neural networks, placing AC, which often deals with uncountably infinite dimensions or choices, at a different level of abstraction.

C. Theoretical Capabilities: Turing Completeness of Neural Architectures

Despite their finite implementation, some neural network architectures, including Transformers and the Neural GPU, have been shown to be Turing complete under certain idealized conditions.22 Turing completeness means that these architectures have the theoretical computational power equivalent to a universal Turing machine; they can, in principle, simulate any algorithm that a classical computer can execute.

For the Transformer architecture, Turing completeness typically requires the inclusion of positional encodings (as without them, the model is order-invariant and cannot even recognize all regular languages) and, critically, the assumption of arbitrary precision for its internal numerical representations.23 If fixed-precision arithmetic is used (as is the case in all practical implementations), the Transformer is no longer strictly Turing complete in the formal sense, though it remains an extremely powerful computational model.23

The significance of Turing completeness is that it establishes a high theoretical ceiling for the computational capabilities of these models. However, it's important to note that "arbitrary precision" is itself a mathematical idealization. While this doesn't directly invoke the Axiom of Choice, it highlights that even when analyzing the theoretical limits of AI, mathematicians and computer scientists sometimes employ idealizations that approach the concerns of infinitary mathematics where AC is a central axiom. Practical AI systems, however, always operate with finite precision, within a more constrained, albeit powerful, computational framework.

D. Insights for Section III (Recap)

LLMs and AI agents are sophisticated systems built upon algorithmic foundations like the Transformer architecture,

backpropagation for learning, and principles of reinforcement learning. They are characterized by their finite nature in terms_of data, model size, and computation. While architectures like the Transformer can be theoretically Turing complete under idealized assumptions (such as arbitrary precision), their practical operation is always within the bounds of finite, constructive computation. The learning processes themselves, such as backpropagation and RL updates, are explicitly constructive, defining precise, iterative steps to refine model behavior. This algorithmic and constructive core of AI stands in contrast to the non-constructive nature of the Axiom of Choice.

IV. The Axiom of Choice in the Mathematical Arsenal of AI

While the Axiom of Choice (AC) is unlikely to be directly implemented in AI algorithms, its influence might be felt indirectly through the mathematical theories that form the bedrock of AI research and development. These include linear algebra, functional analysis, and topology.

A. Linear Algebra and Vector Spaces

Hamel Bases in (Infinite-Dimensional) Vector Spaces:
The Axiom of Choice, typically via its equivalent Zorn's Lemma, guarantees that every vector space possesses a Hamel basis.2 A Hamel basis is a set of linearly independent vectors such that any vector in the space can be expressed as a finite linear combination of these basis vectors.
For finite-dimensional vector spaces, such as Rn, bases are readily constructible (e.g., the standard basis vectors) without any need for AC. However, for infinite-dimensional vector spaces—examples include the space of all real sequences, Lp spaces common in functional analysis, or the set of real numbers R considered as a vector space over the field of rational numbers Q—the existence of a Hamel basis is a non-constructive result. One cannot, for instance, explicitly write down or define a Hamel basis for R over Q; AC merely asserts its existence.10
Relevance to Embeddings and High-Dimensional Representations in LLMs:
LLMs and other machine learning models routinely represent data—such as words, sentences, or images—as vectors in high-dimensional, but crucially, finite-dimensional vector spaces.14 The dimensionality of these embedding spaces, while large (often hundreds or thousands of dimensions), is finite. Standard linear algebra operations performed in these spaces (e.g., dot products for attention scores, matrix multiplications in neural network layers, dimensionality reduction techniques like PCA) rely on the well-understood properties of finite-dimensional linear algebra, where bases can be explicitly constructed and manipulated.
Therefore, the AC-dependent existence of Hamel bases for general infinite-dimensional vector spaces has no direct operational bearing on the computation of these finite-dimensional embeddings or the algorithms that process them.24 The "ideal object" that is a Hamel basis for an arbitrary infinite-dimensional space, guaranteed by AC, remains in the abstract mathematical realm and is not a tool used in the practical construction or manipulation of embeddings in current AI systems.

B. Functional Analysis

Functional analysis, which extends methods of linear algebra and calculus to infinite-dimensional vector spaces

(often function spaces), provides a theoretical framework for many concepts in machine learning.

The Hahn-Banach Theorem (HBT):
The Hahn-Banach Theorem is a central result in functional analysis. In one of its common forms, it states that if M is a linear subspace of a normed vector space X, any bounded linear functional defined on M can be extended to a bounded linear functional on the entire space X without increasing its norm.27 The proof of the Hahn-Banach theorem in its general form, particularly for arbitrary normed spaces (which can be infinite-dimensional), typically relies on Zorn's Lemma, and hence on the Axiom of Choice.27
The HBT also has important geometric formulations, often referred to as hyperplane separation theorems. These theorems essentially state that if two non-empty, disjoint, convex sets are given in a topological vector space, then there exists a continuous linear functional (defining a hyperplane) that separates them.28
Applications in Machine Learning:

1. **Support Vector Machines (SVMs):** The theory underlying SVMs, a powerful class of supervised learning models, is conceptually linked to the separation theorems derived from HBT.[28] SVMs aim to find an optimal hyperplane that separates data points belonging to different classes in a high-dimensional feature space. The existence of such a separating hyperplane for linearly separable data, or more generally, the mathematical framework for finding a maximal margin hyperplane, draws on concepts of convex geometry and functional analysis where HBT guarantees the existence of such separating functionals in general settings.
2. **Optimization Theory:** The Hahn-Banach theorem plays a role in optimization theory, particularly in establishing duality theorems (e.g., in linear programming) and in proving the existence of solutions to certain variational problems and optimal control problems.[27]
3. **Reproducing Kernel Hilbert Spaces (RKHS):** Kernel methods are widely used in machine learning (e.g., in SVMs, Gaussian Processes, kernel PCA). These methods implicitly map data into a high-dimensional (potentially infinite-dimensional) feature space, an RKHS, where linear methods can be applied. Functional analysis, including theorems like HBT and the Riesz Representation Theorem (whose general forms can also depend on AC), provides the rigorous mathematical foundation for RKHS theory.[31]

While the general Hahn-Banach Theorem relies on AC, its application in machine learning contexts like SVMs involves algorithms that operate on finite datasets. The SVM algorithm, for instance, solves a specific quadratic programming problem to find a particular separating hyperplane for the given data. AC's role is in guaranteeing the *general possibility* of such separations or extensions in abstract spaces, thereby providing confidence in the theoretical framework. The algorithms themselves are constructive for the specific, finite problems they address. The AC-dependent aspects provide a broad theoretical assurance, but the actual computational work in AI is typically done in settings where these general existence proofs are not directly invoked for the algorithm's execution.

C. Topology

General topology provides the language for discussing concepts like continuity, convergence, and compactness, which are relevant to learning theory and the analysis of data.

Tychonoff's Theorem and Compactness Arguments:
Tychonoff's Theorem states that the Cartesian product of any collection of compact topological spaces is itself compact in the product topology.4 Compactness is a vital

topological property; for example, a continuous real-valued function on a non-empty compact set is bounded and attains its maximum and minimum values. The proof of Tychonoff's theorem for an arbitrary (potentially uncountable) product of compact spaces is a classical application of the Axiom of Choice (often via Zorn's Lemma or equivalent statements like the Alexander Subbase Theorem).12

Potential Relevance for Learning Theory and Data Manifold Analysis:

1. **Learning Theory:** Theoretical analyses of machine learning algorithms sometimes involve arguments about the compactness of hypothesis spaces (the set of possible functions a model can learn) or parameter spaces. If such a space is compact, it can guarantee, for instance, the existence of an optimal hypothesis or that a learning algorithm will converge.
2. **Data Manifold Analysis:** In many machine learning applications, high-dimensional data is assumed to lie on or near a lower-dimensional manifold embedded in the ambient space. Topological Data Analysis (TDA) is a field that aims to infer the topological and geometric structure of these data manifolds.[32] Concepts like compactness of these manifolds could be theoretically relevant.

If the spaces considered in these theoretical analyses are infinite products of compact spaces, or if general theorems about compactness in abstract topological spaces are invoked, then Tychonoff's theorem, and thus AC, might be an implicit part of the foundational mathematics being used. However, practical machine learning applications invariably deal with finite datasets and models with a finite number of parameters. Data manifolds, if considered, are typically assumed to be locally Euclidean and are often approximated by finite structures (e.g., simplicial complexes in TDA). Thus, the direct impact of AC via Tychonoff's theorem on specific machine learning algorithms is minimal. Its role is more in ensuring the foundational soundness of the general topological concepts that might be used to frame theoretical arguments.

D. Table 2: The Axiom of Choice in Mathematical Tools Relevant to AI

Mathematical Theorem/Concept	Reliance on AC for General Proof (Infinite/Arbitrary Case)	AI-Relevant Application Area	Nature of AC's Impact on AI Application
Existence of Hamel Basis (for any vector space)	Yes, via Zorn's Lemma	Theoretical understanding of (infinite-dimensional) vector spaces; linear algebra foundations.	Foundational for the general theory of vector spaces. Not directly used for the finite-dimensional embeddings common in AI, where bases are constructive. The non-constructive Hamel bases for infinite-dimensional spaces are not algorithmically accessible.
Hahn-Banach Theorem	Yes, via Zorn's Lemma	Theory of Support Vector Machines (SVMs); Optimization theory; Foundations of Reproducing Kernel Hilbert Spaces (RKHS).	Guarantees existence of separating linear functionals or extensions of functionals in general normed spaces. AI algorithms like SVMs constructively find such hyperplanes for specific finite datasets. AC underpins the general theory, but algorithmic implementations are concrete and constructive.

Tychonoff's Theorem	Yes, for arbitrary (infinite) products	Foundational topology for learning theory (e.g., compactness of hypothesis/parameter spaces); Theoretical aspects of data manifold analysis.	Underpins general topological results about compactness. Practical AI applications deal with finite data, finite-dimensional parameter spaces, or finite approximations of manifolds, where the full power of Tychonoff's theorem (and thus AC) is not directly invoked by algorithms. It provides a robust theoretical background.
Every vector space is a projective object in **Set**	Yes	Abstract algebra, category theory (foundational).	Highly abstract; unlikely to have direct implications for current AI algorithms but contributes to the general framework of abstract mathematics that might be used in theoretical computer science.

E. Insights for Section IV (Recap)

Across these mathematical domains, a consistent pattern is evident: the Axiom of Choice is indispensable for proving the existence of certain mathematical objects or the validity of certain properties in their most general, often infinite-dimensional or arbitrarily complex, settings. However, AI applications tend to operate in finite-dimensional spaces or employ constructive methods where AC is not directly invoked by the algorithms themselves. AC provides the assurance that the broader mathematical landscape is coherent and complete in certain ways, but AI algorithms navigate specific, computable paths within this landscape. Many AC-dependent objects, like a Hamel basis for an infinite-dimensional Banach space or a non-principal ultrafilter (related to proofs of Tychonoff's theorem), are "ideal objects"—known to exist but not explicitly constructible. AI, being inherently algorithmic, must work with objects that are, at some level, constructible or finitely representable. The relationship is often: AC enables a general theorem; this general theorem provides a conceptual framework; this framework inspires algorithms for specific, often finite-dimensional and constructive cases, where AC's direct role recedes.

V. Constructivism, Computability, and the Axiom of Choice: An AI Perspective

The relationship between the Axiom of Choice and AI becomes particularly salient when considering the philosophical and practical distinctions between constructive and non-constructive mathematics, especially in light of AI's inherent reliance on computability.

A. The Constructive vs. Non-Constructive Divide in Mathematics

Mathematical constructivism is a philosophical stance asserting that to prove the existence of a mathematical object, one must provide an explicit method or algorithm for constructing that object.10 This approach emphasizes the process of creation. Consequently, constructive mathematics often rejects or restricts principles that lead to non-constructive existence proofs, such as the full Axiom of Choice and, in some interpretations (like intuitionism), the unrestricted Law of the Excluded Middle (which states that any proposition is either true or

false).13
A non-constructive proof, in contrast, establishes the existence of an object without necessarily providing a way to find or build it. Such proofs might proceed by contradiction (showing that the non-existence of the object leads to an absurdity) or by invoking powerful existence axioms like AC.10 Historically, most mathematical proofs were implicitly constructive until the late 19th and early 20th centuries, with the work of Georg Cantor on infinite sets and David Hilbert's formalist program leading to a greater acceptance and use of non-constructive methods.13
The connection between constructive proofs and algorithms is profound: a constructive proof of a statement like "there exists an x such that property P(x) holds" often directly translates into an algorithm for finding such an x.33 This algorithmic content is a key reason for the interest in constructive mathematics within computer science.

B. AI as an Inherently Constructive and Algorithmic Discipline

Artificial intelligence, in its current form, is fundamentally an algorithmic and constructive discipline. LLMs and AI agents are built upon algorithms—finite sequences of well-defined instructions that are executed by a computer.[6] Their learning processes, such as training a neural network via backpropagation [17] or an agent learning a policy through reinforcement learning updates [19], are inherently algorithmic. These are iterative, step-by-step procedures that modify the AI system's parameters or internal state in a precisely defined manner. The outputs, decisions, and actions generated by AI systems are the direct result of these computational, constructive processes. Every aspect of an AI model's operation, from data ingestion and preprocessing to inference and action generation, is ultimately reducible to a sequence of computable steps.

C. Assessing the Impact of AC's Non-Constructivity on AI Theory and Practice

Given AI's constructive nature, the non-constructive essence of AC raises questions about its relevance and impact.

Distinguishing Existence Proofs from Algorithmic Solutions:
Mathematical theorems proven using the Axiom of Choice often guarantee the existence of certain objects or solutions—for example, the existence of a winning strategy in certain infinite games, a Hamel basis for an infinite-dimensional vector space, or a well-ordering for any set. These are powerful theoretical statements, but they are not, in themselves, algorithms.34
AI systems, on the other hand, require algorithms to find solutions, make decisions, or construct representations. A proof that a solution exists can be valuable information for an AI researcher (e.g., it confirms that searching for a solution is not inherently futile). However, such an existence proof is insufficient for an AI agent that needs to compute or construct that solution.34 AI needs the "how-to," which AC often does not provide for its more abstract consequences.
Can Non-Constructive Mathematics (relying on AC) Still Inform AI Research?
Yes, but primarily indirectly. Non-constructive proofs, even those relying on AC, can play a role in guiding AI research by:

1. **Establishing Theoretical Limits or Possibilities:** They can demonstrate that certain types of solutions or structures are theoretically possible (or impossible) in general settings, which can help define the boundaries

of what AI might achieve.

2. **Characterizing Problem Spaces:** They can provide insights into the nature of complex problem spaces, even if the specific objects whose existence is proven are not themselves constructible.
3. **Motivating the Search for Constructive Algorithms:** A non-constructive proof of existence can motivate the search for constructive algorithms that work in specific, practical sub-cases or approximations relevant to AI.[34] For instance, knowing from the Hahn-Banach theorem (AC-dependent in general) that separating hyperplanes exist for disjoint convex sets provides a theoretical underpinning for developing algorithms like SVMs, which then constructively find such hyperplanes for given finite datasets. However, a direct translation of an AC-dependent proof into a practical AI algorithm is generally not feasible due to AC's non-constructive character. The "objects" whose existence AC guarantees (like a well-ordering of R or the non-measurable sets in the Banach-Tarski paradox) are often not algorithmically accessible.

The "Finite Character" Connection:
Some mathematical principles equivalent to AC, such as Tukey's Lemma, deal with collections of sets having "finite character".9 A property has finite character if a set possesses it if and only if all its finite subsets possess it. Proofs using such principles, while ultimately relying on AC for their full generality (especially for infinite sets), have a certain "computational flavor" due to their reference to finite substructures. Similarly, results like the Robertson-Seymour theorem in graph theory (which is non-constructive in parts of its proof concerning the existence of a finite set of forbidden minors) can lead to polynomial-time algorithms for certain graph properties, even if the exact constants or the forbidden minors themselves are not explicitly known.35 These scenarios illustrate a nuanced interaction where non-constructive mathematics can still point towards algorithmic possibilities, albeit sometimes without yielding a fully explicit algorithm.

The inherent non-constructivity of the Axiom of Choice makes it largely a "red herring" for direct application in the design or implementation of AI *algorithms*. AI's core operational paradigm is constructive. While AC is foundational for parts of the mathematical language and theory that AI researchers might use to describe or analyze systems at a high level of abstraction, its direct utility in crafting the step-by-step procedures that constitute AI programs is minimal. The focus of AI is on definable functions and computable processes, whereas AC is most potent precisely when such definitions or processes are not apparent for infinite collections.

Interestingly, the principles of constructive mathematics, which often *reject* AC, align well with the goals of formal verification and the development of trustworthy AI. Systems like Coq and Agda, based on constructive type theory, allow for the development of machine-checked proofs and the extraction of certified algorithms.[36] This suggests that a preference for AC-free or explicitly constructive mathematical frameworks might be more directly relevant to future research in AI safety, reliability, and explainability than the acceptance of the full, non-constructive Axiom of Choice.

Furthermore, the "Bitter Lesson" in AI, as articulated by Rich Sutton, suggests that general-purpose methods that leverage computation and learn from data tend to outperform approaches that rely heavily on human-engineered knowledge or complex, handcrafted features.[37] The Axiom of Choice is a highly abstract piece of human-derived axiomatic knowledge concerning the behavior of infinite sets. While mathematics in general provides a powerful language for abstractly framing concepts that AI can then learn [37], AC itself might be too detached from direct computational realization to be directly useful under the "Bitter Lesson" philosophy for building learning systems. Its influence remains primarily in shaping the abstract mathematical language available to theorists.

D. Insights for Section V (Recap)

AI is an inherently constructive and algorithmic field, contrasting with the non-constructive nature of the Axiom of Choice. While AC-dependent mathematics can provide theoretical context and guarantees of existence that might inspire AI research, AC itself is not directly used in AI algorithms. The need for explicit construction in AI means that the abstract existence proofs afforded by AC are generally insufficient for algorithmic implementation. The philosophical alignment of AI with constructivism suggests that principles of constructive mathematics might be more pertinent to future AI development, particularly in areas concerning reliability and verification.

VI. Theoretical Horizons: Indirect Influences and Speculative Connections

While the direct operational impact of the Axiom of Choice on current AI is limited, its existence and properties touch upon several theoretical and speculative areas that could have long-term, albeit indirect, relevance.

A. Foundational Systems and Limits of AI Reasoning

The Axiom of Choice is independent of the Zermelo-Fraenkel (ZF) axioms of set theory.3 This independence implies that there are mathematical statements that are provable in ZFC (ZF + AC) but unprovable in ZF alone, and conversely, statements provable in ZF + ¬AC (ZF plus the negation of AC) that are unprovable in ZF. This situation is reminiscent of Gödel's incompleteness theorems, which demonstrate fundamental limitations of formal axiomatic systems.

If AI systems are to evolve to the point of engaging in high-level mathematical reasoning or automated theorem proving within foundational systems like set theory, the choice of axioms (e.g., ZF vs. ZFC vs. various constructive systems) would become critically important. The axiomatic basis would define the scope of what such an AI could formally "know" or "prove." For example, an AI system based solely on ZF axioms would not be able to prove the Well-Ordering Theorem, while one based on ZFC could.

Currently, LLMs that exhibit mathematical reasoning capabilities do so primarily by recognizing patterns and structures learned from vast corpora of mathematical text and code, rather than by performing formal logical deductions from first principles like AC.37 However, a future AI dedicated to exploring the frontiers of mathematical foundations would inevitably grapple with such independence results and the role of axioms like AC in defining different "mathematical universes."

B. Paradoxical Ramifications of AC

The Axiom of Choice is famously associated with certain paradoxical or highly counterintuitive results, the most prominent being the Banach-Tarski paradox.

The Banach-Tarski Paradox: This theorem states that a solid 3-dimensional ball can be decomposed into a finite number of disjoint point sets, which can then be rearranged using only rotations and translations to form two complete, solid balls, each identical in volume and shape to the original.[3] The "pieces" involved in this decomposition are extraordinarily complex; they are non-measurable sets, meaning they cannot be assigned a volume in the standard way, and they are not physically constructible.

Relevance (or Lack Thereof) to AI Models of Physical or Data Spaces:

1. **Physical Reality:** The Banach-Tarski paradox does not apply to physical objects. Matter is not infinitely divisible in the way mathematical points in R3 are, nor can physical objects be decomposed into the kinds of pathological sets required by the theorem. AI models that interact with or simulate the physical world would operate under physical laws and constraints where such paradoxes do not manifest.
2. **Data Spaces:** Data processed by AI systems are typically discrete (e.g., pixels, tokens) or are represented as points in finite-dimensional Euclidean spaces (e.g., embedding spaces). While these mathematical spaces are continuous, the operations AI performs (such as clustering, classification, regression) do not involve the kind of arbitrary dissections and isometries over arbitrary subsets that lead to the Banach-Tarski paradox. The "sets" of data points that AI algorithms manipulate are generally "well-behaved" (e.g., definable by predicates, often finite or at least measurable). The paradox highlights the bizarre nature of the collection of *all possible* subsets of R3 if AC is assumed, particularly the existence of non-measurable sets. AI, however, typically works with subsets that are simple enough to be defined and manipulated algorithmically. Thus, the Banach-Tarski paradox, while a fascinating consequence of AC, has little to no direct bearing on the practical operation or theoretical understanding of current AI systems. It underscores that AC's primary domain of impact for such extreme results is within abstract, infinitary mathematics, distinct from the operational world of AI.

C. Hypercomputation and the Role of AC

Hypercomputation is a theoretical branch of computer science that explores models of computation exceeding the capabilities of Turing machines (as defined by the Church-Turing thesis). Such models aim to "compute" functions that are Turing-uncomputable, such as solving the Halting Problem.
Some theoretical models of hypercomputers explicitly invoke the Axiom of Choice or its consequences:

1. **Infinite Choices or Operations:** Certain hypercomputational models might assume the ability to perform infinitely many choices or computational steps in a finite amount of time (supertasks). AC could be invoked to justify the selection of an infinite sequence of actions or states.[38]
2. **Well-Orderings for Transfinite Computation:** Some models involve computation over transfinite ordinals. The Well-Ordering Theorem (equivalent to AC) guarantees that any set can be well-ordered, which could then, in theory, be used to guide a transfinite computational process. For example, enumerating an infinite set of computational states or inputs might rely on AC to establish a well-ordering for the enumeration.[38]
3. **Infinitely Long Programs or Oracles:** An infinitely long program could potentially "hard-code" an infinite number of choices, effectively bypassing the need for AC as an "oracle" but embodying its selective power within the program itself.[38]

Disconnect from Current AI:
It is crucial to emphasize that hypercomputation is a purely theoretical and highly speculative field. Current LLMs and AI agents are implemented on physical hardware that conforms to the standard model of Turing computation (or models computationally equivalent or weaker). They do not possess the kinds of infinitary capabilities hypothesized in hypercomputational models. Therefore, while AC plays a role in some of these theoretical explorations of the ultimate limits of computation, this has no direct consequence for the design, capabilities, or limitations of present-day AI systems. These "what if" scenarios for computation show AC's utility in thought experiments but do not reflect its use in current AI.

D. Insights for Section VI (Recap)

The Axiom of Choice's more profound and paradoxical consequences, such as the Banach-Tarski paradox, and its

role in speculative areas like hypercomputation, primarily concern mathematical operations or objects that are not physically or algorithmically realizable within standard computational frameworks. This reinforces the idea that AC's main sphere of influence is abstract, infinitary mathematics, which is distinct from the operational domain of current AI technologies. The independence of AC from ZF also points to fundamental limits of formal axiomatic systems, a concept that could become relevant if AI systems were to engage in deep meta-mathematical research. For now, these connections remain largely theoretical and do not impinge on the practical aspects of AI development.

VII. Conclusion: Reconciling the Abstract Infinite with Algorithmic Intelligence

The exploration of the effects of the Axiom of Choice (AC) on Large Language Models (LLMs) and AI agents reveals a nuanced relationship, characterized more by indirect foundational support than by direct operational impact.

A. Summary of Effects: Direct vs. Indirect, Theoretical vs. Practical

The direct, operational effects of AC on the algorithms, architectures, and practical implementations of current LLMs and AI agents are negligible to non-existent. These AI systems are fundamentally finite, algorithmic, and constructive in their operations. They process finite data, are represented by finite parameter sets, and execute finite sequences of computational steps.

The primary relevance of AC to the field of AI is indirect. It lies in AC's role as a foundational axiom in ZFC set theory, which underpins much of modern mathematics. Specific theorems in areas like functional analysis (e.g., the Hahn-Banach Theorem), abstract linear algebra (e.g., the existence of a Hamel basis for any vector space), and general topology (e.g., Tychonoff's Theorem) rely on AC for their proofs in the most general, often infinite-dimensional or arbitrarily complex, settings. These mathematical fields, in turn, provide the theoretical language, concepts, and tools that AI researchers and theorists may use to frame problems, develop models, and analyze system properties.

However, a critical distinction exists between the theoretical scope of these AC-dependent theorems and their practical application in AI. While AC guarantees the existence of certain mathematical objects or properties in highly general contexts (often non-constructively), AI applications typically utilize constructive methods, operate in finite-dimensional spaces, or deal with specific instances where the full, AC-dependent power of the general theorem is not required for the particular algorithm at hand. AC helps erect the vast scaffolding of modern mathematics, ensuring its structural integrity in the most general sense. AI practitioners then build specific, concrete structures (algorithms, models) within this framework, typically using tools and techniques that do not require ascending to the most abstract, AC-dependent parts of that scaffolding.

B. The Prevailing Disconnect: Infinitary Non-Constructivism vs. AI's Computational Paradigm

A fundamental disconnect persists between the core nature of the Axiom of Choice and the computational paradigm of AI. AC is most powerful and most relevant when dealing with infinite collections of sets where no explicit rule for selection can be formulated; its essence is often tied to non-constructive existence. AI, conversely, thrives on definable functions, computable procedures, and algorithmic construction. LLMs and AI agents, in their current and foreseeable instantiations, operate firmly within the domain of the finite and the algorithmically constructible. The "effects" of AC are thus largely confined to the theoretical underpinnings of the mathematical language AI researchers speak, rather than the computational sentences they write or the algorithms they

implement.

C. Future Considerations: Potential Shifts in Relevance?

Looking ahead, the relevance of AC to AI could potentially shift under certain, currently speculative, circumstances:

1. **AI in Foundational Mathematics:** If AI systems evolve to a stage where they can engage in original research in foundational mathematics, exploring the limits of axiomatic systems or proposing new axioms, then the implications of AC, its independence, and its alternatives would become directly pertinent to the AI's reasoning processes and discoveries.
2. **Hypercomputational AI:** Should AI research ever move towards realizing theoretical models of hypercomputation that explicitly leverage infinitary principles, then AC's role in such models would warrant significant discussion. However, this remains far from current technological capabilities.
3. **Constructive AI and Verifiability:** Perhaps more plausibly, the increasing demand for robust, verifiable, explainable, and trustworthy AI systems might lead to a greater emphasis on mathematical frameworks that are inherently constructive. Constructive mathematics and type theories, which often avoid or replace AC with weaker, constructive choice principles [36], offer pathways to developing AI systems with stronger formal guarantees. In this sense, the *absence* or *careful handling* of AC-like principles might become more relevant than the uncritical adoption of ZFC as the default background.

For the present, the mathematical tools directly employed in AI are predominantly those that, even if they are special cases of broader AC-dependent theories, possess constructive or finite-dimensional counterparts that are algorithmically tractable. The Axiom of Choice remains a powerful principle governing the abstract landscape of the infinite, while AI continues its journey of mastering the complexities of the finite and computable. The two domains intersect primarily at the level of foundational language, not operational methodology.

Works cited

1. Programmer's Guide To Theory - The Algorithm of Choice, accessed June 9, 2025, https://www.i-programmer.info/programming/theory/6156-axiom-of-choice-the-programmers-guide-.html
2. Axiom of Choice - GeeksforGeeks, accessed June 9, 2025, https://www.geeksforgeeks.org/axiom-of-choice/
3. Axiom Of Choice | Brilliant Math & Science Wiki, accessed June 9, 2025, https://brilliant.org/wiki/axiom-of-choice/
4. Axiom of choice - Wikipedia, accessed June 9, 2025, https://en.wikipedia.org/wiki/Axiom_of_choice
5. www.geeksforgeeks.org, accessed June 9, 2025, https://www.geeksforgeeks.org/axiom-of-choice/#:~:text=The%20Axiom%20of%20Choice%20can,to%20one%20of%20its%20elements.&text=In%20other%20words%2C%20there%20exists,from%20each%20set%20in%20X.
6. Autonomous AI Agents: Leveraging LLMs for Adaptive Decision ..., accessed June 9, 2025, https://www.computer.org/publications/tech-news/community-voices/autonomous-ai-agents
7. An Easy Introduction to LLM Reasoning, AI Agents, and Test Time ..., accessed June 9, 2025,

https://developer.nvidia.com/blog/an-easy-introduction-to-llm-reasoning-ai-agents-and-test-time-scaling/
8. Well-ordering theorem - Wikipedia, accessed June 9, 2025, https://en.wikipedia.org/wiki/Well-ordering_theorem
9. THE AXIOM OF CHOICE AND ITS IMPLICATIONS Contents 1. Introduction 1 2. The Axiom of Choice and Its Equivalents 1 2.1. The Axiom, accessed June 9, 2025, https://math.uchicago.edu/~may/REU2014/REUPapers/Barnum.pdf
10. Non-constructive proofs vs. efficient algorithms - MathOverflow, accessed June 9, 2025, https://mathoverflow.net/questions/117574/non-constructive-proofs-vs-efficient-algorithms
11. www.numberanalytics.com, accessed June 9, 2025, https://www.numberanalytics.com/blog/mastering-tychonoffs-theorem#:~:text=The%20theorem%20states%20that%20the,compact%20spaces%20from%20known%20ones.
12. Mastering Tychonoff's Theorem - Number Analytics, accessed June 9, 2025, https://www.numberanalytics.com/blog/mastering-tychonoffs-theorem
13. Constructive proof - Wikipedia, accessed June 9, 2025, https://en.wikipedia.org/wiki/Constructive_proof
14. Transformer (deep learning architecture) - Wikipedia, accessed June 9, 2025, https://en.wikipedia.org/wiki/Transformer_(deep_learning_architecture)
15. How Transformers Work: A Detailed Exploration of Transformer ..., accessed June 9, 2025, https://www.datacamp.com/tutorial/how-transformers-work
16. How do Transformers work? - Hugging Face LLM Course, accessed June 9, 2025, https://huggingface.co/learn/llm-course/chapter1/4
17. Backpropagation - Wikipedia, accessed June 9, 2025, https://en.wikipedia.org/wiki/Backpropagation
18. What is Backpropagation? | IBM, accessed June 9, 2025, https://www.ibm.com/think/topics/backpropagation
19. What is Reinforcement Learning? - AWS, accessed June 9, 2025, https://aws.amazon.com/what-is/reinforcement-learning/
20. Reinforcement Learning: What It Is, Algorithms, Types and Examples - Turing, accessed June 9, 2025, https://www.turing.com/kb/reinforcement-learning-algorithms-types-examples
21. Reinforcement Learning Basics - SmythOS, accessed June 9, 2025, https://smythos.com/developers/agent-development/reinforcement-learning/
22. www.jmlr.org, accessed June 9, 2025, https://www.jmlr.org/papers/volume22/20-302/20-302.pdf
23. On the Turing Completeness of Modern Neural Network Architectures, accessed June 9, 2025, https://arxiv.org/abs/1901.03429
24. A Bird's Eye View of Linear Algebra: The Basics | Towards Data ..., accessed June 9, 2025, https://towardsdatascience.com/a-birds-eye-view-of-linear-algebra-the-basics/
25. What Are Embeddings in Machine Learning? - Debut Infotech, accessed June 9, 2025,

https://www.debutinfotech.com/blog/what-are-embeddings-in-machine-learning
26. Understanding and Utilizing Embedding in ML: A Brief Overview - MarkovML, accessed June 9, 2025, https://www.markovml.com/blog/embedding-machine-learning
27. Hahn-Banach Theorem: Real Analysis - Number Analytics, accessed June 9, 2025, https://www.numberanalytics.com/blog/hahn-banach-theorem-real-analysis
28. The Ultimate Hahn-Banach Theorem Guide - Number Analytics, accessed June 9, 2025, https://www.numberanalytics.com/blog/ultimate-hahn-banach-theorem-guide
29. The Hahn–Banach theorem, accessed June 9, 2025, https://www.ucl.ac.uk/~ucahad0/3103_handout_6.pdf
30. Hahn–Banach theorem - Wikipedia, accessed June 9, 2025, https://en.wikipedia.org/wiki/Hahn%E2%80%93Banach_theorem
31. Key Concepts in Functional Analysis to Know for Functional Analysis ..., accessed June 9, 2025, https://library.fiveable.me/lists/key-concepts-in-functional-analysis
32. Topological Methods in Machine Learning: A Tutorial for Practitioners, accessed June 9, 2025, https://arxiv.org/html/2409.02901v1
33. Mastering Constructive Logic - Number Analytics, accessed June 9, 2025, https://www.numberanalytics.com/blog/mastering-constructive-logic
34. Do constructivists believe that non-constructive proofs may be false and need to be “confirmed”, or is constructivism simply an exercise in reformulating proofs in a more useful or more interesting way? : r/math - Reddit, accessed June 9, 2025, https://www.reddit.com/r/math/comments/1dn7he7/do_constructivists_believe_that_nonconstructive/
35. Non-constructive algorithm existence proofs - Wikipedia, accessed June 9, 2025, https://en.wikipedia.org/wiki/Non-constructive_algorithm_existence_proofs
36. Advanced Constructive Math Techniques - Number Analytics, accessed June 9, 2025, https://www.numberanalytics.com/blog/advanced-constructive-math-techniques
37. How Math Reveals the Bleeding Edge of AI | Feature | PNNL, accessed June 9, 2025, https://www.pnnl.gov/news-media/how-math-reveals-bleeding-edge-ai
38. Computation and Hypercomputation - MDPI, accessed June 9, 2025, https://www.mdpi.com/2227-7390/10/6/997
39. The Axiom of Choice and Type Theory - Stanford Encyclopedia of Philosophy, accessed June 9, 2025, https://plato.stanford.edu/entries/axiom-choice/choice-and-type-theory.html