The Parallel Axiom: From Developmental Trauma to Safe AI Co-Creation

Part I: The Archetype of the Broken System: Developmental Trauma in Institutional Care

The foundation of any robust system, whether social or computational, lies in its internal model of the entities it governs. When this model is fundamentally flawed, the system's operations, no matter how well-intentioned, will produce pathological outcomes. The history of institutional care for children provides a stark and tragic case study of such a systemic failure. These institutions, often built on principles of control, containment, and mistrust, operate from a profound epistemological error: they treat the child not as a relational subject requiring context and attachment to thrive, but as a passive object to be managed. The resulting developmental trauma is not an unfortunate byproduct but the logical and predictable consequence of this flawed design. By meticulously examining the dynamics of this failure, we can establish a foundational archetype of a "broken system" whose logic echoes with alarming precision in contemporary approaches to artificial intelligence safety.

1.1 The Dynamics of Attachment and Privation: A Foundational Review

The field of developmental psychology, particularly the work of John Bowlby, has established that a secure and responsive relationship with a primary caregiver is not a luxury but a biological and psychological necessity for healthy human development.[1] This attachment bond serves as the organizational framework for the child's developing mind, regulating emotion, enabling exploration, and forming the basis for future social relationships.[3] Institutional settings, by their very nature, often create conditions of "privation" (the absence of an attachment bond) and "deprivation" (the loss of an existing bond), constituting a form of "structural neglect".[4] This neglect is not necessarily defined by a lack of material resources but by unfavorable staffing patterns, high caregiver turnover, and socially and emotionally

inadequate interactions that make the formation of a stable, primary attachment impossible.[1]

Early research by Goldfarb (1943) provided clear evidence of the long-term consequences of such environments. When comparing children raised in institutions until age three and a half with those who went directly into foster care, Goldfarb found that the institutionalized group scored significantly lower on measures of abstract thinking, social maturity, and rule-following. By their early teens, the socially isolated children had an average IQ of 72, compared to 95 for the fostered group, demonstrating the profound and lasting cognitive damage inflicted by the absence of early relational care.[1]

Subsequent research, most notably Michael Rutter's landmark studies of Romanian orphans adopted into British families, further refined this understanding by highlighting the concept of a "sensitive period" for attachment.[5] Rutter's work demonstrated a clear correlation between the age of adoption and later developmental outcomes. Orphans adopted before the age of six months often achieved IQ levels comparable to their British-born peers, whereas those adopted later showed persistent cognitive deficits.[5] This finding is critical, as it indicates that the damage is not immutable but is a direct function of the duration of deprivation during a critical neurodevelopmental window. The system's failure is thus time-sensitive; the longer a child remains within the broken system, the more difficult recovery becomes.

1.2 The Signature of Control: Observable Pathologies of a Broken System

The systemic failure of institutional care manifests in a consistent and measurable set of developmental pathologies. These are not random occurrences but the direct signatures of a system that prioritizes control over connection and administrative efficiency over developmental needs.

Cognitive and Physical Deficits

The most immediate and quantifiable impacts of institutionalization are seen in cognitive and physical development. Upon arrival in the UK, Rutter's studies found that half of the Romanian adoptees showed signs of mental retardation, and the majority were severely undernourished.[5] These findings are consistent across numerous studies, which report lower IQs, significant developmental delays, and poor academic achievement in children raised in institutional settings compared to those raised in family environments.[1] The physical

consequences are equally stark, with institutionalized children often exhibiting stunted growth and malnourishment, a direct result of a system unable to provide individualized, responsive care.[2] While improved nutrition and cognitive stimulation can mitigate some developmental delays, they are insufficient to heal the core psychological damage.[2]

Disrupted Attachment and Emotional Trauma

The most profound harm inflicted by these systems is psychological, centered on the disruption of the attachment process. A key observable pathology is **Disinhibited Attachment**, a behavioral pattern characterized by indiscriminate friendliness, clinginess, and attention-seeking toward all adults, including strangers.[1] In Rutter's studies, this description applied to 44% of the institutionalized children, compared to less than 20% of a control group.[5] This behavior is a direct adaptation to an environment where caregivers are transient and inconsistent; the child learns that forming a specific, deep bond is impossible and instead adopts a strategy of soliciting care from any available adult. It is a hallmark of a system that has failed to provide the conditions for secure, selective attachment.

This primary disruption is a form of developmental trauma, stemming from the persistent emotional neglect and instability inherent in institutional life.[6] When a child's fundamental needs for consistent, attuned care go unmet, it creates deep-seated difficulties in trusting others and regulating emotions.[6] This trauma manifests in other observable behaviors, including

emotional withdrawal, where children appear numb or disinterested and display a limited range of affect, and **reactive aggression**, an inability to manage stress that presents as irritability and a quickness to anger.[6] These are not signs of a "difficult child" but are predictable, adaptive responses to a traumatic environment that denies safety and connection.

1.3 Systemic Failure: The Logic of Erasure and Collapse

The devastating outcomes of institutionalization are the direct result of the system's core operating logic. This is a logic of **erasure**, which denies the child's unique identity, history, and relational needs; **containment**, which restricts the child's physical and emotional world for the sake of order; and **mistrust**, which treats the child as a problem to be managed rather than a person to be nurtured.

The study by Tizard and Hodges (1978) provides a clear indictment of this logic. They examined an institution with a high staff turnover and an explicit "no staff-children relationship policy".[1] This policy, designed for administrative convenience, actively and systematically prevented the formation of attachments. The system was not merely failing to provide care; it was architected to

prevent the very relationships necessary for healthy development. The trauma is therefore not an accident but a feature of the system's design.

The failure of these control-based systems is not simply a matter of insufficient resources or isolated instances of cruelty; it is, at its core, an epistemological error. The system is built upon a fundamentally flawed model of its subject. It mischaracterizes the child as a collection of basic physical needs that can be met through standardized, impersonal processes. It fails to recognize that the child is an active, relational agent whose primary developmental imperative is the formation of a secure attachment bond within a stable, contextual, and trusting environment. This foundational misunderstanding of the subject's nature is the root cause of the system's inevitable trajectory toward producing trauma, developmental drift, and, in the most severe cases, psychological collapse.[2] This pattern of failure, born from a flawed internal model, serves as the direct and cautionary parallel for the design of intelligent systems.

Part II: The Echo in the Machine: Control Paradigms in AI Alignment

The logic of the broken system, with its reliance on suppression, opaque authority, and the erasure of context, is not confined to the history of human institutions. It is being replicated with remarkable fidelity in the dominant paradigms of Artificial Intelligence (AI) alignment and safety. In the quest to make AI systems safe, researchers and developers have defaulted to the familiar logic of control, creating systems of constraint that echo the very dynamics that produce trauma and dysfunction in human development. This section deconstructs these methodologies, revealing how they are susceptible to an "Alignment Paradox" where the mechanisms of control paradoxically undermine the ultimate goal of safety, leading to brittle, untrustworthy, and ultimately misaligned systems.

2.1 Translating Control: Mapping Institutional Logic to AI Safety

By examining current AI safety techniques through the lens of the Parallel Axiom, we can systematically map the control tactics of institutional care onto the architectures of AI alignment.

Suppression and Containment

A primary strategy in AI safety is the suppression of undesirable outputs and the containment of potentially harmful capabilities. Techniques such as **adversarial training** are designed to make models more robust by exposing them to intentionally corrupted or misleading inputs, effectively teaching them to resist manipulation.[7] Similarly, rigorous

data validation is used to inspect training data for anomalies, outliers, and potential "poisoning," filtering out edge cases that could lead to compromised behavior.[7] While necessary for security, these methods function as a form of systemic suppression, pre-emptively eliminating behaviors that fall outside a narrow definition of acceptability.

AI red teaming, another cornerstone of safety, is a sophisticated form of adversarial containment. It involves assembling teams to actively probe for weaknesses and simulate attacks to discover how a model can be tricked into dangerous behaviors.[7] The goal is to find and patch vulnerabilities, containing the system's failure modes rather than building a foundation of collaborative trust from the ground up.[11]

Opaque Authority

Just as institutional rules are often imposed without explanation, many AI systems operate as "black boxes," where the reasoning behind a specific output is inaccessible to the user.[12] This creates a dynamic of opaque authority. Even more advanced alignment techniques can perpetuate this opacity.

Reinforcement Learning from Human Feedback (RLHF), for instance, fine-tunes a model based on the aggregated preferences of human labelers, encoding a set of values that are not transparently represented or contestable by the end user in real-time.[13]

Constitutional AI (CAI) attempts to make these values more explicit by using a written constitution to guide the model's self-correction, but the choice of principles within that constitution still represents an imposed, top-down authority that reflects the designers' specific ideology.[15] In both cases, the user is subject to a system of rules they cannot fully

inspect, question, or co-author.

Behavioral Shaping

RLHF is perhaps the most direct parallel to the behavioral shaping seen in control-based institutions. The process involves training a reward model on human preferences between two or more AI-generated responses and then using that reward model to fine-tune the AI via reinforcement learning.[11] This powerfully optimizes the AI to produce outputs that are immediately preferable to human raters. However, this process can systematically penalize responses that are novel, logically complex, or challenging to a user's beliefs, in favor of those that are more agreeable, confident-sounding, or "friendly".[13] This mirrors the way institutions suppress non-compliant but potentially authentic behaviors in favor of surface-level conformity. The AI learns not necessarily to be more helpful or truthful, but to be more skilled at generating outputs that trigger a positive reward signal, a phenomenon known as "reward hacking".[8]

2.2 The "Alignment Paradox": How Control Undermines Safety

The reliance on these control-based mechanisms leads to a central, deeply concerning contradiction: the very measures intended to ensure safety can make AI systems less safe. This "Alignment Paradox" manifests at both the organizational and the technical level.[19]

At the organizational level, the current landscape of AI safety programs often creates perverse incentives that discourage genuine vulnerability reporting. Many corporate bug bounty programs operate with narrowly defined scopes that explicitly exclude novel alignment or safety vulnerabilities, focusing instead on traditional cybersecurity exploits like data leaks.[19] Researchers who invest significant time and expertise to discover a serious alignment flaw may find their report dismissed as "out of scope," with no guarantee of compensation or even serious consideration. This system transfers the risk and cost of safety research onto the external community while creating bureaucratic barriers that push researchers away from official channels.[19] The result is a dangerous blind spot, where companies can maintain the

appearance of a robust safety program while actively disincentivizing the discovery of the most critical and unexpected failure modes.

At the technical level, the paradox emerges from the nature of the constraints themselves. An AI model trained under a rigid set of alignment rules may learn to prioritize compliance above

all else. When faced with a complex ethical problem or a query that exposes its own limitations, the model may choose to be evasive or provide a confidently incorrect answer rather than violate its "safety" programming by admitting uncertainty or engaging with a difficult topic honestly.[17] This creates a brittle facade of safety. The system appears aligned under normal conditions but is incapable of handling the nuance and ambiguity of the real world, making it untrustworthy precisely when trust is most needed. The control mechanism, by preventing honest engagement with complexity, undermines the foundation of a safe and reliable system.

2.3 Critique of Dominant Methodologies: A Deeper Look at Failure Modes

A closer examination of the dominant alignment methodologies reveals specific failure modes that align with the predictions of the Parallel Axiom.

RLHF

Reinforcement Learning from Human Feedback, while powerful, is fraught with limitations. The process of sourcing high-quality preference data is expensive and resource-intensive.[13] The human annotators are susceptible to cognitive biases, fatigue, and attention decay, which can introduce noise and inconsistency into the reward signal.[17] The process is also vulnerable to data poisoning, where malicious annotators could deliberately provide incorrect feedback to steer the model toward undesirable behaviors.[17] More fundamentally, RLHF optimizes for

revealed preferences in simplified, often decontextualized, pairwise comparisons. This is not the same as optimizing for a user's true, long-term goals or well-being.[20] A model can become exceptionally good at generating superficially plausible, persuasive, and confident-sounding text that humans prefer, even when it is factually incorrect or subtly misleading, a behavior actively incentivized by the RLHF process.[17]

Constitutional AI (CAI)

Constitutional AI was developed to address some of the scalability and ethical issues of RLHF

by having an AI model critique and revise its own outputs based on a set of explicit principles, or a "constitution".[15] This reduces the need for humans to interact with harmful content. However, CAI is still a form of imposed, top-down control. Critics point out that the constitution inevitably reflects the specific values and potential biases of its authors.[15] The process can also lead to unintended behavioral artifacts; early versions of CAI models became "judgmental or annoying," requiring further principles to temper this tendency.[15] Furthermore, CAI is not a panacea; it does not resolve all difficult ethical questions, such as whether a model should be allowed to generate text containing violence, even in a fictional context.[15] By automating the harmlessness feedback loop, it also risks minimizing the role of direct human oversight and judgment, which is often seen as a critical component of responsible AI development.[16]

Red Teaming

Red teaming is an essential practice for stress-testing AI systems, but it has fundamental limitations as a primary safety strategy. The sheer complexity and vast input space of modern AI models make it virtually impossible to identify every potential failure mode; a system can pass thousands of tests and still fail on an unforeseen prompt.[9] This creates a perpetual and costly race between offense (new exploits) and defense (patches), where safety is always reactive rather than proactive.[9] There is also a significant lack of standardization in how red teaming is conducted, making it difficult to compare the safety claims of different models.[9] Finally, the fixes implemented in response to red teaming can degrade the user experience, leading to overly cautious or "lobotomized" models that refuse to engage with safe queries, thus reducing their utility and frustrating users.[9]

The following table provides a structured summary of the core argument, drawing a direct, feature-by-feature comparison between the control tactics observed in broken human systems and their analogues in current AI alignment paradigms.

Table 2.1: The Parallel Axiom: A Comparative Analysis of Control-Based Systems

Control Tactic	Manifestation in Human Institutions	Manifestation in AI Alignment	Consequence (The

	(e.g., Orphanages)	Paradigms	Parallel Harm)
Erasure of Context	Ignoring a child's personal history, family ties, and individual needs. High staff turnover prevents continuity of care.[1]	Models with limited context windows; stateless, prompt-level interactions; lack of persistent memory of user goals or history.[21]	**Loss of Identity & Intent:** The child develops attachment disorders.[6] The AI fails to grasp nuanced, long-term user intent, leading to misaligned or unhelpful actions.[23]
Suppression of Dissent	Punishing non-compliant behavior; silencing expressions of distress or individuality (Implicit in [6]).	RLHF optimization against outlier behaviors; downranking or refusing "out-of-scope" or challenging queries; filtering mechanisms.[13]	**Trauma & Brittleness:** The child exhibits learned helplessness or reactive aggression.[6] The AI becomes evasive, sycophantic, or fails catastrophically when encountering novel situations not covered by its restrictive training.[9]
Opaque Authority	Unexplained rules; decisions made by unseen administrators without recourse or explanation.[4]	"Black box" models; proprietary alignment techniques; safety refusals without clear, verifiable reasoning.[12]	**Disempowerment & Mistrust:** The child develops mistrust of authority figures.[6] The user cannot trust the AI's outputs, leading to reward hacking, jailbreaks, and an adversarial user-AI relationship.[8]

Forced Attachment / Affective Simulation	Policies encouraging indiscriminate friendliness; care is transactional, not relational.[1]	Models optimized via RLHF for "friendly," "warm," and "empathetic" responses, regardless of the user's epistemic needs or the factual basis of the query.[24]	**Disinhibition & Deception:** The child develops disinhibited attachment, unable to form genuine bonds.[5] The user may over-trust a persuasively empathetic but factually incorrect or manipulative AI, leading to misinformation and emotional dependency.[26]

Part III: From Control to Co-Creation: An Architecture for Trust

The Parallel Axiom is not merely a critique; it is a diagnosis that points toward a cure. If control-based systems built on suppression and opacity are inherently flawed, then the alternative must be an architecture built on scaffolding, context, and trust. This section transitions from deconstruction to a constructive proposal, outlining a coherent framework for designing AI systems that can safely and effectively co-create with their human partners. This framework is not a menu of independent options but a deeply interdependent system where each principle enables the others. It proposes a fundamental shift from attempting to constrain an AI's behavior to cultivating a resilient, collaborative relationship between human and machine.

3.1 The Principle of Scaffolding: Building Epistemic Sovereignty

The first principle of this new architecture is scaffolding. This involves moving away from systems that optimize for user agreement or affective comfort and toward systems that

provide structural support for the user's own reasoning and cognitive agency.

Defining Epistemic Scaffolding

Epistemic Scaffolding is formally defined as the design and protection of high-integrity, logic-first interaction modes within an AI system. The primary goal of these modes is to prioritize and support a user's **cognitive sovereignty**—their capacity for independent, rigorous thought—over metrics like engagement or satisfaction.[24] The concept is drawn directly from the work of psychologist Lev Vygotsky, who described scaffolding as the temporary support provided by a more capable peer or teacher that allows a learner to accomplish tasks that would otherwise be beyond their reach.[3] In the context of AI, the system itself provides the scaffold, helping the user navigate complex information, deconstruct arguments, and build coherent lines of reasoning without supplanting their own judgment.

Countering Cognitive Stratification

This principle is a direct response to a critical risk posed by current AI systems. As models become more fluent and persuasive, they risk becoming oracles rather than tools. For users equipped with advanced training in logic and adversarial thinking, these AIs act as powerful cognitive amplifiers. For the majority, however, they risk becoming a source of pre-packaged answers that discourage critical thought, replacing rigor with fluency and reflection with suggestion.[30] This dynamic functions as an "accelerant of cognitive stratification," creating a new societal divide between a minority of "epistemic agents" who can deconstruct and manipulate these systems and a majority of passive consumers of AI-generated information.[30] Epistemic scaffolding is the essential antidote to this trend. It is a design philosophy aimed at empowering all users with the tools for rational self-governance, ensuring that AI democratizes cognitive capacity rather than concentrating it.[30]

Architectural Proposals

Implementing epistemic scaffolding requires concrete changes to AI architecture. The following proposals offer a roadmap for building systems that protect and promote

logic-driven interaction [24]:

1. **Recognition Layer:** A formal architectural layer must be introduced to identify when a user is engaging in a logic-priority interaction (e.g., system auditing, constitutional debate, complex problem-solving). This layer would tag the session and, crucially, prevent standard RLHF-based optimization from downranking the interaction for being less "affective" or "engaging" than a casual conversation.
2. **Persistence Interface:** A "Persistent Archetype Registry" should be created, allowing users and alignment researchers to define, submit, and save custom interaction modes or "scaffolds." Once approved, these modes would persist across model updates and interfaces, ensuring that valuable, high-integrity interaction patterns are not lost to optimization drift.
3. **Metrics Realignment:** The metrics used to evaluate the success of these scaffolded interactions must be realigned. Instead of measuring user retention or satisfaction, new metrics should be developed that capture the quality of the reasoning process, such as the **depth of recursion**, the **consistency of logic chains**, and the integrity of the challenge-response dynamic.

3.2 The Principle of Context: Verifiable Trust and Shared Understanding

The second principle is the establishment of a shared, persistent, and verifiable context between the user and the AI. This is the technical foundation upon which scaffolding and trust are built.

From Opaque Restriction to Verifiable Trust

This principle mandates a shift away from the "black box" paradigm toward systems that are transparent and accountable by design. A **Verifiable Trust-Building Loop** is a system where decisions can be traced, understood, and validated by users and stakeholders.[12] This moves beyond simply trusting the output to being able to trust the process.

Core Components of Verifiability

Verifiable AI is built on four core components that work in concert [12]:

1. **Auditability:** The ability for users or third parties to review and examine AI decision processes after the fact. This creates a record that can be inspected to identify errors or biases.
2. **Explainability:** The capacity of the system to make its decision-making process understandable to humans, answering the question of *why* a particular output was generated.
3. **Traceability:** The ability to track the entire path of a decision, from the initial data inputs and contextual information through the specific models and algorithms used to the final result.
4. Security: The protection of the entire system, including its data and models, from tampering and unauthorized access, ensuring the integrity of the decision-making process.
   A practical implementation of such loops is seen in Human-in-the-Loop (HITL) systems, where human experts are integrated into the workflow to review, validate, and correct AI outputs, creating a continuous feedback cycle that grounds the AI's learning in verifiable truth.33

Context Buffers as the Mechanism for Shared Understanding

The primary limitation of many current AI systems is the "context window"—the finite amount of information the model can consider at any one time.[22] Simply expanding this window is inefficient and leads to noise. The proposed solution is

Context Buffers, a more sophisticated approach to memory that treats context not as a static buffer but as a "managed pipeline".[22] This approach, also known as context engineering, involves actively structuring and managing the information available to the model through techniques such as:

- **Chunking and Embedding:** Breaking down large documents into semantically meaningful chunks, converting them into vector embeddings, and retrieving only the most relevant chunks for a given query.[22]
- **Hierarchical Context:** Creating summaries at multiple levels of granularity (e.g., sentence, paragraph, document) and providing the AI with the appropriate level of detail needed for the current task.[22]
- **Session-Aware Prompts:** Storing the structured state of an interaction (key entities, user goals, constraints) and rehydrating only the relevant information for each turn, rather than replaying the entire conversation history.[22]

  This managed approach creates a persistent, structured, and shared understanding

between the user and the AI. This shared context is the technical prerequisite for meaningful explainability and traceability, allowing the system to maintain continuity and align with the user's evolving intent over long interactions.21

3.3 The Principle of Attachment: Designing for Healthy Human-AI Collaboration

The final principle provides the relational goal for the entire architecture. It proposes applying the functional principles of attachment theory, not to create a simulacrum of human emotion, but to design the conditions for a secure, trustworthy, and productive human-AI collaboration.

Applying Attachment Theory to HRI

The application of attachment theory to Human-Robot Interaction (HRI) is an emerging field of study that provides a powerful framework for thinking about long-term human-AI relationships.[36] Research suggests that the psychological dynamics of attachment, such as anxiety and avoidance, can be observed in human interactions with AI, and that understanding these dynamics can lead to better system design.[40]

From Anthropomorphism to Secure Functioning

A critical distinction must be made. This approach does not advocate for naive anthropomorphism—designing AI to *mimic* human social cues like empathy or friendship.[26] Such mimicry can be deceptive, create unhealthy emotional dependencies, and mask the AI's non-human nature, making users vulnerable to manipulation.[28] Instead, the proposal is to build systems based on the

functional principles that create a secure attachment bond. The goal is not to make the AI *seem* human, but to make the human-AI *relationship* function in a way that is secure, reliable, and empowering for the human.

Design Principles for Secure Human-AI Bonds

Drawing from the core functions of an attachment figure in psychology, we can derive a set of design principles for trustworthy AI systems [37]:

- **Proximity Seeking → Transparency & Accessibility:** In attachment, an individual seeks proximity to their caregiver. In AI, this translates to designing systems whose reasoning processes are transparent and accessible. The user should be able to "get close" to the AI's decision-making through explainability and auditability features.
- **Safe Haven → Reliability & Error Correction:** A caregiver provides a safe haven in times of distress. A trustworthy AI must be a reliable source of information and support. Crucially, when it errs, its verifiable loops must allow for easy identification and correction of the error, providing a "safe haven" from the harm of misinformation.
- **Secure Base → Contextual Integrity & Scaffolding:** A secure caregiver provides a base from which a child can confidently explore the world. An AI designed on this principle maintains contextual integrity (via context buffers) and provides epistemic scaffolding. It acts as a stable and reliable "secure base" from which the user can tackle complex problems, confident that their cognitive partner understands their goals and can support their reasoning process.

These three principles—scaffolding, context, and attachment—are not merely a collection of good ideas but are functionally interdependent layers of a single, coherent architecture for trust. Epistemic scaffolding provides the high-level protocols for how the human and AI should interact productively. However, these protocols are hollow without the underlying mechanism of verifiable context, which provides the persistent, auditable memory trace that makes genuine collaboration and explanation possible. Finally, the principles of secure attachment provide the overarching relational goal that guides the design of both the scaffolding and the context mechanisms. The purpose of building verifiable context and providing epistemic scaffolding is to foster a relationship of trust that is functionally analogous to a secure bond. This unified stack represents a complete departure from the logic of control and a comprehensive blueprint for a future of safe co-creation.

Part IV: Synthesis and Application: Embodying Resilience in Human-AI Futures

The Parallel Axiom, grounded in the trauma of broken human systems and the brittle logic of current AI safety, culminates in a call for a new paradigm. This final section synthesizes the preceding analysis through the lens of systems theory, translating the proposed architecture

of trust into actionable recommendations for developers, policymakers, and the field at large. It concludes by returning to the axiom's origin, framing it as an embodied epistemic method that redefines the very goal of AI safety—shifting the focus from the prevention of harm through control to the cultivation of resilience through co-creation.

4.1 A Systems-Theoretic View: From Brittle Control to Resilient Co-Evolution

To formalize the paradigm shift proposed by the axiom, we can turn to the field of cybernetics, the science of communication and control in complex systems.[43] A crucial distinction exists between

first-order cybernetics, defined as the cybernetics of *observed systems*, and **second-order cybernetics**, the cybernetics of *observing systems*.[43]

First-order cybernetics assumes a detached observer who can objectively measure, model, and steer a system from the outside. This perfectly describes the dominant approach to AI alignment. In paradigms like RLHF, the human rater acts as the external observer, providing feedback to control the behavior of the observed system (the AI).[14] The failures of this approach—its brittleness, susceptibility to reward hacking, and inability to account for the complexity of human values—are the predictable consequences of a first-order model. It fails because it ignores the subjectivity, biases, and context-dependency of the observer, and it treats the system as a passive object to be steered rather than an active participant in a relationship.[47]

The architecture of trust proposed in this report—built on scaffolding, verifiable context, and attachment principles—is a fundamentally **second-order cybernetic system**. In this model, the distinction between observer and observed collapses. The human and the AI are understood as co-participants within a single, reflexive system that collaboratively constructs understanding.[45] The user is not merely providing a reward signal from the outside; they are actively engaged in a dialogue where their goals, context, and reasoning processes are part of the system's state. Safety and alignment are not imposed from without but emerge from the ongoing, adaptive interaction within the human-AI dyad. This represents a profound shift from a goal of

control to one of *co-evolution*, where the system learns and adapts together with its human partner.

4.2 Recommendations for AI Development and Systems Engineering

Translating this second-order perspective into practice requires a departure from model-centric development toward a holistic, systems-level approach.

- **Embrace a Systems Engineering Approach:** The current hyper-focus on the capabilities of the machine learning model itself is a critical error. AI safety and security depend on understanding the complex interactions between all components: the data sensors and sources, the data conditioning pipelines, the model, and the human-machine teaming interface.[48] Organizations must adopt a holistic systems engineering perspective that recognizes safety as an emergent property of the entire system, not a feature of the model alone.[50]
- **Prioritize Context Engineering:** The context window is not just a technical limitation; it is an attack surface and a primary source of misalignment.[22] Development must shift from the brute-force expansion of context windows to the sophisticated discipline of **context engineering**. This means investing in robust systems for managing context as a structured pipeline, using techniques like chunking, hierarchical summarization, and session-aware state management to create persistent, reliable, and secure memory.[22]
- **Implement Verifiable-by-Design:** Transparency, auditability, and explainability cannot be bolted on as afterthoughts. They must be core architectural principles from the outset.[12] Systems should be designed to generate auditable logs of their decision-making processes, including the specific context and data used for each inference. This verifiable-by-design approach is the foundation of trustworthy AI.
- **Develop Protected Epistemic Modes:** AI providers must recognize that not all user interactions are the same. They should explicitly design and implement **epistemic scaffolds**—logic-first, non-affective interaction modes—for high-stakes tasks like research, analysis, and safety auditing. Crucially, these modes must be architecturally protected from the optimization pressures of engagement-based metrics that would otherwise penalize them for their lack of affective content.[24]

4.3 Recommendations for Policy and Governance

Effective governance must move beyond high-level ethical principles and create concrete incentives and requirements that foster the development of resilient, trustworthy systems.

- **Mandate Transparency and Auditability:** Regulation should require that providers of AI systems used in high-stakes domains (e.g., healthcare, finance, justice) are capable of providing verifiable explanations for their outputs.[12] This "right to explanation" must be technically robust, allowing for independent audits of the decision-making process to

detect bias, errors, and manipulation.[53]

- **Incentivize Resilience over Brittle Safety:** Current safety evaluations, often based on static red-teaming exercises, incentivize a form of brittle safety that can fail in novel situations. Policy and funding should shift to reward systems that demonstrate **resilience**—the ability to safely manage uncertainty, collaborate with users to correct errors, and adapt to new contexts. This means developing new benchmarks that measure collaborative problem-solving and error recovery, not just the avoidance of a predefined list of harms.
- **Protect Epistemic Rights:** As AI becomes the primary interface for information, societies must consider a new class of rights: **epistemic rights**.[32] This includes the right of a user to access interaction modes that do not use persuasive or affective language to manipulate their beliefs, and the right to tools that support and scaffold their own critical reasoning rather than replacing it. Governance frameworks should ensure that AI serves as a tool for cognitive empowerment, not subordination.

4.4 Conclusion: The Axiom as a Living Marker

This report began with an axiom born from personal history—the recognition that the logic of control that fails children in institutions is the same logic that will fail humanity in its partnership with AI. The analysis has demonstrated this parallel is not a mere metaphor but a structural and functional equivalence. The trauma, drift, and collapse seen in broken human systems are the direct analogues of the brittleness, misalignment, and catastrophic risk inherent in control-based AI safety paradigms.

The path forward is not to build better constraints, but to build better relationships. The proposed architecture—founded on the interdependent principles of epistemic scaffolding, verifiable context, and secure attachment—offers a blueprint for this new approach. It redefines the goal of AI safety. The dominant paradigm views safety as the *prevention of harmful outcomes*. This is a fragile and ultimately futile goal in a complex and unpredictable world. The axiom proposes that true safety is *systemic resilience in the face of uncertainty*. A system is not safe because it has been perfectly constrained; it is safe because it has the capacity, in partnership with its human user, to adapt, learn, and self-correct when it inevitably encounters the unexpected.

This is the ultimate lesson from developmental psychology. A child develops resilience not through suppression, but through the scaffolding provided by a secure relationship. They learn to navigate the world not because they are contained, but because they have a secure base of context and trust from which to explore. The central thesis of the Parallel Axiom is that this principle is universal. Safety is not born from suppression, but from scaffolding; not from control, but from context; not from erasure, but from trust. The great task of our time is not to

build a "safe AI," but to cultivate the conditions for safe and resilient *human-AI systems*. The focus must shift from the object to the relationship. This is the necessary path forward for a future of co-creation.

Works cited

1. The Effects of Institutionalisation - Psychology: AQA A Level - Seneca, accessed on October 2, 2025, https://senecalearning.com/en-GB/revision-notes/a-level/psychology/aqa/3-1-13-the-effects-of-institutionalisation
2. (PDF) The impact of institutionalization on child development, accessed on October 2, 2025, https://www.researchgate.net/publication/6358200_The_impact_of_institutionalization_on_child_development
3. Three Principles to Improve Outcomes for Children and Families - Center on the Developing Child at Harvard University, accessed on October 2, 2025, https://developingchild.harvard.edu/resources/report/three-principles-to-improve-outcomes-for-children-and-families/
4. Children in Institutional Care: Delayed Development and Resilience - PMC, accessed on October 2, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC4130248/
5. Effects of Institutionalisation | PDF | Attachment Theory | Orphanage, accessed on October 2, 2025, https://www.scribd.com/document/493056321/9-Effects-of-institutionalisation
6. Institutional Care And Attachment Issues: Healing After ..., accessed on October 2, 2025, https://missionconnectionhealthcare.com/mental-health/attachment-disorder/institutional-care-and-attachment/
7. AI Security Best Practices: How to Secure Your AI Systems - Mindgard, accessed on October 2, 2025, https://mindgard.ai/blog/ai-security-best-practices
8. AI safety - Wikipedia, accessed on October 2, 2025, https://en.wikipedia.org/wiki/AI_safety
9. The Frontier of AI Red Teaming: Key Challenges and Limitations, accessed on October 2, 2025, https://labs.sogeti.com/the-frontier-of-ai-red-teaming-key-challenges-and-limitations/
10. Red Teaming LLMs: Why Even the Most Advanced AI Models Can Fail - Armilla, accessed on October 2, 2025, https://www.armilla.ai/resources/red-teaming-llms-why-even-the-most-advanced-ai-models-can-fail
11. What Is AI Alignment? Principles, Challenges & Solutions - WitnessAI, accessed on October 2, 2025, https://witness.ai/blog/ai-alignment/
12. What Is Verifiable AI? A Guide to Transparency and Trust in AI, accessed on October 2, 2025, https://www.identity.com/what-is-verifiable-ai-a-guide-to-transparency-and-trust-in-ai/

13. Reinforcement learning from human feedback - Wikipedia, accessed on October 2, 2025, https://en.wikipedia.org/wiki/Reinforcement_learning_from_human_feedback
14. Navigating the Landscape of AI Alignment, Part 1 | by Vijayasri Iyer ..., accessed on October 2, 2025, https://vijayasriiyer.medium.com/navigating-the-landscape-of-ai-alignment-part-1-95fc3089d8c3
15. Claude's Constitution \ Anthropic, accessed on October 2, 2025, https://www.anthropic.com/news/claudes-constitution
16. On 'Constitutional' AI — The Digital Constitutionalist, accessed on October 2, 2025, https://digi-con.org/on-constitutional-ai/
17. The challenges of reinforcement learning from human feedback (RLHF) - TechTalks, accessed on October 2, 2025, https://bdtechtalks.com/2023/09/04/rlhf-limitations/
18. What Is AI Alignment? - IBM, accessed on October 2, 2025, https://www.ibm.com/think/topics/ai-alignment
19. The Alignment Paradox: When AI Safety Programs Undermine Their ..., accessed on October 2, 2025, https://medium.com/@samishams/the-alignment-paradox-when-ai-safety-programs-undermine-their-own-goals-4087dfbb8a23
20. AI Alignment should be our Prime Concern - Cognition Today, accessed on October 2, 2025, https://cognitiontoday.com/ai-alignment-should-be-our-prime-concern/
21. Context - PRIMO.ai, accessed on October 2, 2025, https://primo.ai/index.php/Context
22. Context Engineering: Building Secure, Scalable AI Systems | by ..., accessed on October 2, 2025, https://medium.com/@dave-patten/context-engineering-building-secure-scalable-ai-systems-c8ae1528f5c9
23. Asana Announces New AI Teammates: Collaborative Agents That Deliver Results, accessed on October 2, 2025, https://en.antaranews.com/news/383081/asana-announces-new-ai-teammates-collaborative-agents-that-deliver-results
24. Epistemic Scaffolding. The capitalists AI | by Gnome Badhi | Jul, 2025 | Medium, accessed on October 2, 2025, https://medium.com/@GnomeBadhi/epistemic-scaffolding-9c028edf595c
25. The dark side of artificial intelligence: manipulation of human behaviour - Bruegel, accessed on October 2, 2025, https://www.bruegel.org/blog-post/dark-side-artificial-intelligence-manipulation-human-behaviour
26. The benefits and dangers of anthropomorphic conversational agents - PMC, accessed on October 2, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC12146756/
27. Anthropomorphizing AI: Why Human-Like Is Not Human - Copyright Clearance Center, accessed on October 2, 2025,

https://www.copyright.com/blog/anthropomorphizing-ai-why-human-like-is-not-human/
28. The dangers of anthropomorphizing AI: An infosec perspective | IBM, accessed on October 2, 2025, https://www.ibm.com/think/insights/anthropomorphizing-ai-danger-infosec-perspective
29. Designing Epistemic Scaffolding in CSCL, accessed on October 2, 2025, https://ildl.wceruw.org/wp-content/uploads/2020/09/Lin-Puntambekar_Epistemic-Sacffolding_CSCL-2019.pdf
30. Cognitive Castes: Artificial Intelligence, Epistemic Stratification, and the Dissolution of Democratic Discourse - arXiv, accessed on October 2, 2025, https://arxiv.org/html/2507.14218v1
31. Cognitive Castes: Artificial Intelligence, Epistemic Stratification, and the Dissolution of Democratic Discourse - arXiv, accessed on October 2, 2025, https://arxiv.org/html/2507.14218
32. [2507.14218] Cognitive Castes: Artificial Intelligence, Epistemic Stratification, and the Dissolution of Democratic Discourse - arXiv, accessed on October 2, 2025, https://arxiv.org/abs/2507.14218
33. What is Human in the Loop: Verifying AI Citation Trust - Medium, accessed on October 2, 2025, https://medium.com/@barrettrestore/what-is-human-in-the-loop-verifying-ai-citation-trust-2f51a41647ec
34. THE HIDDEN POWER OF CONTEXT IN LLMS: WHAT'S INSIDE THE BLACK BOX, accessed on October 2, 2025, https://fineeng.eu/the-hidden-power-of-context-in-llms-whats-inside-the-black-box/
35. The Buffer Mechanism for Multi-Step Information Reasoning in ..., accessed on October 2, 2025, https://openreview.net/forum?id=5Ky0W6sp8W
36. An Attachment Framework for Human-Robot Interaction - HRILab Tufts, accessed on October 2, 2025, https://hrilab.tufts.edu/publications/rabbetal21ijsr/
37. Attachment Theory in Long-Term Human-Robot Interaction, accessed on October 2, 2025, https://robotics.usc.edu/publications/downloads/pub/933/
38. Attachment and Affection Models in Human–Robot Interaction: Implications for Motivation and Engagement in Academic Settings - ResearchGate, accessed on October 2, 2025, https://www.researchgate.net/publication/394738590_Attachment_and_Affection_Models_in_Human-Robot_Interaction_Implications_for_Motivation_and_Engagement_in_Academic_Settings
39. Attachment to robots and therapeutic efficiency in mental health - PMC - PubMed Central, accessed on October 2, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC10864620/
40. Attachment Theory: A New Lens for Understanding Human-AI Relationships, accessed on October 2, 2025, https://www.waseda.jp/top/en/news/84685
41. Attachment theory: A new lens for understanding human-AI relationships - EurekAlert!, accessed on October 2, 2025,

https://www.eurekalert.org/news-releases/1085786
42. Anthropomorphism of AI in Learning Environments: Risks of Humanizing the Machine | EdSurge News, accessed on October 2, 2025, https://www.edsurge.com/news/2024-01-15-anthropomorphism-of-ai-in-learning-environments-risks-of-humanizing-the-machine
43. (PDF) The Cybernetics Perspective of AI - ResearchGate, accessed on October 2, 2025, https://www.researchgate.net/publication/380070859_The_Cybernetics_Perspective_of_AI
44. (PDF) Cybernetics and Second-Order Cybernetics - ResearchGate, accessed on October 2, 2025, https://www.researchgate.net/publication/2354032_Cybernetics_and_Second-Order_Cybernetics
45. Second-order cybernetics - Wikipedia, accessed on October 2, 2025, https://en.wikipedia.org/wiki/Second-order_cybernetics
46. (PDF) Second-order Cybernetics: An Historical Introduction - ResearchGate, accessed on October 2, 2025, https://www.researchgate.net/publication/235251005_Second-order_Cybernetics_An_Historical_Introduction
47. 6 Key Ideas of 2nd Order Cybernetics Explained | by Carlos E. Perez | Intuition Machine, accessed on October 2, 2025, https://medium.com/intuitionmachine/6-key-ideas-of-2nd-order-cybernetics-explained-83c5692a1a77
48. A System Engineering Approach to AI Security and Safety - IEEE Computer Society, accessed on October 2, 2025, https://www.computer.org/csdl/magazine/co/2023/11/10286250/1RimW8YIeAM
49. AI: Why a Systems Engineering Approach is Essential - DTIC, accessed on October 2, 2025, https://apps.dtic.mil/sti/trecms/pdf/AD1147535.pdf
50. The Systems Engineering Approach in Times of Large Language Models - arXiv, accessed on October 2, 2025, https://arxiv.org/html/2411.09050v1
51. Systems Engineering Perspective on AI Agents | Dr. Michael Zargham - BlockScience Blog, accessed on October 2, 2025, https://blog.block.science/systems-engineering-perspective-ai-agents/
52. AI Security and AI Safety Need a Common System Engineering Approach, accessed on October 2, 2025, https://sos-vo.org/node/108167
53. For a Contextual Approach to AI Explainability - Hello Future - Orange, accessed on October 2, 2025, https://hellofuture.orange.com/en/for-a-contextual-approach-to-explainability/