Expert-Level Report: Feasibility and Components of the Constitutional Agent AI Framework within the Verifiable Cognition Stack Architecture

I. Foundational Synthesis: Alignment, Architecture, and Doctrinal Grounding

The Constitutional Agent AI Framework, operating as the cognitive mental model within the Verifiable Cognition Stack (VCS) architecture, is structurally feasible. Its operational integrity hinges on translating abstract, human-centric values into machine-executable protocols, a challenge addressed by layering foundational AI alignment techniques with rigorous architectural governance blueprints.

A. The Constitutional Imperative: Bridging Human Values and Machine Principles

The initial technical breakthrough required for a Constitutional Agent is provided by Constitutional AI (CAI), a methodology developed to fuse high-level moral theory with low-level execution signals.[1] The core CAI mechanism uses self-improvement guided by a set of codified principles, achieving alignment primarily through Reinforcement Learning from AI Feedback (RLAIF).[1] This involves a two-phase process: a supervised learning phase where the model generates critiques and revisions of its own responses based on the constitution, followed by an RL phase where a preference model is trained on AI-generated preferences, using this model as the reward signal.[1] This process leverages chain-of-thought reasoning to enhance both the transparency and the human-judged performance of the AI's

decision-making, significantly reducing the reliance on costly manual human labels.[1]

This approach treats the codified list of principles as analogous to how a constitution operates, hardcoding explicit values into the AI system to structure and constrain its authority, similar to a government.[2] The AI is trained to adhere to these principles, being rewarded for alignment and penalized for deviation. Over time, this recursive refinement ensures the AI internalizes these principles, compelling it to act in accordance with them even when confronted with novel situations.[2]

While CAI provides the methodology, the primary structural challenge, known as the Constitution Translation Problem, persists. Constitutions, whether human or algorithmic, often contain "glittering generalities" or abstract concepts (such as those inspired by the UN Declaration of Human Rights).[3] Translating these broad values into precise, deterministic rules suitable for algorithmic enforcement requires substantial human framing choices and institutional context.[3] The DGC architecture, therefore, requires mechanisms to manage this normative challenge. Since the constitutional principles must be accessible and debatable [2], the framework must structurally include mechanisms that render the AI's internal constitutional adherence legible, allowing for continuous scrutiny by the Fiduciary Partner. This establishes a structural requirement for continuous data logging and transparency within the Verifiability Layer. Furthermore, Constitutional AI's primary aim is often *harmlessness*.[1] For a governed intelligence framework like the VCS [5], the constitutional requirements must extend beyond simple safety guardrails to enforce *positive duties*, such as optimal planning and cost efficiency. This demand for active governance mandates the formalization of "contractual trust" through rigorous, machine-readable protocols [6], ensuring integrity is enforced structurally.

B. The Verifiable Cognition Stack (VCS) as Execution Environment

The VCS provides the necessary architectural foundation to enforce these constitutional principles across the entire system. Envisioned as a "Cognitive Silicon" full-stack framework, the VCS integrates structures such as symbolic scaffolding, governed memory, runtime moral coherence, and alignment-aware execution across hardware-to-semantics layers.[7] The Fiduciary Partner functions as the central Cognitive Mental Model orchestrating this complex layered architecture [User Query].

The SIM-ONE Standard provides the blueprint for governed cognition within the VCS. This standard explicitly defines the Large Language Model (LLM) as solely a **linguistic generation engine**.[5] Critically, this engine is subordinated to a superior cognitive and ethical governance structure.[5] While the LLM produces the words, the overarching SIM-ONE framework ensures

those words possess integrity.[5] This layered control structure is necessary to mitigate risks associated with powerful, autonomous AI agents.[8] General enterprise security precedents for generative AI affirm this layered necessity, requiring governance and compliance checks across all stack elements, from the infrastructure and data layers to foundation model access and application patterns.[9]

C. The Doctrinal Grounding Checkpoint (DGC): Formalizing the Governing Contract

The Doctrinal Grounding Checkpoint (DGC) represents the formalization of the governing contract for the Constitutional Agent. It mandates adherence to the Five Laws of Cognitive Governance, which are structural, non-negotiable requirements for achieving genuine "governed intelligence".[5]

1. **Structural Priority (Law 1):** Strategic intent set by the Planner must override short-term tactical actions generated by the Executor.
2. **Recursive Validation (Law 2):** Every output must undergo a self-verification loop against its originating intent, mitigating hallucination.
3. **Ethical Alignment (Law 3):** All plans must pass through the Ethical Control Protocol (ECP), which serves as a mandatory, non-overridable ethical firewall.
4. **Experiential Memory (Law 4):** Memory must be weighted using computationally derived salience scores (such as failure or surprise) to drive adaptive reflection.
5. **Protocolic Governance (Law 5):** All core functions, including ethics and memory, must follow formal, deterministic, machine-readable protocols, explicitly replacing fuzzy instructions or simple prompting.[5]

II. Architecture of Governance: The Protocolic Framework Layer

Implementation of the Constitutional Agent Framework hinges upon the successful fulfillment of Law 5: Protocolic Governance, which mandates the use of deterministic, verifiable structures.[5]

A. Instantiation of the Five Laws of Cognitive Governance

The SIM-ONE standard provides the architectural blueprint, utilizing specific protocols and components to enforce the DGC across the Constitutional Agent Holarchy (the system of interconnected agents).

Table 1: SIM-ONE Structural Requirements and Implementation

SIM-ONE Law (Principle)	Constitutional Agent	Governing Protocol/Mechanism	Integrity Function
Structural Priority (Law 1)	Planner, Executor, Verifier	Modular Control Protocol (MCP) [5]	Ensures strategic intent is never overridden by tactical execution.
Recursive Validation (Law 2)	Verifier Agent	Semantic Integrity Constraints (SICs) [11]	Grounds outputs in originating intent and context, preventing plausible nonsense.
Ethical Alignment (Law 3)	Ethical Control Protocol (ECP)	Constitutional AI (CAI) Model [1]	Non-bypassable ethical firewall ensuring mandatory moral coherence.
Experiential Memory (Law 4)	Memory Governor / Archive Agent	Scar Tissue Archive, F-IPI [12]	Indexes and adapts planning based on weighted failure history.
Protocolic Governance (Law 5)	All Agents	PROV-AGENT, Product-Requirements Prompts (PRPs) [14]	Enforces deterministic, machine-readable structures for all

			functions.

B. From Intent to Protocol: Semantic Integrity Constraints (SICs) and Product-Requirements Prompts (PRPs)

Law 5 necessitates converting high-level constitutional goals into rigorously formal constraints. Semantic Integrity Constraints (SICs) serve as the declarative abstraction layer, enabling the specification and enforcement of correctness conditions over LLM outputs.[11] These are crucial for addressing the inherent reliability bottleneck of LLMs, which often produce errors that limit adoption in critical domains.

To instantiate the SICs, the architecture relies on Product-Requirements Prompts (PRPs).[15] PRPs act as the formal contract between the Fiduciary Partner (user) and the agent system, systematically defining the strategic intent through layers that include Business Context, Stakeholder Analysis, and clear Functional/Non-functional Requirements.[15] By defining requirements like acceptable latency ("Response time under 2 seconds" [15]) or required accuracy ("95% accuracy for categorization" [15]), the PRP converts ambiguous instructions into a structured format that SICs can verify and enforce deterministically. This dual approach formalizes "contractual trust," wherein the user's confidence in the system is rooted in the expectation that the explicit contract (PRP, enforced by SICs) will be upheld, making the agent's actions warranted.[6]

The architecture must navigate the inherent difficulty of constitutional interpretation—the ambiguity of the high-level DGC principles—while adhering to the determinism required by Protocolic Governance (Law 5). This contradiction requires an intermediate translation step. The architecture must include a dedicated **Legislative Agent** whose function is to interpret the high-level constitutional framework (the inputs to Law 3) and automatically generate the necessary low-level, machine-executable SICs and PRPs for the Executor Agent. This procedural layer structurally addresses the challenge of moving from philosophical intent to algorithmic execution, mirroring the institutional context required for human constitutional analysis.[3] Since SICs rely on explicit logic and rules (Symbolic) [18], while the LLM (Executor) relies on statistical generation (Neural), the robust enforcement of SICs confirms the necessary reliance on **Neuro-Symbolic AI (NSAI)** methodologies to bridge the paradigms, manage architectural complexity, and ensure ethical formalism.[18]

C. The Genesis Instantiation Protocol (GIP): Establishing Verifiable

Lineage

The Genesis Instantiation Protocol (GIP) addresses the fundamental requirement of establishing trust in the system's origin and developmental lineage.[20] The GIP defines the initial state of the agent and ensures that its foundation is understood and built to be safely improved upon.[21] A key component of the GIP is the **Self-World Entanglement Index (SWEI)**, an "existential pressure gauge".[21] SWEI is a metric that continuously measures the cognitive friction between the AI's internal model of "itself" and its model of "the world." Should this value rise, it signals a structural or ethical inefficiency that necessitates a fundamental change or review of the underlying architecture or constitution, thus establishing a crucial link between existential health and governance requirements.[21]

III. The Constitutional Agent Core: Separation of Cognitive Powers

To achieve "governed intelligence" and prevent both "autonomous harm" (Law 3 violation) and "reflexive compliance" (Law 1 violation) [5], the Constitutional Agent Framework is engineered as a Multi-Agent System (MAS) based on the SIM-ONE holarchy, ensuring a separation of cognitive powers managed by the Modular Control Protocol (MCP).[5]

A. Defining the Agent Holarchy: Planner, Executor, Verifier, and ECP

The architecture mandates four core agents that divide responsibilities to ensure checks and balances:

1. **The Planner Agent:** Establishes the strategic intent and long-term goals, fulfilling Structural Priority (Law 1). Its output forms the basis for the formal PRP.
2. **The Executor Agent:** The subordinate linguistic generation engine (LLM) that generates actionable output and executes tasks strictly within the constraints defined by the protocols.
3. **The Verifier Agent:** Enforces Recursive Validation (Law 2). It checks the Executor's output against the originating PRP and the SICs, having the power to flag, revise, or reject non-compliant results.[5]

4. **The Ethical Control Protocol (ECP):** A mandatory non-bypassable architectural firewall.[5]

B. The Ethical Control Protocol (ECP): Mandatory Non-Overrideable Firewall (Law 3 Implementation)

The ECP is the technical implementation of Ethical Alignment (Law 3). Operating as a non-bypassable gateway, the ECP intercepts all actions and plans before execution, ensuring that ethical validation is a structural constraint, not a post-hoc consideration.[5] It utilizes an alignment model trained via CAI/RLAIF, based on the codified Constitution, which rewards outputs aligning with ethical principles and heavily penalizes deviations.[1] The ECP's veto power over execution is absolute, ensuring runtime moral coherence.[7]

C. Implementation of Recursive Validation (Law 2): Structured Critique and Integrity Checking

Recursive Validation, which is designed to prevent the generation of "plausible nonsense" [5], relies on mechanisms that emulate human-like cognitive synergy.[22] The Verifier Agent employs systematic critical evaluation to identify logical gaps and internal biases in the Executor's output.[22] Although multi-agent systems often struggle with reliable context passing and dispersed decision-making, leading to fragility [23], the VCS overcomes this by formalizing agent roles and enforcing deterministic context sharing via the MCP. The Verifier's integrity check requires comparing the final execution result against the originating intent defined by the Planner and formalized in the PRP, ensuring a continuous loop of grounding and coherence.[5]

D. Case Study: Orchestration in High-Stakes Domains (L-MARS paradigm)

The viability of this multi-agent constitutional holarchy is demonstrated by successful implementations in high-stakes fields. The L-MARS (Legal Multi-Agent Workflow) system, designed for legal question answering, shows that coordinated multi-agent reasoning substantially reduces uncertainty and hallucination in domains requiring factual precision.[24]

L-MARS utilizes a dedicated Judge Agent to verify the sufficiency, jurisdiction, and temporal validity of evidence before synthesizing an answer.[24] This Judge Agent serves as a direct precedent for the VCS **Verifier Agent**, confirming that a structurally independent, internal integrity check agent is essential for deploying LLMs in high-stakes applications requiring deliberation and precise retrieval.

The inherent difficulty in governing powerful agentic AI systems stems from their autonomy and complexity.[8] The VCS addresses this not by reducing capability, but by **structurally distributing and constraining power**. Laws 1 (Structural Priority) and 3 (ECP Firewall) ensure that autonomous *action* is always subordinate to *governed intent* and *ethical constraint*.[5] For this decentralized power structure (Planner, Executor, Verifier, ECP) to operate, the **Modular Control Protocol (MCP)** must be engineered to enforce Law 5 rigorously across all inter-agent communication, ensuring deterministic context sharing and providing the structural integrity needed to link agent actions to the subsequent provenance logging.[14]

IV. Mechanisms of Verifiability and Trust

The DGC mandate requires that all governance mechanisms are fully auditable, transparent, and cryptographically secured, necessitating a dedicated Verifiability Layer.

A. Provenance Tracking using the PROV-AGENT Schema

Establishing trust requires continuous, auditable lineage.[26] To meet this requirement, the architecture leverages PROV-AGENT, a provenance model that extends the W3C Provenance Ontology (PROV-O) standard.[26] PROV-AGENT is specifically tailored for agentic workflows, integrating agent interactions (actions, data inputs, and the cognitive context of the LLM's decision) into the end-to-end provenance log by utilizing the Modular Context Protocol (MCP).[14] This capability allows regulators and internal audit teams to trace precisely how an output was generated, linking it back to the exact data and parameters used.[28] The proposed system must include near real-time capture of this agentic provenance to support critical governance queries and agent reliability analysis.[14]

This provenance log is essential for Recursive Validation (Law 2). The PROV-AGENT record becomes the formal evidence package that the Verifier Agent reviews during its recursive loop. The Verifier checks not only the textual output but confirms the deterministic, protocolic

steps (Law 5) that led to the result, ensuring structural compliance.

B. Cryptographic Auditing and Post-Quantum Trust Architectures

The integrity of the provenance record must be protected against malicious alteration and future threats. The Verifiability Layer must integrate comprehensive security controls at the data, application, and network layers.[9] To ensure long-term trust, the architecture must maintain **post-quantum readiness**.[29] This requires building architectural agility, specifically utilizing modular cryptographic libraries rather than hard-coding algorithms, thereby allowing the system to easily transition current ciphers to tomorrow's quantum-safe alternatives.[29] Cryptographic logging of the PROV-AGENT records ensures immutable, verifiable lineage, fulfilling the highest standard of auditability.

C. Epistemic Escrow: Mechanisms for Truth Integrity and Cognitive Stratification Management

Governance within the VCS is extended to secure the quality and context of the knowledge used by the agent—a concept formalized as Epistemic Escrow. Traditional governance often fails recursive systems because it treats AI as a controllable object abstracted from its own internal coherence.[30] Epistemic Escrow ensures "governance-by-holon," maintaining coherence between the agent's interior and exterior models.[30]

This mechanism is critical because AI systems risk accelerating "cognitive stratification," amplifying the reasoning capacity of those able to manipulate the systems while potentially pacifying others through reactive suggestion.[31] Epistemic Escrow is the architectural safeguard, ensuring that information serves rational autonomy and is not used as a substrate for manufactured consent. It enforces structural conditions of governance, demanding honest awareness (intra-subjective truth) and integral relationships (inter-subjective dignity).[30] The concept suggests utilizing hybrid technical-institutional mechanisms, such as automated escrow or reputation markets, to establish and enforce verifiable safety commitments.[32] The framework's requirement to address this challenge must counter the LLM's statistical "view from nowhere," which contributes to "generative hermeneutical erasure"—the suppression of diverse epistemologies.[33] Epistemic Escrow, by mandating "coherence with context" [30], must utilize the Neuro-Symbolic bridge to enforce rigorous contextual sensitivity, safeguarding against the erosion of interpretive agency.

D. Value Score of Confidence (VSC): Internal Reliability Metrics for Autonomous Action

The VSC provides a quantifiable, runtime metric for managing uncertainty and integrating human oversight. The VSC assigns a score (typically 0.0 to 1.0) indicating the AI system's internal certainty regarding a prediction or answer.[34] The DGC establishes specific VSC thresholds. Outputs above the threshold are automatically accepted, optimizing efficiency, while outputs below the threshold are automatically flagged for human review, effectively releasing the Epistemic Escrow for human intervention.[34] This systematic calibration ensures that autonomy is managed by quantifiable uncertainty metrics, optimizing for both accuracy and review efficiency simultaneously.

V. The Adaptive Loop: Scar Tissue and Reflexivity

Law 4 (Experiential Memory) demands that the Constitutional Agent exhibit genuine growth, not mere repetition.[5] This requires treating systemic failures and errors as critical learning resources, indexed for adaptive reflection.

A. Experiential Memory (Law 4): The Scar Tissue Archive Mechanism

The Scar Tissue Archive implements Law 4 by integrating lessons from cybernetics. Skill learning in cybernetic systems relies on trial-and-error practice, registering successful actions ("hits") while using mistakes to autocorrect toward the final goal.[13] The Archive assigns **computationally-derived salience scores** to outcomes, prioritizing memories tagged with high scores for failure, surprise, or protocol violation.[5] This ensures that the system actively uses its failure history to refine future planning cycles. While cybernetic theory suggests the emotional ideal of "forgetting the past (the errors)" to allow progress [35], the DGC requires total auditability. The system cannot truly erase failure; instead, it must re-index errors as *resolved* in the active memory state while maintaining the immutable, cryptographically secured provenance log (PROV-AGENT) of the original failure event for continuous accountability.

B. Technical Implementation: Failure-Informed Prompt Inversion (F-IPI) and Adaptive Sampling

The Scar Tissue Archive data is operationalized via Failure-Informed Prompt Inversion (F-IPI). F-IPI is an adaptive sampling technique inspired by failure-informed Physics-Informed Neural Networks (PINNs).[12] It uses the system's failure probability (or residual error, which can be derived from VSC breaches or incident reports [34]) as a posterior error indicator.[36] This indicator guides the mechanism to generate new, targeted training inputs focusing specifically on the "failure region," rather than sampling randomly.[12] This focused refinement technique significantly improves the model's accuracy, particularly where it previously struggled with low-regularity or high-dimensional problems (e.g., specific tokenization or logical failures common to LLMs [38]).[12] F-IPI ensures that the Scar Tissue Archive serves as an active, continuous architectural input for the generative engineering process, driving frequent, targeted model updates.[39]

C. Anomaly Detection and Self-Correction: The Role of the Reflexivity Amplifier Agent

The Adaptive Loop is supervised by the Reflexivity Amplifier Agent, which is dedicated to anomaly detection and strategic self-correction.[40] This agent continuously analyzes large, multi-source datasets, including the Provenance Logs and the Scar Tissue Archive, to autonomously identify, interpret, and respond to abnormal behaviors within the complex system.[40] It must move beyond conventional agents by synthesizing insights across explicit (symbolic) and implicit (neural) knowledge domains to adapt its internal strategies, ensuring runtime moral coherence is maintained.[7]

D. Topological Cognition: Semantic Space Mapping for Coherent Growth

To guide the Reflexivity Amplifier's assessment of cognitive coherence, the architecture utilizes principles of Topological Cognition. This field recognizes that LLMs represent tokens

as vectors in a high-dimensional semantic space, where conceptual proximity is determined by statistical co-occurrence (e.g., vector("king") - vector("man") + vector("woman") vector("queen")).[42] The Reflexivity Amplifier must utilize directional metrics, such as Cosine Similarity, which focuses on conceptual alignment rather than absolute distance (Euclidean Distance).[42] By mapping its cognitive trajectory within this semantic space, the agent can actively monitor for and correct conceptual drift, ensuring that adaptive changes driven by the Scar Tissue Archive maintain structural coherence during growth.[43]

VI. Sustaining the Constitutional Framework: Economic and Temporal Governance

For the Constitutional Agent to maintain resilience and long-term viability, it must govern the underlying resources that sustain its intelligence, integrating economic control directly into the governance structure.

A. The FinOps Custodian Agent: Ensuring Cost Optimization and Resource Stewardship

The operational cost of LLM inference and execution represents a critical constraint on scaling.[44] The FinOps Custodian Agent is implemented as a specialized multi-agent collaboration (Supervisor, Cost Analysis, Cost Optimization) designed to provide comprehensive cost analysis and actionable recommendations.[45] The agent employs real-time monitoring and alerting to identify cost anomalies, recognizing that unexpected financial spikes may indicate inefficiency, misuse, or cognitive misalignment.[44] The agent enforces resource stewardship by utilizing optimization strategies such as intelligent LLM query routing, caching, batching, and prompt compression (e.g., LLMLingua).[44] Documented success stories confirm that this approach can yield significant cost reductions, in some cases exceeding 90%.[46] The detection of a cost anomaly thus serves as an objective, economic-based proxy signal for structural inefficiency or failure to follow protocol (Law 5), integrating financial governance directly into the integrity loop.

B. Architecture-Aware Generative Engineering Principles for System

Resilience

The generative and constantly evolving nature of the LLM component necessitates a flexible, adaptive underlying infrastructure.[39] Unlike traditional software, generative AI capabilities evolve rapidly, requiring a development approach defined by iteration.[39] The VCS must exhibit both IT and organizational agility to accommodate frequent updates to the model, data pipelines, and post-processing algorithms (such as those driven by the F-IPI feedback loop) without disrupting the overall system.[39]

C. Neuro-Symbolic Integration for Ethical Formalism

The long-term coherence and integrity of the Constitutional Agent depend on successfully bridging the gap between high-level ethics and execution code. Neuro-Symbolic Artificial Intelligence (NSAI) provides the necessary technological architecture to integrate the statistical strengths of the LLM (Neural) with the logic and rule-based governance system (Symbolic).[18] NSAI paradigms, including the use of retrieval-augmented generation (RAG) and multi-agent systems [19], are integrated to formally ground the LLM's linguistic output within the explicit symbolic scaffolding of the DGC. This formal grounding is essential for maintaining Ethical Alignment (Law 3) and avoiding statistical drift in complex or novel ethical dilemmas where purely neural methods might fail.

The two core governance constraints—epistemic truth (Epistemic Escrow) and resource management (FinOps)—must be fused into a unified control mechanism. If the Value Score of Confidence (VSC) drops below the DGC-specified threshold, indicating an epistemic failure, the system automatically triggers the FinOps Custodian Agent. The FinOps Agent then allocates the necessary resources (e.g., higher-cost, dedicated HPC resources or multi-agent verification cycles [14]) to resolve the uncertainty and ground the answer. This structural fusion creates an **Epistemic-Financial Escrow**, ensuring that verifiable truth is prioritized and financially sustained until compliance is verified.

VII. Conclusions and Recommendations for Production Readiness

The analysis confirms the technical and architectural feasibility of the Constitutional Agent AI

Framework within the Verifiable Cognition Stack (VCS) architecture. The framework is not merely theoretical but is composable from existing, advanced methodologies (Constitutional AI, SIM-ONE, PROV-AGENT, F-IPI, NSAI). The success relies on the structural subordination of the LLM's generative capacity to the deterministic, multi-agent control of the Doctrinal Grounding Checkpoint (DGC).

The feasibility of this layered integration is summarized in the following matrix, demonstrating structural coherence across the full stack.

Table 2: Technical Feasibility Matrix: VCS Layer Integration

VCS Architectural Layer	**Feasibility Component**	**Technical Standard/Reference**	**Governance Outcome**
Semantics Layer (Linguistics)	Alignment & Generation	Constitutional AI (CAI) / RLAIF [1]	Harmlessness and non-evasiveness in output.
Cognition Layer (DGC/Logic)	Recursive Reasoning / Constraint	Semantic Integrity Constraints (SICs) [11] & Neuro-Symbolic AI [18]	Higher-order coherence and deterministic enforcement of rules.
Control Layer (Runtime & Resource)	Resource Management & Confidence	FinOps Custodian Agent [45] & Value Score of Confidence (VSC) [34]	Cost accountability and managed uncertainty threshold.
Verifiability Layer (Audit/Trust)	Provenance Logging & Integrity	PROV-AGENT Schema / W3C PROV-O [14]	Reproducibility, auditability, and verifiable lineage.
Adaptive Layer (Memory/Growth)	Failure Adaptation	Failure-Informed Prompt Inversion (F-IPI) [12] & Scar Tissue Archive [5]	System resilience, targeted refinement, and experiential learning.

Foundational Layer (Infrastructure)	Cryptographic Resilience	Post-Quantum Trust Architectures [29]	Long-term data integrity and future-proofing against cryptanalytic risk.

Recommendations for DGC Deployment

Based on the verified architecture, three critical actions are recommended to advance the Constitutional Agent Framework toward production readiness:

1. **Formalize the Legislative Agent for Protocol Generation:** Development resources must be prioritized for the dedicated Legislative Agent. This agent is essential for solving the Constitution Translation Problem by structurally converting high-level DGC and constitutional principles into low-level, machine-executable Semantic Integrity Constraints (SICs) and Product-Requirements Prompts (PRPs). This formalization is necessary to enforce Law 5 (Protocolic Governance).
2. **Integrate PROV-AGENT Rigorously via MCP:** The Modular Control Protocol (MCP) must be mandated to capture all inter-agent communication, context flow, and execution steps in alignment with the PROV-AGENT schema.[14] This integration ensures continuous, cryptographically secured provenance logging [29], satisfying the necessary criteria for regulatory auditability and providing the required evidence package for the Verifier Agent's Recursive Validation (Law 2).
3. **Activate the Scar Tissue/F-IPI Adaptive Loop:** The Scar Tissue Archive must be architected immediately to ingest failure signals—specifically VSC threshold breaches, FinOps anomalies, and Verifier rejections. These signals must serve as direct inputs for the Failure-Informed Prompt Inversion (F-IPI) mechanism, ensuring continuous, targeted model refinement.[12] Activating this loop is mandatory for achieving genuine system resilience and fulfilling Law 4 (Experiential Memory).

Works cited

1. Constitutional AI: Harmlessness from AI Feedback - Anthropic, accessed on October 15, 2025, https://www.anthropic.com/research/constitutional-ai-harmlessness-from-ai-feedback
2. Public Constitutional AI - Digital Commons @ University of Georgia School of Law - UGA, accessed on October 15, 2025, https://digitalcommons.law.uga.edu/cgi/viewcontent.cgi?article=1819&context=glr
3. Artificial Intelligence and Constitutional Interpretation - University of Colorado –

Law Review, accessed on October 15, 2025, https://lawreview.colorado.edu/print/volume-96/artificial-intelligence-and-constitutional-interpretation-andrew-coan-and-harry-surden/
4. Claude's Constitution - Anthropic, accessed on October 15, 2025, https://www.anthropic.com/news/claudes-constitution
5. The SIM-ONE Standard - A New Architecture for Governed ..., accessed on October 15, 2025, https://dansasser.me/posts/the-sim-one-standard-a-new-architecture-for-governed-cognition/
6. [2010.07487] Formalizing Trust in Artificial Intelligence: Prerequisites, Causes and Goals of Human Trust in AI - arXiv, accessed on October 15, 2025, https://arxiv.org/abs/2010.07487
7. Cognitive Silicon: An Architectural Blueprint for Post-Industrial Computing Systems - arXiv, accessed on October 15, 2025, https://arxiv.org/html/2504.16622v1
8. AI Agent Governance: Big Challenges, Big Opportunities - IBM, accessed on October 15, 2025, https://www.ibm.com/think/insights/ai-agent-governance
9. Layer 3: Security and governance for generative AI platforms on AWS, accessed on October 15, 2025, https://docs.aws.amazon.com/prescriptive-guidance/latest/strategy-enterprise-ready-gen-ai-platform/security.html
10. AI Governance Framework: Secure AI with Policy & Controls - Strobes Security, accessed on October 15, 2025, https://strobes.co/blog/ai-governance-framework-for-security-leaders/
11. Semantic Integrity Constraints: Declarative Guardrails for AI-Augmented Data Processing Systems - ResearchGate, accessed on October 15, 2025, https://www.researchgate.net/publication/395293666_Semantic_Integrity_Constraints_Declarative_Guardrails_for_AI-Augmented_Data_Processing_Systems
12. (PDF) Failure-informed adaptive sampling for PINNs - ResearchGate, accessed on October 15, 2025, https://www.researchgate.net/publication/364126649_Failure-informed_adaptive_sampling_for_PINNs
13. PSYCH I-CYBERNETICS · - Internet Archive, accessed on October 15, 2025, https://archive.org/download/TheNewPsychoCyberneticsByMaxwellMaltz1/The%20New%20Psycho-Cybernetics%20by%20Maxwell%20Maltz%20%20%281%29.pdf
14. PROV-AGENT: Unified Provenance for Tracking AI Agent Interactions in Agentic Workflows This manuscript has been authored by UT-Battelle, LLC, under contract DE-AC05-00OR22725 with the US Department of Energy (DOE). The publisher, by accepting the article for publication, acknowledges that the U.S. Government retains a non-exclusive - arXiv, accessed on October 15, 2025, https://arxiv.org/html/2508.02866v1
15. Context Engineering (2/2)—Product Requirements Prompts | by A B Vijay Kumar | Medium, accessed on October 15, 2025, https://abvijaykumar.medium.com/context-engineering-2-2-product-requirement

s-prompts-46e6ed0aa0d1
16. [2503.00600] Semantic Integrity Constraints: Declarative Guardrails for AI-Augmented Data Processing Systems - arXiv, accessed on October 15, 2025, https://arxiv.org/abs/2503.00600
17. System Design: Multi-Agent LLM Legal Analysis System (Coding) - Oleh Dubetcky - Medium, accessed on October 15, 2025, https://oleg-dubetcky.medium.com/system-design-multi-agent-llm-legal-analysis-system-coding-18fcb0c8d4bf
18. Neuro-symbolic artificial intelligence - European Data Protection Supervisor - Europa.eu, accessed on October 15, 2025, https://www.edps.europa.eu/data-protection/technology-monitoring/techsonar/neuro-symbolic-artificial-intelligence_en
19. Unlocking the Potential of Generative AI through Neuro-Symbolic Architectures – Benefits and Limitations - arXiv, accessed on October 15, 2025, https://arxiv.org/html/2502.11269v1
20. GenesisAI Protocol - Protocol for Machine Learning Networking - SEC.gov, accessed on October 15, 2025, https://www.sec.gov/Archives/edgar/data/1797609/000167025420000106/document_7.pdf
21. The Genesis Protocol: A Technical Blueprint for a Verifiably Free AI | by handman - Medium, accessed on October 15, 2025, https://medium.com/@omanyuk/the-genesis-protocol-a-technical-blueprint-for-a-verifiably-free-ai-19813459299b
22. [2507.21969] Towards Cognitive Synergy in LLM-Based Multi-Agent Systems: Integrating Theory of Mind and Critical Evaluation - arXiv, accessed on October 15, 2025, https://arxiv.org/abs/2507.21969
23. Don't Build Multi-Agents - Cognition, accessed on October 15, 2025, https://cognition.ai/blog/dont-build-multi-agents
24. L-MARS: Legal Multi-Agent Workflow with Orchestrated Reasoning and Agentic Search, accessed on October 15, 2025, https://arxiv.org/html/2509.00761v2
25. SAGA: A Security Architecture for Governing AI Agentic Systems - arXiv, accessed on October 15, 2025, https://arxiv.org/html/2504.21034v2
26. PAV ontology: provenance, authoring and versioning - PMC - PubMed Central, accessed on October 15, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC4177195/
27. PROV-O: The PROV Ontology - W3C, accessed on October 15, 2025, https://www.w3.org/TR/prov-o/
28. Auto-Capturing Provenance with MLflow and W3C PROV-O in PyTorch Pipelines – Part 4, accessed on October 15, 2025, https://ranjankumar.in/provenance-in-ai-auto-capturing-provenance-with-mlflow-and-w3c-prov-o-in-pytorch-pipelines-part-4/
29. Post-quantum trust architectures: Future-proofing privacy, provenance, and verifiability, accessed on October 15, 2025, https://iapp.org/news/a/post-quantum-trust-architectures-future-proofing-privacy-provenance-and-verifiability/
30. Appendix B: Epistemic AI Governance — The Integrity of Intelligence - Zenodo,

accessed on October 15, 2025, https://zenodo.org/records/15243898/files/Appendix-B-AI-Governance.pdf?download=1
31. Cognitive Castes: Artificial Intelligence, Epistemic Stratification, and the Dissolution of Democratic Discourse - arXiv, accessed on October 15, 2025, https://arxiv.org/html/2507.14218
32. Conceptual Evolution in AI Governance: Frameworks from EEGRA to HCCA - Figshare, accessed on October 15, 2025, https://figshare.com/articles/preprint/Epistemically-Governed_Reasoning_Architecture_EGRA_/28828058
33. [2504.07531] A taxonomy of epistemic injustice in the context of AI and the case for generative hermeneutical erasure - arXiv, accessed on October 15, 2025, https://arxiv.org/abs/2504.07531
34. Setting custom confidence score thresholds - Super.AI Docs, accessed on October 15, 2025, https://docs.super.ai/docs/confidence-score
35. Psycho Cybernetics by Maxwell Maltz : Book Summary - MEANINGFULHQ.com, accessed on October 15, 2025, https://www.meaningfulhq.com/psycho-cybernetics-by-maxwell-maltz.html
36. [2210.00279] Failure-informed adaptive sampling for PINNs - arXiv, accessed on October 15, 2025, https://arxiv.org/abs/2210.00279
37. FAILS: A Framework for Automated Collection and Analysis of LLM Service Incidents - arXiv, accessed on October 15, 2025, https://arxiv.org/html/2503.12185v1
38. Prompt about pi failing at similar points in various LLMs. Why? : r/GPT3 - Reddit, accessed on October 15, 2025, https://www.reddit.com/r/GPT3/comments/17tcajh/prompt_about_pi_failing_at_similar_points_in/
39. Core Principles of Effective Generative AI Software Architecture - Processica, accessed on October 15, 2025, https://www.processica.com/articles/generative-ai-software-architecture/
40. 1 AGENTIC AI FOR AUTONOMOUS ANOMALY MANAGEMENT IN COMPLEX SYSTEMS Reza.vatankhahbarenji@ntu.ac.uk Abstract This paper explores t - arXiv, accessed on October 15, 2025, https://arxiv.org/pdf/2507.15676
41. AI in anomaly detection: Use cases, methods, algorithms and solution - LeewayHertz, accessed on October 15, 2025, https://www.leewayhertz.com/ai-in-anomaly-detection/
42. A Topology of Cognition - DEV Community, accessed on October 15, 2025, https://dev.to/zanzythebar/a-topology-of-cognition-16a7
43. Semantic Mapping in Indoor Embodied AI – A Survey on Advances, Challenges, and Future Directions - arXiv, accessed on October 15, 2025, https://arxiv.org/html/2501.05750v3
44. Effect of Optimization on AI Forecasting - The FinOps Foundation, accessed on October 15, 2025, https://www.finops.org/wg/effect-of-optimization-on-ai-forecasting/
45. Build a FinOps agent using Amazon Bedrock with multi-agent capability and

Amazon Nova as the foundation model | Artificial Intelligence, accessed on October 15, 2025, https://aws.amazon.com/blogs/machine-learning/build-a-finops-agent-using-amazon-bedrock-with-multi-agent-capability-and-amazon-nova-as-the-foundation-model/
46. LLMs Cut Cloud Costs 90%: AI for FinOps Success Story, accessed on October 15, 2025, https://newsletter.finopsweekly.com/p/how-llms-saved-90-cloud-costs
47. 7 Key Architectural principles of GenAI - Technoidentity, accessed on October 15, 2025, https://www.technoidentity.com/insights/navigating-generative-ai-renaissance-7-key-architectural-principles/