Legitimacy Debt: The Unseen Liability Shaping the Future of Artificial Intelligence

Executive Summary

This report introduces and analyzes the concept of **Legitimacy Debt**, a structural burden accrued when artificial intelligence (AI) systems are developed and scaled on a foundation of unconsented, unverified, and poorly governed data. Analogous to technical debt in software engineering, Legitimacy Debt represents a trade-off where short-term acceleration in capability is achieved at the expense of long-term stability and trust. While technical debt manifests as internal friction, Legitimacy Debt is an external-facing liability owed to society, data creators, and the legal system.

The core thesis of this analysis is that the AI industry's prevailing ethos of rapid, unrestrained scaling has generated a significant and systemic Legitimacy Debt. This debt is now coming due, manifesting as a global wave of high-stakes litigation, intensifying regulatory scrutiny, and a quantifiable erosion of public trust. The consequences are not merely financial or reputational; this debt creates a structural drag on the entire AI ecosystem, favoring market centralization by erecting liability barriers that only the largest corporations can surmount, thereby slowing the democratization and broader adoption of AI.

Key findings of this report include:

- **A Distinct Liability:** Legitimacy Debt is a unique form of risk, distinct from but interconnected with technical, ethical, and verification debt. It is the foundational liability, arising from the violation of social contracts regarding data, which in turn seeds other forms of debt throughout the AI development lifecycle.
- **Primary Manifestations:** The "interest payments" on this debt are now materializing. They include a surge in copyright and privacy lawsuits from creators and individuals, severe economic damage from data scandals that undermine consumer confidence, and a market structure where the immense liability risk concentrates power among a few dominant technology firms.
- **Pathways to Repayment:** A multi-layered strategy is required to manage and repay this

debt. This involves implementing robust governance frameworks like the EU AI Act, adopting technical transparency layers such as the C2PA standard and AI watermarking, fostering the creation of trusted and ethically sourced data ecosystems, and developing new economic models for data compensation, including formal licensing markets and agentic micropayments.

The strategic recommendation for organizations is to transition from a reactive, risk-mitigation posture to a proactive strategy of building **"Legitimacy by Design."** This approach requires treating data provenance, consent, and ethical governance not as compliance afterthoughts but as core, indispensable components of innovation, competitive advantage, and long-term value creation in the age of AI.

Part 1: Defining Legitimacy Debt: A New Framework for AI Risk

1.1 Introduction: Beyond Technical Debt

The concept of "technical debt" has been a cornerstone of software engineering for decades. Coined by Ward Cunningham, it uses a financial metaphor to describe the implied cost of rework caused by choosing an easy, expedient solution in the present instead of a more robust approach that would take longer.[1] This "debt" accrues "interest" over time, manifesting as increased maintenance overhead, slower development cycles, and a higher propensity for bugs, ultimately hindering a system's evolution.[1]

The advent of the AI era has supercharged this dynamic. AI-powered coding assistants have made it unprecedentedly easy to generate vast quantities of code, yet they have also made it easier to accumulate technical debt at an alarming rate.[3] These tools can produce code that is redundant, logically flawed, or contains hidden security vulnerabilities, often because they lack a holistic understanding of the broader system architecture.[3] A recent Stanford study found that developers using AI assistants wrote "significantly less secure code" yet were "more likely to believe they wrote secure code," highlighting a dangerous combination of accelerated debt accrual and false confidence.[3]

However, the industry's intense focus on the internal liabilities of code and system design has obscured a far greater, more foundational risk: **Legitimacy Debt**. This is not a debt of

suboptimal code but a debt of illegitimate foundations. It is accrued when the very raw material of modern AI—data—is acquired and used in ways that violate legal rights, ethical norms, and the social contract. While technical debt compromises a system from within, Legitimacy Debt compromises its right to exist from without. It is the structural liability incurred by prioritizing speed and scale over the painstaking work of ensuring data is ethically sourced, properly licensed, and transparently governed.

1.2 The Anatomy of Debt: Legitimacy vs. Technical, Ethical, and Verification Debt

To manage risk effectively, it is crucial to develop a precise vocabulary. The term "debt" is often used loosely in technology circles. Legitimacy Debt is a distinct concept with unique origins, consequences, and stakeholders. A comparative analysis clarifies its specific nature.

- **Technical Debt** is an *internal-facing* liability concerning the health of a codebase or system architecture. It arises from pragmatic, often necessary, shortcuts taken during development, such as using inconsistent coding standards, creating duplicate code snippets, or failing to refactor legacy systems.[1] The "interest" is paid by the organization in the form of increased maintenance costs, slower feature velocity, and higher bug rates.[1] The debt is owed primarily by the development team to its future self and to the business.
- **Verification Debt** is a related concept that arises from prioritizing speed over the thorough verification of AI-generated outputs, whether it be code, academic papers, or other artifacts.[4] Like technical debt, it creates long-term costs, but its impact is broader, affecting trustworthiness and accuracy. A critical feature of verification debt is the introduction of
  moral hazard: the party incurring the debt (e.g., a developer who blindly accepts AI-generated code) may not be the same party that pays the cost (e.g., the end-user who suffers from a security flaw, or society at large).[4]
- **Ethical Debt** is a broader, societal-level liability. It is incurred when AI systems are "created and deployed without fully examining and addressing the potential ethical consequences".[5] This debt encompasses a wide range of harms, including algorithmic bias that perpetuates discrimination, the environmental costs of training large models, and the socio-economic impact of job displacement.[6] The cost of ethical debt is borne by society, and disproportionately by marginalized communities who may be unfairly treated by an algorithm without their knowledge or recourse.[8]
- **Legitimacy Debt**, in contrast, is an *external-facing, structural* liability rooted in the *provenance and legal standing* of the foundational assets used to build an AI system. It is incurred not by writing bad code or failing to anticipate downstream harms, but by building powerful systems on a foundation of data acquired without clear consent, legal

right, or transparent governance. It stems directly from the AI industry's foundational "grab-first-ask-questions-later" approach to data acquisition, which has relied heavily on the mass scraping of web content.[9] This debt is a direct violation of the social contract regarding data ownership, privacy, and intellectual property. The "interest" is paid externally through high-stakes litigation, crippling regulatory fines, and the systemic erosion of public trust. This debt is owed by the AI developer to data creators, rights holders, regulators, and the public at large.

These forms of debt are not isolated; they exist in a cascading relationship. The initial act of accruing Legitimacy Debt—for instance, by training a model on a vast, uncurated corpus of scraped internet data—is the "original sin" that seeds other liabilities throughout the system. Such a dataset is likely to contain copyrighted material, private information, and reflections of historical societal biases.[3] A model trained on this illegitimate foundation will, in turn, generate outputs—be it code, text, or images—that may contain security vulnerabilities (Technical Debt), perpetuate harmful stereotypes (Ethical Debt), or simply be inaccurate (creating Verification Debt). A developer who then uses these outputs without sufficient scrutiny compounds the problem. Therefore, Legitimacy Debt is the foundational layer of risk. Addressing algorithmic bias or code quality without resolving the underlying illegitimacy of the training data is akin to treating symptoms while ignoring the disease. A truly robust governance framework must address risk at its source, starting with the legitimacy of its data assets.

Table 1: A Comparative Framework of AI-Related Debts

Debt Type	Core Definition	Who Incurs It?	Who Pays the Cost?	Manifestation of "Interest"	Example
Technical Debt	The implied future cost of choosing an easy, suboptimal technical solution now instead of a better, more robust one.[1]	Development Team / Organization	Development Team / Organization	Slower development, increased maintenance, more bugs, system instability.[1]	Using a quick coding hack instead of refactoring a module, leading to future integration problems.
Verificatio	The	AI User /	Downstrea	Inaccurate	A student

n Debt	accumulated risk from prioritizing speed over the thorough verification of AI-generated outputs, leading to a decline in trustworthiness.[4]	Developer	m Users / Society	outputs, flawed mental models, trust cascades on shaky foundations, moral hazard.[4]	using an LLM to write an essay without fact-checking, submitting a paper with fabricated citations.
Ethical Debt	The societal cost of deploying AI systems without adequately considering and mitigating potential negative consequences and societal harms.[5]	AI Developer / Deploying Organization	Society / Marginalized Groups	Algorithmic bias, discrimination, job displacement, environmental damage, erosion of autonomy.[7]	Deploying a hiring algorithm trained on historical data that perpetuates gender bias against female candidates.
Legitimacy Debt	The structural burden created by building and scaling AI on data acquired without clear	AI Developer / Deploying Organization	Data Creators / Rights Holders / The Public	Litigation, regulatory fines, public mistrust, reputational damage, market centralization.[10]	Training a large language model on billions of web pages containing copyrighted books and personal

	consent, legal rights, provenance, or ethical governance.				blogs without permission.

1.3 The Causal Chain: How Legitimacy Debt Accumulates

Legitimacy Debt is not an accidental byproduct of innovation; it is the predictable outcome of a specific developmental methodology. Its accumulation follows a clear causal chain, moving from unrestrained technical ambition to systemic market dysfunction.

1. **Acceleration without Restraint:** The generative AI boom was catalyzed by a race for scale. The prevailing belief was that larger models trained on more data would yield greater capabilities. This led to a "grab-first-ask-questions-later" mentality, where automated web scraping became the primary method for amassing petabyte-scale datasets.[9] The imperative was to build the most powerful model as quickly as possible, with considerations of data rights and provenance being secondary.
2. **The Legitimacy Gap:** This approach created a profound legitimacy gap. The act of scraping publicly accessible data was often conflated with having a legal and ethical right to use that data for commercial AI training. However, as numerous legal and ethical analyses have pointed out, "publicly available" does not mean "consented for use".[9] This practice bypassed fundamental ethical principles of consent, transparency, and fairness.[14] The scraped corpora were inherently illegitimate, containing vast amounts of copyrighted material, sensitive personal information, and unfiltered societal biases, all ingested without the knowledge or permission of the original creators and subjects.[9]
3. **The Litigation Load:** The legitimacy gap inevitably became a legal battleground. As the commercial value of AI models skyrocketed, creators, publishers, and individuals whose data was used without permission began to seek recourse through the legal system. This has resulted in a global wave of litigation, with lawsuits alleging copyright infringement, privacy violations, and breach of contract serving as the primary counter-balance to the industry's unrestrained data acquisition practices.[10] These lawsuits represent the first major "interest payments" on the accrued Legitimacy Debt.
4. **Structural Drag & Centralization:** The immense financial and legal liability associated with this litigation load creates a powerful structural drag on the AI ecosystem. The potential for multi-billion dollar damages and the high cost of legal defense and compliance create a formidable barrier to entry. Only large, well-resourced corporations possess the capital reserves, legal teams, and insurance coverage necessary to operate

in such a high-risk environment.[12] This dynamic stifles competition from startups and smaller innovators, concentrating AI development power in the hands of a few dominant players who can afford to absorb the liability.

5. **Damaged Reality of AI:** The cumulative effect of high-profile lawsuits, recurring privacy scandals, and the public's growing awareness of unethical data harvesting is a deep and widening trust deficit. A 2025 Pew Research survey found that 50% of Americans are now more concerned than excited about the increased use of AI, a significant jump from 37% in 2021.[21] This public skepticism is not a soft, intangible risk; it is a direct impediment to progress. It slows the adoption of beneficial AI tools, discourages public and private investment in experimentation, and fuels calls for potentially heavy-handed regulation that could stifle innovation.[22] The damaged public perception of AI's legitimacy becomes a self-fulfilling prophecy, constraining its potential.

Part 2: Manifestations and Consequences of Legitimacy Debt

The accumulation of Legitimacy Debt is no longer a theoretical risk. Its consequences are now visible and measurable across legal, economic, and market domains. The "interest" on this debt is being paid through a cascade of legal challenges, tangible financial losses, and a distortion of the competitive landscape.

2.1 The Litigation Load: Copyright, Privacy, and the Legal Backlash

The most direct manifestation of Legitimacy Debt is the surge of litigation aimed at the heart of the generative AI business model: its training data. This legal backlash is proceeding on multiple fronts.

- **Copyright Infringement:** The central claim in dozens of lawsuits is that AI companies have engaged in mass, unauthorized reproduction of copyrighted works to train their commercial models.[10] Prominent examples include Getty Images' suit against Stability AI for allegedly copying over 12 million photographs [10]; class-action lawsuits brought by authors and artists against OpenAI, Meta, and Anthropic [10]; and legal challenges from major news organizations like The New York Times, the Toronto Star, and Postmedia, who argue their journalistic content was systematically scraped and used to build products that now compete with them.[19] AI developers have typically countered with a "fair use" defense, arguing that training a model is a transformative act. However, this argument is

facing increasing skepticism. The U.S. Copyright Office, in a 2025 report, stated that arguments analogizing AI training to human learning are "mistaken," as AI models absorb "the essence of linguistic expression" and can generate outputs that directly dilute the market for the original works, a key factor weighing against fair use.[27]

- **Privacy Violations and Data Scraping:** A parallel track of litigation focuses on the violation of privacy rights through the large-scale, unconsented scraping of personal data. A putative class action filed in California alleges that OpenAI and Microsoft violated the privacy of millions by scraping their personal data from social media and other internet sources.[18] These cases draw on the precedent set by scandals like that of Clearview AI, which built a massive facial recognition database by scraping three billion photos from platforms like Facebook and Twitter without user consent, an act that regulators deemed "unethical and illegal".[9] These lawsuits challenge the fundamental premise that data publicly visible on the internet is free for any commercial use.
- **Breach of Contract:** A third legal strategy bypasses the complexities of copyright and privacy law by focusing on platforms' terms of service. Reddit's lawsuit against Anthropic, for example, does not claim copyright ownership over user posts but instead alleges that Anthropic's scraping activities constituted a direct breach of Reddit's User Agreement, which explicitly prohibits unauthorized commercial exploitation of its content.[17] This approach reframes the issue as a contractual violation, providing another avenue to hold AI developers accountable for their data acquisition methods.

2.2 Reputational and Economic Fallout: The High Cost of Eroding Public Trust

Beyond direct legal costs, Legitimacy Debt imposes severe and lasting reputational and economic damage. A foundation of illegitimate data practices makes an organization vulnerable to scandals that destroy public trust, a critical asset in the digital economy.

- **Erosion of Public Trust:** Public sentiment is turning increasingly wary. A 2025 Gallup poll revealed that 80% of Americans believe the government should prioritize AI safety rules even if it slows down development, a clear indicator of low public trust.[28] A Rutgers survey highlighted a growing societal divide, with trust in AI significantly higher among those with higher income and education levels, suggesting that the technology's benefits and perceived legitimacy are not being distributed evenly.[29] This trust deficit is a direct liability, as it can lead to user boycotts, employee activism, and a general reluctance to adopt or engage with AI-powered services.
- **Reputational Damage from Data Practices:** History is replete with examples of companies that suffered immense reputational harm due to their data practices. The Facebook-Cambridge Analytica scandal remains a canonical case of how the misuse of

user data can trigger a global crisis of confidence.[30] Other significant incidents include DeepMind's use of 1.6 million UK patient records without explicit consent for app development, which led to regulatory rebuke [30]; Amazon's AI recruiting tool, which had to be scrapped after it was found to be systematically biased against women [31]; and a 2025 discovery that the AI photo editing app Wondershare RepairIt was leaking sensitive user images due to poor security, directly contradicting its privacy policy.[11] These events create lasting public memory and reinforce the narrative that AI companies cannot be trusted with personal data.

- **Direct and Indirect Economic Costs:** The financial consequences of these failures are substantial and multi-faceted.
  - **Direct Costs:** Regulatory fines for data privacy violations can be staggering. The EU's GDPR has enabled significant penalties, such as the €530 million fine levied against TikTok.[11] The global average cost of a single data breach reached a record high of $4.88 million in 2024, an amount that can be existential for smaller firms.[32] These costs include forensic investigations, legal fees, and patient/customer notification expenses.
  - **Indirect Costs:** Often exceeding the direct costs, these long-tail impacts include customer churn, damage to brand equity, increased cyber insurance premiums, loss of strategic partnerships, and a negative impact on company valuation.[11] A 2019 Verizon study found that 69% of consumers actively avoid businesses that have suffered a data breach, and nearly a third never return.[32] For small businesses, the consequences can be fatal; over 60% of small businesses close within six months of a significant cyberattack.[32]

2.3 Structural Drag: How Liability Risk Drives Market Centralization

Perhaps the most insidious long-term consequence of Legitimacy Debt is its effect on the very structure of the AI market. The immense and unpredictable liability associated with AI development creates a powerful force for centralization, undermining the goal of a diverse and competitive ecosystem.

- **The Liability Barrier:** The potential for multi-million or even billion-dollar lawsuits and regulatory fines acts as a formidable barrier to entry.[12] Startups and small-to-medium-sized enterprises (SMEs) typically lack the capital to withstand a single major lawsuit, let alone the ongoing costs of navigating complex global compliance regimes and securing adequate insurance coverage.[12] This risk premium makes it difficult for them to attract investment and compete on a level playing field.
- **Concentration of Power:** This dynamic disproportionately benefits large, well-resourced incumbents. Big Tech firms can afford to maintain large in-house legal teams, pay for expensive insurance policies, absorb the costs of litigation and fines as a cost of doing

business, and invest heavily in compliance infrastructure.[12] This allows them to operate in the high-liability environment created by Legitimacy Debt, while smaller competitors are pushed out. The result is the concentration of AI development and deployment in the hands of a few dominant companies, which limits consumer choice, stifles innovation, and reduces competitive pressure to improve ethical practices.[12]

- **Systemic Financial Risk:** This concentration carries macroeconomic risks. In critical sectors like finance, a "monoculture" could emerge where a small number of dominant AI models, likely trained on similar massive datasets, are used across the industry.[35] SEC Chair Gary Gensler has warned that the data-hungry nature of AI could lead to a convergence on a few key data providers and AI-as-a-Service companies.[35] This homogeneity could lead to correlated behaviors, herding, and amplified market volatility during periods of stress, increasing systemic risk for the entire financial system.[35]

This situation reveals a deep paradox. Legitimacy Debt, which was born from a *lack* of rules and a culture of "acceleration without restraint," has triggered a powerful backlash demanding accountability. This backlash, in turn, is leading to the creation of new, complex liability regimes and regulations, such as the EU AI Act and the Colorado AI Act.[12] While these regulations are well-intentioned and necessary, their high compliance costs can inadvertently reinforce the market power of the very incumbents whose actions necessitated the regulation in the first place. Startups and smaller innovators, who may be more inclined to adopt ethical, provenance-first approaches, can be priced out of the market by the high cost of demonstrating their legitimacy. If not designed carefully, the regulatory response to Legitimacy Debt risks becoming a tool of market consolidation, entrenching the dominance of the original debtors and creating a vicious cycle of centralization.

Part 3: Counter-Balance Systems: Pathways to Legitimacy Debt Repayment

Addressing Legitimacy Debt requires a fundamental shift from a reactive posture of risk mitigation to a proactive strategy of building trust and legitimacy by design. This involves a multi-layered approach that combines robust governance, technical transparency, the creation of trusted data ecosystems, and new economic models that align incentives. These counter-balance systems represent the pathways to "repaying" the debt and establishing a sustainable foundation for the future of AI.

3.1 Governance and Restraint: The Foundational Layer

The first step in repaying any debt is to stop accumulating more of it. This requires establishing clear governance frameworks and protocols of restraint at both the organizational and regulatory levels.

- **Comprehensive Governance Frameworks:** Organizations must move beyond ad-hoc ethics committees to embed governance into their operational DNA. The **Hourglass Model of Organizational AI Governance**, developed by researchers at the Turku School of Economics, offers a structured approach.[38] It visualizes governance as a process of translating high-level "environmental" inputs (such as laws, ethical principles, and stakeholder expectations) through a narrowing "organizational" layer (which defines strategy, roles, and capabilities) into the broad "AI system" layer, where these principles are implemented as concrete operational practices for data management, algorithm monitoring, and risk assessment.[40] This model provides a practical roadmap for turning abstract ethical commitments into auditable, day-to-day actions.
- **Regulatory Guardrails: The EU AI Act:** On a macro level, landmark legislation like the European Union's AI Act provides a powerful template for mandated restraint. The Act implements a risk-based framework that is fundamentally changing the calculus for AI developers.[37]
    - It explicitly **bans "unacceptable risk" AI systems**, including those used for social scoring and the untargeted scraping of facial images from the internet to create databases—practices that are primary sources of Legitimacy Debt.[37]
    - For **"high-risk" systems** (e.g., in employment, credit, or law enforcement), it imposes stringent obligations, including requirements for high-quality, representative training data; detailed documentation of data provenance; robust human oversight; and measures to detect and mitigate bias.[37]
    - For **generative AI models**, it introduces critical transparency requirements, mandating that providers disclose that content is AI-generated and publish summaries of the copyrighted data used for training.[42]

      By codifying these requirements, the AI Act effectively shifts the burden of proof, compelling developers to demonstrate legitimacy from the outset rather than dealing with the consequences after the fact.

3.2 Transparency Layers: Building Trust Through Provenance

Governance frameworks must be supported by technical mechanisms that make their implementation verifiable. Data provenance—the creation of a detailed, auditable record of a data asset's origin and history—is the cornerstone of technical transparency and a direct

antidote to the opacity that fuels Legitimacy Debt.[46]

- **Data Provenance Standards (C2PA):** The Coalition for Content Provenance and Authenticity (C2PA), a consortium including Adobe, Microsoft, and other major tech firms, has developed an open technical standard for creating verifiable "Content Credentials" for digital media.[49]
  - **Mechanism:** C2PA functions like a digital "nutrition label".[49] When a C2PA-enabled device or software application creates or modifies a piece of content, it generates a manifest containing assertions about that action (e.g., creator identity, tool used, timestamp, whether AI was involved). This manifest is cryptographically signed and embedded within the file. Any subsequent alteration to the content or its metadata will invalidate the cryptographic signature, making tampering evident.[50]
  - **Adoption:** This standard is gaining traction. OpenAI, for example, is now embedding C2PA metadata in images generated by its DALL·E 3 model, allowing any user with a verification tool to confirm that the image is AI-generated.[52] This provides a crucial, verifiable layer of context.
  - **Limitations:** The primary limitation of C2PA is that it is an opt-in system, and the metadata can be deliberately or accidentally stripped from a file. However, its absence on a piece of content may itself become a signal of potential manipulation as adoption becomes more widespread.[52]
- **AI Watermarking:** A complementary approach is watermarking, which involves embedding a persistent, often invisible, signal directly into AI-generated content to identify its origin.[53]
  - **Technology:** These techniques are becoming increasingly sophisticated. Google's SynthID, for instance, makes subtle modifications to the pixels of an AI-generated image to embed a digital watermark that is designed to survive common transformations like compression, filtering, and cropping.[53] For text, research is underway on methods that train language models to subtly favor certain word choices or sentence structures, creating a statistical signature that can identify the text's origin.[53]
  - **Market Growth:** The demand for these solutions is creating a significant new market. The global AI watermarking market was valued at approximately $500-600 million in 2024-2025 and is projected to grow to over $3 billion by the early 2030s, with a compound annual growth rate (CAGR) of around 25%.[54] This growth is driven by regulatory pressures, such as the labeling requirements in the EU AI Act, and the urgent enterprise need to combat deepfakes and misinformation.[56]
  - **Challenges:** Despite its promise, widespread adoption remains a challenge. A 2025 empirical analysis found that only 38% of leading AI image generators were implementing adequate watermarking practices.[58] Key hurdles include the need for greater interoperability between different watermarking systems and improving the robustness of watermarks against sophisticated adversarial attacks designed to remove them.[54]

3.3 Trusted Data Spaces: Creating Legitimate Data Ecosystems

The most fundamental way to address Legitimacy Debt is to shift the data supply chain away from the unpermissioned scraping of the open web and toward the creation of curated, rights-managed, and ethically governed data ecosystems.

- **The Concept:** Trusted Data Spaces are secure and standardized digital environments designed to facilitate data sharing between organizations based on the principles of data sovereignty, interoperability, and trust.[59] In these spaces, data providers retain control over their assets and can define the precise conditions under which they can be accessed and used, creating a direct counter-model to the "Wild West" of web scraping.
- **European Initiatives:** The European Union has made the development of such spaces a strategic priority. The European Commission's **AI Continent Action Plan** is a multi-pronged strategy to build a trustworthy AI ecosystem in Europe.[61]
    - **AI Factories and Data Labs:** This initiative aims to establish the necessary high-performance computing infrastructure ("AI Factories") and pair it with "Data Labs" that will bring together and curate large, high-quality, and ethically sourced datasets from diverse public and private sources. The goal is to provide European startups and researchers with legitimate, robust data to train trustworthy AI models.[61]
    - **Data Union Strategy:** A comprehensive Data Union Strategy, set to launch in 2025, will seek to create a true single market for data within the EU, allowing legitimate AI solutions to scale across borders on a foundation of trusted data.[61]
    - **Sector-Specific Projects:** Numerous EU-funded projects are already building these data spaces for specific industries, including healthcare (the VELES project), agriculture (DIVINE), and urban logistics (IKIGAI), all designed around common standards for trustworthy data sharing.[60]

3.4 Economic Alignment: The Future of Data Compensation

Ultimately, repaying Legitimacy Debt requires a move beyond technical and legal compliance to fundamental economic realignment. New models are emerging that aim to fairly compensate creators for the value their data provides, transforming the data relationship from extractive to transactional.

- **Data Licensing Markets:** In direct response to the wave of copyright litigation, a formal market for licensing data for AI training is rapidly taking shape.[26]

- **Market Growth:** The scale of this emerging market is substantial. Forecasts project the global AI training dataset market to grow to between $17 billion and $27 billion by 2032.[62] The sub-market for academic research data alone is projected to reach nearly $1.6 billion by 2030, up from $381.8 million in 2024.[63]
- **Key Players and Deals:** Major AI developers are now actively pursuing licensing agreements. OpenAI, for example, has signed multi-million dollar deals with content providers like The Associated Press and Axel Springer to license their archives for training.[26] This shift is creating an entire ecosystem of intermediaries, such as Human Native AI and Corpus, which act as marketplaces to connect rights holders with AI companies seeking legitimate training data.[62]
- **Challenges:** The primary obstacle to the growth of this market remains the legal ambiguity around fair use. As long as some AI companies maintain that they do not need to license data, citing untested fair use defenses, the incentive to participate in the licensing market is diminished.[27]

- **Agentic Micropayments:** A more futuristic but potentially transformative model for repayment involves the use of autonomous AI agents to execute granular, real-time payments for data consumption.[64]
  - **Mechanism:** This model envisions a future where an AI agent, when performing a task like training or generating content, would automatically execute a micropayment—perhaps fractions of a cent—for each piece of data it accesses or uses. This creates a direct, transactional link between data usage and creator compensation, moving beyond bulk licensing to a true pay-per-use economy.[66]
  - **Enabling Technologies:** This vision requires new financial infrastructure. A leading example is the **x402 protocol**, a new standard being developed by a foundation co-founded by Coinbase and Cloudflare.[68] It repurposes the long-dormant HTTP "402 Payment Required" status code to embed instant, stablecoin-based payments directly into web interactions.[70] This would allow one AI agent to automatically pay another for a data crawl, an API call, or access to a piece of content, with the transaction settled programmatically in seconds.[66]
  - **Implication for Repayment:** Agentic micropayments represent a powerful mechanism for repaying Legitimacy Debt. They would create a continuous, automated revenue stream for data creators, ensuring they are compensated for the value their work provides to the AI ecosystem. This model aligns economic incentives with ethical principles, fostering a sustainable and legitimate data supply chain where value is exchanged, not just extracted.[67]

Part 4: Strategic Implications and Future Outlook

Understanding Legitimacy Debt is not merely an academic exercise; it is a strategic imperative. Its accumulation and repayment will have profound implications for the long-term economic trajectory of AI, its societal acceptance, and the very structure of the digital world. Connecting this concept to broader critiques of technology reveals its true significance.

4.1 Legitimacy Debt in Context: Algorithmic Fairness, Data Colonialism, and Surveillance Capitalism

Legitimacy Debt serves as a unifying framework that connects several critical discourses about the societal impact of technology.

- **Algorithmic Fairness:** The field of algorithmic fairness has largely focused on identifying and mitigating bias in the outputs of AI models, using various statistical metrics to measure parity across different demographic groups.[72] However, this approach often treats the symptoms rather than the cause. Legitimacy Debt highlights that if an algorithm is trained on data reflecting decades of systemic discrimination—such as historical redlining in mortgage lending—it will inevitably learn and perpetuate those biases, no matter how sophisticated the debiasing techniques applied to its outputs.[74] The "unfairness" of the algorithm is a manifestation of a deeper Legitimacy Debt owed by the societal system (e.g., the criminal justice or financial system) due to historical injustice.[72] Therefore, achieving true fairness requires not just tweaking algorithms, but addressing the illegitimacy of the foundational data and the societal systems from which it is drawn.
- **Surveillance Capitalism:** Coined by Shoshana Zuboff, surveillance capitalism describes an economic logic based on the "unilateral claiming of private human experience as free raw material" to be converted into behavioral data for prediction and profit.[76] Legitimacy Debt is the direct and predictable liability that arises from this business model. The "grab-first-ask-questions-later" data acquisition that defines surveillance capitalism is the primary engine for accumulating Legitimacy Debt.[78] The legal and social backlash against unpermissioned data scraping is, in essence, a rejection of the core tenets of surveillance capitalism.
- **Data Colonialism:** This concept extends the critique of surveillance capitalism to a global scale, arguing that the data practices of Big Tech mirror historical colonial dynamics.[79] In this view, powerful entities, predominantly from the Global North, extract a new form of "raw material"—data derived from human life—from global populations, often without fair compensation or respect for local sovereignty.[25] The mass scraping of the global internet to train proprietary AI models, the benefits of which flow primarily to a handful of corporations and nations, is a textbook example of data colonialism in action. The resulting Legitimacy Debt is therefore not just a corporate liability but a geopolitical issue, reflecting and reinforcing global power imbalances.

4.2 Long-Term Economic Impact: The Macro View of a Trust Deficit

If left unaddressed, the systemic accumulation of Legitimacy Debt poses a significant threat to the long-term economic promise of AI.

- **A Potential Productivity Paradox:** Economists project that AI could drive a major, sustained surge in productivity, potentially adding trillions to the global economy and reversing decades of slowing growth.[82] However, this optimistic scenario assumes widespread and efficient diffusion of the technology. The structural drag created by Legitimacy Debt—including constant litigation, a patchwork of restrictive regulations, and deep public mistrust—could severely hamper this diffusion. If businesses and consumers are reluctant to adopt AI tools due to perceived risks and a lack of legitimacy, the technology's full productivity potential may never be realized, leading to a disappointing gap between hype and economic reality.[84]
- **Erosion of Social Capital and Institutional Trust:** The economic impact extends beyond simple productivity metrics. When critical societal functions—such as credit allocation, hiring, and even public policy—are delegated to AI systems that are perceived as illegitimate, it erodes trust not only in the technology itself but also in the institutions that deploy it.[84] In a liberal democracy, institutional legitimacy is derived from processes of public deliberation, consent, and accountability. Outsourcing moral and political judgments to opaque, technically optimized systems that lack this legitimacy can destabilize society and undermine the very social contract upon which a healthy economy depends.[84]
- **The Escalating Cost of Inaction:** The economic consequences of data privacy scandals and security breaches are already well-documented and severe, encompassing direct regulatory fines, loss of market value, customer abandonment, and increased operational costs.[11] As AI systems become more deeply integrated into the core functions of the economy, the potential scale of these damages will grow exponentially. A 2025 IBM report found that ungoverned AI systems are not only more likely to be breached but are also significantly more costly when a breach occurs.[85] Inaction is not a neutral stance; it is an active choice to accept an escalating and potentially systemic economic risk.

4.3 Conclusion: From Debt Accumulation to Principled Innovation

The central axiom of this analysis—"When acceleration outruns legitimacy, trust collapses into litigation"—has been validated by the events of recent years. The AI industry, fueled by an

ambition that outpaced its commitment to governance, has accrued a formidable Legitimacy Debt. It now stands at a critical inflection point. The path of continued debt accumulation leads to a future of perpetual legal conflict, stifled innovation, market concentration, and a public alienated from a technology that could otherwise bring immense benefits.

The alternative is to begin the process of repayment. This requires a fundamental reorientation of priorities, moving from a paradigm of unrestrained acceleration to one of principled innovation. The counter-balance systems outlined in this report—comprehensive governance, technical transparency, trusted data ecosystems, and fair economic models—provide the necessary tools. The path forward involves embracing Legitimacy Debt Repayment not as a reactive compliance burden but as a proactive design principle.

This is where practical applications, such as the envisioned WordPress distribution centre framed as a **"Legitimacy Debt Repayment Node,"** become critically important. Such a node represents a microcosm of a healthier AI ecosystem. It would function as a "Trusted Space," embodying restraint by curating a specific, high-quality corpus of data rather than indiscriminately scraping the web. It would champion provenance by ensuring every piece of content has a clear, verifiable history. It would operate on principles of transparency and fair compensation, potentially leveraging new economic models to reward creators. By doing so, it would not only mitigate its own risk but also build a foundation of trust that becomes its core value proposition.

Ultimately, repaying Legitimacy Debt is not about slowing down innovation; it is about making it sustainable. By building on a foundation of legitimacy, the AI industry can move beyond the vicious cycle of mistrust and litigation and begin to realize its true, transformative potential. This approach transforms the repayment of debt from a cost center into a source of profound competitive differentiation and lasting value creation.

Works cited

1. What Is Technical Debt in AI Codes & How to Manage It - GAP, accessed on September 27, 2025, https://www.growthaccelerationpartners.com/blog/what-is-technical-debt-in-ai-generated-codes-how-to-manage-it
2. Tech Debt vs. Feature Velocity: How to Find the Right Balance - CTO Magazine, accessed on September 27, 2025, https://ctomagazine.com/tech-debt-vs-feature-velocity-balance/
3. AI vs Technical Debt: Is This A Race to the Bottom? - DeepDocs, accessed on September 27, 2025, https://deepdocs.dev/ai-and-technical-debt-a-race-to-the-bottom/
4. Verification debt is the AI era's technical debt - Kevin Browne, accessed on September 27, 2025, https://www.kevinbrowne.ca/verification-debt-is-the-ai-eras-technical-debt/
5. Who pays for ethical debt in AI? | Request PDF - ResearchGate, accessed on

September 27, 2025, https://www.researchgate.net/publication/348194864_Who_pays_for_ethical_debt_in_AI
6. AI Has Social Consequences, But Who Pays the Price? Tech Companies' Problem with 'Ethical Debt' - Government Technology, accessed on September 27, 2025, https://www.govtech.com/artificial-intelligence/ai-has-social-consequences-but-who-pays-the-price-tech-companies-problem-with-ethical-debt
7. Navigating the Minefield: Technical, Ethical, and Governance Debt in AI/ML and Generative AI Models | by Mo Meskarian | Medium, accessed on September 27, 2025, https://medium.com/@m.a.meskarian/navigating-the-minefield-technical-ethical-and-governance-debt-in-ai-ml-and-generative-ai-models-fd7fa83861e7
8. Overcoming 'Ethical Debt' - Johns Hopkins Whiting School of Engineering, accessed on September 27, 2025, https://engineering.jhu.edu/news/overcoming-ethical-debt/
9. Between Idealism and Reality — Ethically Sourced Data in AI | by Markus Brinsa - Medium, accessed on September 27, 2025, https://medium.com/@markus_brinsa/between-idealism-and-reality-ethically-sourced-data-in-ai-9138446d2a5c
10. Copyright lawsuits highlight potential risks for agents using AI tools in marketing, accessed on September 27, 2025, https://insurancenewsnet.com/innarticle/copyright-lawsuits-highlight-potential-risks-for-agents-using-ai-tools-in-marketing
11. AI-Powered App Exposes User Data, Creates Risk of Supply Chain ..., accessed on September 27, 2025, https://www.trendmicro.com/en_us/research/25/i/ai-powered-app-exposes-user-data.html
12. The Risks of AI-Specific Liability Regimes | American Enterprise ..., accessed on September 27, 2025, https://www.aei.org/technology-and-innovation/the-risks-of-ai-specific-liability-regimes/
13. Strategies for Ethical and Effective Web Data Collection in AI - RTS Labs, accessed on September 27, 2025, https://rtslabs.com/web-data-collection-in-ai
14. www.alation.com, accessed on September 27, 2025, https://www.alation.com/blog/data-ethics-in-ai-6-key-principles-for-responsible-machine-learning/#:~:text=The%20six%20fundamental%20data%20ethics,sampling%20(ensuring%20representative%20and%20unbiased
15. Data Ethics in AI: 6 Key Principles for Responsible Machine Learning - Alation, accessed on September 27, 2025, https://www.alation.com/blog/data-ethics-in-ai-6-key-principles-for-responsible-machine-learning/
16. Ethical Data – Why Your AI Models Need It - Appen, accessed on September 27, 2025, https://www.appen.com/blog/ethical-data-why-your-ai-models-needs-it
17. Reddit's Lawsuit Over Data-Scraping Could Reshape the Future of ..., accessed on September 27, 2025,

https://www.findlaw.com/legalblogs/courtside/reddits-lawsuit-over-data-scraping-could-reshape-the-future-of-ai/
18. Data Scraping, Privacy Law, and the Latest Challenge to the Generative AI Business Model, accessed on September 27, 2025, https://www.afslaw.com/perspectives/privacy-counsel/data-scraping-privacy-law-and-the-latest-challenge-the-generative-ai
19. Scraping the Surface: OpenAI Sued for Data Scraping in Canada - American Bar Association, accessed on September 27, 2025, https://www.americanbar.org/groups/business_law/resources/business-law-today/2025-february/openai-sued-data-scraping-canada/
20. Significant Liability Risks For Companies That Utilize AI - MehaffyWeber, accessed on September 27, 2025, https://www.mehaffyweber.com/news/significant-liability-risks-for-companies-that-utilize-ai/
21. How Americans View AI and Its Impact on Human Abilities, Society ..., accessed on September 27, 2025, https://www.pewresearch.org/science/2025/09/17/how-americans-view-ai-and-its-impact-on-people-and-society/
22. Stanford researchers use AI to push the boundaries of what's possible, accessed on September 27, 2025, https://news.stanford.edu/stories/2025/09/ai-applications-campus-robotics-mining-nueroscience-law
23. Reputational pitfalls in the age of artificial intelligence - Edelman, accessed on September 27, 2025, https://www.edelman.com/uk/insights/reputational-pitfalls-artificial-intelligence
24. Technical and ethical debt in the AI fair use crisis - ResearchGate, accessed on September 27, 2025, https://www.researchgate.net/publication/394315158_Technical_and_ethical_debt_in_the_AI_fair_use_crisis
25. Data Colonialism and Data Sets - Harvard Law Review, accessed on September 27, 2025, https://harvardlawreview.org/blog/2023/06/data-colonialism-and-data-sets/
26. The Emergence of Data Licensing Ecosystems for AI Training: Legal Outcomes Shaping the Future | by Trent V. Bolar, Esq. | Medium, accessed on September 27, 2025, https://medium.com/@trentice.bolar/the-emergence-of-data-licensing-ecosystems-for-ai-training-legal-outcomes-shaping-the-future-55e99abed915
27. Copyright Office Weighs In on AI Training and Fair Use | Skadden, Arps, Slate, Meagher & Flom LLP, accessed on September 27, 2025, https://www.skadden.com/insights/publications/2025/05/copyright-office-report
28. Americans Prioritize AI Safety and Data Security - Gallup News, accessed on September 27, 2025, https://news.gallup.com/poll/694685/americans-prioritize-safety-data-security.aspx
29. Survey Highlights an Emerging Divide Over Artificial Intelligence in the U.S.,

accessed on September 27, 2025, https://comminfo.rutgers.edu/news/survey-highlights-emerging-divide-over-artificial-intelligence-us
30. The Reputational Risks of AI | California Management Review, accessed on September 27, 2025, https://cmr.berkeley.edu/2022/01/the-reputational-risks-of-ai/
31. Top 50 AI Scandals [2025] - DigitalDefynd, accessed on September 27, 2025, https://digitaldefynd.com/IQ/top-ai-scandals/
32. Exploring the Financial Impacts of Data Breaches: Direct Costs and Long-Term Effects on Companies' Bottom Lines | Simbo AI, accessed on September 27, 2025, https://www.simbo.ai/blog/exploring-the-financial-impacts-of-data-breaches-direct-costs-and-long-term-effects-on-companies-bottom-lines-3549721/
33. Economic and Financial Consequences of Corporate Cyberattacks | NBER, accessed on September 27, 2025, https://www.nber.org/digest/jun18/economic-and-financial-consequences-corporate-cyberattacks
34. Understanding AI Liability in Marketing: Risks and Mitigation for Small Businesses - Tish.Law, accessed on September 27, 2025, https://tish.law/blog/understanding-ai-liability-in-marketing-risks-and-mitigation-for-small-businesses/
35. Artificial Intelligence in Financial Markets: Systemic Risk and Market ..., accessed on September 27, 2025, https://www.sidley.com/en/insights/newsupdates/2024/12/artificial-intelligence-in-financial-markets-systemic-risk-and-market-abuse-concerns
36. The rise of artificial intelligence: benefits and risks for financial stability, accessed on September 27, 2025, https://www.ecb.europa.eu/press/financial-stability-publications/fsr/special/html/ecb.fsrart202405_02~58c3ce5246.en.html
37. AI Act | Shaping Europe's digital future - European Union, accessed on September 27, 2025, https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai
38. AI Governance Framework, accessed on September 27, 2025, https://ai-governance.eu/
39. Putting AI Ethics into Practice: The Hourglass Model of Organizational AI Governance, accessed on September 27, 2025, https://www.researchgate.net/publication/361023169_Putting_AI_Ethics_into_Practice_The_Hourglass_Model_of_Organizational_AI_Governance
40. Ensuring Trust and Compliance with the Hourglass Model of AI ..., accessed on September 27, 2025, https://aiethix.com/blog/ensuring-trust-and-compliance-with-the-hourglass-model-of-ai-governance
41. The hourglass model - AI Governance Framework, accessed on September 27, 2025, https://ai-governance.eu/ai-governance-framework/the-hourglass-model/
42. EU AI Act: first regulation on artificial intelligence | Topics - European Parliament, accessed on September 27, 2025, https://www.europarl.europa.eu/topics/en/article/20230601STO93804/eu-ai-act-f

irst-regulation-on-artificial-intelligence
43. Article 5: Prohibited AI Practices | EU Artificial Intelligence Act, accessed on September 27, 2025, https://artificialintelligenceact.eu/article/5/
44. Article 10: Data and Data Governance | EU Artificial Intelligence Act, accessed on September 27, 2025, https://artificialintelligenceact.eu/article/10/
45. Key Issue 5: Transparency Obligations - EU AI Act, accessed on September 27, 2025, https://www.euaiact.com/key-issue/5
46. Establishing Data Provenance for Responsible Artificial Intelligence Systems., accessed on September 27, 2025, https://pure.itu.dk/en/publications/establishing-data-provenance-for-responsible-artificial-intellige
47. What is Data Provenance? | IBM, accessed on September 27, 2025, https://www.ibm.com/think/topics/data-provenance
48. Data Provenance in AI | LEARN - Data Foundation, accessed on September 27, 2025, https://datafoundation.org/news/reports/697/697-Data-Provenance-in-AI
49. C2PA | Verifying Media Content Sources, accessed on September 27, 2025, https://c2pa.org/
50. How it works - Content Authenticity Initiative, accessed on September 27, 2025, https://contentauthenticity.org/how-it-works
51. C2PA and Content Credentials Explainer :: C2PA Specifications, accessed on September 27, 2025, https://spec.c2pa.org/specifications/specifications/2.2/explainer/Explainer.html
52. C2PA in ChatGPT Images - OpenAI Help Center, accessed on September 27, 2025, https://help.openai.com/en/articles/8912793-c2pa-in-chatgpt-images
53. AI Output Disclosures: Use, Provenance, Adverse Incidents, accessed on September 27, 2025, https://www.ntia.gov/issues/artificial-intelligence/ai-accountability-policy-report/developing-accountability-inputs-a-deeper-dive/information-flow/ai-output-disclosures
54. AI Watermarking Market | Global Market Analysis Report - 2035, accessed on September 27, 2025, https://www.futuremarketinsights.com/reports/ai-watermarking-market
55. AI Watermarking Market Size & Share | Industry Report, 2033 - Grand View Research, accessed on September 27, 2025, https://www.grandviewresearch.com/industry-analysis/ai-watermarking-market-report
56. AI Watermarking Market Size, Share & Analysis Report, 2034, accessed on September 27, 2025, https://www.gminsights.com/industry-analysis/ai-watermarking-market
57. AI Watermarking Market Size and YoY Growth Rate, 2025-2032, accessed on September 27, 2025, https://www.coherentmarketinsights.com/industry-reports/ai-watermarking-market
58. Adoption of Watermarking Measures for AI-Generated content and Implications under the EU AI Act - arXiv, accessed on September 27, 2025,

https://arxiv.org/html/2503.18156v2
59. The Power of Data Provenance in Building Trustworthy and Compliant Digital Ecosystems, accessed on September 27, 2025, https://www.fiware.org/2024/11/21/the-power-of-data-provenance-in-building-trustworthy-and-compliant-digital-ecosystems/
60. Projects - International Data Spaces, accessed on September 27, 2025, https://internationaldataspaces.org/make/projects/
61. The European AI Office: Driving innovation with open and trusted ..., accessed on September 27, 2025, https://data.europa.eu/en/news-events/news/european-ai-office-driving-innovation-open-and-trusted-data
62. The generative AI data licensing market - Squarespace, accessed on September 27, 2025, https://static1.squarespace.com/static/5cc5785816b6406e50258c5c/t/6787f6e8ea6436205b66940f/1736963826323/The+generative+AI+data+licensing+market.pdf
63. AI Datasets & Licensing For Academic Research And Publishing Market Report 2030, accessed on September 27, 2025, https://www.grandviewresearch.com/industry-analysis/ai-datasets-licensing-academic-research-publishing-market-report
64. Agentic payments: what are they, what are the legal risks and what's next? - Linklaters, accessed on September 27, 2025, https://www.linklaters.com/insights/blogs/fintechlinks/2025/august/agentic-payments
65. What are agentic payments, and what do you need to know? | Very Good Security, accessed on September 27, 2025, https://www.verygoodsecurity.com/blog/posts/what-are-agentic-payments-and-what-do-you-need-to-know
66. Google Agentic Payments Protocol + x402: Agents Can Now Actually Pay Each Other, accessed on September 27, 2025, https://www.coinbase.com/developer-platform/discover/launches/google_x402
67. Cloudflare Launches NET Dollar Stablecoin for AI Agent Micropayments - WebProNews, accessed on September 27, 2025, https://www.webpronews.com/cloudflare-launches-net-dollar-stablecoin-for-ai-agent-micropayments/
68. Coinbase and Cloudflare Will Launch the x402 Foundation: Building ..., accessed on September 27, 2025, https://www.coinbase.com/en-gb/blog/coinbase-and-cloudflare-will-launch-x402-foundation
69. Coinbase and Cloudflare to Launch x402 Foundation to Pioneer AI-Driven Payments, accessed on September 27, 2025, https://www.leaprate.com/cryptocurrency/coinbase-and-cloudflare-launch-x402-foundation-to-pioneer-ai-driven-payments/
70. Digital Assets Exchange Coinbase And Cloudflare To Launch X402 Foundation, accessed on September 27, 2025,

https://www.crowdfundinsider.com/2025/09/252788-digital-assets-exchange-coinbase-and-cloudflare-to-launch-x402-foundation/
71. Coinbase and Cloudflare Launch x402 Foundation for AI Payment Standard - CoinCentral, accessed on September 27, 2025, https://coincentral.com/coinbase-and-cloudflare-launch-x402-foundation-for-ai-payment-standard/
72. Algorithmic fairness: challenges to building an effective regulatory ..., accessed on September 27, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC12426052/
73. Algorithmic Fairness in Lending: How Should We Measure Fairness? - Open Banker, accessed on September 27, 2025, https://openbanker.beehiiv.com/p/algorithmicfairness
74. Algorithmic Fairness: Are computer-aided decisions actually fair? - Boston University, accessed on September 27, 2025, https://www.bu.edu/articles/2018/algorithmic-fairness/
75. When Algorithms Judge Your Credit: Understanding AI Bias in Lending Decisions - Accessible Law - UNT Dallas, accessed on September 27, 2025, https://www.accessiblelaw.untdallas.edu/post/when-algorithms-judge-your-credit-understanding-ai-bias-in-lending-decisions
76. Harvard professor says surveillance capitalism is undermining democracy, accessed on September 27, 2025, https://news.harvard.edu/gazette/story/2019/03/harvard-professor-says-surveillance-capitalism-is-undermining-democracy/
77. Surveillance Capitalism: Why Ethical Data Analytics Matters to Preserve Human Experience, accessed on September 27, 2025, https://chibili.medium.com/the-dark-side-of-data-analytics-surveillance-capitalism-20364ee2e387
78. The Ethics And Implications Of Surveillance Capitalism ... - Consensus, accessed on September 27, 2025, https://consensus.app/questions/ethics-implications-surveillance-capitalism/
79. Unveiling the role of Artificial Intelligence in colonial dynamics ..., accessed on September 27, 2025, https://unipd-centrodirittiumani.it/en/topics/unveiling-the-role-of-artificial-intelligence-in-colonial-dynamics
80. Are emerging technologies colonial? - Faculty of Health Sciences - University of Cape Town, accessed on September 27, 2025, https://health.uct.ac.za/ethics-lab/articles/2023-11-29-recap-are-emerging-technologies-colonial
81. The Threat of Data Colonialism | Berkman Klein Center, accessed on September 27, 2025, https://cyber.harvard.edu/story/2024-10/threat-data-colonialism
82. The Economic Impact of Artificial Intelligence - The Concord Coalition, accessed on September 27, 2025, https://www.concordcoalition.org/facing-the-future/the-economic-impact-of-artificial-intelligence/
83. AI's Promise for the Global Economy - International Monetary Fund (IMF), accessed on September 27, 2025,

https://www.imf.org/en/Publications/fandd/issues/2024/09/AIs-promise-for-the-global-economy-Michael-Spence
84. Beyond the Boom: Understanding AI's True Economic Impact and ..., accessed on September 27, 2025, https://www.fairobserver.com/world-news/beyond-the-boom-understanding-ais-true-economic-impact-and-hidden-financial-danger/
85. Cost of a Data Breach Report 2025 - IBM, accessed on September 27, 2025, https://www.ibm.com/reports/data-breach